

Crawl Report for: <https://docs.ansible.com/>

Content from: <https://docs.ansible.com/>

Ansible Documentation
Ansible documentation
Toggle navigation
Products
Blog
Community
Webinars and training
Try it now
Toggle navigation
Join the community
Users
Developers
Maintainers
Ansible core
Ansible ecosystem
Red Hat Ansible Automation Platform
Ansible community documentation
Ansible offers open-source automation that is simple, flexible, and powerful.
Got thoughts or feedback on this site? We want to hear from you!
Join us in the
Ansible Forum
or open a
GitHub issue
in the docsite repository.
Quicklinks
Ansible package docs home
Collection index
Modules and plugins index
Get started with Ansible
Understand the fundamentals of Ansible automation
Install the Ansible package
Run your first ad hoc command in a few easy steps
Users
Start writing Ansible playbooks
Learn about Ansible modules
Build inventory files to manage multiple hosts
Continue the Ansible user journey
Developers
Set up your development environment
Learn how Ansible works
Write custom modules or plugins
Continue the Ansible developer journey
Maintainers
Review community maintainer responsibilities
Understand Ansible contributor paths
Explore ways to grow community
Continue the Ansible community maintainer journey
CC BY-SA 4.0 |

Crawl Report for: <https://docs.ansible.com/>

Content from: <https://docs.ansible.com/>

Ansible Documentation
Ansible documentation
Toggle navigation
Products
Blog
Community
Webinars and training
Try it now
Toggle navigation
Join the community
Users
Developers
Maintainers
Ansible core
Ansible ecosystem
Red Hat Ansible Automation Platform
Ansible community documentation
Ansible offers open-source automation that is simple, flexible, and powerful.
Got thoughts or feedback on this site? We want to hear from you!
Join us in the
Ansible Forum
or open a
GitHub issue
in the docsite repository.
Quicklinks
Ansible package docs home
Collection index
Modules and plugins index
Get started with Ansible
Understand the fundamentals of Ansible automation
Install the Ansible package
Run your first ad hoc command in a few easy steps
Users
Start writing Ansible playbooks
Learn about Ansible modules
Build inventory files to manage multiple hosts
Continue the Ansible user journey
Developers
Set up your development environment
Learn how Ansible works
Write custom modules or plugins
Continue the Ansible developer journey
Maintainers
Review community maintainer responsibilities
Understand Ansible contributor paths
Explore ways to grow community
Continue the Ansible community maintainer journey
CC BY-SA 4.0 |

Content from: <https://docs.ansible.com/>

[Ansible Documentation](#)

[Ansible documentation](#)

[Toggle navigation](#)

[Products](#)

[Blog](#)

[Community](#)

[Webinars and training](#)

[Try it now](#)

[Toggle navigation](#)

[Join the community](#)

[Users](#)

[Developers](#)

[Maintainers](#)

[Ansible core](#)

[Ansible ecosystem](#)

[Red Hat Ansible Automation Platform](#)

[Ansible community documentation](#)

Ansible offers open-source automation that is simple, flexible, and powerful.

Got thoughts or feedback on this site? We want to hear from you!

[Join us in the](#)

[Ansible Forum](#)

[or open a](#)

[GitHub issue](#)

[in the docsite repository.](#)

[Quicklinks](#)

[Ansible package docs home](#)

[Collection index](#)

[Modules and plugins index](#)

[Get started with Ansible](#)

[Understand the fundamentals of Ansible automation](#)

[Install the Ansible package](#)

[Run your first ad hoc command in a few easy steps](#)

[Users](#)

[Start writing Ansible playbooks](#)

[Learn about Ansible modules](#)

[Build inventory files to manage multiple hosts](#)

[Continue the Ansible user journey](#)

[Developers](#)

[Set up your development environment](#)

[Learn how Ansible works](#)

[Write custom modules or plugins](#)

[Continue the Ansible developer journey](#)

[Maintainers](#)

[Review community maintainer responsibilities](#)

[Understand Ansible contributor paths](#)

[Explore ways to grow community](#)

[Continue the Ansible community maintainer journey](#)

[CC BY-SA 4.0 |](#)

[Privacy policy |](#)

[Sponsored by](#)

Content from: https://docs.ansible.com/ansible-core/2.11_ja/index.html

Ansible Core ?????? ? Ansible Core Documentation
AnsibleFest
Products
Community
Webinars & Training
Blog
Documentation
Ansible
2.11_ja
????????????????????
?????????
Ansible Core ?????
Ansible Core ???
????????
Ansible Core ?????
Ansible ??????????
Ansible ???
??????
????????
?????????????
????????????????????
Playbook ?????
???
Ansible ?????
Ansible ??????: ?????????
YAML ??
Python 3 ?????
?????????????
????????????????
?????????????
????????
???????? (FAQ)
???
Ansible ?????: ??????????????
?????
Red Hat Ansible Automation Platform
Ansible Automation Hub
Ansible ?????????
????????
ansible-core ??????
Ansible
Ansible Core ??????
GitHub ???
You are reading an unmaintained version of the Ansible documentation. Unmaintained Ansible versions can contain unfixed security vulnerabilities (CVE). Please upgrade to a maintained version. See the latest Ansible documentation
.
Ansible Core ??????
?
ansible-core ???
?
Ansible ? IT ??? IT ??? (??)

Ansible ??????????

?????
Ansible ???
??????
????????????????????
????????????
Ansible ???????
??? Ansible ??????????????
????????????
Ansible ??? Python 3
????????
????????????
Windows ??????????
Cisco ACI ??????
????????
Ansible ???
Ansible ??????????????
????????
????????????
ansible-core
???
Ansible ?????????
Python API
????????
????????????????
????????
?????Galaxy ?????????
????? Galaxy ??????
?????Galaxy ?????????
Ansible ??????
?????
????????
????????????????
Playbook ?????
???
Ansible ???
Ansible ??????: ??????
YAML ??
Python 3 ???
????????
????????????
????????
?????
????? (FAQ)
???
Ansible ?????: ?????????
?????
Red Hat Ansible Automation Platform
Ansible Automation Hub
Ansible ??????
?????
ansible-core ?????
Ansible-core 2.11
Ansible-base 2.10
Next

????????????????

Ansible ??????????

?????
Ansible ???
??????
????????????????????
????????????
Ansible ???????
??? Ansible ??????????????
????????????
Ansible ??? Python 3
????????
????????????
Windows ??????????
Cisco ACI ??????
????????
Ansible ???
Ansible ??????????????
????????
????????????
ansible-core
???
Ansible ?????????
Python API
????????
????????????????
????????
?????Galaxy ?????????
?????? Galaxy ??????
Ansible ??????
??????
Collection Index
????????????????
Playbook ?????
???
Ansible ???
Ansible ?????: ??????
YAML ??
Python 3 ???
????????
????????????
????????
?????
????? (FAQ)
???
Ansible ?????: ?????????
?????
Red Hat Ansible Automation Platform
Ansible Automation Hub
Ansible ??????
?????
ansible-core ?????
Ansible-core 2.12
Ansible-core 2.11
Ansible-base 2.10
Next

Content from: https://docs.ansible.com/ansible-core/2.13_ja/index.html

Ansible Core ?????? ? Ansible Core Documentation
AnsibleFest
Products
Community
Webinars & Training
Blog
Documentation
Ansible
2.13_ja
Select version:
2.15_ja
2.14_ja
2.13_ja
Search docs:
Ansible ???????
Ansible ????????
????????????????????
?????????
Ansible Core ?????
Ansible Core ???
??????
Ansible Core ?????
Ansible ??????????
ansible-core ??????????
????????????????
Ansible ??????????????
Ansible ???
??????
Ansible Galaxy
Galaxy ???????
Galaxy ??????
??????
Collection Index
????????????????????
Playbook ?????
???
Ansible ?????
Ansible ??????: ?????????
YAML ??
Python 3 ?????
????????????
????????????
????????????
??????
?????? (FAQ)
???
Ansible ?????: ?????????????
?????
Red Hat Ansible Automation Platform
Ansible Automation Hub
Ansible ???????
??????

Ansible-base 2.10 ?????
Ansible Core ???
???????
??????????? Playbook ???
???????????
???????
Playbook ???
???????????
??
Ansible Core ????
Ansible ??????????
????
????????????
ansible-core ??????????
????????????
Ansible ??????????
Ansible ?????????
????????????????
????????????
?????
Playbook ??????????
???????
Ansible ?????????
????????????????
????????????????
????????????????
GitHub ???
Ansible ?????????????
?????????
reStructuredText ??????
????????????
????
Ansible ???
?????
????????????????????
????????????
?????????
??? Ansible ?????????????
????????????
Ansible ??? Python 3
?????????
????????????
Windows ?????????
Cisco ACI ??????
????????
Ansible ???
Ansible ?????????????
????????
????????????
ansible-core
???
Ansible ?????????????
Python API
?????????

????????????????????
?????????
?????Galaxy ??????????????
?????? Galaxy ??????
Ansible ??????
Ansible Galaxy
Galaxy ??????
Galaxy ?????????
?????????????
Galaxy ??????
Galaxy ?????????????
Galaxy ??????
Galaxy ?????????
Galaxy ??????
??????
Collection Index
????????????????????
Playbook ?????
??
Ansible ???
Ansible ?????: ??????
YAML ??
Python 3 ???
?????????
?????????????
?????????
?????
????? (FAQ)
??
Ansible ????: ??????????
?????
Red Hat Ansible Automation Platform
Ansible Automation Hub
Ansible ??????
?????
ansible-core ?????
Ansible-core 2.13
Ansible-core 2.12
Ansible-core 2.11
Ansible-base 2.10
Next
© Copyright Ansible project contributors.
????: ?? 19, 2023

Content from: https://docs.ansible.com/ansible-core/2.14_ja/index.html

Ansible Core ?????? ? Ansible Core Documentation
AnsibleFest
Products
Community
Webinars & Training
Blog
Documentation
Ansible
2.14_ja
2.15_ja
2.14_ja
2.13_ja
Ansible ??????
MySQL ??????
????????????????????
?????????
Ansible Core ?????
Ansible Core ???
Ansible ?????????
Ansible ?????????????
Ansible Playbook ???
Ansible Vault ?????????????
Ansible ?????????????
Ansible ?????????
Ansible ? Windows ??? BSD ??????
Ansible ??????
Ansible Core ?????
Ansible ?????????
ansible-core ?????????
????????????
Ansible ?????????
Ansible ???
?????
Ansible Galaxy
Galaxy ??????
Galaxy ??????
?????
Collection Index
????????????????????
Playbook ?????
???
Ansible ?????
Ansible ??????: ??????
YAML ??
Python 3 ?????
????????
????????
????????
?????
?????
????? (FAQ)
???
Ansible ?????: ?????????

Ansible ????????

????????????????? Ansible ???????

Ansible ???

Ansible Core ?????

Ansible-core 2.14 ?????

Ansible-core 2.13 ?????

Ansible-core 2.12 ?????

Ansible-core 2.11 ?????

Ansible-base 2.10 ?????

Ansible Core ???

Ansible ??????????

?????????????

?????????????

?????: ??????????????????

?????????

Ansible ??????????????

?????????????

?????????????

Ansible CLI ???????

Ansible Playbook ???

Ansible Playbook

Playbook ???

Playbook ???

??? Playbook ???

??????

Ansible Vault ??????????????

Ansible Vault

Vault ????????

Ansible Vault ??????????????

?????????????????

?????????????????????????????

?????????????????

Ansible Vault ??????????????

Ansible ??????????????????

?????????

?????????????????

?????????

?????????

?????????????????????????

Ansible ??????????

?????????????????

?????????????????

?????????????

?????????

Playbook ??????????????

?????????????????

Ansible ? Windows ??? BSD ??????????

Windows ??????????

Ansible ??? Windows ???

Windows ???????

Desired State Configuration

Windows ????????

Windows ?????????? (FAQ)

Ansible ?????? BSD ???????

Ansible ????????

??????
Playbook ???
?????????
?????
Ansible ???
Ansible Core ???
Ansible ??????
????
?????????
ansible-core ??????
?????????
Ansible ??????
Ansible ??????
?????????
?????????
?????
Playbook ??????
?????
Ansible ??????
?????????
?????????
?????????
?????????
GitHub ??
Ansible ??????
?????????
reStructuredText ??????
?????????
????
Ansible ??
?????
?????????
?????????
??? Ansible ??????
?????????
Ansible ??? Python 3
?????????
?????????
?? YAML ??????
Windows ??????
?????????
Ansible ???
Ansible ??????
?????????
?????????
ansible-core
??
Ansible ??????
Python API
?????????
?????????
?????????
?????Galaxy ??????
????? Galaxy ??????

Ansible ???????
Ansible Galaxy
Galaxy ???????
Galaxy ?????????
?????????????
Galaxy ???????
Galaxy ??????????
Galaxy ??????
Galaxy ?????????
Galaxy ??????
??????
Collection Index
????????????????????
Playbook ?????
???
Ansible ????
Ansible ??????: ???????
YAML ??
Python 3 ????
????????????
?????????????
????????????
??????
?????? (FAQ)
???
Ansible ?????: ??????????
?????
Red Hat Ansible Automation Platform
Ansible Automation Hub
Ansible ??????
?????
ansible-core ??????
Ansible-core 2.14
Ansible-core 2.13
Ansible-core 2.12
Ansible-core 2.11
Ansible-base 2.10
Next
© Copyright Ansible project contributors.
????: ?? 19, 2023

Content from: https://docs.ansible.com/ansible-core/2.15_ja/index.html

Ansible Core ?????? ? Ansible Core Documentation

AnsibleFest

Products

Community

Webinars & Training

Blog

Documentation

Ansible

2.15_ja

Select version:

2.15_ja

2.14_ja

2.13_ja

Search docs:

Ansible getting started

MySQL ????????

????????????????????

??????????

Ansible Core ?????

Ansible Core ???

Ansible ??????????

Ansible ?????????????

Ansible Playbook ???

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Ansible ?????????

Ansible ? Windows ??? BSD ????????

Ansible tips and tricks

Ansible Core ?????

Ansible ?????????

ansible-core ?????????????

????????????????

Ansible ?????????????

Ansible ???

??????

Ansible Galaxy

Galaxy ???????

Galaxy ??????

??????

Collection Index

????????????????????

Playbook ?????

???

Ansible ?????

Ansible ??????: ????????

YAML ??

Python 3 ?????

????????????

????????????????

????????????

??????

?????? (FAQ)

???

Ansible ????: ?????????????

?????

Red Hat Ansible Automation Platform

Ansible Automation Hub

Ansible ???????

??????

ansible-core ??????

Ansible

Ansible Core ??????

Ansible Core ??????

?

ansible-core ???

?

Ansible ? IT ??? IT ??? (???)
????????????????????????????????

Ansible core, or
ansible-core

is the main building block and architecture for Ansible, and includes:

CLI tools such as
ansible-playbook

,

ansible-doc

. and others for driving and interacting with automation.

The Ansible language that uses YAML to create a set of rules for developing Ansible Playbooks and includes functions
such as conditionals, blocks, includes, loops, and other Ansible imperatives.

An architectural framework that allows extensions through Ansible collections.

Ansible ??? OpenSSH ???
(??)
??IT ??Ansible
??
?????
AnsibleFest
(Red Hat ??????Ansible ???) ??????????????AnsibleFest
??
Ansible
??OpenS
SH ??? 1 ??Ansible ???????????????? OS
???Ansible ??Kerberos?LDAP????????????????????????????????????
This documentation covers the version of
ansible-core
noted in the upper left corner of this page. We maintain multiple versions of
ansible-core
and of the documentation, so please be sure you are using the version of the documentation that covers the version of
Ansible you're using. For recent features, we note the version of Ansible where the feature was added.
ansible-core
?????? 2 ??? 2.10
??
Ansible getting started
MySQL ??????????
?????????????
Playbook ???
Ansible ???
????????????????????????????????

?????????
Ansible ??????
????????????????? Ansible ??????
Ansible ???
Ansible Core ?????
Ansible-core 2.15 Porting Guide
Ansible-core 2.14 ?????
Ansible-core 2.13 ?????
Ansible-core 2.12 ?????
Ansible-core 2.11 ?????
Ansible-base 2.10 ?????
Ansible Core ???
Ansible ??????????
?????????????
?????????????
?????: ??????????????????
?????????
Ansible ??????????????
?????????????
?????????????
Ansible CLI ??????
Ansible Playbook ???
Ansible Playbook
Playbook ???
Playbook ???
??? Playbook ???
?????
Protecting sensitive data with Ansible vault
Ansible Vault
Managing vault passwords
Encrypting content with Ansible Vault
Using encrypted variables and files
Configuring defaults for using encrypted content
When are encrypted files made visible?
Format of files encrypted with Ansible Vault
Using Ansible modules and plugins
Introduction to modules
Module maintenance and support
Rejecting modules
?????????
Modules and plugins index
Ansible ??????????
?????????????????
?????????????????
?????????????????
?????????????
?????????????
Playbook ??????????????
?????????????????
Ansible ? Windows ??? BSD ??????????
Windows ??????????????
Ansible ??? Windows ???
Windows ?????????
Desired State Configuration
Windows ?????????

Windows ?????????? (FAQ)
Ansible ????? BSD ?????
Ansible tips and tricks
General tips
Playbook tips
Inventory tips
Execution tricks
Sample Ansible setup
Ansible Core ???
Ansible ??????????
????
????????????
ansible-core ?????????????
????????????
Ansible ??????????
Ansible ??????????
????????????????
Ansible ??????????
????????????????
????????????????
????????????????
????????????????
GitHub ???
Ansible ?????????????
?????????
reStructuredText ??????
????????????????
????
Ansible ???
?????
????????????????????
????????????
?????????
??? Ansible ?????????????
????????????
Ansible ??? Python 3
?????????
????????????????
?? YAML ??????????
Windows ?????????????
????????????
Ansible ???
Ansible ?????????????????
?????????
????????????
ansible-core
???
Ansible ?????????????
Python API
?????????
????????????????
?????????
?????Galaxy ?????????????
????? Galaxy ??????
Ansible ????????

Ansible Galaxy
Galaxy ???????
Galaxy ??????????
?????????????
Galaxy ????????
Galaxy ?????????????
Galaxy ??????
Galaxy ??????????
Galaxy ??????
???????
Collection Index
?????????????????????????
Playbook ??????
???
Ansible ?????
Ansible ??????: ????????
YAML ??
Python 3 ?????
?????????????
?????????????
?????????????
?????????????
???????
?????? (FAQ)
???
Ansible ?????: ??????????????
?????
Red Hat Ansible Automation Platform
Ansible Automation Hub
Ansible ???????
??????
ansible-core ??????
Ansible-core 2.15
Ansible-core 2.14
Ansible-core 2.13
Ansible-core 2.12
Ansible-core 2.11
Ansible-base 2.10
Next
© Copyright Ansible project contributors.
?????: ?? 25, 2023

Content from: https://docs.ansible.com/ansible-core/devel/dev_guide/testing_running_locally.html

Testing Ansible and Collections ? Ansible Core Documentation

Blog

Ansible community forum

Documentation

Ansible Core Documentation

Ansible Core

Select version:

2.19

2.18

2.17

devel

Search docs:

Ansible getting started

Getting started with Ansible

Installation, Upgrade & Configuration

Installation Guide

Ansible Core Porting Guides

Using Ansible Core

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible Core

Ansible Community Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery

Releases and maintenance

Testing Strategies

Sanity Tests

Frequently Asked Questions

Glossary

Ansible Reference: Module Utilities

Special Variables

Red Hat Ansible Automation Platform

Ansible Automation Hub

Logging Ansible output

Roadmaps

ansible-core Roadmaps

Ansible Core

Testing Ansible and Collections

Edit on GitHub

Testing Ansible and Collections

?

This document describes how to run tests using
ansible-test

.

Setup

Testing an Ansible Collection

Testing

ansible-core

Commands

Environments

Containers

Choosing a container

Custom containers

Docker and SELinux

Docker Desktop with WSL2

Configuration requirements

Setup instructions

Linux cgroup configuration

Podman

Remote virtual machines

Python virtual environments

Composite environment arguments

Additional Requirements

Environment variables

Interactive shell

Code coverage

Setup

?

Before running

ansible-test

, set up your environment for

Testing an Ansible Collection

or

Testing ansible-core

, depending on which scenario applies to you.

Warning

If you use

git

for version control, make sure the files you are working with are not ignored by

git

.

If they are,

ansible-test

will ignore them as well.

Testing an Ansible Collection

?

If you are testing an Ansible Collection, you need a copy of the collection, preferably a Git clone.

For example, to work with the
community.windows

collection, follow these steps:

Clone the collection you want to test into a valid collection root:

```
git
```

```
clone
```

```
https://github.com/ansible-collections/community.windows
```

```
~/dev/ansible_collections/community/windows
```

Important

The path must end with

```
/ansible_collections/{collection_namespace}/{collection_name}
```

where

```
{collection_namespace}
```

is the namespace of the collection and

```
{collection_name}
```

is the collection name.

Clone any collections on which the collection depends:

```
git
```

```
clone
```

```
https://github.com/ansible-collections/ansible.windows
```

```
~/dev/ansible_collections/ansible/windows
```

Important

If your collection has any dependencies on other collections, they must be in the same collection root, since
ansible-test

will not use your configured collection roots (or other Ansible configuration).

Note

See the collection's

```
galaxy.yml
```

for a list of possible dependencies.

Switch to the directory where the collection to test resides:

```
cd
```

```
~/dev/ansible_collections/community/windows
```

Testing

```
ansible-core
```

?

If you are testing

```
ansible-core
```

itself, you need a copy of the

```
ansible-core
```

source code, preferably a Git clone.

Having an installed copy of

```
ansible-core
```

is not sufficient or required.

For example, to work with the

```
ansible-core
```

source cloned from GitHub, follow these steps:

Clone the

```
ansible-core
```

repository:

```
git
```

```
clone
```

<https://github.com/ansible/ansible>
~/dev/ansible

Switch to the directory where the
ansible-core
source resides:

```
cd  
~/dev/ansible
```

Add
ansible-core
programs to your
PATH

```
:  
source  
hacking/env-setup
```

Note

You can skip this step if you only need to run
ansible-test
, and not other
ansible-core
programs.

In that case, simply run
bin/ansible-test
from the root of the
ansible-core
source.

Caution

If you have an installed version of
ansible-core
and are trying to run
ansible-test
from your
PATH

,
make sure the program found by your shell is the one from the
ansible-core
source:

which
ansible-test

Commands

?

The most commonly used test commands are:

ansible-test

sanity

- Run sanity tests (mostly linters and static analysis).

ansible-test

integration

- Run integration tests.

ansible-test

units

- Run unit tests.

Run

ansible-test

--help

to see a complete list of available commands.

Note

For detailed help on a specific command, add the

`--help`

option after the command.

Environments

?

Most

`ansible-test`

commands support running in one or more isolated test environments to simplify testing.

Containers

?

Containers are recommended for running sanity, unit and integration tests, since they provide consistent environments.

Unit tests will be run with network isolation, which avoids unintentional dependencies on network resources.

The

`--docker`

option runs tests in a container using either Docker or Podman.

Note

If both Docker and Podman are installed, Docker will be used.

To override this, set the environment variable

`ANSIBLE_TEST_PREFER_PODMAN`

to any non-empty value.

Choosing a container

?

Without an additional argument, the

`--docker`

option uses the

default

container.

To use another container, specify it immediately after the

`--docker`

option.

Note

The

default

container is recommended for all sanity and unit tests.

To see the list of supported containers, use the

`--help`

option with the

`ansible-test`

command you want to use.

Note

The list of available containers is dependent on the

`ansible-test`

command you are using.

You can also specify your own container.

When doing so, you will need to indicate the Python version in the container with the

`--python`

option.

Custom containers

?

When building custom containers, keep in mind the following requirements:

The

`USER`

should be

root

.

Use an

init

process, such as

systemd

.

Include

sshd

and accept connections on the default port of

22

.

Include a POSIX compatible

sh

shell which can be found on

PATH

.

Include a

sleep

utility which runs as a subprocess.

Include a supported version of Python.

Avoid using the

VOLUME

statement.

Docker and SELinux

?

Using Docker on a host with SELinux may require setting the system in permissive mode.

Consider using Podman instead.

Docker Desktop with WSL2

?

These instructions explain how to use

ansible-test

with WSL2 and Docker Desktop

without

systemd

support.

Note

If your WSL2 environment includes

systemd

support, these steps are not required.

Configuration requirements

?

Open Docker Desktop and go to the

Settings

screen.

On the the

General

tab:

Uncheck the

Start Docker Desktop when you log in

checkbox.

Check the

Use the WSL 2 based engine

checkbox.

On the
Resources
tab under the
WSL Integration
section:

Enable distros you want to use under the
Enable integration with additional distros
section.

Click
Apply and restart
if changes were made.

Setup instructions
?

Note
If all WSL instances have been stopped, these changes will need to be re-applied.
Verify Docker Desktop is properly configured (see
Configuration requirements
).

Quit Docker Desktop if it is running:

Right click the
Docker Desktop
taskbar icon.

Click the
Quit Docker Desktop
option.

Stop any running WSL instances with the command:
wsl

--shutdown

Verify all WSL instances have stopped with the command:
wsl

-l

-v

Start a WSL instance and perform the following steps as
root

:

Verify the
systemd
subsystem is not registered:

Check for the
systemd
cgroup hierarchy with the following command:

grep
systemd
/proc/self/cgroup

If any matches are found, re-check the
Configuration requirements
and follow the
Setup instructions
again.

Mount the
systemd
cgroup hierarchy with the following commands:

mkdir
/sys/fs/cgroup/systemd

```
mount
cgroup
-t
cgroup
/sys/fs/cgroup/systemd
-o
none,name
=
```

```
systemd,xattr
```

Start Docker Desktop.

You should now be able to use

```
ansible-test
```

with the

```
--docker
```

option.

Linux cgroup configuration

?

Note

These changes will need to be re-applied each time the container host is booted.

For certain container hosts and container combinations, additional setup on the container host may be required.

In these situations

```
ansible-test
```

will report an error and provide additional instructions to run as

```
root
```

```
:
```

```
mkdir
```

```
/sys/fs/cgroup/systemd
```

```
mount
```

```
cgroup
```

```
-t
```

```
cgroup
```

```
/sys/fs/cgroup/systemd
```

```
-o
```

```
none,name
```

```
=
```

```
systemd,xattr
```

If you are using rootless Podman, an additional command must be run, also as

```
root
```

```
.
```

Make sure to substitute your user and group for

```
{user}
```

and

```
{group}
```

respectively:

```
chown
```

```
-R
```

```
{
```

```
user
```

```
}
```

```
:
```

```
{
```

```
group
```

```
}
```

```
/sys/fs/cgroup/systemd
```

Podman

?

When using Podman, you may need to stop existing Podman processes after following the Linux cgroup configuration

instructions. Otherwise Podman may be unable to see the new mount point.

You can check to see if Podman is running by looking for

podman

and

catatonit

processes.

Remote virtual machines

?

Remote virtual machines are recommended for running integration tests not suitable for execution in containers.

The

--remote

option runs tests in a cloud hosted ephemeral virtual machine.

Note

An API key is required to use this feature, unless running under an approved Azure Pipelines organization.

To see the list of supported systems, use the

--help

option with the

ansible-test

command you want to use.

Note

The list of available systems is dependent on the

ansible-test

command you are using.

Python virtual environments

?

Python virtual environments provide a simple way to achieve isolation from the system and user Python environments.

They are recommended for unit and integration tests when the

--docker

and

--remote

options cannot be used.

The

--venv

option runs tests in a virtual environment managed by

ansible-test

.

Requirements are automatically installed before tests are run.

Composite environment arguments

?

The environment arguments covered in this document are sufficient for most use cases.

However, some scenarios may require the additional flexibility offered by composite environment arguments.

The

--controller

and

--target

options are alternatives to the

--docker

,

--remote

and

--venv
options.

Note

When using the
shell

command, the

--target

option is replaced by three platform specific options.

Add the

--help

option to your

ansible-test

command to learn more about the composite environment arguments.

Additional Requirements

?

Some

ansible-test

commands have additional requirements.

You can use the

--requirements

option to automatically install them.

Note

When using a test environment managed by

ansible-test

the

--requirements

option is usually unnecessary.

Environment variables

?

When using environment variables to manipulate tests there some limitations to keep in mind. Environment variables are:

Not propagated from the host to the test environment when using the

--docker

or

--remote

options.

Not exposed to the test environment unless enabled in

test/lib/ansible_test/_internal/util.py

in the

common_environment

function.

Example:

ANSIBLE_KEEP_REMOTE_FILES=1

can be set when running

ansible-test

integration

--venv

. However, using the

--docker

option would

require running

ansible-test

shell

to gain access to the Docker environment. Once at the shell prompt, the environment variable could be set

and the tests executed. This is useful for debugging tests inside a container by following the Debugging modules

instructions.

Interactive shell

?

Use the

ansible-test

shell

command to get an interactive shell in the same environment used to run tests. Examples:

ansible-test

shell

--docker

- Open a shell in the default docker container.

ansible-test

shell

--venv

--python

3.10

- Open a shell in a Python 3.10 virtual environment.

Code coverage

?

Code coverage reports make it easy to identify untested code for which more tests should be written. Online reports are available but only cover the

devel

branch (see

Testing Ansible

). For new code local reports are needed.

Add the

--coverage

option to any test command to collect code coverage data. If you

aren't using the

--venv

or

--docker

options which create an isolated python

environment then you may have to use the

--requirements

option to ensure that the

correct version of the coverage module is installed:

ansible-test

coverage

erase

ansible-test

units

--coverage

apt

ansible-test

integration

--coverage

aws_lambda

ansible-test

coverage

html

Reports can be generated in several different formats:

ansible-test

coverage

report

- Console report.

ansible-test

coverage

html

- HTML report.

ansible-test

coverage

xml

- XML report.

To clear data between test runs, use the

ansible-test

coverage

erase

command.

© Copyright Ansible project contributors.

Last updated on Oct 19, 2025.

Content from: <https://docs.ansible.com/ansible-core/devel/index.html>

[Ansible Core Documentation](#) ? [Ansible Core Documentation](#)

[Blog](#)

[Ansible community forum](#)

[Documentation](#)

[Ansible Core Documentation](#)

[Ansible Core](#)

Select version:

[2.19](#)

[2.18](#)

[2.17](#)

[devel](#)

Search docs:

[Ansible getting started](#)

[Getting started with Ansible](#)

[Installation, Upgrade & Configuration](#)

[Installation Guide](#)

[Ansible Core Porting Guides](#)

[Using Ansible Core](#)

[Building Ansible inventories](#)

[Using Ansible command line tools](#)

[Using Ansible playbooks](#)

[Protecting sensitive data with Ansible vault](#)

[Using Ansible modules and plugins](#)

[Using Ansible collections](#)

[Using Ansible on Windows, BSD, and z/OS UNIX](#)

[Ansible tips and tricks](#)

[Contributing to Ansible Core](#)

[Ansible Community Guide](#)

[ansible-core Contributors Guide](#)

[Advanced Contributor Guide](#)

[Ansible documentation style guide](#)

[Extending Ansible](#)

[Developer Guide](#)

[Ansible Galaxy](#)

[Galaxy User Guide](#)

[Galaxy Developer Guide](#)

[Reference & Appendices](#)

[Collection Index](#)

[Indexes of all modules and plugins](#)

[Playbook Keywords](#)

[Return Values](#)

[Ansible Configuration Settings](#)

[Controlling how Ansible behaves: precedence rules](#)

[YAML Syntax](#)

[Python 3 Support](#)

[Interpreter Discovery](#)

[Releases and maintenance](#)

[Testing Strategies](#)

[Sanity Tests](#)

[Frequently Asked Questions](#)

[Glossary](#)

[Ansible Reference: Module Utilities](#)

Special Variables

Red Hat Ansible Automation Platform

Ansible Automation Hub

Logging Ansible output

Roadmaps

ansible-core Roadmaps

Ansible Core

Ansible Core Documentation

Ansible Core Documentation

?

Ansible Core, or

ansible-core

is the main building block and architecture for Ansible, and includes:

CLI tools such as

ansible-playbook

,

ansible-doc

. and others for driving and interacting with automation.

The Ansible language that uses YAML to create a set of rules for developing Ansible Playbooks and includes functions such as conditionals, blocks, includes, loops, and other Ansible imperatives.

An architectural framework that allows extensions through Ansible collections.

This documentation covers the version of

ansible-core

noted in the upper left corner of this page.

We maintain multiple versions of

ansible-core

and of the documentation, so please be sure you are using the version of the documentation that covers the version of Ansible you're using.

For recent features, we note the version of Ansible where the feature was added.

ansible-core

releases a new major release approximately twice a year.

The core application evolves somewhat conservatively, valuing simplicity in language design and setup.

Contributors develop and change modules and plugins, hosted in collections, much more quickly.

Ansible getting started

Getting started with Ansible

Introduction to Ansible

Start automating with Ansible

Building an inventory

Creating a playbook

Ansible concepts

Installation, Upgrade & Configuration

Installation Guide

Installing Ansible

Installing Ansible on specific operating systems

Configuring Ansible

Ansible Core Porting Guides

Ansible-core 2.19 Porting Guide

Ansible-core 2.18 Porting Guide

Ansible-core 2.17 Porting Guide

Ansible-core 2.16 Porting Guide

Ansible-core 2.15 Porting Guide

Ansible-core 2.14 Porting Guide

Ansible-core 2.13 Porting Guide

Ansible-core 2.12 Porting Guide

- Ansible-core 2.11 Porting Guide
- Ansible-base 2.10 Porting Guide
- Using Ansible Core
- Building Ansible inventories
 - How to build your inventory
 - Working with dynamic inventory
- Patterns: targeting hosts and groups
- Connection methods and details
- Using Ansible command line tools
- Introduction to ad hoc commands
- Working with command line tools
- Ansible CLI cheatsheet
- Using Ansible playbooks
 - Ansible playbooks
 - Working with playbooks
 - Executing playbooks
 - Advanced playbook syntax
- Manipulating data
 - Protecting sensitive data with Ansible vault
- Ansible Vault
 - Managing vault passwords
 - Encrypting content with Ansible Vault
 - Using encrypted variables and files
 - Configuring defaults for using encrypted content
 - When are encrypted files made visible?
 - Format of files encrypted with Ansible Vault
- Using Ansible modules and plugins
 - Introduction to modules
 - Boolean variables
 - Module maintenance and support
 - Rejecting modules
 - Working with plugins
 - Modules and plugins index
- Using Ansible collections
 - Installing collections
 - Removing a collection
 - Downloading collections
 - Listing collections
 - Verifying collections
 - Using collections in a playbook
 - Collections index
- Using Ansible on Windows, BSD, and z/OS UNIX
 - Managing BSD hosts with Ansible
 - Managing Windows hosts with Ansible
 - Managing z/OS UNIX hosts with Ansible
- Ansible tips and tricks
 - General tips
 - Playbook tips
 - Inventory tips
 - Execution tricks
- Sample Ansible setup
- Contributing to Ansible Core
- Ansible Community Guide
- Getting started

Contributor path
ansible-core Contributors Guide
Reporting bugs and requesting features
Contributing to the Ansible Documentation
The Ansible Development Cycle
Other Tools and Programs
Working with the Ansible repo
Advanced Contributor Guide
Committers Guidelines
Release Manager Guidelines
GitHub Admins
Ansible Ecosystem Project Development Resources
Ansible documentation style guide
Linguistic guidelines
reStructuredText guidelines
Markdown guidelines
Accessibility guidelines
More resources
Extending Ansible
Developer Guide
Adding modules and plugins locally
Should you develop a module?
Developing modules
Contributing your module to an existing Ansible collection
Conventions, tips, and pitfalls
Ansible and Python 3
Debugging modules
Module format and documentation
Ansible markup
Adjacent YAML documentation files
Windows module development walkthrough
Creating a new collection
Testing Ansible
The lifecycle of an Ansible module or plugin
Developing plugins
Developing dynamic inventory
Developing
ansible-core
Ansible module architecture
Python API
Rebasing a pull request
Using and developing module utilities
Ansible collection creator path
Developing collections
Migrating Roles to Roles in Collections on Galaxy
Collection Galaxy metadata structure
Ansible architecture
Ansible Galaxy
Galaxy User Guide
Finding collections on Galaxy
Finding roles on Galaxy
Installing roles from Galaxy
Galaxy Developer Guide
Creating collections for Galaxy

[Creating roles for Galaxy](#)
[Reference & Appendices](#)
[Collection Index](#)
[Indexes of all modules and plugins](#)
[Playbook Keywords](#)
[Return Values](#)
[Ansible Configuration Settings](#)
[Controlling how Ansible behaves: precedence rules](#)
[YAML Syntax](#)
[Python 3 Support](#)
[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[ansible-core Roadmaps](#)
[Ansible-core 2.20](#)
[Ansible-core 2.19](#)
[Ansible-core 2.18](#)
[Ansible-core 2.17](#)
[Ansible-core 2.16](#)
[Ansible-core 2.15](#)
[Ansible-core 2.14](#)
[Ansible-core 2.13](#)
[Ansible-core 2.12](#)
[Ansible-core 2.11](#)
[Ansible-base 2.10](#)
[Next](#)
[© Copyright Ansible project contributors.](#)
[Last updated on Oct 19, 2025.](#)

Content from: <https://docs.ansible.com/ansible-prior-versions.html>

[Ansible community documentation archive](#) | [Ansible Documentation](#)

[Ansible documentation](#)

[Toggle navigation](#)

[Products](#)

[Blog](#)

[Community](#)

[Webinars and training](#)

[Try it now](#)

[Toggle navigation](#)

[Join the community](#)

[Users](#)

[Developers](#)

[Maintainers](#)

[Ansible core](#)

[Ansible ecosystem](#)

[Red Hat Ansible Automation Platform](#)

[Ansible community documentation archive](#)

Archive page for the Ansible community documentation that provides older versions of content.

Got thoughts or feedback on this site? We want to hear from you!

[Join us in the](#)

[Ansible Forum](#)

[or open a](#)

[GitHub issue](#)

[in the docsite repository.](#)

[Ansible community documentation for version 2.7](#)

[Visit the documentation](#)

[Ansible community documentation for version 2.6](#)

[Visit the documentation](#)

[Ansible community documentation for version 2.5](#)

[Visit the documentation](#)

[Ansible community documentation for version 2.4](#)

[Visit the documentation](#)

[Ansible community documentation for version 2.3](#)

[Visit the documentation](#)

[CC BY-SA 4.0](#) |

[Privacy policy](#) |

[Sponsored by](#)

Content from: <https://docs.ansible.com/ansible/2.9/plugins/inventory/virtualbox.html>

virtualbox ? virtualbox inventory source ? Ansible Documentation

AnsibleFest

Products

Community

Webinars & Training

Blog

Documentation

Ansible

2.9

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

User Guide

Ansible Quickstart Guide

Ansible concepts

Getting Started

How to build your inventory

Working with dynamic inventory

Patterns: targeting hosts and groups

Introduction to ad-hoc commands

Connection methods and details

Working with command line tools

Working With Playbooks

Understanding privilege escalation: become

Ansible Vault

Working With Modules

Working With Plugins

Action Plugins

Become Plugins

Cache Plugins

Callback Plugins

Cliconf Plugins

Connection Plugins

Httpapi Plugins

Inventory Plugins

Enabling inventory plugins

Using inventory plugins

Plugin List

Lookup Plugins

Netconf Plugins

Shell Plugins

Strategy Plugins

Vars Plugins

Filters

Tests

Plugin Filter Configuration

Ansible and BSD

Windows Guides

Using collections

Contributing to Ansible

Ansible Community Guide

Extending Ansible
Developer Guide
Common Ansible Scenarios
Public Cloud Guides
Network Technology Guides
Virtualization and Containerization Guides
Ansible for Network Automation
Ansible for Network Automation
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Module Index
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Release and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Tower
Logging Ansible output
Roadmaps
Ansible Roadmap
Ansible
Docs

»

User Guide

»

Working With Playbooks

»

Advanced Playbooks Features

»

Working With Plugins

»

Inventory Plugins

»

virtualbox ? virtualbox inventory source

Edit on GitHub

For community users, you are reading an unmaintained version of the Ansible documentation. Unmaintained Ansible versions can contain unfixed security vulnerabilities (CVE). Please upgrade to a maintained version. See the latest Ansible community documentation

. For Red Hat customers, see the Red Hat AAP platform lifecycle

.

virtualbox ? virtualbox inventory source

¶

Synopsis

Parameters

Examples

Status

Synopsis

¶

Get inventory hosts from the local virtualbox installation.

Uses a YAML configuration file that ends with virtualbox.(yml|yaml) or vbox.(yml|yaml).

The inventory_hostname is always the ?Name? of the virtualbox instance.

Parameters

¶

Parameter

Choices/

Defaults

Configuration

Comments

cache

boolean

Choices:

no

?

yes

ini entries:

[inventory]

cache = no

env:ANSIBLE_INVENTORY_CACHE

Toggle to enable/disable the caching of the inventory's source data, requires a cache plugin setup to work.

cache_connection

string

ini entries:

[defaults]

fact_caching_connection = VALUE

[inventory]

cache_connection = VALUE

env:ANSIBLE_CACHE_PLUGIN_CONNECTION

env:ANSIBLE_INVENTORY_CACHE_CONNECTION

Cache connection data or path, read cache plugin documentation for specifics.

cache_plugin

string

Default:

"memory"

ini entries:

[defaults]

fact_caching = memory

[inventory]

cache_plugin = memory

env:ANSIBLE_CACHE_PLUGIN

env:ANSIBLE_INVENTORY_CACHE_PLUGIN

Cache plugin to use for the inventory's source data.

cache_prefix

-

Default:

"ansible_inventory_"

ini entries:

[default]
fact_caching_prefix = ansible_inventory_
[inventory]
cache_prefix = ansible_inventory_
env:ANSIBLE_CACHE_PLUGIN_PREFIX
env:ANSIBLE_INVENTORY_CACHE_PLUGIN_PREFIX
Prefix to use for cache plugin files/tables
cache_timeout
integer
Default:
3600
ini entries:
[defaults]
fact_caching_timeout = 3600
[inventory]
cache_timeout = 3600
env:ANSIBLE_CACHE_PLUGIN_TIMEOUT
env:ANSIBLE_INVENTORY_CACHE_TIMEOUT
Cache duration in seconds
compose
dictionary
Default:
{}
Create vars from jinja2 expressions.
groups
dictionary
Default:
{}
Add hosts to group based on Jinja2 conditionals.
keyed_groups
list
Default:
[]
Add hosts to group based on the values of a variable.
network_info_path
-
Default:
"/VirtualBox/GuestInfo/Net/0/V4/IP"
property path to query for network information (ansible_host)
plugin
-
/
required
Choices:
virtualbox
token that ensures this is a source file for the 'virtualbox' plugin
query
dictionary
Default:
{}
create vars from virtualbox properties
running_only
boolean
Default:

"no"

toggles showing all vms vs only those currently running

settings_password_file

-

provide a file containing the settings password (equivalent to --settingspwfile)

strict

boolean

Choices:

no

?

yes

If

yes

make invalid entries a fatal error, otherwise skip and continue.

Since it is possible to use facts in the expressions they might not always be available and we ignore those errors by default.

Examples

¶

file must be named vbox.yaml or vbox.yml

simple_config_file

:

plugin

:

virtualbox

settings_password_file

:

/etc/virtualbox/secrets

query

:

logged_in_users

:

/VirtualBox/GuestInfo/OS/LoggedInUsersList

compose

:

ansible_connection

:

('indows' in vbox_Guest_OS)|ternary('winrm', 'ssh')

add hosts (all match with minishift vm) to the group container if any of the vms are in ansible_inventory'

plugin

:

virtualbox

groups

:

container

:

"minis"

in

(inventory_hostname)"

Status

¶

This inventory is not guaranteed to have a backwards compatible interface.

[preview]

This inventory is

maintained by the Ansible Community

.

[community]

Authors



UNKNOWN

Hint

If you notice any issues in this documentation, you can
edit this document
to improve it.

Hint

Configuration entries for each entry type have a low to high priority order. For example, a variable that is lower in the list
will override a variable that is higher up.

Next

Previous

© Copyright 2019 Red Hat, Inc.

Last updated on May 27, 2022.

Content from: https://docs.ansible.com/ansible/2.9/reference_appendices/module_utils.html

Ansible Reference: Module Utilities ? Ansible Documentation

AnsibleFest

Products

Community

Webinars & Training

Blog

Documentation

Ansible

2.9

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

User Guide

Contributing to Ansible

Ansible Community Guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Public Cloud Guides

Network Technology Guides

Virtualization and Containerization Guides

Ansible for Network Automation

Ansible for Network Automation

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Module Index

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery

Release and maintenance

Testing Strategies

Sanity Tests

Frequently Asked Questions

Glossary

Ansible Reference: Module Utilities

AnsibleModule

Basic

Special Variables

Red Hat Ansible Tower

Logging Ansible output

Roadmaps

Ansible Roadmap

Ansible

Docs

»

Ansible Reference: Module Utilities

[Edit on GitHub](#)

For community users, you are reading an unmaintained version of the Ansible documentation. Unmaintained Ansible versions can contain unfixed security vulnerabilities (CVE). Please upgrade to a maintained version. See the latest Ansible community documentation

. For Red Hat customers, see the Red Hat AAP platform lifecycle

.

Ansible Reference: Module Utilities



This page documents utilities intended to be helpful when writing Ansible modules in Python.

AnsibleModule



To use this functionality, include

from

ansible.module_utils.basic

import

AnsibleModule

in your module.

class

ansible.module_utils.basic.

AnsibleModule

(

argument_spec

,

bypass_checks

=

False

,

no_log

=

False

,

check_invalid_arguments

=

None

,

mutually_exclusive

=

None

,

required_together

=

None

,

required_one_of

=

None

,

add_file_common_args

=

False

,

```
supports_check_mode
```

```
=
```

```
False
```

```
,
```

```
required_if
```

```
=
```

```
None
```

```
,
```

```
required_by
```

```
=
```

```
None
```

```
)
```

```
¶
```

Common code for quickly building an ansible module in Python

(although you can write modules with anything that can return JSON).

See

[Ansible module development: getting started](#)

for a general introduction

and

[Ansible module architecture](#)

for more detailed explanation.

```
add_path_info
```

```
(
```

```
kwargs
```

```
)
```

```
¶
```

for results that are files, supplement the info about the file

in the return path with stats about the file path.

```
atomic_move
```

```
(
```

```
src
```

```
,
```

```
dest
```

```
,
```

```
unsafe_writes
```

```
=
```

```
False
```

```
)
```

```
¶
```

atomically move src to dest, copying attributes from dest, returns true on success

it uses `os.rename` to ensure this as it is an atomic operation, rest of the function is

to work around limitations, corner cases and ensure selinux context is saved if possible

```
backup_local
```

```
(
```

```
fn
```

```
)
```

```
¶
```

make a date-marked backup of the specified file, return True or False on success or failure

boolean

```
(
```

```
arg
```

```
)
```

```
¶
```

Convert the argument to a boolean

digest_from_file

```
(  
filename  
,  
algorithm  
)
```

¶

Return hex digest of local file for a digest_method specified by name, or None if file is not present.

exit_json

```
(  
**  
kwargs  
)
```

¶

return from the module, without error

fail_json

```
(  
**  
kwargs  
)
```

¶

return from the module, with an error message

get_bin_path

```
(  
arg  
,  
required  
=  
False  
,  
opt_dirs  
=  
None  
)
```

¶

Find system executable in PATH.

Parameters

arg

? The executable to find.

required

? if executable is not found and required is

True

, fail_json

opt_dirs

? optional list of directories to search in addition to

PATH

Returns

if found return full path; otherwise return None

is_executable

```
(  
path  
)
```

¶

is the given path executable?

Parameters

path

? The path of the file to check.

Limitations:

Does not account for FSACLs.

Most times we really want to know ?Can the current user execute this file?. This function does not tell us that, only if any execute bit is set.

is_special_selinux_path

```
(  
path  
)
```

¶¶
Returns a tuple containing (True, selinux_context) if the given path is on a NFS or other ?special? fs mount point, otherwise the return will be (False, None).

load_file_common_arguments

```
(  
params  
)
```

¶¶
many modules deal with files, this encapsulates common options that the file module accepts such that it is directly available to all modules and they can share code.

md5

```
(  
filename  
)
```

¶¶
Return MD5 hex digest of local file using digest_from_file().
Do not use this function unless you have no other choice for:
Optional backwards compatibility
Compatibility with a third party protocol
This function will not work on systems complying with FIPS-140-2.
Most uses of this function can use the module.sha1 function instead.

preserved_copy

```
(  
src  
,  
dest  
)
```

¶¶
Copy a file with preserved ownership, permissions and context

run_command

```
(  
args  
,  
check_rc  
=  
False  
,  
close_fds  
=  
True  
,  
executable
```

```
=
None
,
data
=
None
,
binary_data
=
False
,
path_prefix
=
None
,
cwd
=
None
,
use_unsafe_shell
=
False
,
prompt_regex
=
None
,
environ_update
=
None
,
umask
=
None
,
encoding
=
'utf-8'
,
errors
=
'surrogate_or_strict'
,
expand_user_and_vars
=
True
,
pass_fds
=
None
,
before_communicate_callback
=
None
```


)

¶

Execute a command, returns rc, stdout, and stderr.

Parameters

args

? is the command to run

* If args is a list, the command will be run with shell=False.

* If args is a string and use_unsafe_shell=False it will split args to a list and run with shell=False

* If args is a string and use_unsafe_shell=True it runs with shell=True.

Kw check_rc

Whether to call fail_json in case of non zero RC.

Default False

Kw close_fds

See documentation for subprocess.Popen(). Default True

Kw executable

See documentation for subprocess.Popen(). Default None

Kw data

If given, information to write to the stdin of the command

Kw binary_data

If False, append a newline to the data. Default False

Kw path_prefix

If given, additional path to find the command in.

This adds to the PATH environment variable so helper commands in the same directory can also be found

Kw cwd

If given, working directory to run the command inside

Kw use_unsafe_shell

See

args

parameter. Default False

Kw prompt_regex

Regex string (not a compiled regex) which can be used to detect prompts in the stdout which would otherwise cause the execution to hang (especially if no input data is specified)

Kw environ_update

dictionary to

update

os.environ with

Kw umask

Umask to be used when running the command. Default None

Kw encoding

Since we return native strings, on python3 we need to know the encoding to use to transform from bytes to text. If you want to always get bytes back, use encoding=None. The default is ?utf-8?. This does not affect transformation of strings given as args.

Kw errors

Since we return native strings, on python3 we need to transform stdout and stderr from bytes to text. If the bytes are undecodable in the

encoding

specified, then use this error

handler to deal with them. The default is

surrogate_or_strict

which means that the bytes will be decoded using the surrogateescape error handler if available (available on all python3 versions we support) otherwise a UnicodeError traceback will be raised. This does not affect transformations of strings given as args.

Kw expand_user_and_vars

When

use_unsafe_shell=False

this argument

dictates whether

~

is expanded in paths and environment variables are expanded before running the command. When

True

a string such as

\$SHELL

will be expanded regardless of escaping. When

False

and

use_unsafe_shell=False

no path or variable expansion will be done.

Kw pass_fds

When running on Python 3 this argument

dictates which file descriptors should be passed

to an underlying

Popen

constructor. On Python 2, this will

set

close_fds

to False.

Kw before_communicate_callback

This function will be called

after

Popen

object will be created

but before communicating to the process.

(

Popen

object will be passed to callback as a first argument)

Returns

A 3-tuple of return code (integer), stdout (native string), and stderr (native string). On python2, stdout and stderr are both byte strings. On python3, stdout and stderr are text strings converted according to the encoding and errors parameters. If you want byte strings on python3, use encoding=None to turn decoding to text off.

sha1

(

filename

)

¶

Return SHA1 hex digest of local file using digest_from_file().

sha256

(

filename

)

¶

Return SHA-256 hex digest of local file using `digest_from_file()`.

Basic

¶

To use this functionality, include

`import`

`ansible.module_utils.basic`

in your module.

`exception`

`ansible.module_utils.basic.`

`AnsibleFallbackNotFound`

¶

`ansible.module_utils.basic.`

`env_fallback`

(

*

`args`

,

**

`kwargs`

)

¶

Load value from environment

`ansible.module_utils.basic.`

`get_all_subclasses`

(

`cls`

)

¶

Deprecated

: Use `ansible.module_utils.common._utils.get_all_subclasses` instead

`ansible.module_utils.basic.`

`get_platform`

(

)

¶

Deprecated

Use

`platform.system()`

directly.

Returns

Name of the platform the module is running on in a native string

Returns a native string that labels the platform (?Linux?, ?Solaris?, etc). Currently, this is

the result of calling

`platform.system()`

.

`ansible.module_utils.basic.`

`heuristic_log_sanitize`

(

`data`

,

`no_log_values`

=

None

)

¶

Remove strings that look like passwords from log messages

ansible.module_utils.basic.

load_platform_subclass

(

cls

,

*

args

,

**

kwargs

)

¶

Deprecated

: Use `ansible.module_utils.common.sys_info.get_platform_subclass` instead

ansible.module_utils.basic.

remove_values

(

value

,

no_log_strings

)

¶

Remove strings in `no_log_strings` from `value`. If `value` is a container type, then remove a lot more.

Use of `deferred_removals` exists, rather than a pure recursive solution, because of the potential to hit the maximum recursion depth when dealing with large amounts of data (see issue #24560).

ansible.module_utils.basic.

sanitize_keys

(

obj

,

no_log_strings

,

ignore_keys

=

frozenset({})

)

¶

Sanitize the keys in a container object by removing `no_log` values from key names.

This is a companion function to the

`remove_values()`

function. Similar to that function,

we make use of `deferred_removals` to avoid hitting maximum recursion depth in cases of large data structures.

Parameters

obj

? The container object to sanitize. Non-container objects are returned unmodified.

no_log_strings

? A set of string values we do not want logged.

`ignore_keys`

? A set of string values of keys to not sanitize.

Returns

An object with sanitized keys.

Next

Previous

© Copyright 2019 Red Hat, Inc.

Last updated on May 27, 2022.

Content from: <https://docs.ansible.com/ansible/devel/community/communication.html>

Communicating with the Ansible community ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Getting started

Community Code of Conduct

Developer Certificate Of Origin

Communicating with the Ansible community

Code of Conduct

Forum

Real-time chat

Working groups

Ansible Community Topics

Ansible Automation Platform support questions

How can I help?

Other ways to get involved

Contributor path

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Community Guide
Getting started
Communicating with the Ansible community
Edit on GitHub
Communicating with the Ansible community
?
Code of Conduct
Forum
The Bullhorn
Regional and Language-specific channels
Real-time chat
General channels
Working group-specific channels
Meetings on Matrix
Working groups
Requesting a forum group
Requesting a community collection repository
Ansible Community Topics
Ansible Automation Platform support questions
Code of Conduct
?
All interactions within the Ansible Community are governed by our
Community Code of Conduct
. Please read and understand it before participating.
Forum
?

The Ansible Forum is the default and recommended starting point for most community interactions. It's ideal for:

- Asking questions and seeking help.
- Participating in development discussions.
- Learning about events and news.

To get started:

Register:
Sign up to create an account and join the community.

Explore topics:
Browse by categories and tags to discover discussions, or simply start a new topic of your own.

Stay updated:
Subscribe to specific categories or tags that interest you. Just click the bell icon in the top-right corner of the relevant category or tag page and select your notification preference.

Explore forum groups that match your interests. Joining a group often automatically subscribes you to related posts.

The Bullhorn ?

The Bullhorn is our community's weekly newsletter, published directly in the Forum:

Subscribe:
Click the bell button under the Bullhorn category description, then select Watching.

.
Submit News:
See the About the Newsletter category post for submission guidelines.
Questions about the newsletter:
Ask us in the Ansible Social room on Matrix.

.
Regional and Language-specific channels ?
Communicate in your preferred language by visiting the International Communities forum category.

. Current subcategories include:

- Deutsche (German)
- Español (Spanish)
- Français (French)
- Italiano (Italian)

Norsk (Norwegian)
Português (Portuguese)

Join an
Ansible Meetup
near you.

For details on requesting a new language subcategory, see the
About the International Communities category post

Real-time chat

?

For real-time interactions, the Ansible community uses the
Matrix protocol

Note
The
Forum
is our default communication platform. We recommend engaging there before considering other options like Matrix.
To join the community on Matrix:

Get a Matrix account:

From
Matrix.org
or any other Matrix homeserver.

Choose a Matrix client:

We recommend
Element Webchat

Join rooms:

Use the links in the

General channels

or

Working groups

to join specific rooms.

For more information, see the community-hosted
Matrix FAQ

You can add Matrix shields to your repository?

README.md

using the shield in the
community-topics
repository as a template.

Note

IRC channels are no longer official communication channels. Use the Forum and Matrix instead.

General channels

?

The clickable links below take you directly to the Matrix room in your browser. Room/channel information is also given
for use in other clients:

Community social room & posting news for the Bullhorn newsletter

General usage & support questions

Developer & code-related topics

Community & collections related topics

Working group-specific channels

?

Many working groups have dedicated chat channels. See the
Working groups

for details.

Meetings on Matrix

?

The Ansible community holds regular meetings on Matrix. All interested individuals are invited to participate.

Check the

meeting schedule and agenda page

for more information.

Working groups

?

Working Groups enable Ansible community members to self-organize around specific interests.

Find a complete list of groups and their communication channels within the

Forum groups

.

Requesting a forum group

?

To request a new working group:

First, check if there is no appropriate

Forum group

you can join instead of starting a new one.

Review the

things you can ask for post

regarding working groups.

Submit your request in the

forum topic

.

If a Matrix chat channel is also needed, consult the

Ansible Community Matrix FAQ

.

Requesting a community collection repository

?

Working groups are often built around Ansible community collections. You can use a repository under your organization

or request one under

ansible-collections

on the forum. Create a topic in the

Project Discussions category and the ?coll-repo-request? tag

.

Ansible Community Topics

?

The

Ansible Community Steering Committee

uses the

Forum

for asynchronous discussions and voting on community topics.

For more information, see:

Creating new policy proposals & inclusion requests

Community topics workflow

Community topics on the Forum

Ansible Automation Platform support questions

?

Red Hat Ansible

Automation Platform

is a subscription service providing support, certified content, and tooling for Ansible, including content management, a controller, UI and REST API.

For questions related to Ansible Automation Platform, visit

Red Hat support
instead of community communication platforms.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 19, 2025.

Content from: https://docs.ansible.com/ansible/devel/community/contributions_collections.html

[Ansible Collections Contributor Guide](#) ? [Ansible Community Documentation](#)

[Blog](#)

[Ansible community forum](#)

[Documentation](#)

[Ansible Community Documentation](#)

[Ansible](#)

Select version:

[latest](#)

[11](#)

[devel](#)

Search docs:

[Ansible getting started](#)

[Getting started with Ansible](#)

[Getting started with Execution Environments](#)

[Installation, Upgrade & Configuration](#)

[Installation Guide](#)

[Ansible Porting Guides](#)

[Using Ansible](#)

[Building Ansible inventories](#)

[Using Ansible command line tools](#)

[Using Ansible playbooks](#)

[Protecting sensitive data with Ansible vault](#)

[Using Ansible modules and plugins](#)

[Using Ansible collections](#)

[Using Ansible on Windows, BSD, and z/OS UNIX](#)

[Ansible tips and tricks](#)

[Contributing to Ansible](#)

[Ansible Community Guide](#)

[Ansible Collections Contributor Guide](#)

[The Ansible Collections Development Cycle](#)

[Requesting changes to a collection](#)

[Creating your first collection pull request](#)

[Testing Collection Contributions](#)

[Review checklist for collection PRs](#)

[Ansible community package collections requirements](#)

[Guidelines for collection maintainers](#)

[Contributing to Ansible-maintained Collections](#)

[Ansible Community Steering Committee](#)

[Contributing to the Ansible Documentation](#)

[Other Tools and Programs](#)

[Working with the Ansible collection repositories](#)

[ansible-core Contributors Guide](#)

[Advanced Contributor Guide](#)

[Ansible documentation style guide](#)

[Extending Ansible](#)

[Developer Guide](#)

[Common Ansible Scenarios](#)

[Legacy Public Cloud Guides](#)

[Network Automation](#)

[Network Getting Started](#)

[Network Advanced Topics](#)

[Network Developer Guide](#)

[Ansible Galaxy](#)
[Galaxy User Guide](#)
[Galaxy Developer Guide](#)
[Reference & Appendices](#)
[Collection Index](#)
[Indexes of all modules and plugins](#)
[Playbook Keywords](#)
[Return Values](#)
[Ansible Configuration Settings](#)
[Controlling how Ansible behaves: precedence rules](#)
[YAML Syntax](#)
[Python 3 Support](#)
[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[ansible-core Roadmaps](#)
[Ansible](#)
[Ansible Collections Contributor Guide](#)
[Edit on GitHub](#)
[Ansible Collections Contributor Guide](#)
[?](#)
[The Ansible Collections Development Cycle](#)
[Macro development: roadmaps, releases, and projects](#)
[Micro development: the lifecycle of a PR](#)
[Making your PR merge-worthy](#)
[Requesting changes to a collection](#)
[Reporting a bug](#)
[Requesting a feature](#)
[Creating your first collection pull request](#)
[Prepare your environment](#)
[Change the code](#)
[Fix the bug](#)
[Test your changes](#)
[Submit a pull request](#)
[Testing Collection Contributions](#)
[How to test a collection PR](#)
[Add unit tests to a collection](#)
[Adding integration tests to a collection](#)
[Review checklist for collection PRs](#)
[Reviewing bug reports](#)
[Reviewing suggested changes](#)
[Review tests in the PR](#)
[Review for merge commits and breaking changes](#)
[Ansible community package collections requirements](#)

Overview

Feedback and communications

Keeping informed

Communication and Working Groups

Collection infrastructure

Python Compatibility

Standards for developing module and plugin utilities

Repository structure requirements

Contributor Workflow

Naming

Collection licensing requirements

Contributor License Agreements

Repository management

CI Testing

When moving modules between collections

Development conventions

Collection Dependencies

Other requirements

Guidelines for collection maintainers

Maintainer responsibilities

Expanding the collection community

Maintaining good collection documentation

Ansible Collection Maintenance and Workflow

Stepping down as a collection maintainer

Releasing collections

Contributing to Ansible-maintained Collections

Ansible-maintained collections

Deciding where your contribution belongs

Requirements to merge your PR

Ansible Community Steering Committee

Steering Committee mission and responsibilities

Steering Committee membership guidelines

Steering Committee past members

Community topics workflow

Contributing to the Ansible Documentation

Editing docs directly on GitHub

Reviewing or solving open issues

Reviewing open PRs

Opening a new issue and/or PR

Verifying your documentation PR

Joining the documentation working group

Other Tools and Programs

Popular editors

Tools for validating playbooks

Collection development tools

Other tools

If you have a specific Ansible interest or expertise (for example, VMware, Linode, and so on), consider joining a working group

.

Working with the Ansible collection repositories

?

How can I find

editors, linters, and other tools

that will support my Ansible development efforts?

Where can I find guidance on
coding in Ansible

?

How do I
create a collection

?

How do I
rebase my PR

?

How do I learn about Ansible's
testing (CI) process

?

How do I
deprecate a module

?

See

Collection developer tutorials

for a quick introduction on how to develop and test your collection contributions.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 19, 2025.

Content from: https://docs.ansible.com/ansible/devel/community/contributor_path.html

Contributor path ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Getting started

Contributor path

Determine your area of interest

Find the corresponding project

Learn

Specific knowledge for code developers

Making your first contribution

Continue to contribute

Teach others

Become a collection maintainer

Become a steering committee member

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Community Guide
Contributor path
Edit on GitHub
Contributor path
?

This section describes the contributor's journey from the beginning to becoming a leader who helps shape the future of Ansible. You can use this path as a roadmap for your long-term participation.

Any contribution to the project, even a small one, is very welcome and valuable. Any contribution counts, whether it is feedback on an issue, a pull request, a topic or documentation change, or a coding contribution. When you contribute regularly, your proficiency and judgment in the related area increase and, along with this, the importance of your presence in the project.

Determine your area of interest

Find the corresponding project

Learn

Specific knowledge for code developers

Making your first contribution

Continue to contribute

Teach others

Become a collection maintainer

Become a steering committee member

Determine your area of interest

?

First, determine areas that are interesting to you. Consider your current experience and what you'd like to gain. For example, if you use a specific collection, have a look there. See

How can I help?

for more ideas on how to help.

Find the corresponding project

?

These are multiple community projects in the Ansible ecosystem you could contribute to:

Ansible Core

Collections

AWX

Galaxy

ansible-lint

Molecule

Learn

?

The required skillset depends on the area of interest and the project you'll be working on. Remember that the best way to learn is by doing.

Specific knowledge for code developers

?

Code development requires the most technical knowledge. Let's sort out what an Ansible developer should learn.

You should understand at least the

basics

of the following tools:

Python programming language

Git

GitHub collaborative development model through forks and pull requests

You can learn these tools more in-depth when working on your first contributions.

Each Ansible project has its own set of contributor guidelines. Familiarize yourself with these as you prepare your first contributions.

Ansible Core development

.

Ansible collection development

and the collection-level contributor guidelines in the collection repository.

Making your first contribution

?

You can find some ideas on how you can contribute in

How can I help?

.

If you are interested in contributing to collections, take a look at

collection contributions

and the

collection repository

's

README

and

CONTRIBUTING

files. To make your first experience as smooth as possible, read the repository documentation carefully, then ask the repository maintainers for guidance if you have any questions.

Take a look at GitHub issues labeled with the

easyfix

and

good_first_issue

labels for:

Ansible collections repositories

All other Ansible projects

Issues labeled with the

docs

label in

Ansible collections

and

other

Ansible projects can be also good to start with.

When you choose an issue to work on, add a comment directly on the GitHub issue to say you are looking at it and let others know to avoid conflicting work.

You can also ask for help in a comment if you need it.

Continue to contribute

?

We don't expect everybody to know everything. Start small, think big. When you contribute regularly, your proficiency and judgment in the related area will improve quickly and, along with this, the importance of your presence in the project.

See

Communicating with the Ansible community

for ways to communicate and engage with the Ansible community, including working group meetings, accessing the Bullhorn news bulletin, and upcoming contributor summits.

Teach others

?

Share your experience with other contributors through

improving documentation

, answering questions from other contributors and users on

Matrix/Libera.Chat IRC

, giving advice on issues and pull requests, and discussing topics on the

Forum

.

Become a collection maintainer

?

If you are a code contributor to a collection, you can get extended permissions in the repository and become a maintainer. A collection maintainer is a contributor trusted by the community who makes significant and regular contributions to the project and showed themselves as a specialist in the related area. See

Guidelines for collection maintainers

for details.

For some collections that use the

collection bot

, such as

community.general

and

community.network

, you can have different levels of access and permissions:

File-level permissions: the stage prior to becoming a collection maintainer. The file is usually a module or plugin. File maintainers have indirect commit rights.

Supershipit permissions: similar to being a file maintainer but the scope where a maintainer has the indirect commit is the whole repository.

Triage access to the repository: allows contributors to manage issues and pull requests.

Write access to the repository also known as

commit

: allows contributors to merge pull requests to the development branch as well as perform all the other activities listed in the

Guidelines for collection maintainers

.

For information about permission levels, see the

GitHub official documentation

.

Become a steering committee member

?

Note

You do NOT have to be a programmer to become a steering committee member.

The

Steering Committee

member status reflects the highest level of trust and allows contributors to lead the project by making important decisions for the Ansible project. The Committee members are community leaders who shape the project's future and the future of automation in the IT world in general.

To reach the status, as the current Committee members did before getting it, along with the things mentioned in this document, you should:

Subscribe to, comment on, and vote on the
community topics<creating_community_topic>

.

Propose your topics.

If time permits, join the

Community meetings

. Note this is

NOT

a requirement.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 19, 2025.

Content from: https://docs.ansible.com/ansible/devel/community/steering/community_steering_con

Steering Committee mission and responsibilities ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

The Ansible Collections Development Cycle

Requesting changes to a collection

Creating your first collection pull request

Testing Collection Contributions

Review checklist for collection PRs

Ansible community package collections requirements

Guidelines for collection maintainers

Contributing to Ansible-maintained Collections

Ansible Community Steering Committee

Steering Committee mission and responsibilities

Steering Committee responsibilities

Current Steering Committee members

Creating new policy proposals & inclusion requests

Community Working Group meetings

Steering Committee membership guidelines

Steering Committee past members

Community topics workflow

Contributing to the Ansible Documentation

Other Tools and Programs

Working with the Ansible collection repositories

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible
Developer Guide
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Collections Contributor Guide
Ansible Community Steering Committee
Steering Committee mission and responsibilities
Edit on GitHub
Steering Committee mission and responsibilities
?

The Steering Committee mission is to provide continuity, guidance, and suggestions to the Ansible community to ensure the delivery and high quality of the Ansible package. In addition, the committee helps decide the technical direction of the Ansible project. It is responsible for approving new proposals and policies in the community, package, and community collections world, new community collection-inclusion requests, and other technical aspects regarding inclusion and packaging.

The Committee should reflect the scope and breadth of the Ansible community.

Steering Committee responsibilities

?

The Committee:

Designs policies and procedures for the community collections world.

Votes on approval changes to established policies and procedures.

Reviews community collections for compliance with the policies.

Helps create and define roadmaps for our deliverables such as the

ansible

package, major community collections, and documentation.

Reviews community collections submitted for inclusion in the Ansible package and decides whether to include them or not.

Review other proposals of importance that need the Committee's attention and provide feedback.

Current Steering Committee members

?

The following table lists the current Steering Committee members. See

Steering Committee past members

for a list of past members.

The Steering Committee itself is also a group in the forum (

@SteeringCommittee

).

Current Steering committee members

?

Name

GitHub

Forum

Start year

Alexei Znamensky

russoz

russoz

2022

Alicia Cozine

acozine

acozine

2021

Andrew Klychkov

Andersson007

Andersson007

2021

Brian Scholer

briantist

briantist

2022

Core Team repr.

See Forum

2024

Felix Fontein

felixfontein

felixfontein

2021

James Cassell

jamescassell

cassel

2021

John Barker

gundalow

gundalow

2021

Mario Lenz

mariolenz

mariolenz

2022

Markus Bergholz

markuman

markuman

2022

Maxwell G

gotmax23

gotmax23

2022

Sorin Sbarnea

ssbarnea

ssbarnea

2021

John Barker (

gundalow

) has been elected by the Committee as its

Chairperson

.

Committee members are selected based on their active contribution to the Ansible Project and its community. See [Steering Committee membership guidelines](#)

to learn details.

Creating new policy proposals & inclusion requests

?

The Committee uses the

Ansible Forum

to asynchronously discuss with the Community and vote on the proposals in corresponding

community topics

.

You can

create a community topic

(make sure you use the

Project

Discussions

category and the

community-wg

tag) if you want to discuss an idea that impacts any of the following:

Ansible Community

Community collection best practices and requirements

Community collection inclusion/exclusion policy and workflow

The Community governance

Other proposals of importance that need the Committee's or overall Ansible community attention

To request changes to the inclusion policy and

Ansible community package collections requirements

:

Submit a new pull request against the document by clicking the

Edit

button on its web page.

Create a

community topic

containing the rationale for the proposed changes.

To submit new collections for inclusion into the Ansible package:

Submit the new collection inclusion requests through a new discussion in the

ansible-inclusion

repository.

Depending on a topic you want to discuss with the Community and the Committee, as you prepare your proposal,

please consider the requirements established by:
Community Code of Conduct

- Ansible community package collections requirements

- Ansible Collection Inclusion Checklist

- Ansible Community Package Collections Removal Process

- Community topics workflow

?

The Committee uses the
Community topics workflow
to asynchronously discuss and vote on the
community topics

- The quorum, the minimum number of Committee members who must vote on a topic in order for a decision to be officially made, is half of the whole number of the Committee members. If the quorum number contains a fractional part, it is rounded up to the next whole number. For example, if there are thirteen members currently in the committee, the quorum will be seven.

Votes must always have ?no change? as an option.

In case of equal numbers of votes for and against a topic, the chairperson?s vote will break the tie. For example, if there are six votes for and six votes against a topic, and the chairperson?s vote is among those six which are for the topic, the final decision will be positive. If the chairperson has not voted yet, other members ask them to vote.

For votes with more than two options, one choice must have at least half of the votes. If two choices happen to both have half of the votes, the chairperson?s vote will break the tie. If no choice has at least half of the votes, the vote choices have to be adjusted so that a majority can be found for a choice in a new vote.

Community topics triage

?

The Committee conducts a triage of
community topics
periodically (every three to six months).

The triage goals are:

Sparking interest for forgotten topics.

Identifying and closing irrelevant topics, for example, when the reason of the topic does not exist anymore or the topic is out of the Committee responsibilities scope.

Identifying and closing topics that the Community are not interested in discussing. As indicators, it can be absence of comments or no activity in comments, at least, for the last six months.

Identifying and closing topics that were solved and implemented but not closed (in this case, such a topic can be closed on the spot with a comment that it has been implemented).

Identifying topics that have been in pending state for a long time, for example, when it is waiting for actions from someone for several months or when the topics were solved but not implemented.

A person starting the triage:

Identifies the topics mentioned above.

Creates a special triage topic containing an enumerated list of the topics-candidates for closing.

Establishes a vote date considering a number of topics, their complexity and comment-history size giving the Community sufficient time to go through and discuss them.

The Community and the Committee vote on each topic-candidate listed in the triage topic whether to close it or keep it open.

Collection inclusion requests workflow

?

When reviewing community collection
inclusion requests

, the Committee members check if a collection adheres to the

Ansible community package collections requirements

.

Note

The Steering Committee can reject a collection inclusion request or exclude a collection from the Ansible package even when the collection satisfies the requirements if the Steering Committee agrees that presence of the collection will significantly deteriorate the Ansible package user experience or the package build process. In case of rejection/removal, the collection maintainers receive comprehensive feedback from the Committee explaining the reasons of starting the process. In case the reasons are fixable, the feedback will contain information what the maintainers need to change.

A Committee member who conducts the inclusion review copies the

Ansible community collection checklist

into a corresponding

discussion

.

In the course of the review, the Committee member marks items as completed or leaves a comment saying whether the reviewer expects an issue to be addressed or whether it is optional (for example, it could be

MUST FIX:

<what> or

SHOULD FIX:

<what> under an item).

For a collection to be included in the Ansible community package, the collection:

MUST be reviewed and approved as compliant with the requirements by at least two Steering Committee members.

At least one of the reviews checks compliance with the entire checklist.

All subsequent reviews can focus only on compliance with documentation and development conventions.

Reviewers must not be involved significantly in development of the collection. They MUST declare any potential conflict of interest (for example, being friends/relatives/coworkers of the maintainers/authors, being users of the collection, or having contributed to that collection recently or in the past).

After the collection gets two Committee member approvals, a Committee member creates a community topic

linked to the corresponding inclusion request. The issue's description says that the collection has been approved by the Committee and establishes a date (a week by default) when the inclusion decision will be considered made.

The inclusion automatically gets suspended if the Committee members raise concerns or start another inclusion review within this time period.

When there are no more objections or ongoing inclusion reviews, the inclusion date gets prolonged for another week.

If the inclusion has not been suspended by the established date, the inclusion request is considered successfully resolved. In this case, a Committee member:

Declares the decision in the topic and in the inclusion request.

Moves the request to the

Resolved

reviews

category.

Adds the collection to the

ansible.in

file in a corresponding directory of the

ansible-build-data repository

.

Announces the inclusion through the

Bullhorn newsletter

.

Closes the topic.

Collection exclusion workflow

?

The Committee uses the

Ansible Community Package Collections Removal Process

to remove collections not satisfying the

Ansible community package collections requirements
from the Ansible package.

Community Working Group meetings

?

See the Community Working Group meeting
schedule

. Meeting summaries are posted on the
Forum

.

Note

Participation in the Community Working Group meetings is optional for Committee members. Decisions on
community topics

are made asynchronously in the topics themselves.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 19, 2025.

Content from: https://docs.ansible.com/ansible/devel/dev_guide/developing_collections_distributing_collections.html

Distributing collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

- Interpreter Discovery
- Releases and maintenance
- Testing Strategies
- Sanity Tests
- Frequently Asked Questions
- Glossary
- Ansible Reference: Module Utilities
- Special Variables
- Red Hat Ansible Automation Platform
- Ansible Automation Hub
- Logging Ansible output
- Roadmaps
- Ansible Roadmap
- ansible-core Roadmaps
- Ansible
- Developer Guide
- Developing collections
- Distributing collections
- Edit on GitHub
- Distributing collections
- ?

A collection is a distribution format for Ansible content. A typical collection contains modules and other plugins that address a set of related use cases. For example, a collection might automate administering a particular database. A collection can also contain roles and playbooks.

Note

Before distributing your collection, ensure you have updated the `galaxy.yml` file.

See

Collection structure
for details.

To distribute your collection and allow others to use it, you can publish your collection on one or more distribution server

. Distribution servers include:

Distribution server

Collections accepted

Ansible Galaxy

All collections

Pulp 3 Galaxy

All collections, supports signed collections

Red Hat Automation Hub

Only collections certified by Red Hat, supports signed collections

Privately hosted Automation Hub

Collections authorized by the owners

Distributing collections involves four major steps:

Initial configuration of your distribution server or servers

Building your collection tarball

Preparing to publish your collection

Publishing your collection

Initial configuration of your distribution server or servers

Creating a namespace

Getting your API token

Specifying your API token and distribution server

Building your collection tarball

Ignoring files and folders
Signing a collection
Preparing to publish your collection
Installing your collection locally
Reviewing your collection
Understanding collection versioning
Publishing your collection
Publishing a collection from the command line
Publishing a collection from the website
Initial configuration of your distribution server or servers
?

Configure a connection to one or more distribution servers so you can publish collections there. You only need to configure each distribution server once. You must repeat the other steps (building your collection tarball, preparing to publish, and publishing your collection) every time you publish a new collection or a new version of an existing collection.

Create a namespace on each distribution server you want to use.

Get an API token for each distribution server you want to use.

Specify the API token for each distribution server you want to use.

Creating a namespace

?

You must upload your collection into a namespace on each distribution server. If you have a login for Ansible Galaxy, your Ansible Galaxy username is usually also an Ansible Galaxy namespace.

Warning

Namespaces on Ansible Galaxy cannot include hyphens. If you have a login for Ansible Galaxy that includes a hyphen, your Galaxy username is not also a Galaxy namespace. For example,

awesome-user

is a valid username for Ansible Galaxy, but it is not a valid namespace.

You can create additional namespaces on Ansible Galaxy if you choose. For Red Hat Automation Hub and private Automation Hub you must create a namespace before you can upload your collection. To create a namespace:

To create a namespace on Galaxy, see

Galaxy namespaces

on the Galaxy docsite for details.

To create a namespace on Red Hat Automation Hub, see the

Red Hat automation content documentation

.

Specify the namespace in the

galaxy.yml

file for each collection. For more information on the

galaxy.yml

file, see

Collection Galaxy metadata structure

.

Getting your API token

?

An API token authenticates your connection to each distribution server. You need a separate API token for each distribution server. Use the correct API token to connect to each distribution server securely and protect your content.

To get your API token:

To get an API token for Galaxy, see

the Galaxy documentation

.

To get an API token for Automation Hub, see the

Red Hat automation content documentation

.

Specifying your API token and distribution server

?

Each time you publish a collection, you must specify the API token and the distribution server to create a secure connection. You have two options for specifying the token and distribution server:

You can configure the token in configuration, as part of a `galaxy_server_list`

entry in your

`ansible.cfg`

file. Using configuration is the most secure option.

You can pass the token at the command line as an argument to the

`ansible-galaxy`

command. If you pass the token at the command line, you can specify the server at the command line, by using the default setting, or by setting the server in configuration. Passing the token at the command line is insecure, because typing secrets at the command line may expose them to other users on the system.

Specifying the token and distribution server in configuration

?

By default, Ansible Galaxy is configured as the only distribution server. You can add other distribution servers and specify your API token or tokens in configuration by editing the

`galaxy_server_list`

section of your

`ansible.cfg`

file. This is the most secure way to manage authentication for distribution servers. Specify a URL and token for each server. For example:

```
[galaxy]
```

```
server_list
```

```
=
```

```
release_galaxy
```

```
[galaxy_server.release_galaxy]
```

```
url
```

```
=
```

```
https://galaxy.ansible.com/
```

```
token
```

```
=
```

```
abcdefghijklmnopqrstuvwxyz
```

You cannot use

`apt-key`

with any servers defined in your

`galaxy_server_list`

. See

Configuring the `ansible-galaxy` client

for complete details.

Specifying the token at the command line

?

You can specify the API token at the command line using the

`--token`

argument of the

`ansible-galaxy`

command. There are three ways to specify the distribution server when passing the token at the command line:

using the

`--server`

argument of the

`ansible-galaxy`

command

relying on the default (

`https://galaxy.ansible.com`

)

setting a server in configuration by creating a

GALAXY_SERVER

setting in your

ansible.cfg

file

For example:

ansible-galaxy

collection

publish

path/to/my_namespace-my_collection-1.0.0.tar.gz

--token

abcdefghijklmnopqrstuvwxyz

Warning

Using the

--token

argument is insecure. Passing secrets at the command line may expose them to others on the system.

Building your collection tarball

?

After configuring one or more distribution servers, build a collection tarball. The collection tarball is the published artifact, the object that you upload and other users download to install your collection. To build a collection tarball:

Review all settings in your

galaxy.yml

file. See

Collection Galaxy metadata structure

for details. Ensure you have updated the version number. Each time you publish your collection, it must have a new version number. You cannot make changes to existing versions of your collection on a distribution server. If you try to upload the same collection version more than once, the distribution server returns the error

Code:

conflict.collection_exists

. Collections follow semantic versioning rules. For more information on versions, see

Understanding collection versioning

. For more information on the

galaxy.yml

file, see

Collection Galaxy metadata structure

.

Run

ansible-galaxy

collection

build

from inside the top-level directory of the collection. For example:

collection_dir#>

ansible-galaxy

collection

build

This command builds a tarball of the collection in the current directory, which you can upload to your selected distribution server:

my_collection/

???

galaxy.yml

???

...

???

my_namespace-my_collection-1.0.0.tar.gz

???

...

Note

To reduce the size of collections, certain files and folders are excluded from the collection tarball by default. See

Ignoring files and folders

if your collection directory contains other files you want to exclude.

The current Galaxy maximum tarball size is 2 MB.

You can upload your tarball to one or more distribution servers. You can also distribute your collection locally by copying the tarball to install your collection directly on target systems.

Ignoring files and folders

?

You can exclude files from your collection with either

build_ignore

or

Manifest Directives

. For more information on the

galaxy.yml

file, see

Collection Galaxy metadata structure

.

Include all, with explicit ignores

?

By default, the build step includes all the files in the collection directory in the tarball except for the following:

galaxy.yml

*.pyc

*.retry

tests/output

previously built tarballs in the root directory

various version control directories such as

.git/

To exclude other files and folders from your collection tarball, set a list of file glob-like patterns in the

build_ignore

key in the collection?s

galaxy.yml

file. These patterns use the following special characters for wildcard matching:

*

: Matches everything

?

: Matches any single character

[seq]

: Matches any character in sequence

[!seq]

:Matches any character not in sequence

For example, to exclude the

sensitive

folder within the

playbooks

folder as well any

.tar.gz

archives, set the following in your

galaxy.yml

file:

build_ignore

:
-
playbooks/sensitive
-
'*.tar.gz'
Note
The
build_ignore
feature is only supported with
ansible-galaxy
collection
build
in Ansible 2.10 or newer.
Manifest Directives
?
New in version 2.14.
The
galaxy.yml
file supports manifest directives that are historically used in Python packaging, as described in
MANIFEST.in commands
.
Note
The use of
manifest
requires installing the optional
distlib
Python dependency.
Note
The
manifest
feature is only supported with
ansible-galaxy
collection
build
in
ansible-core
2.14 or newer, and is mutually exclusive with
build_ignore
.
For example, to exclude the
sensitive
folder within the
playbooks
folder as well as any
.tar.gz
archives, set the following in your
galaxy.yml
file:
manifest
:
directives
:
-
recursive-exclude playbooks/sensitive **

-

global-exclude *.tar.gz

By default, the

MANIFEST.in

style directives would exclude all files by default, but there are default directives in place. Those default directives are described below. To see the directives in use during build, pass

-vvv

with the

ansible-galaxy

collection

build

command.

include meta/*.yml

include *.txt *.md *.rst COPYING LICENSE

recursive-include tests **

recursive-include docs *.rst *.yml *.yaml *.json *.j2 *.txt

recursive-include roles *.yml *.yaml *.json *.j2

recursive-include playbooks *.yml *.yaml *.json

recursive-include changelogs *.yml *.yaml

recursive-include plugins */*.py

recursive-include plugins/become *.yml *.yaml

recursive-include plugins/cache *.yml *.yaml

recursive-include plugins/callback *.yml *.yaml

recursive-include plugins/cliconf *.yml *.yaml

recursive-include plugins/connection *.yml *.yaml

recursive-include plugins/filter *.yml *.yaml

recursive-include plugins/httpapi *.yml *.yaml

recursive-include plugins/inventory *.yml *.yaml

recursive-include plugins/lookup *.yml *.yaml

recursive-include plugins/netconf *.yml *.yaml

recursive-include plugins/shell *.yml *.yaml

recursive-include plugins/strategy *.yml *.yaml

recursive-include plugins/test *.yml *.yaml

recursive-include plugins/vars *.yml *.yaml

recursive-include plugins/modules *.ps1 *.yml *.yaml

recursive-include plugins/module_utils *.ps1 *.psm1 *.cs

manifest.directives from galaxy.yml inserted here

exclude galaxy.yml galaxy.yaml MANIFEST.json FILES.json <namespace>-<name>*.tar.gz

recursive-exclude tests/output **

global-exclude /* /__pycache__

Note

<namespace>-<name>*.tar.gz

is expanded with the actual

namespace

and

name

.

The

manifest.directives

supplied in

galaxy.yml

are inserted after the default includes and before the default excludes.

To enable the use of manifest directives without supplying your own, insert either

manifest:

```
{}  
or  
manifest:  
  null  
in the  
galaxy.yml  
file and remove any use of  
build_ignore
```

.

If the default manifest directives do not meet your needs, you can set `manifest.omit_default_directives` to a value of `true` in `galaxy.yml`.

You then must specify a full compliment of manifest directives in `galaxy.yml`.

The defaults documented above are a good starting point.

Below is an example where the default directives are not included.

```
manifest  
:  
  directives  
  :  
  -  
  include meta/runtime.yml  
  -  
  include README.md LICENSE  
  -  
  recursive-include plugins /**.py  
  -  
  exclude galaxy.yml MANIFEST.json FILES.json <namespace>-<name>*.tar.gz  
  -  
  recursive-exclude tests/output **  
omit_default_directives  
:  
  true
```

Signing a collection

?

You can include a GnuPG signature with your collection on a

Pulp 3 Galaxy

server. See

Enabling collection signing

for details.

You can manually generate detached signatures for a collection using the

gpg

CLI using the following step. This step assume you have generated a GPG private key, but do not cover this process.

```
ansible-galaxy
```

```
collection
```

```
build
```

```
tar
```

```
-Oxzf
```

```
namespace-name-1.0.0.tar.gz
```

```
MANIFEST.json
```

```
|
```

```
gpg
--output
namespace-name-1.0.0.asc
--detach-sign
--armor
--local-user
[email protected]
-
```

Preparing to publish your collection

?

Each time you publish your collection, you must create a new version

on the distribution server. After you publish a version of a collection, you cannot delete or modify that version. To avoid unnecessary extra versions, check your collection for bugs, typos, and other issues locally before publishing:

Install the collection locally.

Review the locally installed collection before publishing a new version.

Installing your collection locally

?

You have two options for installing your collection locally:

Install your collection locally from the tarball.

Install your collection locally from your Git repository.

Installing your collection locally from the tarball

?

To install your collection locally from the tarball, run

```
ansible-galaxy
```

```
collection
```

```
install
```

and specify the collection tarball. You can optionally specify a location using the

```
-p
```

flag. For example:

```
collection_dir#>
```

```
ansible-galaxy
```

```
collection
```

```
install
```

```
my_namespace-my_collection-1.0.0.tar.gz
```

```
-p
```

```
./collections
```

Install the tarball into a directory configured in

COLLECTIONS_PATHS

so Ansible can easily find and load the collection. If you do not specify a path value,

```
ansible-galaxy
```

```
collection
```

```
install
```

installs the collection in the first path defined in

COLLECTIONS_PATHS

.

Installing your collection locally from a Git repository

?

To install your collection locally from a Git repository, specify the repository and the branch you want to install:

```
collection_dir#>
```

```
ansible-galaxy
```

```
collection
```

```
install
```

```
git+https://github.com/org/repo.git,devel
```

You can install a collection from a git repository instead of from Galaxy or Automation Hub. As a developer, installing from a git repository lets you review your collection before you create the tarball and publish the collection. As a user, installing from a git repository lets you use collections or versions that are not in Galaxy or Automation Hub yet. This functionality is meant as a minimal shortcut for developers of content as previously described, and git repositories may not support the full set of features from the

ansible-galaxy

CLI. In complex cases, a more flexible option may be to

git

clone

the repository into the correct file structure of the collection installation directory.

The repository must contain a

galaxy.yml

or

MANIFEST.json

file. This file provides metadata such as the version number and namespace of the collection.

Installing a collection from a git repository at the command line

?

To install a collection from a git repository at the command line, use the URI of the repository instead of a collection name or path to a

tar.gz

file. Use the prefix

git+

, unless you're using SSH authentication with the user

git

(for example,

git@github.com:ansible-collections/ansible.windows.git

). You can specify a branch, commit, or tag using the comma-separated

git commit-ish

syntax.

For example:

Install a collection in a repository using the latest commit on the branch 'devel'

ansible-galaxy

collection

install

git+https://github.com/organization/repo_name.git,devel

Install a collection from a private GitHub repository

ansible-galaxy

collection

install

:organization/repo_name.git

Install a collection from a local git repository

ansible-galaxy

collection

install

git+file:///home/user/path/to/repo_name.git

Warning

Embedding credentials into a git URI is not secure. Use safe authentication options to prevent your credentials from being exposed in logs or elsewhere.

Use

SSH

authentication

Use

netrc

authentication

Use

http.extraHeader

in your git configuration

Use

url.<base>.pushInsteadOf

in your git configuration

Specifying the collection location within the git repository

?

When you install a collection from a git repository, Ansible uses the collection

galaxy.yml

or

MANIFEST.json

metadata file to build the collection. By default, Ansible searches two paths for collection

galaxy.yml

or

MANIFEST.json

metadata files:

The top level of the repository.

Each directory in the repository path (one level deep).

If a

galaxy.yml

or

MANIFEST.json

file exists in the top level of the repository, Ansible uses the collection metadata in that file to install an individual collection.

??? galaxy.yml

??? plugins/

? ??? lookup/

? ??? modules/

? ??? module_utils/

???? README.md

If a

galaxy.yml

or

MANIFEST.json

file exists in one or more directories in the repository path (one level deep), Ansible installs each directory with a metadata file as a collection. For example, Ansible installs both collection1 and collection2 from this repository structure by default:

??? collection1

? ??? docs/

? ??? galaxy.yml

? ??? plugins/

? ??? inventory/

? ??? modules/

??? collection2

??? docs/

??? galaxy.yml

??? plugins/

| ??? filter/

| ??? modules/

??? roles/

If you have a different repository structure or only want to install a subset of collections, you can add a fragment to the end of your URI (before the optional comma-separated version) to indicate the location of the metadata file or files. The

path should be a directory, not the metadata file itself. For example, to install only collection2 from the example repository with two collections:

```
ansible-galaxy collection install git+https://github.com/organization/repo_name.git#/collection2/
```

In some repositories, the main directory corresponds to the namespace:

```
namespace/
```

```
??? collectionA/
```

```
|   ??? docs/
```

```
|   ??? galaxy.yml
```

```
|   ??? plugins/
```

```
|   ?   ??? README.md
```

```
|   ?   ??? modules/
```

```
|   ??? README.md
```

```
|   ??? roles/
```

```
??? collectionB/
```

```
    ??? docs/
```

```
    ??? galaxy.yml
```

```
    ??? plugins/
```

```
    ?   ??? connection/
```

```
    ?   ??? modules/
```

```
    ??? README.md
```

```
    ??? roles/
```

You can install all collections in this repository, or install one collection from a specific commit:

```
# Install all collections in the namespace
```

```
ansible-galaxy
```

```
collection
```

```
install
```

```
git+https://github.com/organization/repo_name.git#/namespace/
```

```
# Install an individual collection using a specific commit
```

```
ansible-galaxy
```

```
collection
```

```
install
```

```
git+https://github.com/organization/repo_name.git#/namespace/collectionA/7b60ddc245bc416b72d8ea6ed7b799885110f5e5
```

```
Reviewing your collection
```

```
?
```

Review the collection:

Run a playbook that uses the modules and plugins in your collection. Verify that new features and functionality work as expected. For examples and more details see

Using collections

```
.
```

Check the documentation for typos.

Check that the version number of your tarball is higher than the latest published version on the distribution server or servers.

If you find any issues, fix them and rebuild the collection tarball.

Understanding collection versioning

```
?
```

The only way to change a collection is to release a new version. The latest version of a collection (by highest version number) is the version displayed everywhere in Galaxy and Automation Hub. Users can still download older versions.

Follow semantic versioning when setting the version for your collection. In summary:

Increment the major version number,

```
x
```

```
of
```

```
x.y.z
```

, for an incompatible API change.

Increment the minor version number,

y

of

x.y.z

, for new functionality in a backwards compatible manner (for example new modules/plugins, parameters, return values).

Increment the patch version number,

z

of

x.y.z

, for backwards compatible bug fixes.

Read the official

Semantic Versioning

documentation for details and examples.

Publishing your collection

?

The last step in distributing your collection is publishing the tarball to Ansible Galaxy, Red Hat Automation Hub, or a privately hosted Automation Hub instance. You can publish your collection in two ways:

from the command line using the

`ansible-galaxy`

`collection`

`publish`

command

from the website of the distribution server (Galaxy, Automation Hub) itself

Publishing a collection from the command line

?

To upload the collection tarball from the command line using

`ansible-galaxy`

:

`ansible-galaxy`

`collection`

`publish`

`path/to/my_namespace-my_collection-1.0.0.tar.gz`

Note

This `ansible-galaxy` command assumes you have retrieved and stored your API token in configuration. See

Specifying your API token and distribution server

for details.

The

`ansible-galaxy`

`collection`

`publish`

command triggers an import process, just as if you uploaded the collection through the Galaxy website. The command waits until the import process completes before reporting the status back. If you want to continue without waiting for the import result, use the

`--no-wait`

argument and manually look at the import progress in your

My Imports

page.

Publishing a collection from the website

?

See the

Galaxy documentation

to learn how to publish your collection directly on the Galaxy website.

See also

Using Ansible collections

[Learn how to install and use collections.](#)

[Collection Galaxy metadata structure](#)

[Table of fields used in the](#)

[galaxy.yml](#)

[file](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 19, 2025.

Content from: https://docs.ansible.com/ansible/devel/getting_started_ee/index.html

Getting started with Execution Environments ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Introduction to Execution Environments

Setting up your environment

Building your first Execution Environment

Running your EE

Running Ansible with the community EE image

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

[Return Values](#)
[Ansible Configuration Settings](#)
[Controlling how Ansible behaves: precedence rules](#)
[YAML Syntax](#)
[Python 3 Support](#)
[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)

[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[ansible-core Roadmaps](#)

[Ansible](#)
[Getting started with Execution Environments](#)
[Edit on GitHub](#)
[Getting started with Execution Environments](#)
[?](#)

You can run Ansible automation in containers, like any other modern software application.

Ansible uses container images known as Execution Environments (EE) that act as control nodes.

EEs remove complexity to scale out automation projects and make things like deployment operations much more straightforward.

An Execution Environment image contains the following packages as standard:

- `ansible-core`
- `ansible-runner`
- Python

Ansible content dependencies

In addition to the standard packages, an EE can also contain:

- one or more Ansible collections and their dependencies
- other custom components

This getting started guide shows you how to build and test a simple Execution Environment.

The resulting container image represents an Ansible control node that contains:

- standard EE packages
- `community.postgresql` collection
- `psycpg2-binary`

Python package

[Introduction to Execution Environments](#)

[Setting up your environment](#)

[Building your first Execution Environment](#)

[Running your EE](#)

[Running Ansible with the community EE image](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 19, 2025.

Content from: <https://docs.ansible.com/ansible/index.html>

[Ansible Documentation](#) ? [Ansible Community Documentation](#)

[Blog](#)

[Ansible community forum](#)

[Documentation](#)

[Ansible Community Documentation](#)

[Ansible](#)

Select version:

[latest](#)

[11](#)

[devel](#)

Search docs:

[Ansible getting started](#)

[Getting started with Ansible](#)

[Getting started with Execution Environments](#)

[Installation, Upgrade & Configuration](#)

[Installation Guide](#)

[Ansible Porting Guides](#)

[Using Ansible](#)

[Building Ansible inventories](#)

[Using Ansible command line tools](#)

[Using Ansible playbooks](#)

[Protecting sensitive data with Ansible vault](#)

[Using Ansible modules and plugins](#)

[Using Ansible collections](#)

[Using Ansible on Windows, BSD, and z/OS UNIX](#)

[Ansible tips and tricks](#)

[Contributing to Ansible](#)

[Ansible Community Guide](#)

[Ansible Collections Contributor Guide](#)

[ansible-core Contributors Guide](#)

[Advanced Contributor Guide](#)

[Ansible documentation style guide](#)

[Extending Ansible](#)

[Developer Guide](#)

[Common Ansible Scenarios](#)

[Legacy Public Cloud Guides](#)

[Network Automation](#)

[Network Getting Started](#)

[Network Advanced Topics](#)

[Network Developer Guide](#)

[Ansible Galaxy](#)

[Galaxy User Guide](#)

[Galaxy Developer Guide](#)

[Reference & Appendices](#)

[Collection Index](#)

[Indexes of all modules and plugins](#)

[Playbook Keywords](#)

[Return Values](#)

[Ansible Configuration Settings](#)

[Controlling how Ansible behaves: precedence rules](#)

[YAML Syntax](#)

[Python 3 Support](#)

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Documentation
Ansible Documentation
?

Welcome to Ansible community documentation!

This documentation covers the version of Ansible noted in the upper left corner of this page.

We maintain multiple versions of Ansible and of the documentation, so please be sure you are using the version of the documentation that covers the version of Ansible you're using.

For recent features, we note the version of Ansible where the feature was added.

Ansible releases a new major release approximately twice a year.

The core application evolves somewhat conservatively, valuing simplicity in language design and setup.

Contributors develop and change modules and plugins, hosted in collections, much more quickly.

Ansible getting started

Getting started with Ansible

Introduction to Ansible

Start automating with Ansible

Building an inventory

Creating a playbook

Ansible concepts

Getting started with Execution Environments

Introduction to Execution Environments

Setting up your environment

Building your first Execution Environment

Running your EE

Running Ansible with the community EE image

Installation, Upgrade & Configuration

Installation Guide

Installing Ansible

Installing Ansible on specific operating systems

Configuring Ansible

Ansible Porting Guides

Ansible 12 Porting Guide

Ansible 11 Porting Guide

Ansible 10 Porting Guide

Ansible 9 Porting Guide

Ansible 8 Porting Guide

Ansible 7 Porting Guide

Ansible 6 Porting Guide

Ansible 5 Porting Guide

Ansible 4 Porting Guide

- Ansible 3 Porting Guide
- Ansible 2.10 Porting Guide
- Ansible 2.9 Porting Guide
- Ansible 2.8 Porting Guide
- Ansible 2.7 Porting Guide
- Ansible 2.6 Porting Guide
- Ansible 2.5 Porting Guide
- Ansible 2.4 Porting Guide
- Ansible 2.3 Porting Guide
- Ansible 2.0 Porting Guide
- Using Ansible
 - Building Ansible inventories
 - How to build your inventory
 - Working with dynamic inventory
 - Patterns: targeting hosts and groups
 - Connection methods and details
 - Using Ansible command line tools
 - Introduction to ad hoc commands
 - Working with command line tools
 - Ansible CLI cheatsheet
 - Using Ansible playbooks
 - Ansible playbooks
 - Working with playbooks
 - Executing playbooks
 - Advanced playbook syntax
 - Manipulating data
 - Protecting sensitive data with Ansible vault
 - Ansible Vault
 - Managing vault passwords
 - Encrypting content with Ansible Vault
 - Using encrypted variables and files
 - Configuring defaults for using encrypted content
 - When are encrypted files made visible?
 - Format of files encrypted with Ansible Vault
 - Using Ansible modules and plugins
 - Introduction to modules
 - Boolean variables
 - Module maintenance and support
 - Rejecting modules
 - Working with plugins
 - Modules and plugins index
 - Using Ansible collections
 - Installing collections
 - Removing a collection
 - Downloading collections
 - Listing collections
 - Verifying collections
 - Using collections in a playbook
 - Collections index
 - Using Ansible on Windows, BSD, and z/OS UNIX
 - Managing BSD hosts with Ansible
 - Managing Windows hosts with Ansible
 - Managing z/OS UNIX hosts with Ansible
 - Ansible tips and tricks

General tips
Playbook tips
Inventory tips
Execution tricks
Sample Ansible setup
Contributing to Ansible
Ansible Community Guide
Getting started
Contributor path
Ansible Collections Contributor Guide
The Ansible Collections Development Cycle
Requesting changes to a collection
Creating your first collection pull request
Testing Collection Contributions
Review checklist for collection PRs
Ansible community package collections requirements
Guidelines for collection maintainers
Contributing to Ansible-maintained Collections
Ansible Community Steering Committee
Contributing to the Ansible Documentation
Other Tools and Programs
Working with the Ansible collection repositories
ansible-core Contributors Guide
Reporting bugs and requesting features
Contributing to the Ansible Documentation
The Ansible Development Cycle
Other Tools and Programs
Working with the Ansible repo
Advanced Contributor Guide
Committers Guidelines
Release Manager Guidelines
GitHub Admins
Ansible Ecosystem Project Development Resources
Ansible documentation style guide
Linguistic guidelines
reStructuredText guidelines
Markdown guidelines
Accessibility guidelines
More resources
Extending Ansible
Developer Guide
Adding modules and plugins locally
Should you develop a module?
Developing modules
Contributing your module to an existing Ansible collection
Conventions, tips, and pitfalls
Ansible and Python 3
Debugging modules
Module format and documentation
Ansible markup
Adjacent YAML documentation files
Windows module development walkthrough
Creating a new collection
Testing Ansible

The lifecycle of an Ansible module or plugin
Developing plugins
Developing dynamic inventory
Developing
ansible-core
Ansible module architecture
Python API
Rebasing a pull request
Using and developing module utilities
Ansible collection creator path
Developing collections
Migrating Roles to Roles in Collections on Galaxy
Collection Galaxy metadata structure
Ansible architecture
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Basic Concepts
How Network Automation is Different
Run Your First Command and Playbook
Build Your Inventory
Use Ansible network roles
Beyond the basics
Working with network connection options
Resources and next steps
Network Advanced Topics
Network Resource Modules
Ansible Network Examples
Parsing semi-structured text with Ansible
Validate data against set criteria with Ansible
Network Debug and Troubleshooting Guide
Working with command output and prompts in network modules
Ansible Network FAQ
Platform Options
Network Developer Guide
Developing network resource modules
Developing network plugins
Documenting new network platforms
Ansible Galaxy
Galaxy User Guide
Finding collections on Galaxy
Finding roles on Galaxy
Installing roles from Galaxy
Galaxy Developer Guide
Creating collections for Galaxy
Creating roles for Galaxy
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules

[YAML Syntax](#)
[Python 3 Support](#)
[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[Ansible project 12.0](#)
[Ansible project 11.0](#)
[Ansible project 10.0](#)
[Ansible project 9.0](#)
[Ansible project 8.0](#)
[Ansible project 7.0](#)
[Ansible project 6.0](#)
[Ansible project 5.0](#)
[Ansible project 4.0](#)
[Ansible project 3.0](#)
[Ansible project 2.10](#)
[Older Roadmaps](#)
[ansible-core Roadmaps](#)
[Ansible-core 2.19](#)
[Ansible-core 2.18](#)
[Ansible-core 2.17](#)
[Ansible-core 2.16](#)
[Ansible-core 2.15](#)
[Ansible-core 2.14](#)
[Ansible-core 2.13](#)
[Ansible-core 2.12](#)
[Ansible-core 2.11](#)
[Ansible-base 2.10](#)
[Next](#)
[© Copyright Ansible project contributors.](#)
[Last updated on Oct 08, 2025.](#)

Content from: https://docs.ansible.com/ansible/latest/collections/all_plugins.html

Indexes of all modules and plugins ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Index of all Become Plugins

Index of all Cache Plugins

Index of all Callback Plugins

Index of all Cliconf Plugins

Index of all Connection Plugins

Index of all Filter Plugins

[Index of all Httpapi Plugins](#)
[Index of all Inventory Plugins](#)
[Index of all Lookup Plugins](#)
[Index of all Modules](#)
[Index of all Netconf Plugins](#)
[Index of all Roles](#)
[Index of all Shell Plugins](#)
[Index of all Strategy Plugins](#)
[Index of all Test Plugins](#)
[Index of all Vars Plugins](#)
[Index of all deprecated collections and plugins](#)
[Playbook Keywords](#)
[Return Values](#)
[Ansible Configuration Settings](#)
[Controlling how Ansible behaves: precedence rules](#)
[YAML Syntax](#)
[Python 3 Support](#)
[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[ansible-core Roadmaps](#)
[Ansible](#)
[Indexes of all modules and plugins](#)
[Indexes of all modules and plugins](#)
[?](#)
[Plugin indexes](#)
[Index of all Become Plugins](#)
[Index of all Cache Plugins](#)
[Index of all Callback Plugins](#)
[Index of all Cliconf Plugins](#)
[Index of all Connection Plugins](#)
[Index of all Filter Plugins](#)
[Index of all Httpapi Plugins](#)
[Index of all Inventory Plugins](#)
[Index of all Lookup Plugins](#)
[Index of all Modules](#)
[Index of all Netconf Plugins](#)
[Index of all Roles](#)
[Index of all Shell Plugins](#)
[Index of all Strategy Plugins](#)
[Index of all Test Plugins](#)
[Index of all Vars Plugins](#)
[Index of all deprecated collections and plugins](#)
[Previous](#)

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/collections/community/general/merge_variables_lookup

community.general.merge_variables lookup ? Merge variables whose names match a given pattern ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Collections in the Amazon Namespace

Collections in the Ansible Namespace

Collections in the Arista Namespace

Collections in the Awx Namespace

Collections in the Azure Namespace

Collections in the Check_point Namespace

Collections in the Chocolatey Namespace
Collections in the Cisco Namespace
Collections in the Cloud Namespace
Collections in the Cloudscale_ch Namespace
Collections in the Community Namespace
Community.Aws
Community.Ciscosmb
Community.Crypto
Community.Digitalocean
Community.Dns
Community.Docker
Community.General
Description
Communication
Changelog
Guides
Technology Guides
Developer Guides
Plugin Index
Community.Grafana
Community.Hashi_Vault
Community.Hrobot
Community.Library_Inventory_Filtering_V1
Community.Libvirt
Community.Mongodb
Community.Mysql
Community.Okd
Community.Postgresql
Community.Proxmox
Community.Proxysql
Community.Rabbitmq
Community.Routeros
Community.Sap_Libs
Community.Sops
Community.Vmware
Community.Windows
Community.Zabbix
Collections in the Containers Namespace
Collections in the Cyberark Namespace
Collections in the Dellemc Namespace
Collections in the F5networks Namespace
Collections in the Fortinet Namespace
Collections in the Frr Namespace
Collections in the Gluster Namespace
Collections in the Google Namespace
Collections in the Grafana Namespace
Collections in the Hetzner Namespace
Collections in the Hitachivantara Namespace
Collections in the Hpe Namespace
Collections in the Ibm Namespace
Collections in the Ieisystem Namespace
Collections in the Infinidat Namespace
Collections in the Infoblox Namespace
Collections in the Inspur Namespace

Collections in the Junipernetworks Namespace

Collections in the Kaytus Namespace

Collections in the Kubernetes Namespace

Collections in the Kubevirt Namespace

Collections in the Lowlydba Namespace

Collections in the Mellanox Namespace

Collections in the Microsoft Namespace

Collections in the Netapp Namespace

Collections in the Netapp_eseries Namespace

Collections in the Netbox Namespace

Collections in the Nginx_io Namespace

Collections in the Openstack Namespace

Collections in the Openvswitch Namespace

Collections in the Ovirt Namespace

Collections in the Purestorage Namespace

Collections in the Ravendb Namespace

Collections in the Sensus Namespace

Collections in the Servicenow Namespace

Collections in the Splunk Namespace

Collections in the T_systems_mms Namespace

Collections in the Telekom_mms Namespace

Collections in the Theforeman Namespace

Collections in the Vmware Namespace

Collections in the Vultr Namespace

Collections in the Vyos Namespace

Collections in the Wti Namespace

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery

Releases and maintenance

Testing Strategies

Sanity Tests

Frequently Asked Questions

Glossary

Ansible Reference: Module Utilities

Special Variables

Red Hat Ansible Automation Platform

Ansible Automation Hub

Logging Ansible output

Roadmaps

Ansible Roadmap

ansible-core Roadmaps

Ansible

Collection Index

Collections in the Community Namespace

Community.General

community.general.merge_variables lookup ? Merge variables whose names match a given pattern

Edit on GitHub

community.general.merge_variables lookup ? Merge variables whose names match a given pattern

?

Note

This lookup plugin is part of the community.general collection (version 11.4.0).

You might already have this collection installed if you are using the ansible package.

It is not included in ansible-core

.

To check whether it is installed, run
ansible-galaxy
collection
list

.

To install it, use:

ansible-galaxy
collection
install
community.general

.

To use it in a playbook, specify:
community.general.merge_variables

.

New in community.general 6.5.0

Synopsis

Terms

Keyword parameters

Notes

Examples

Return Value

Synopsis

?

This lookup returns the merged result of all variables in scope that match the given prefixes, suffixes, or regular expressions, optionally.

Terms

?

Parameter

Comments

Terms

list

/

elements=string

/

required

Depending on the value of

pattern_type

, this is a list of prefixes, suffixes, or regular expressions that is used to match all variables that should be merged.

Keyword parameters

?

This describes keyword parameters of the lookup. These are the values

key1=value1

,

key2=value2

and so on in the following

examples:

```
lookup('community.general.merge_variables',
```

```
key1=value1,
```

```
key2=value2,
```

```
...)
```

and

```
query('community.general.merge_variables',
```

```
key1=value1,
```

```
key2=value2,
```

```
...)
```

Parameter

Comments

groups

list

/

elements=string

added in community.general 8.5.0

Search for variables across hosts that belong to the given groups. This allows to collect configuration pieces across different hosts (for example a service on a host with its database on another host).

initial_value

any

An initial value to start with.

override

string

Return an error, print a warning or ignore it when a key is overwritten.

The default behavior

error

makes the plugin fail when a key would be overwritten.

When

warn

and

ignore

are used, note that it is important to know that the variables are sorted by name before being merged. Keys for later variables in this order overwrite keys of the same name for variables earlier in this order. To avoid potential confusion, better use

override=error

whenever possible.

Choices:

"error"

? (default)

"warn"

"ignore"

Configuration:

INI entry:

```
[merge_variables_lookup]
```

```
override
```

```
=
```

```
error
```

Environment variable:

ANSIBLE_MERGE_VARIABLES_OVERRIDE

pattern_type

string

Change the way of searching for the specified pattern.

Choices:

"prefix"

"suffix"

"regex"

? (default)

Configuration:

INI entry:

[merge_variables_lookup]

pattern_type

=

regex

Environment variable:

ANSIBLE_MERGE_VARIABLES_PATTERN_TYPE

Note

Configuration entries listed above for each entry type (Ansible variable, environment variable, and so on) have a low to high priority order.

For example, a variable that is lower in the list will override a variable that is higher up.

The entry types are also ordered by precedence from low to high priority order.

For example, an ansible.cfg entry (further up in the list) is overwritten by an Ansible variable (further down in the list).

Notes

?

Note

When keyword and positional parameters are used together, positional parameters must be listed before keyword parameters:

lookup('community.general.merge_variables',

term1,

term2,

key1=value1,

key2=value2)

and

query('community.general.merge_variables',

term1,

term2,

key1=value1,

key2=value2)

Examples

?

Some example variables, they can be defined anywhere as long as they are in scope

test_init_list

:

-

"list

init

item

1"

-

"list

init

item

2"

testa__test_list

:

-

```

"test
a
item
1"
testb__test_list
:
-
"test
b
item
1"
testa__test_dict
:
ports
:
-
1
testb__test_dict
:
ports
:
-
3
# Merge variables that end with '__test_dict' and store the result in a variable 'example_a'
example_a
:
"
{{
lookup
(
'community.general.merge_variables'
,
'__test_dict'
,
pattern_type
=
'suffix'
)
}}
"

# The variable example_a now contains:
# ports:
# - 1
# - 3
# Merge variables that match the '^.__test_list$' regular expression, starting with an initial value and store the
# result in a variable 'example_b'
example_b
:
"
{{
lookup
(
'community.general.merge_variables'
,

```

```
'^.__test_list$'
```

```
,
```

```
initial_value
```

```
=
```

```
test_init_list
```

```
)
```

```
}}
```

```
"
```

```
# The variable example_b now contains:
```

```
# - "list init item 1"
```

```
# - "list init item 2"
```

```
# - "test a item 1"
```

```
# - "test b item 1"
```

```
Return Value
```

```
?
```

```
Key
```

```
Description
```

```
Return value
```

```
any
```

In case the search matches list items, a list is returned. In case the search matches dicts, a dict is returned.

```
Returned:
```

```
success
```

```
Authors
```

```
?
```

```
Roy Lenferink (@rlenferink)
```

```
Mark Ettema (@m-a-r-k-e)
```

```
Alexander Petrenz (@alpex8)
```

```
Collection links
```

```
?
```

```
Issue Tracker
```

```
Repository (Sources)
```

```
Ask for help
```

```
Submit a bug report
```

```
Request a feature
```

```
Communication
```

```
Previous
```

```
Next
```

```
© Copyright Ansible project contributors.
```

```
Last updated on Oct 08, 2025.
```

Content from: <https://docs.ansible.com/ansible/latest/collections/index.html>

[Collection Index ? Ansible Community Documentation](#)

[Blog](#)

[Ansible community forum](#)

[Documentation](#)

[Ansible Community Documentation](#)

[Ansible](#)

[Select version:](#)

[latest](#)

[11](#)

[devel](#)

[Search docs:](#)

[Ansible getting started](#)

[Getting started with Ansible](#)

[Getting started with Execution Environments](#)

[Installation, Upgrade & Configuration](#)

[Installation Guide](#)

[Ansible Porting Guides](#)

[Using Ansible](#)

[Building Ansible inventories](#)

[Using Ansible command line tools](#)

[Using Ansible playbooks](#)

[Protecting sensitive data with Ansible vault](#)

[Using Ansible modules and plugins](#)

[Using Ansible collections](#)

[Using Ansible on Windows, BSD, and z/OS UNIX](#)

[Ansible tips and tricks](#)

[Contributing to Ansible](#)

[Ansible Community Guide](#)

[Ansible Collections Contributor Guide](#)

[ansible-core Contributors Guide](#)

[Advanced Contributor Guide](#)

[Ansible documentation style guide](#)

[Extending Ansible](#)

[Developer Guide](#)

[Common Ansible Scenarios](#)

[Legacy Public Cloud Guides](#)

[Network Automation](#)

[Network Getting Started](#)

[Network Advanced Topics](#)

[Network Developer Guide](#)

[Ansible Galaxy](#)

[Galaxy User Guide](#)

[Galaxy Developer Guide](#)

[Reference & Appendices](#)

[Collection Index](#)

[Collections in the Amazon Namespace](#)

[Collections in the Ansible Namespace](#)

[Collections in the Arista Namespace](#)

[Collections in the Awx Namespace](#)

[Collections in the Azure Namespace](#)

[Collections in the Check_point Namespace](#)

[Collections in the Chocolatey Namespace](#)

Collections in the Cisco Namespace
Collections in the Cloud Namespace
Collections in the Cloudscale_ch Namespace
Collections in the Community Namespace
Collections in the Containers Namespace
Collections in the Cyberark Namespace
Collections in the Dellemc Namespace
Collections in the F5networks Namespace
Collections in the Fortinet Namespace
Collections in the Frr Namespace
Collections in the Gluster Namespace
Collections in the Google Namespace
Collections in the Grafana Namespace
Collections in the Hetzner Namespace
Collections in the Hitachivantara Namespace
Collections in the Hpe Namespace
Collections in the Ibm Namespace
Collections in the Ieismystem Namespace
Collections in the Infinidat Namespace
Collections in the Infoblox Namespace
Collections in the Inspur Namespace
Collections in the Junipernetworks Namespace
Collections in the Kaytus Namespace
Collections in the Kubernetes Namespace
Collections in the Kubevirt Namespace
Collections in the Lowlydba Namespace
Collections in the Mellanox Namespace
Collections in the Microsoft Namespace
Collections in the Netapp Namespace
Collections in the Netapp_eseries Namespace
Collections in the Netbox Namespace
Collections in the Nginx_io Namespace
Collections in the Openstack Namespace
Collections in the Openvswitch Namespace
Collections in the Ovirt Namespace
Collections in the Purestorage Namespace
Collections in the Ravendb Namespace
Collections in the Sensus Namespace
Collections in the Servicenow Namespace
Collections in the Splunk Namespace
Collections in the T_systems_mms Namespace
Collections in the Telekom_mms Namespace
Collections in the Theforeman Namespace
Collections in the Vmware Namespace
Collections in the Vultr Namespace
Collections in the Vyos Namespace
Collections in the Wti Namespace
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Collection Index
Collection Index
?

These are the collections with docs hosted on
docs.ansible.com

.

amazon.aws
ansible.builtin
ansible.netcommon
ansible.posix
ansible.utils
ansible.windows
arista.eos
awx.awx
azure.azcollection
check_point.mgmt
chocolatey.chocolatey
cisco.aci
cisco.dnac
cisco.intersight
cisco.ios
cisco.iosxr
cisco.meraki
cisco.mso
cisco.nxos
cisco.ucs
cloudscale_ch.cloud
community.aws
community.ciscosmb
community.crypto
community.digitalocean
community.dns
community.docker
community.general
community.grafana
community.hashi_vault
community.hrobot
community.library_inventory_filtering_v1
community.libvirt

community.mongodb
community.mysql
community.okd
community.postgresql
community.proxmox
community.proxysql
community.rabbitmq
community.routeros
community.sap_libs
community.sops
community.vmware
community.windows
community.zabbix
containers.podman
cyberark.conjur
cyberark.pas
dell EMC.enterprise_sonic
dell EMC.openmanage
dell EMC.powerflex
dell EMC.unity
f5networks.f5_modules
fortinet.fortimanager
fortinet.fortios
google.cloud
grafana.grafana
hetzner.hcloud
hitachivantara.vspone_block
hitachivantara.vspone_object
ibm.qradar
DEPRECATED: REMOVED IN 13
ibm.storage_virtualize
ieisystem.inmanage
infinidat.infinibox
infoblox.nios_modules
inspur.ispim
junipernetworks.junos
kaytus.ksmanage
kubernetes.core
kubevirt.core
lowlydba.sqlserver
microsoft.ad
microsoft.iis
netapp.cloudmanager
netapp.ontap
netapp.storagegrid
netapp_eseries.santricity
netbox.netbox
ngine_io.cloudstack
openstack.cloud
ovirt.ovirt
purestorage.flasharray
purestorage.flashblade
ravendb.ravendb
splunk.es

[telekom_mms.icinga_director](#)

[theforeman.foreman](#)

[vmware.vmware](#)

[vmware.vmware_rest](#)

[vultr.cloud](#)

[vyos.vyos](#)

[wti.remote](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/collections_guide/index.html

Using Ansible collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Installing collections

Removing a collection

Downloading collections

Listing collections

Verifying collections

Using collections in a playbook

Collections index

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

- Indexes of all modules and plugins
- Playbook Keywords
- Return Values
- Ansible Configuration Settings
- Controlling how Ansible behaves: precedence rules
- YAML Syntax
- Python 3 Support
- Interpreter Discovery
- Releases and maintenance
- Testing Strategies
- Sanity Tests
- Frequently Asked Questions
- Glossary
- Ansible Reference: Module Utilities
- Special Variables
- Red Hat Ansible Automation Platform
- Ansible Automation Hub
- Logging Ansible output
- Roadmaps
- Ansible Roadmap
- ansible-core Roadmaps
- Ansible
- Using Ansible collections
- Edit on GitHub
- Using Ansible collections
- ?

Note

Making Open Source More Inclusive

Red Hat is committed to replacing problematic language in our code, documentation, and web properties. We are beginning with these four terms: master, slave, blacklist, and whitelist. We ask that you open an issue or pull request if you come upon a term that we have missed. For more details, see our CTO Chris Wright's message

Welcome to the Ansible guide for working with collections.

Collections are a distribution format for Ansible content that can include playbooks, roles, modules, and plugins.

You can install and use collections through a distribution server, such as Ansible Galaxy, or a Pulp 3 Galaxy server.

Installing collections

Installing collections in containers

Installing collections with

ansible-galaxy

Installing collections with signature verification

Installing an older version of a collection

Install multiple collections with a requirements file

Downloading a collection for offline use

Installing collections adjacent to playbooks

Installing a collection from source files

Installing a collection from a Git repository

Configuring the

ansible-galaxy

client

Removing a collection

Downloading collections

Listing collections

Verifying collections

Verifying collections with
ansible-galaxy

Verifying signed collections

Using collections in a playbook

Simplifying module names with the
collections

keyword

Using
collections

in roles

Using
collections

in playbooks

Using a playbook from a collection

Collections index

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/community/code_of_conduct.html

Community Code of Conduct ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Getting started

Community Code of Conduct

Anti-harassment policy

Policy violations

Developer Certificate Of Origin

Communicating with the Ansible community

How can I help?

Other ways to get involved

Contributor path

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Community Guide
Getting started
Community Code of Conduct
Edit on GitHub
Community Code of Conduct
?

Topics
Community Code of Conduct
Anti-harassment policy
Policy violations

Every community can be strengthened by a diverse variety of viewpoints, insights, opinions, skillsets, and skill levels. However, with diversity comes the potential for disagreement and miscommunication. The purpose of this Code of Conduct is to ensure that disagreements and differences of opinion are conducted respectfully and on their own merits, without personal attacks or other behavior that might create an unsafe or unwelcoming environment.

These policies are not designed to be a comprehensive set of Things You Cannot Do. We ask that you treat your fellow community members with respect and courtesy, and in general, Don't Be A Jerk. This Code of Conduct is meant to be followed in spirit as much as in letter and is not exhaustive.

All Ansible events and participants therein are governed by this Code of Conduct and anti-harassment policy. We expect organizers to enforce these guidelines throughout all events, and we expect attendees, speakers, sponsors, and volunteers to help ensure a safe environment for our whole community. Specifically, this Code of Conduct covers participation in all Ansible-related forums and mailing lists, code and documentation contributions, public chat (Matrix, IRC), private correspondence, and public meetings.

Ansible community members are?

Considerate

Contributions of every kind have far-ranging consequences. Just as your work depends on the work of others, decisions you make surrounding your contributions to the Ansible

community will affect your fellow community members. You are strongly encouraged to take those consequences into account while making decisions.

Patient

Asynchronous communication can come with its own frustrations, even in the most responsive of communities. Please remember that our community is largely built on volunteered time, and that questions, contributions, and requests for support may take some time to receive a response. Repeated ?bumps? or ?reminders? in rapid succession are not good displays of patience. Additionally, it is considered poor manners to ping a specific person with general questions. Pose your question to the community as a whole, and wait patiently for a response.

Respectful

Every community inevitably has disagreements, but remember that it is possible to disagree respectfully and courteously. Disagreements are never an excuse for rudeness, hostility, threatening behavior, abuse (verbal or physical), or personal attacks.

Kind

Everyone should feel welcome in the Ansible community, regardless of their background. Please be courteous, respectful and polite to fellow community members. Do not make or post offensive comments related to skill level, gender, gender identity or expression, sexual orientation, disability, physical appearance, body size, race, or religion. Sexualized images or imagery, real or implied violence, intimidation, oppression, stalking, sustained disruption of activities, publishing the personal information of others without explicit permission to do so, unwanted physical contact, and unwelcome sexual attention are all strictly prohibited. Additionally, you are encouraged not to make assumptions about the background or identity of your fellow community members.

Inquisitive

The only stupid question is the one that does not get asked. We encourage our users to ask early and ask often. Rather than asking whether you can ask a question (the answer is always yes!), instead, simply ask your question. You are encouraged to provide as many specifics as possible. Code snippets in the form of Gists or other paste site links are almost always needed in order to get the most helpful answers. Refrain from pasting multiple lines of code directly into the chat channels - instead use gist.github.com or another paste site to provide code snippets.

Helpful

The Ansible community is committed to being a welcoming environment for all users, regardless of skill level. We were all beginners once upon a time, and our community cannot grow without an environment where new users feel safe and comfortable asking questions. It can become frustrating to answer the same questions repeatedly; however, community members are expected to remain courteous and helpful to all users equally, regardless of skill or knowledge level. Avoid providing responses that prioritize snideness and snark over useful information. At the same time, everyone is expected to read the provided documentation thoroughly. We are happy to answer questions, provide strategic guidance, and suggest effective workflows, but we are not here to do your job for you.

Anti-harassment policy

?

Harassment includes (but is not limited to) all of the following behaviors:

Offensive comments related to gender (including gender expression and identity), age, sexual orientation, disability, physical appearance, body size, race, and religion

Derogatory terminology including words commonly known to be slurs

Posting sexualized images or imagery in public spaces

Deliberate intimidation

Stalking

Posting others' personal information without explicit permission

Sustained disruption of talks or other events

Inappropriate physical contact

Unwelcome sexual attention

Participants asked to stop any harassing behavior are expected to comply immediately.

Sponsors are also subject to the anti-harassment policy. In particular, sponsors should not use sexualized images, activities, or other material. Meetup organizing staff and other volunteer organizers should not use sexualized attire or otherwise create a sexualized environment at community events.

In addition to the behaviors outlined above, continuing to behave a certain way after you have been asked to stop also constitutes harassment, even if that behavior is not specifically outlined in this policy. It is considerate and respectful to stop doing something after you have been asked to stop, and all community members are expected to comply with such requests immediately.

Policy violations

?

Instances of abusive, harassing, or otherwise unacceptable behavior may be reported by contacting

codeofconduct

@

ansible

.

com

, to anyone with administrative power in community chat (Admins or Moderators on Matrix, ops on IRC), or to the local organizers of an event. Meetup

organizers are encouraged to prominently display points of contact for reporting unacceptable behavior at local events.

If a participant engages in harassing behavior, the meetup organizers may take any action they deem appropriate. These actions may include but are not limited to warning the offender, expelling the offender from the event, and barring the offender from future community events.

Organizers will be happy to help participants contact security or local law enforcement, provide escorts to an alternate location, or otherwise assist those experiencing harassment to feel safe for the duration of the meetup. We value the safety and well-being of our community members and want everyone to feel welcome at our events, both online and offline.

We expect all participants, organizers, speakers, and attendees to follow these policies at all of our event venues and event-related social events.

The Ansible Community Code of Conduct is licensed under the Creative Commons Attribution-Share Alike 3.0 license. Our Code of Conduct was adapted from Codes of Conduct of other open source projects, including:

Contributor Covenant

Elastic

The Fedora Project

OpenStack

Puppet Labs

Ubuntu

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/community/collection_contributors/collection_contributors.html

Releasing collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

The Ansible Collections Development Cycle

Requesting changes to a collection

Creating your first collection pull request

Testing Collection Contributions

Review checklist for collection PRs

Ansible community package collections requirements

Guidelines for collection maintainers

Maintainer responsibilities

Expanding the collection community

Maintaining good collection documentation

Ansible Collection Maintenance and Workflow

Stepping down as a collection maintainer

Releasing collections

Preparing to release a collection

Collection versioning and deprecation

Collection changelogs

Options for releasing a collection

Contributing to Ansible-maintained Collections

Ansible Community Steering Committee

Contributing to the Ansible Documentation

Other Tools and Programs

Working with the Ansible collection repositories

ansible-core Contributors Guide

Advanced Contributor Guide
Ansible documentation style guide
Extending Ansible
Developer Guide
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Collections Contributor Guide
Guidelines for collection maintainers
Releasing collections
Edit on GitHub
Releasing collections
?
Collection maintainers release all supported stable versions of the collections regularly,
provided that there have been enough changes merged to release.
Preparing to release a collection
Collection versioning and deprecation
Collection changelogs
Options for releasing a collection
Releasing without release branches
Hybrid approach
Releasing with release branches
Preparing to release a collection
?

The collections under the
ansible-collections organization
follow
semantic versioning
when releasing. See
Collection versioning and deprecation
for details.

To prepare for a release, a collection must have:

A publicly available policy of releasing, versioning, and deprecation. This can be, for example, written in its README or in a dedicated pinned issue.

A pinned issue when its release managers inform the community about planned or completed releases. This can be combined with the release policy issue mentioned above.

A
changelog

Releases of the collection tagged in the collection's repository.

CI pipelines up and running. This can be implemented by using GitHub Actions, Azure Pipelines, Zuul.

All CI tests running against a commit that releases the collection. If they do not pass, the collection **MUST NOT** be released.

See

Including a collection in Ansible

if you plan on adding a new collection to the Ansible package.

Note

Your collection must pass

ansible-test

sanity

tests. See

Testing collections

for details.

Collection versioning and deprecation

?

Note

Collections **MUST** adhere to

semantic versioning

To preserve backward compatibility for users, every Ansible minor version series (5.1.x, 5.2.x, and so on) will keep the major version of a collection constant. For example, if Ansible 5.0.0 includes

community.general

4.0.2, then each Ansible 5.X.x release will include the latest

community.general

4.y.z release available at build time. Ansible 5.x.x will

never

include a

community.general

5.y.x release, even if it is available. Major collection version changes will be included in the next Ansible major release (6.0.0 in this case).

Ensure that the current major release of your collection included in 6.0.0 receives at least bugfixes as long as new Ansible 6.X.X releases are produced.

Since new minor releases are included, you can include new features, modules and plugins. You must make sure that you do not break backward compatibility. See

semantic versioning

for more details. This means in particular:

You can fix bugs in

patch

releases but not add new features or deprecate things.

You can add new features and deprecate things in

minor

releases, but not remove things or change the behavior of existing features.

You can only remove things or make breaking changes in

major

releases.

Ensure that if a deprecation is added in a collection version that is included in 5.x.y, the removal itself will only happen in a collection version included in 7.0.0 or later.

Ensure that the policy of releasing, versioning, and deprecation is announced to contributors and users in some way.

For an example of how to do this, see

the announcement in `community.general`

. You could also do this in the collection README file.

Collection changelogs

?

Collections **MUST** include a changelog. To give a consistent feel for changelogs across collections and ensure changelogs exist for collections included in the

`ansible`

package, we suggest you use

`antsibull-changelog`

to maintain and generate this.

Before releasing, verify the following for your changelogs:

All merged pull requests since the last release, except ones related to documentation and new modules/plugins, have changelog fragments

.

New module and plugin pull requests, except `jinja2` test and filter plugins, do

not

need a changelog fragment, they are auto-detected by the changelog generator by their

`version_added`

value.

All the fragments follow the

changelog entry format

.

Options for releasing a collection

?

There are several approaches on how to release a collection. If you are not aware of which approach to use, ask in the `#ansible-community`

IRC channel or the

community

Matrix channel.

This section assumes that publishing the collection is done with

`Zuul`

and that

`antsibull-changelog`

is used for the changelog.

Releasing without release branches

?

Use releasing without release branches when:

There are no prior major releases of the collection.

There are no breaking changes introduced since the

1.0.0

release of the collection.

See

Releasing collections without release branches

for details.

When there is a need to introduce breaking changes, you can switch to the next approach.

Hybrid approach

?

In this approach, releases for the current major version are made from the main

branch, while new releases for older major versions are made from release branches for these versions.

Releasing with release branches

?

Use releasing with release branches when breaking changes have been introduced. This approach is usually only used by the large community collections,

community.general

and

community.network

.

See

Releasing collections with release branches

for details.

Releasing collections without release branches

Releasing collections with release branches

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: <https://docs.ansible.com/ansible/latest/community/communication.html>

Communicating with the Ansible community ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Getting started

Community Code of Conduct

Developer Certificate Of Origin

Communicating with the Ansible community

Code of Conduct

Forum

Real-time chat

Working groups

Ansible Community Topics

Ansible Automation Platform support questions

How can I help?

Other ways to get involved

Contributor path

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Community Guide
Getting started
Communicating with the Ansible community
Edit on GitHub
Communicating with the Ansible community
?
Code of Conduct
Forum
The Bullhorn
Regional and Language-specific channels
Real-time chat
General channels
Working group-specific channels
Meetings on Matrix
Working groups
Requesting a forum group
Requesting a community collection repository
Ansible Community Topics
Ansible Automation Platform support questions
Code of Conduct
?
All interactions within the Ansible Community are governed by our
Community Code of Conduct
. Please read and understand it before participating.
Forum
?

The Ansible Forum is the default and recommended starting point for most community interactions. It's ideal for:

- Asking questions and seeking help.
- Participating in development discussions.
- Learning about events and news.

To get started:

Register:
Sign up to create an account and join the community.

Explore topics:
Browse by categories and tags to discover discussions, or simply start a new topic of your own.

Stay updated:
Subscribe to specific categories or tags that interest you. Just click the bell icon in the top-right corner of the relevant category or tag page and select your notification preference.

Explore forum groups that match your interests. Joining a group often automatically subscribes you to related posts.

The Bullhorn?

The Bullhorn is our community's weekly newsletter, published directly in the Forum:

Subscribe:
Click the bell button under the Bullhorn category description, then select Watching.

.
Submit News:
See the About the Newsletter category post for submission guidelines.
Questions about the newsletter:
Ask us in the Ansible Social room on Matrix.

.
Regional and Language-specific channels?

Communicate in your preferred language by visiting the International Communities forum category.

. Current subcategories include:

- Deutsche (German)
- Español (Spanish)
- Français (French)
- Italiano (Italian)

Norsk (Norwegian)
Português (Portuguese)

Join an
Ansible Meetup
near you.

For details on requesting a new language subcategory, see the
About the International Communities category post

Real-time chat

?

For real-time interactions, the Ansible community uses the
Matrix protocol

Note
The
Forum
is our default communication platform. We recommend engaging there before considering other options like Matrix.
To join the community on Matrix:

Get a Matrix account:

From
Matrix.org
or any other Matrix homeserver.

Choose a Matrix client:

We recommend
Element Webchat

Join rooms:

Use the links in the

General channels

or

Working groups

to join specific rooms.

For more information, see the community-hosted
Matrix FAQ

You can add Matrix shields to your repository?

README.md

using the shield in the
community-topics
repository as a template.

Note

IRC channels are no longer official communication channels. Use the Forum and Matrix instead.

General channels

?

The clickable links below take you directly to the Matrix room in your browser. Room/channel information is also given
for use in other clients:

Community social room & posting news for the Bullhorn newsletter

General usage & support questions

Developer & code-related topics

Community & collections related topics

Working group-specific channels

?

Many working groups have dedicated chat channels. See the
Working groups

for details.

Meetings on Matrix

?

The Ansible community holds regular meetings on Matrix. All interested individuals are invited to participate.

Check the

meeting schedule and agenda page

for more information.

Working groups

?

Working Groups enable Ansible community members to self-organize around specific interests.

Find a complete list of groups and their communication channels within the

Forum groups

.

Requesting a forum group

?

To request a new working group:

First, check if there is no appropriate

Forum group

you can join instead of starting a new one.

Review the

things you can ask for post

regarding working groups.

Submit your request in the

forum topic

.

If a Matrix chat channel is also needed, consult the

Ansible Community Matrix FAQ

.

Requesting a community collection repository

?

Working groups are often built around Ansible community collections. You can use a repository under your organization

or request one under

ansible-collections

on the forum. Create a topic in the

Project Discussions category and the ?coll-repo-request? tag

.

Ansible Community Topics

?

The

Ansible Community Steering Committee

uses the

Forum

for asynchronous discussions and voting on community topics.

For more information, see:

Creating new policy proposals & inclusion requests

Community topics workflow

Community topics on the Forum

Ansible Automation Platform support questions

?

Red Hat Ansible

Automation Platform

is a subscription service providing support, certified content, and tooling for Ansible, including content management, a controller, UI and REST API.

For questions related to Ansible Automation Platform, visit

Red Hat support
instead of community communication platforms.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/community/contributions_collections.html

[Ansible Collections Contributor Guide](#) ? [Ansible Community Documentation](#)

[Blog](#)

[Ansible community forum](#)

[Documentation](#)

[Ansible Community Documentation](#)

[Ansible](#)

Select version:

[latest](#)

[11](#)

[devel](#)

Search docs:

[Ansible getting started](#)

[Getting started with Ansible](#)

[Getting started with Execution Environments](#)

[Installation, Upgrade & Configuration](#)

[Installation Guide](#)

[Ansible Porting Guides](#)

[Using Ansible](#)

[Building Ansible inventories](#)

[Using Ansible command line tools](#)

[Using Ansible playbooks](#)

[Protecting sensitive data with Ansible vault](#)

[Using Ansible modules and plugins](#)

[Using Ansible collections](#)

[Using Ansible on Windows, BSD, and z/OS UNIX](#)

[Ansible tips and tricks](#)

[Contributing to Ansible](#)

[Ansible Community Guide](#)

[Ansible Collections Contributor Guide](#)

[The Ansible Collections Development Cycle](#)

[Requesting changes to a collection](#)

[Creating your first collection pull request](#)

[Testing Collection Contributions](#)

[Review checklist for collection PRs](#)

[Ansible community package collections requirements](#)

[Guidelines for collection maintainers](#)

[Contributing to Ansible-maintained Collections](#)

[Ansible Community Steering Committee](#)

[Contributing to the Ansible Documentation](#)

[Other Tools and Programs](#)

[Working with the Ansible collection repositories](#)

[ansible-core Contributors Guide](#)

[Advanced Contributor Guide](#)

[Ansible documentation style guide](#)

[Extending Ansible](#)

[Developer Guide](#)

[Common Ansible Scenarios](#)

[Legacy Public Cloud Guides](#)

[Network Automation](#)

[Network Getting Started](#)

[Network Advanced Topics](#)

[Network Developer Guide](#)

[Ansible Galaxy](#)
[Galaxy User Guide](#)
[Galaxy Developer Guide](#)
[Reference & Appendices](#)
[Collection Index](#)
[Indexes of all modules and plugins](#)
[Playbook Keywords](#)
[Return Values](#)
[Ansible Configuration Settings](#)
[Controlling how Ansible behaves: precedence rules](#)
[YAML Syntax](#)
[Python 3 Support](#)
[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[ansible-core Roadmaps](#)
[Ansible](#)
[Ansible Collections Contributor Guide](#)
[Edit on GitHub](#)
[Ansible Collections Contributor Guide](#)
[?](#)
[The Ansible Collections Development Cycle](#)
[Macro development: roadmaps, releases, and projects](#)
[Micro development: the lifecycle of a PR](#)
[Making your PR merge-worthy](#)
[Requesting changes to a collection](#)
[Reporting a bug](#)
[Requesting a feature](#)
[Creating your first collection pull request](#)
[Prepare your environment](#)
[Change the code](#)
[Fix the bug](#)
[Test your changes](#)
[Submit a pull request](#)
[Testing Collection Contributions](#)
[How to test a collection PR](#)
[Add unit tests to a collection](#)
[Adding integration tests to a collection](#)
[Review checklist for collection PRs](#)
[Reviewing bug reports](#)
[Reviewing suggested changes](#)
[Review tests in the PR](#)
[Review for merge commits and breaking changes](#)
[Ansible community package collections requirements](#)

Overview

Feedback and communications

Keeping informed

Communication and Working Groups

Collection infrastructure

Python Compatibility

Standards for developing module and plugin utilities

Repository structure requirements

Contributor Workflow

Naming

Collection licensing requirements

Contributor License Agreements

Repository management

CI Testing

When moving modules between collections

Development conventions

Collection Dependencies

Other requirements

Guidelines for collection maintainers

Maintainer responsibilities

Expanding the collection community

Maintaining good collection documentation

Ansible Collection Maintenance and Workflow

Stepping down as a collection maintainer

Releasing collections

Contributing to Ansible-maintained Collections

Ansible-maintained collections

Deciding where your contribution belongs

Requirements to merge your PR

Ansible Community Steering Committee

Steering Committee mission and responsibilities

Steering Committee membership guidelines

Steering Committee past members

Community topics workflow

Contributing to the Ansible Documentation

Editing docs directly on GitHub

Reviewing or solving open issues

Reviewing open PRs

Opening a new issue and/or PR

Verifying your documentation PR

Joining the documentation working group

Other Tools and Programs

Popular editors

Tools for validating playbooks

Collection development tools

Other tools

If you have a specific Ansible interest or expertise (for example, VMware, Linode, and so on), consider joining a working group

.

Working with the Ansible collection repositories

?

How can I find

editors, linters, and other tools

that will support my Ansible development efforts?

Where can I find guidance on
coding in Ansible

?

How do I
create a collection

?

How do I
rebase my PR

?

How do I learn about Ansible's
testing (CI) process

?

How do I
deprecate a module

?

See

Collection developer tutorials

for a quick introduction on how to develop and test your collection contributions.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/community/contributor_path.html

Contributor path ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Getting started

Contributor path

Determine your area of interest

Find the corresponding project

Learn

Specific knowledge for code developers

Making your first contribution

Continue to contribute

Teach others

Become a collection maintainer

Become a steering committee member

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Community Guide
Contributor path
Edit on GitHub
Contributor path
?

This section describes the contributor's journey from the beginning to becoming a leader who helps shape the future of Ansible. You can use this path as a roadmap for your long-term participation.

Any contribution to the project, even a small one, is very welcome and valuable. Any contribution counts, whether it is feedback on an issue, a pull request, a topic or documentation change, or a coding contribution. When you contribute regularly, your proficiency and judgment in the related area increase and, along with this, the importance of your presence in the project.

Determine your area of interest

Find the corresponding project

Learn

Specific knowledge for code developers

Making your first contribution

Continue to contribute

Teach others

Become a collection maintainer

Become a steering committee member

Determine your area of interest

?

First, determine areas that are interesting to you. Consider your current experience and what you'd like to gain. For example, if you use a specific collection, have a look there. See

How can I help?

for more ideas on how to help.

Find the corresponding project

?

These are multiple community projects in the Ansible ecosystem you could contribute to:

Ansible Core

Collections

AWX

Galaxy

ansible-lint

Molecule

Learn

?

The required skillset depends on the area of interest and the project you'll be working on. Remember that the best way to learn is by doing.

Specific knowledge for code developers

?

Code development requires the most technical knowledge. Let's sort out what an Ansible developer should learn.

You should understand at least the

basics

of the following tools:

Python programming language

Git

GitHub collaborative development model through forks and pull requests

You can learn these tools more in-depth when working on your first contributions.

Each Ansible project has its own set of contributor guidelines. Familiarize yourself with these as you prepare your first contributions.

Ansible Core development

.

Ansible collection development

and the collection-level contributor guidelines in the collection repository.

Making your first contribution

?

You can find some ideas on how you can contribute in

How can I help?

.

If you are interested in contributing to collections, take a look at

collection contributions

and the

collection repository

's

README

and

CONTRIBUTING

files. To make your first experience as smooth as possible, read the repository documentation carefully, then ask the repository maintainers for guidance if you have any questions.

Take a look at GitHub issues labeled with the

easyfix

and

good_first_issue

labels for:

Ansible collections repositories

All other Ansible projects

Issues labeled with the

docs

label in

Ansible collections

and

other

Ansible projects can be also good to start with.

When you choose an issue to work on, add a comment directly on the GitHub issue to say you are looking at it and let others know to avoid conflicting work.

You can also ask for help in a comment if you need it.

Continue to contribute

?

We don't expect everybody to know everything. Start small, think big. When you contribute regularly, your proficiency and judgment in the related area will improve quickly and, along with this, the importance of your presence in the project.

See

Communicating with the Ansible community

for ways to communicate and engage with the Ansible community, including working group meetings, accessing the Bullhorn news bulletin, and upcoming contributor summits.

Teach others

?

Share your experience with other contributors through

improving documentation

, answering questions from other contributors and users on

Matrix/Libera.Chat IRC

, giving advice on issues and pull requests, and discussing topics on the

Forum

.

Become a collection maintainer

?

If you are a code contributor to a collection, you can get extended permissions in the repository and become a maintainer. A collection maintainer is a contributor trusted by the community who makes significant and regular contributions to the project and showed themselves as a specialist in the related area. See

Guidelines for collection maintainers

for details.

For some collections that use the

collection bot

, such as

community.general

and

community.network

, you can have different levels of access and permissions:

File-level permissions: the stage prior to becoming a collection maintainer. The file is usually a module or plugin. File maintainers have indirect commit rights.

Supershipit permissions: similar to being a file maintainer but the scope where a maintainer has the indirect commit is the whole repository.

Triage access to the repository: allows contributors to manage issues and pull requests.

Write access to the repository also known as

commit

: allows contributors to merge pull requests to the development branch as well as perform all the other activities listed in the

Guidelines for collection maintainers

.

For information about permission levels, see the

GitHub official documentation

.

Become a steering committee member

?

Note

You do NOT have to be a programmer to become a steering committee member.

The

Steering Committee

member status reflects the highest level of trust and allows contributors to lead the project by making important decisions for the Ansible project. The Committee members are community leaders who shape the project's future and the future of automation in the IT world in general.

To reach the status, as the current Committee members did before getting it, along with the things mentioned in this document, you should:

Subscribe to, comment on, and vote on the
community topics<creating_community_topic>

.

Propose your topics.

If time permits, join the

Community meetings

. Note this is

NOT

a requirement.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/community/how_can_I_help.html

How can I help? ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Getting started

Community Code of Conduct

Developer Certificate Of Origin

Communicating with the Ansible community

How can I help?

Become a power user

Ask and answer questions online

Review, fix, and maintain the documentation

Participate in your local meetup

File and verify issues

Review and submit pull requests

Become a collection maintainer

Join a working group

Teach Ansible to others

Social media

Other ways to get involved

Contributor path

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides
Network Automation
Network Getting Started
Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Community Guide
Getting started
How can I help?
Edit on GitHub
How can I help?
?
Become a power user
Ask and answer questions online
Review, fix, and maintain the documentation
Participate in your local meetup
File and verify issues
Review and submit pull requests
Become a collection maintainer
Join a working group
Teach Ansible to others
Social media
Thanks for being interested in helping the Ansible project!
There are many ways to help the Ansible project?but first, please read and understand the
Community Code of Conduct
.
Become a power user
?

A great way to help the Ansible project is to become a power user:

Use Ansible everywhere you can

Take tutorials and classes

Read the

official documentation

Study some of the

many excellent books

about Ansible

Get certified

.

When you become a power user, your ability and opportunities to help the Ansible project in other ways will multiply quickly.

Ask and answer questions online

?

There are many online platforms where Ansible users and contributors ask and answer questions including the

Forum

.

Reach out and communicate with your fellow Ansible enthusiasts.

You can find the official

Ansible communication channels

.

Review, fix, and maintain the documentation

?

Typos are everywhere, even in the Ansible documentation. We work hard to keep the documentation up-to-date, but you may also find outdated examples. We offer easy ways to

report and/or fix documentation errors

.

Participate in your local meetup

?

There are Ansible meetups

all over the world

. Join your local meetup. Attend regularly. Ask good questions. Volunteer to give a presentation about how you use Ansible.

If there is no meetup near you, we are happy to help you start one

.

File and verify issues

?

All software has bugs, and Ansible is no exception. When you find a bug, you can help tremendously by telling us about it:

Filing

issues for ansible-core

.

Filing

issues for collections

.

If the bug you found already exists in an issue, you can help by verifying the behavior of the reported bug with a comment in that issue, or by reporting any additional information.

Review and submit pull requests

?

As you become more familiar with how Ansible works, you may be able to fix issues or develop new features yourself. If you think you have a fix for a bug in Ansible, or if you have a new feature that you would like to share with millions of Ansible users, read all about the development process

to learn how to get your code accepted into Ansible.

You can also get started with solving GitHub issues labeled with the

[easyfix](#)

and

[good_first_issue](#)

labels for:

[Ansible collections](#)

[All other Ansible projects](#)

When you choose an issue to work on, add a comment directly on the GitHub issue to say you are looking at it and let others know to avoid conflicting work.

You can also ask for help in a comment if you need it.

For collections, refer to the

[Creating your first collection pull request](#)

page to learn how to quickly set up your local environment, test your changes, and submit a ready-for-review pull request.

Another good way to help is to review pull requests that other Ansible users have submitted. Ansible Core keeps a full list of

[open pull requests by file](#)

, so if a particular module or plugin interests you, you can easily keep track of all the relevant new pull requests and provide testing or feedback. Alternatively, you can review the pull requests for any collections that interest you. Click

[Issue tracker](#)

on the collection documentation page to find the issues and PRs for that collection.

[Become a collection maintainer](#)

?

Once you have learned about the development process and have contributed code to a collection, we encourage you to become a maintainer of that collection. There are hundreds of modules in dozens of Ansible collections, and the vast majority of them are written and maintained entirely by members of the Ansible community.

See

[collection maintainer guidelines](#)

to learn more about the responsibilities of being an Ansible collection maintainer.

[Join a working group](#)

?

Working groups are a way for Ansible community members to self-organize around particular topics of interest. We have working groups around various topics. To join or create a working group, please read the

[Ansible Working Groups](#)

.

[Teach Ansible to others](#)

?

We are working on a standardized

[Ansible workshop](#)

that can provide a good hands-on introduction to Ansible usage and concepts.

[Social media](#)

?

If you like Ansible and just want to spread the good word, feel free to share on your social media platform of choice, and let us know by using

[@ansible](#)

or

[#ansible](#)

. We'll be looking for you.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/community/maintainers_guidelines.html

Maintainer responsibilities ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

The Ansible Collections Development Cycle

Requesting changes to a collection

Creating your first collection pull request

Testing Collection Contributions

Review checklist for collection PRs

Ansible community package collections requirements

Guidelines for collection maintainers

Maintainer responsibilities

Collection maintainer definition

Maintainer responsibilities

Becoming a maintainer

Communicating as a maintainer

Establishing working group communication

Participating in community topics

Expanding the collection community

Maintaining good collection documentation

Ansible Collection Maintenance and Workflow

Stepping down as a collection maintainer

Releasing collections

Contributing to Ansible-maintained Collections

Ansible Community Steering Committee

Contributing to the Ansible Documentation

Other Tools and Programs

Working with the Ansible collection repositories

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery

Releases and maintenance

Testing Strategies

Sanity Tests

Frequently Asked Questions

Glossary

Ansible Reference: Module Utilities

Special Variables

Red Hat Ansible Automation Platform

Ansible Automation Hub

Logging Ansible output

Roadmaps

Ansible Roadmap

ansible-core Roadmaps

Ansible

Ansible Collections Contributor Guide

Guidelines for collection maintainers

Maintainer responsibilities

Edit on GitHub

Maintainer responsibilities

?

This document provides guidance for:

Contributors to collections who want to join maintainer teams.

Collection maintainers seeking to understand their roles better.

This document defines the role of an Ansible collection maintainer, outlines their responsibilities, and describes the process for becoming one.

Collection maintainer definition

Maintainer responsibilities

Becoming a maintainer

Communicating as a maintainer

Establishing working group communication

Participating in community topics

Collection maintainer definition

?

An Ansible collection maintainer, or simply maintainer, is a contributor who:

Makes significant and regular contributions to a project.

Demonstrates expertise in the area the collection automates.

Earns the community's trust. To fulfill their duties, maintainers have write

or higher access to the collection.

Maintainer responsibilities

?

Collection maintainers perform the following tasks:

Act in accordance with the

Community Code of Conduct

.

Subscribe to the repository they maintain. In GitHub, click Watch > All activity

.

Keep the

README

, development guidelines, and other general

Maintaining good collection documentation

current.

Review and commit changes from other contributors using the

Review checklist for collection PRs

.

Backport

changes to stable branches.

Plan and perform releases

.

Ensure the collection adheres to the

Ansible community package collections requirements

.

Track changes announced through the

news-for-maintainers

forum tag. Click the

Bell

button to subscribe. Update the collection accordingly.

Build a healthy community

to increase the number of active contributors and maintainers for collections.

Multiple maintainers can divide these responsibilities among themselves.

Becoming a maintainer

?

If you are interested in becoming a maintainer and meet the requirements

, nominate yourself. You can also nominate another person by following these steps:

Create a GitHub issue in the relevant repository.

If you receive no response, message the

Red Hat Ansible Community Engineering Team

on the

Ansible forum

.

Communicating as a maintainer

?

Maintainers communicate with the community through the channels listed in the Ansible communication guide

.

Establishing working group communication

?

Working groups rely on efficient communication. As a maintainer, you can establish communication for your working groups using these techniques:

Find and join an existing

forum group

and use tags that suit your project.

If no suitable options exist,

request them

.

Provide working group details and chat room links in the contributor section of your project's

README.md

.

Encourage contributors to join the forum group and use appropriate tags.

Participating in community topics

?

The Community and the

Steering Committee

discuss and vote on

community topics

asynchronously. These topics impact the entire project or its components, including collections and packaging.

Share your opinion and vote on the topics to help the community make informed decisions.

Expanding the collection community

?

You can expand the community around your collection in the following ways:

Explicitly state in your

README

that the collection welcomes new maintainers and contributors.

Give

newcomers a positive first experience

.

Invite contributors to join forum groups and subscribe to tags related to your project.

Maintain

good documentation

with guidelines for new contributors.

Make people feel welcome personally and individually. Greet and thank them.

Use labels to identify easy fixes and leave non-critical easy fixes to newcomers.

Offer help explicitly.

Include quick ways contributors can help and provide contributor documentation references in your

README

.

Be responsive in issues, pull requests (PRs), and other communication channels.

Conduct PR days regularly.

Maintain a zero-tolerance policy toward behavior that violates the

Community Code of Conduct

.

* Include information about how people can report code of conduct violations in your

README

and

CONTRIBUTING

files.

Look for new maintainers among active contributors.

Maintaining good collection documentation

?

Ensure the collection documentation meets these criteria:

It is up-to-date.

It matches the

Ansible documentation style guide

.

Collection module and plugin documentation adheres to the
Ansible documentation format

.

Collection user guides follow the
Collection documentation format

.

Repository files include at least a
README
and
CONTRIBUTING
file.

The

README

file contains all sections from
collection_template/README.md

.

The

CONTRIBUTING

file includes all details or links to details on how new or continuing contributors can contribute to your collection.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/community/maintainers_workflow.html

Ansible Collection Maintenance and Workflow ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

The Ansible Collections Development Cycle

Requesting changes to a collection

Creating your first collection pull request

Testing Collection Contributions

Review checklist for collection PRs

Ansible community package collections requirements

Guidelines for collection maintainers

Maintainer responsibilities

Expanding the collection community

Maintaining good collection documentation

Ansible Collection Maintenance and Workflow

Managing pull requests

Releasing a collection

Backporting

Including a collection in Ansible

Stepping down as a collection maintainer

Releasing collections

Contributing to Ansible-maintained Collections

Ansible Community Steering Committee

Contributing to the Ansible Documentation

Other Tools and Programs

Working with the Ansible collection repositories

ansible-core Contributors Guide

Advanced Contributor Guide
Ansible documentation style guide
Extending Ansible
Developer Guide
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Collections Contributor Guide
Guidelines for collection maintainers
Ansible Collection Maintenance and Workflow
Edit on GitHub
Ansible Collection Maintenance and Workflow
?
Each collection community can set its own rules and workflows for managing pull requests (PRs), bug reports, documentation issues, feature requests, as well as for adding and replacing maintainers.
Collection maintainers have
write
or higher access to a collection, allowing them to merge pull requests and perform other administrative tasks.
Managing pull requests
?
Maintainers review and merge PRs according to the following guidelines:
Community Code of Conduct
Maintainer responsibilities
Committer guidelines

PR review checklist

Releasing a collection

?

Collection maintainers are responsible for releasing new collection versions. The general release process includes:

Planning and announcement

: Define the release scope and communicate it.

Changelog generation

: Create a comprehensive list of changes.

Git tagging

: Create and push a release Git tag.

Automated publication

: The release tarball is automatically published on

Ansible Galaxy

via the

Zuul dashboard

or a custom GitHub Actions workflow.

Final announcement

: Communicate the successful release.

See

Releasing collections

for more information.

Backporting

?

Collection maintainers backport merged pull requests to stable branches if they exist. This process adheres to the collection's

semantic versioning

and release policies.

The manual backporting process mirrors the

ansible-core backporting guidelines

.

For streamlined backporting, GitHub bots like the

Patchback app

can automate the process through labeling, as implemented in the

community.general

collection.

Including a collection in Ansible

?

To include a collection in the Ansible community package, maintainers create a discussion in the ansible-collections/ansible-inclusion repository

. See the

submission process

and the

Ansible community package collections requirements

for details.

Stepping down as a collection maintainer

?

If you can no longer continue as a collection maintainer, follow these steps:

Inform other maintainers

: Notify your co-maintainers.

Notify the community

: For collections under the

ansible-collections

GitHub organization, inform the relevant

Real-time chat

channels, or email
ansible-community@redhat.com

.

Identify potential replacements

: Look for active contributors within the collection who could become new maintainers. Discuss these candidates with other maintainers or the Ansible community team

.

Announce the need for maintainers (if no replacement is found)

: If you cannot find a replacement, create a pinned issue in the collection repository announcing the need for new maintainers.

Post in the Bullhorn newsletter

: Make the same announcement through the Bullhorn newsletter

.

Engage in candidate discussions

: Be available to discuss potential candidates identified by other maintainers or the community team. Remember, this is a community and you are welcome to rejoin at any time.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/developing_api.html

Python API ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[ansible-core Roadmaps](#)
[Ansible](#)
[Developer Guide](#)
[Python API](#)
[Edit on GitHub](#)
[Python API](#)
[?](#)

[Topics](#)

[Python API](#)

[Attention](#)

The Ansible API is intended for internal Ansible use. Ansible may make changes to this API at any time that could break backward compatibility with older versions of the API. Because of this, external use is not supported by Ansible. If you want to use Python API only for executing playbooks or modules, consider [ansible-runner](#)

first.

If you would like to use Ansible programmatically from a language other than Python, trigger events asynchronously, or have access control and logging demands, please see the [AWX project](#)

.

[See also](#)

[Developing dynamic inventory](#)

[Developing dynamic inventory integrations](#)

[Developing modules](#)

[Getting started on developing a module](#)

[Developing plugins](#)

[How to develop plugins](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/developing_collections.html

Developing collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

- Interpreter Discovery
- Releases and maintenance
- Testing Strategies
- Sanity Tests
- Frequently Asked Questions
- Glossary
- Ansible Reference: Module Utilities
- Special Variables
- Red Hat Ansible Automation Platform
- Ansible Automation Hub
- Logging Ansible output
- Roadmaps
- Ansible Roadmap
- ansible-core Roadmaps
- Ansible
- Developer Guide
- Developing collections
- Edit on GitHub
- Developing collections
- ?

Collections are a distribution format for Ansible content. You can package and distribute playbooks, roles, modules, and plugins using collections. A typical collection addresses a set of related use cases. For example, the `cisco.ios`

collection automates management of Cisco IOS devices.

You can create a collection and publish it to

Ansible Galaxy

or to a private Automation Hub instance. You can publish certified collections to the Red Hat Automation Hub, part of the Red Hat Ansible Automation Platform.

Examine the

Ansible collection creator path

to understand how to go from creating a collection to having it included in the Ansible package distribution.

Developing new collections

Creating collections

Naming your collection

Creating a new collection

Creating a collection from a custom template

Creating collections with `ansible-creator`

Using shared resources in collections

Using documentation fragments in collections

Leveraging optional module utilities in collections

Listing collection dependencies

Testing collections

Testing tools

Distributing collections

Initial configuration of your distribution server or servers

Building your collection tarball

Preparing to publish your collection

Publishing your collection

Documenting collections

Documenting modules and plugins

Documenting roles

Verifying your collection documentation

Build a docsite with `antsibull-docs`

Working with existing collections

[Migrating Ansible content to a different collection](#)

[Migrating content](#)

[Contributing to collections](#)

[Contributing to a collection: community.general](#)

[Generating changelogs and porting guide entries in a collection](#)

[Understanding ansible-changelog](#)

[Including collection changelogs into Ansible](#)

[Collections references](#)

[Collection structure](#)

[Collection directories and files](#)

[Collection Galaxy metadata structure](#)

[Structure](#)

[Examples](#)

[For instructions on developing modules, see](#)

[Developing modules](#)

.

[See also](#)

[Using Ansible collections](#)

[Learn how to install and use collections in playbooks and roles](#)

[Contributing to Ansible-maintained Collections](#)

[Guidelines for contributing to selected collections](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/developing_collections_creating.h

Creating collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing collections
Creating collections
Edit on GitHub
Creating collections
?

To create a collection:

Create a
new collection
, optionally using a custom
collection template
, with the
ansible-galaxy
collection
init
command.

Add modules and other content to the collection.

Build the collection into a collection artifact with
ansible-galaxy collection build

.
Publish the collection artifact to Galaxy with
ansible-galaxy collection publish

.
A user can then install your collection on their systems.

Naming your collection

Creating a new collection

Creating a collection from a custom template

Creating collections with ansible-creator

Naming your collection

?
Collection names consist of a namespace and a name, separated by a period (

.
) . Both namespace and name should be valid Python identifiers. This means that they should consist of ASCII letters, digits, and underscores.

Note

Usually namespaces and names use lower-case letters, digits, and underscores, but no upper-case letters.

You should make sure that the namespace you use is not registered by someone else by checking on

Ansible Galaxy's namespace list

. If you chose a namespace or even a full collection name that collides with another collection on Galaxy, it can happen

that if you or someone else runs

ansible-galaxy

collection

install

with your collection name, you end up with another collection. Even if the namespace currently does not exist, it could be created later by someone else.

If you want to request a new namespace on Ansible Galaxy,

create an issue on github.com/ansible/galaxy

.

There are a few special namespaces:

ansible

:

The

ansible namespace

is owned by Red Hat and reserved for official Ansible collections. Two special members are the synthetic

ansible.builtin

and

ansible.legacy

collections. These cannot be found on Ansible Galaxy, but are built-in into ansible-core.

community

:

The

community namespace

is owned by the Ansible community. Collections from this namespace generally live in the

GitHub ansible-collection organization

. If you want to create a collection in this namespace,

request

it on the forum.

local

:

The

local namespace

does not contain any collection on Ansible Galaxy, and the intention is that this will never change. You can use the

local

namespace for collections that are locally on your machine or locally in your Git repositories, without having to fear collisions with actually existing collections on Ansible Galaxy.

Creating a new collection

?

Create your collection skeleton in a path that includes

ansible_collections

, for example

collections/ansible_collections/

.

To start a new collection, run the following command in your collections directory:

ansible_collections#>

ansible-galaxy

collection

init

my_namespace.my_collection

Note

Both the namespace and collection names use the same strict set of requirements. Both are limited to alphanumeric characters and underscores, must have a minimum length of two characters, and cannot start with an underscore.

It will create the structure

[my_namespace]/[my_collection]/[collection

skeleton]

.

Hint

If Git is used for version control, the corresponding repository should be initialized in the collection directory.

Once the collection exists, you can populate the directories with the content you want inside the collection. See

ansible-collections

GitHub Org to get a better idea of what you can place inside a collection.

Reference: the

ansible-galaxy

collection

command

Currently the

ansible-galaxy

collection

command implements the following sub commands:

init

: Create a basic collection based on the default template included with Ansible or your own template.

build

: Create a collection artifact that can be uploaded to Galaxy or your own repository.

publish

: Publish a built collection artifact to Galaxy.

install

: Install one or more collections.

To learn more about the

ansible-galaxy

command-line tool, see the

ansible-galaxy

man page.

Creating a collection from a custom template

?

The built-in collection template is a simple example of a collection that works with

ansible-core

, but if you want to simplify your development process you may want to create a custom collection template to pass to

ansible-galaxy

collection

init

.

A collection skeleton is a directory that looks like a collection directory but any

.j2

files (excluding those in

templates/

and

roles/*/templates/

) will be templated by

ansible-galaxy

collection

init

. The skeleton's

galaxy.yml.j2

file should use the variables

namespace

and

collection_name

which are derived from

ansible-galaxy
init
namespace.collection_name
, and will populate the metadata in the initialized collection?
galaxy.yml
file. There are a few additional variables available by default (for example,
version
is
1.0.0
) , and these can be supplemented/overridden using
--extra-vars

.
An example
galaxy.yml.j2
file that accepts an optional dictionary variable
dependencies
could look like this:

```
namespace:
{{
namespace
}}
name:
{{
collection_name
}}
version:
{{
(
version
|
quote
)
is
version
(
'0.0.0'
,
operator
=
'gt'
,
version_type
=
'semver'
)|
ternary
(
version
,
undef
(
'version must be a valid semantic version greater than 0.0.0'
))
}}
```

dependencies:

```
{{
dependencies
|
default
({},
true
)
}}
```

To initialize a collection using the new template, pass the path to the skeleton with

ansible-galaxy

collection

init

:

ansible_collections#>

ansible-galaxy

collection

init

--collection-skeleton

/path/to/my/namespace/skeleton

--extra-vars

"@my_vars_file.json"

my_namespace.my_collection

Note

Before

ansible-core

2.17, collection skeleton templating is limited to the few hardcoded variables including

namespace

,

collection_name

, and

version

.

Note

The default collection skeleton uses an internal filter

comment_ify

that isn't accessibly to

--collection-skeleton

. Use

ansible-doc

-t

filter|test

--list

to see available plugins.

Creating collections with ansible-creator

?

ansible-creator

is designed to quickly scaffold an Ansible collection project.

Note

The

Ansible Development Tools

package offers a convenient way to install

ansible-creator

along with a curated set of tools for developing automation content.

After

installing

ansible-creator

you can initialize a project in one of the following ways:

Use the

init

subcommand.

Use

ansible-creator

with the

Ansible extension

in Visual Studio Code.

See also

[Using Ansible collections](#)

[Learn how to install and use collections.](#)

[Collection structure](#)

[Directories and files included in the collection skeleton](#)

[Ansible Development Tools \(ADT\)](#)

[Python package of tools to create and test Ansible content.](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/developing_collections_testing.html

Testing collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing collections
Testing collections
Edit on GitHub
Testing collections
?

Testing your collection ensures that your code works well and integrates well with the rest of the Ansible ecosystem. Your collection should pass the sanity tests for Ansible code. You should also add unit tests to cover the code in your collection and integration tests to cover the interactions between your collection and ansible-core.

Testing tools
Sanity tests
Adding unit tests
Adding integration tests
Testing tools
?

The main tool for testing collections is
ansible-test

, Ansible's testing tool described in

Testing Ansible
and provided by both the
ansible
and
ansible-core
packages.

Use
ansible-test
tool from your collection directory, which must include
ansible_collections
in the path, for example
collections/ansible_collections/community/general
for the
community.general
collection. See
Testing Collection Contributions
and
Testing Ansible and Collections
for testing guidelines.
You can run several sanity tests, as well as run unit and integration tests for plugins using
ansible-test

. When you test collections, test against the ansible-core version(s) you are targeting.

You must always execute

`ansible-test`

from the root directory of a collection.

You can run

`ansible-test`

in Docker containers without installing any special requirements.

The Ansible team uses this approach in Azure Pipelines both in the ansible/ansible GitHub repository and in the large community collections such as

`community.general`

and

`community.network`

to automatically run the tests when pull requests are submitted.

Many collections which do not require running tests on different OS distributions use GitHub Actions as their continuous integration (CI) platform.

The

`collection_template` repository

contains GitHub Actions workflow

templates

that collection developers are free to use to easily set up CI in their collection repositories.

The examples below demonstrate running tests in Docker containers.

Sanity tests

?

To run all sanity tests:

`ansible-test sanity --docker default -v`

See

Sanity Tests

for more information. See the

full list of sanity tests

for details on the sanity tests and how to fix identified issues.

Adding unit tests

?

You must place unit tests in the appropriate

`tests/unit/plugins/`

directory. For example, you would place tests for

`plugins/module_utils/foo/bar.py`

in

`tests/unit/plugins/module_utils/foo/test_bar.py`

or

`tests/unit/plugins/module_utils/foo/bar/test_bar.py`

. For examples, see the

unit tests in `community.general`

.

To run all unit tests for all supported Python versions:

`ansible-test units --docker default -v`

To run all unit tests only for a specific Python version:

`ansible-test units --docker default -v --python 3.6`

To run only a specific unit test:

`ansible-test units --docker default -v --python 3.6 tests/unit/plugins/module_utils/foo/test_bar.py`

You can specify Python requirements in the

`tests/unit/requirements.txt`

file. See

Unit Tests

for more information, especially on fixture files.

Adding integration tests

?

You must place integration tests in the appropriate

tests/integration/targets/

directory. For module integration tests, you can use the module name alone. For example, you would place integration tests for

plugins/modules/foo.py

in a directory called

tests/integration/targets/foo/

. For non-module plugin integration tests, you must add the plugin type to the directory name. For example, you would place integration tests for

plugins/connections/bar.py

in a directory called

tests/integration/targets/connection_bar/

. For lookup plugins, the directory must be called

lookup_foo

, for inventory plugins,

inventory_foo

, and so on.

You can write two different kinds of integration tests:

Ansible role tests run with

ansible-playbook

and validate various aspects of the module. They can depend on other integration tests (usually named

prepare_bar

or

setup_bar

, which prepare a service or install a requirement named

bar

in order to test module

foo

) to set-up required resources, such as installing required libraries or setting up server services.

runme.sh

tests run directly as scripts. They can set up inventory files, and execute

ansible-playbook

or

ansible-inventory

with various settings.

For examples, see the

integration tests in community.general

. See also

Integration tests

for more details.

Since integration tests can install requirements, and set-up, start and stop services, we recommended running them in docker containers or otherwise restricted environments whenever possible. By default,

ansible-test

supports Docker images for several operating systems. See the

list of supported docker images

for all options. Use the

default

image mainly for platform-independent integration tests, such as those for cloud modules. The following examples use the

fedora35

image.

To execute all integration tests for a collection:

`ansible-test integration --docker fedora35 -v`

If you want more detailed output, run the command with

`-vvv`

instead of

`-v`

. Alternatively, specify

`--retry-on-error`

to automatically re-run failed tests with higher verbosity levels.

To execute only the integration tests in a specific directory:

`ansible-test integration --docker fedora35 -v connection_bar`

You can specify multiple target names. Each target name is the name of a directory in `tests/integration/targets/`

.

See also

[Testing Ansible](#)

[More resources on testing Ansible](#)

[Contributing to Ansible-maintained Collections](#)

[Guidelines for contributing to selected collections](#)

[Communication](#)

Got questions? Need help? Want to share your ideas? Visit the [Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/developing_locally.html

Adding modules and plugins locally ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Adding modules and plugins locally
Edit on GitHub
Adding modules and plugins locally

?

You can extend Ansible by adding custom modules or plugins. You can create them from scratch or copy existing ones for local use. You can store a local module or plugin on your Ansible control node and share it with your team or organization. You can also share plugins and modules by including them in a collection, then publishing the collection on Ansible Galaxy.

If you are using a local module or plugin but Ansible cannot find it, this page is all you need.

If you want to create a plugin or a module, see

Developing plugins

,

Developing modules

and

Developing collections

.

Extending Ansible with local modules and plugins offers shortcuts such as:

You can copy other people's modules and plugins.

When writing a new module, you can choose any programming language you like.

You do not have to clone any repositories.

You do not have to open a pull request.

You do not have to add tests (though we recommend that you do!).

Modules and plugins: what is the difference?

Adding modules and plugins in collections

Adding a module or plugin outside of a collection

Adding standalone local modules for all playbooks and roles

Adding standalone local modules for selected playbooks or a single role

Adding a non-module plugin locally outside of a collection

Adding local non-module plugins for all playbooks and roles

Adding standalone local plugins for selected playbooks or a single role

Using

ansible.legacy

to access custom versions of an

ansible.builtin

module

Modules and plugins: what is the difference?

?

If you are looking to add functionality to Ansible, you might wonder whether you need a module or a plugin. Here is a

quick overview to help you understand what you need:

Plugins

extend Ansible's core functionality. Most plugin types execute on the control node within the

/usr/bin/ansible

process. Plugins offer options and extensions for the core features of Ansible: transforming data, logging output, connecting to inventory, and more.

Modules are a type of plugin that execute automation tasks on a ?target? (usually a remote system). Modules work as standalone scripts that Ansible executes in their own process outside of the control node. Modules interface with Ansible mostly with JSON, accepting arguments and returning information by printing a JSON string to stdout before exiting. Unlike the other plugins (which must be written in Python), modules can be written in any language; although Ansible provides modules in Python and Powershell only.

Adding modules and plugins in collections

?

You can add modules and plugins by

creating a collection

. With a collection, you can use custom modules and plugins in any playbook or role. You can share your collection easily at any time through Ansible Galaxy.

The rest of this page describes other methods of using local, standalone modules or plugins.

Adding a module or plugin outside of a collection

?

You can configure Ansible to load standalone local modules or plugins in specific locations and make them available to all playbooks and roles (using configured paths). Alternatively, you can make a non-collection local module or plugin available only to certain playbooks or roles (with adjacent paths).

Adding standalone local modules for all playbooks and roles

?

To load standalone local modules automatically and make them available to all playbooks and roles, use the

DEFAULT_MODULE_PATH

configuration setting or the

ANSIBLE_LIBRARY

environment variable. The configuration setting and environment variable take a colon-separated list, similar to

\$PATH

. You have two options:

Add your standalone local module to one of the default configured locations. See the

DEFAULT_MODULE_PATH

configuration setting for details. Default locations may change without notice.

Add the location of your standalone local module to an environment variable or configuration:

the

ANSIBLE_LIBRARY

environment variable

the

DEFAULT_MODULE_PATH

configuration setting

To view your current configuration settings for modules:

ansible-config dump |grep DEFAULT_MODULE_PATH

After you save your module file in one of these locations, Ansible loads it and you can use it in any local task, playbook, or role.

To confirm that

my_local_module

is available:

type

ansible

localhost

-m

my_local_module

to see the output for that module, or

type

ansible-doc

-t

module

my_local_module

to see the documentation for that module

Note

This applies to all plugin types but requires specific configuration and/or adjacent directories for each plugin type, see below.

Note

The

ansible-doc

command can parse module documentation from modules written in Python or an adjacent YAML file. If you have a module written in a programming language other than Python, you should write the documentation in a Python or YAML file adjacent to the module file.

Adjacent YAML documentation files

Adding standalone local modules for selected playbooks or a single role

?

Ansible automatically loads all executable files from certain directories adjacent to your playbook or role as modules. Standalone modules in these locations are available only to the specific playbook, playbooks, or role in the parent directory.

To use a standalone module only in a selected playbook or playbooks, store the module in a subdirectory called library

in the directory that contains the playbook or playbooks.

To use a standalone module only in a single role, store the module in a subdirectory called library

within that role.

Note

This applies to all plugin types but requires specific configuration and/or adjacent directories for each plugin type, see below.

Warning

Roles contained in collections cannot contain any modules or other plugins. All plugins in a collection must live in the collection

plugins

directory tree. All plugins in that tree are accessible to all roles in the collection. If you are developing new modules, we recommend distributing them in

collections

, not in roles.

Adding a non-module plugin locally outside of a collection

?

You can configure Ansible to load standalone local plugins in a specified location or locations and make them available to all playbooks and roles. Alternatively, you can make a standalone local plugin available only to specific playbooks or roles.

Note

Although modules are plugins, the naming patterns for directory names and environment variables that apply to other plugin types do not apply to modules. See

Adding a module or plugin outside of a collection

.

Adding local non-module plugins for all playbooks and roles

?

To load standalone local plugins automatically and make them available to all playbooks and roles, use the configuration setting or environment variable for the type of plugin you are adding. These configuration settings and environment variables take a colon-separated list, similar to

\$PATH

. You have two options:

Add your local plugin to one of the default configured locations. See

configuration settings

for details on the correct configuration setting for the plugin type. Default locations may change without notice.

Add the location of your local plugin to an environment variable or configuration:

the relevant

ANSIBLE_plugin_type_PLUGINS

environment variable - for example,

\$ANSIBLE_INVENTORY_PLUGINS

or

\$ANSIBLE_VARS_PLUGINS

the relevant

plugin_type_PATH

configuration setting, most of which begin with

DEFAULT_

- for example,

DEFAULT_CALLBACK_PLUGIN_PATH

or

DEFAULT_FILTER_PLUGIN_PATH

or

BECOME_PLUGIN_PATH

To view your current configuration settings for non-module plugins:

ansible-config dump |grep plugin_type_PATH

After your plugin file is added to one of these locations, Ansible loads it and you can use it in any local module, task, playbook, or role. For more information on environment variables and configuration settings, see

Ansible Configuration Settings

.

To confirm that

plugins/plugin_type/my_local_plugin

is available:

type

ansible-doc

-t

<plugin_type>

my_local_lookup_plugin

to see the documentation for that plugin - for example,

ansible-doc

-t

lookup

my_local_lookup_plugin

The

ansible-doc

command works for most plugin types, but not for action, filter, or test plugins. See

ansible-doc

for more details.

Adding standalone local plugins for selected playbooks or a single role

?

Ansible automatically loads all plugins from certain directories adjacent to your playbook or role, loading each type of plugin separately from a directory named for the type of plugin. Standalone plugins in these locations are available only to the specific playbook, playbooks, or role in the parent directory.

To use a standalone plugin only in a selected playbook or playbooks, store the plugin in a subdirectory for the correct plugin_type

(for example,

callback_plugins

or

inventory_plugins

) in the directory that contains the playbooks. These directories must use the

_plugins

suffix. For a full list of plugin types, see

Working with plugins

.

To use a standalone plugin only in a single role, store the plugin in a subdirectory for the correct

plugin_type

(for example,

cache_plugins

or

strategy_plugins

) within that role. When shipped as part of a role, the plugin is available as soon as the role is executed. These directories must use the

_plugins

suffix. For a full list of plugin types, see

Working with plugins

.

Warning

Roles contained in collections cannot contain any plugins. All plugins in a collection must live in the collection

plugins

directory tree. All plugins in that tree are accessible to all roles in the collection. If you are developing new plugins, we recommend distributing them in

collections

, not in roles.

Warning

Some plugin types are needed early during Ansible execution, such as callbacks, inventory, and cache. These plugin types cannot be loaded dynamically and must exist in configured paths or be referenced by FQCN in configuration.

Using

ansible.legacy

to access custom versions of an

ansible.builtin

module

?

If you need to override one of the

ansible.builtin

modules and are using FQCN, you need to use

ansible.legacy

as part of the fully-qualified collection name (FQCN). For example, if you had your own

copy

module, you would access it as

ansible.legacy.copy

. See

Using ansible.legacy to access local custom modules from collections-based roles

for details on how to use custom modules with roles within a collection.

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/developing_modules_checklist.html

Contributing your module to an existing Ansible collection ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible

Developer Guide
Contributing your module to an existing Ansible collection
Edit on GitHub
Contributing your module to an existing Ansible collection

?

If you want to contribute a module to an existing collection, you must meet the community's objective and subjective requirements. Please read the details below, and also review our tips for module development

.

Modules accepted into certain collections are included in every Ansible release on PyPI. However, contributing to one of these collections is not the only way to distribute a module - you can create your own collection, embed modules in roles on Galaxy or simply share copies of your module code for local use

.

Contributing modules: objective requirements

?

To contribute a module to most Ansible collections, you must:
write your module in either Python or Powershell for Windows
use the

AnsibleModule

common code

support Python 2.6 and Python 3.5 - if your module cannot support Python 2.6, explain the required minimum Python version and rationale in the requirements section in

DOCUMENTATION

use proper

Python 3 syntax

follow

PEP 8

Python style conventions - see

pep8

for more information

license your module under the GPL license (GPLv3 or later)

understand the

DCO agreement

, which applies to contributions to the

Ansible Core

and

Ansible Documentation

repositories.

conform to Ansible?

formatting and documentation

standards

include comprehensive

tests

for your module

minimize module dependencies

support

check_mode

if possible

ensure your code is readable

if a module is named

<something>_facts

, it should be because its main purpose is returning

ansible_facts

. Do not name modules that do not do this with

_facts

. Only use

ansible_facts

for information that is specific to the host machine, for example network interfaces and their configuration, which operating system and which programs are installed.

Modules that query/return general information (and not

ansible_facts

) should be named

_info

. General information is non-host specific information, for example information on online/cloud services (you can access different accounts for the same online service from the same host), or information on VMs and containers accessible from the machine.

Additional requirements may apply for certain collections. Review the individual collection repositories for more information.

Please make sure your module meets these requirements before you submit your PR/proposal. If you have questions, visit the

Ansible communication guide

for information on how to reach out to the community.

Contributing to Ansible: subjective requirements

?

If your module meets these objective requirements, collection maintainers will review your code to see if they think it is clear, concise, secure, and maintainable. They will consider whether your module provides a good user experience, helpful error messages, reasonable defaults, and more. This process is subjective, with no exact standards for acceptance. For the best chance of getting your module accepted, follow our

tips for module development

.

Other checklists

?

Tips for module development

.

Windows development checklist

.

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/developing_modules_general.html

Developing modules ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[ansible-core Roadmaps](#)
[Ansible](#)
[Developer Guide](#)
[Developing modules](#)
[Edit on GitHub](#)
[Developing modules](#)
[?](#)

A module is a reusable, standalone script that Ansible runs on your behalf, either locally or remotely. Modules interact with your local machine, an API, or a remote system to perform specific tasks like changing a database password or spinning up a cloud instance. Each module can be used by the Ansible API, or by the `ansible`

or

`ansible-playbook`

programs. A module provides a defined interface, accepts arguments, and returns information to Ansible by printing a JSON string to stdout before exiting.

If you need functionality that is not available in any of the thousands of Ansible modules found in collections, you can easily write your own custom module. When you write a module for local use, you can choose any programming language and follow your own rules. Use this topic to learn how to create an Ansible module in Python. After you create a module, you must add it locally to the appropriate directory so that Ansible can find and execute it. For details about adding a module locally, see

[Adding modules and plugins locally](#)

.

If you are developing a module in a collection

, see those documents instead.

[Preparing an environment for developing Ansible modules](#)

[Creating a standalone module](#)

[Creating a module in a collection](#)

[Creating an info or a facts module](#)

[Verifying your module code](#)

[Verifying your module code locally](#)

[Verifying your module code in a playbook](#)

[Testing your newly-created module](#)

[Contributing back to Ansible](#)

[Communication and development support](#)

[Credit](#)

[Preparing an environment for developing Ansible modules](#)

?

You just need

`ansible-core`

installed to test the module. Modules can be written in any language,

but most of the following guide is assuming you are using Python.

Modules for inclusion in Ansible itself must be Python or Powershell.

One advantage of using Python or Powershell for your custom modules is being able to use the `module_utils`

common code that does a lot of the

heavy lifting for argument processing, logging and response writing, among other things.

Creating a standalone module

?

It is highly recommended that you use a

`venv`

or

`virtualenv`

for Python development.

To create a standalone module:

Create a

library

directory in your workspace. Your test play should live in the same directory.

Create your new module file:

\$

`touch`

`library/my_test.py`

. Or just open/create it with your editor of choice.

Paste the content below into your new module file. It includes the required Ansible format and documentation

, a simple

argument spec for declaring the module options

, and some example code.

Modify and extend the code to do what you want your new module to do. See the programming tips

and

Python 3 compatibility

pages for pointers on writing clean and concise module code.

Creating a module in a collection

?

To create a new module in an existing collection called

`my_namespace.my_collection`

:

Create your new module file:

\$

`touch`

`<PATH_TO_COLLECTION>/ansible_collections/my_namespace/my_collection/plugins/modules/my_test.py`

. Or just create it with your editor of choice.

Paste the content below into your new module file. It includes the required Ansible format and documentation

, a simple

argument spec for declaring the module options

, and some example code.

Modify and extend the code to do what you want your new module to do. See the programming tips

and

Python 3 compatibility

pages for pointers on writing clean and concise module code.

`#!/usr/bin/python`

Copyright: (c) 2018, Terry Jones <

>

GNU General Public License v3.0+ (see COPYING or <https://www.gnu.org/licenses/gpl-3.0.txt>)

from

__future__

import

(

absolute_import

,

division

,

print_function

)

__metaclass__

=

type

DOCUMENTATION

=

r

"""

module: my_test

short_description: This is my test module

If this is part of a collection, you need to use semantic versioning,

i.e. the version is of the form "2.5.0" and not "2.4".

version_added: "1.0.0"

description: This is my longer description explaining my test module.

options:

name:

description: This is the message to send to the test module.

required: true

type: str

new:

description:

- Control to demo if the result of this module is changed or not.

- Parameter description can be a list as well.

required: false

type: bool

Specify this value according to your collection

in format of namespace.collection.doc_fragment_name

extends_documentation_fragment:

- my_namespace.my_collection.my_doc_fragment_name

author:

- Your Name (@yourGitHubHandle)

"""

EXAMPLES

=

r

"""

Pass in a message

- name: Test with a message

my_namespace.my_collection.my_test:

name: hello world

pass in a message and have changed true


```

- name: Test with a message and changed output
my_namespace.my_collection.my_test:
name: hello world
new: true
# fail the module
- name: Test failure of the module
my_namespace.my_collection.my_test:
name: fail me
'''

RETURN

=

r
'''

# These are examples of possible return values, and in general should use other names for return values.
original_message:
description: The original name param that was passed in.
type: str
returned: always
sample: 'hello world'
message:
description: The output message that the test module generates.
type: str
returned: always
sample: 'goodbye'
'''

from
ansible.module_utils.basic
import
AnsibleModule
def
run_module
():
# define available arguments/parameters a user can pass to the module
module_args
=
dict
(
name
=
dict
(
type
=
'str'
,
required
=
True
),
new
=
dict
(
type

```

```

=
'bool'
,
required
=
False
,
default
=
False
)
)
# seed the result dict in the object
# we primarily care about changed and state
# changed is if this module effectively modified the target
# state will include any data that you want your module to pass back
# for consumption, for example, in a subsequent task
result
=
dict
(
changed
=
False
,
original_message
=
"
,
message
=
"
)
# the AnsibleModule object will be our abstraction working with Ansible
# this includes instantiation, a couple of common attr would be the
# args/params passed to the execution, as well as if the module
# supports check mode
module
=
AnsibleModule
(
argument_spec
=
module_args
,
supports_check_mode
=
True
)
# if the user is working with this module in only check mode we do not
# want to make any changes to the environment, just return the current
# state with no modifications
if
module

```

```

.
check_mode
:
module
.
exit_json
(
**
result
)
# manipulate or modify the state as needed (this is going to be the
# part where your module will do what it needs to do)
result
[
'original_message'
]
=
module
.
params
[
'name'
]
result
[
'message'
]
=
'goodbye'
# use whatever logic you need to determine whether or not this module
# made any modifications to your target
if
module
.
params
[
'new'
]:
result
[
'changed'
]
=
True
# during the execution of the module, if there is an exception or a
# conditional state that effectively causes a failure, run
# AnsibleModule.fail_json() to pass in the message and the result
if
module
.
params
[
'name'
]

```

```

==
'fail me'
:
module
.
fail_json
(
msg
=
'You requested this to fail'
,
**
result
)
# in the event of a successful module execution, you will want to
# simple AnsibleModule.exit_json(), passing the key/value results
module
.
exit_json
(
**
result
)
def
main
():
run_module
()
if
__name__
==
'__main__'
:
main
()
```

Creating an info or a facts module
?

Ansible gathers information about the target machines using facts modules, and gathers information on other objects or files using info modules.

If you find yourself trying to add

state:

info

or

state:

list

to an existing module, that is often a sign that a new dedicated

_facts

or

_info

module is needed.

In Ansible 2.8 and onwards, we have two type of information modules, they are

*_info

and

*_facts

- If a module is named
 <something>_facts
 , it should be because its main purpose is returning
 ansible_facts
 . Do not name modules that do not do this with
 _facts
- Only use
 ansible_facts
 for information that is specific to the host machine, for example network interfaces and their configuration, which
 operating system and which programs are installed.
 Modules that query/return general information (and not
 ansible_facts
) should be named
 _info
- General information is non-host specific information, for example information on online/cloud services (you can access
 different accounts for the same online service from the same host), or information on VMs and containers accessible
 from the machine, or information on individual files or programs.
 Info and facts modules, are just like any other Ansible Module, with a few minor requirements:
 They MUST be named
 <something>_info
 or
 <something>_facts
 , where <something> is singular.
 Info
 *_info
 modules MUST return in the form of the
 result dictionary
 so other modules can access them.
 Fact
 *_facts
 modules MUST return in the
 ansible_facts
 field of the
 result dictionary
 so other modules can access them.
 They MUST support
 check_mode
- They MUST NOT make any changes to the system.
 They MUST document the
 return fields
 and
 examples
- You can add your facts into
 ansible_facts
 field of the result as follows:
 module
- exit_json
 (

```
changed
=
False
,
ansible_facts
=
dict
(
my_new_fact
=
value_of_fact
))
```

The rest is just like creating a normal module.

Verifying your module code

?

After you modify the sample code above to do what you want, you can try out your module.

Our

debugging tips

will help if you run into bugs as you verify your module code.

Verifying your module code locally

?

The simplest way is to use

ansible

ad hoc command:

ANSIBLE_LIBRARY

=

./library

ansible

-m

my_test

-a

'name=hello new=true'

remotehost

If your module does not need to target a remote host, you can quickly and easily exercise your code locally like this:

ANSIBLE_LIBRARY

=

./library

ansible

-m

my_test

-a

'name=hello new=true'

localhost

For a module developed in an existing collection called

my_namespace.my_collection

(as mentioned above):

\$

ansible

localhost

-m

my_namespace.my_collection.my_test

-a

'name=hello new=true'

--playbook-dir

```
=
$PWD
If you use
pdb
,
print()
, or some other method of local debugging for faster iteration,
you can avoid going through Ansible by creating an arguments file, which is
a basic JSON config file that passes parameters to your module so that you can run it.
```

Name the arguments file
/tmp/args.json
and add the following content:

```
{
"ANSIBLE_MODULE_ARGS"
:
{
"name"
:
"hello"
,
"new"
:
true
}
}
```

Then the module can be tested locally and directly. This skips the packing steps and uses module_utils files directly:

```
$
python
library/my_test.py
/tmp/args.json
```

You might also need to add your collection's path to the Python path. This tells Python to look for additional module_utils code in the given path.

You can run the module code, as in the following example:

```
$
export
PYTHONPATH
=
PATH_TO_COLLECTIONS:
$PYTHONPATH
$
python
-m
ansible_collections.my_namespace.my_collection.plugins.modules.my_test
/tmp/args.json
```

It should return output like this:

```
{
"changed"
:
true
,
"state"
:
{
"original_message"
```

```

:
"hello"
,
"new_message"
:
"goodbye"
},
"invocation"
:
{
"module_args"
:
{
"name"
:
"hello"
,
"new"
:
true
}}}

```

Verifying your module code in a playbook

?

You can easily run a full test by including it in a playbook, as long as the library

directory is in the same directory as the play:

Create a playbook in any directory:

\$

touch

testmod.yml

Add the following to the new playbook file:

```

-
name
:
test my new module
hosts
:
localhost
tasks
:
-
name
:
run the new module
my_test
:
name
:
'hello'
new
:
true
register
:

```


testout

-

name

:

dump test output

debug

:

msg

:

'{{

testout

}}'

Run the playbook and analyze the output:

\$

ansible-playbook

./testmod.yml

Testing your newly-created module

?

Review our

testing

section for more detailed

information, including instructions for

testing module documentation

, adding

integration tests

, and more.

Note

If contributing to Ansible, every new module and plugin should have integration tests, even if the tests cannot be run on Ansible CI infrastructure.

In this case, the tests should be marked with the

unsupported

alias in

aliases file

.

Contributing back to Ansible

?

If you would like to contribute to

ansible-core

by adding a new feature or fixing a bug,

create a fork

of the ansible/ansible repository and develop against a new feature branch using the

devel

branch as a starting point. When you have a good working code change, you can submit a pull request to the Ansible repository by selecting your feature branch as a source and the Ansible devel branch as a target.

If you want to contribute a module to an

Ansible collection

, review our

submission checklist

,

programming tips

, and

strategy for maintaining Python 2 and Python 3 compatibility

, as well as information about

testing

before you open a pull request.

The

Community Guide

covers how to open a pull request and what happens next.

Communication and development support

?

Visit the

Ansible communication guide

for information on how to join the conversation.

Credit

?

Thank you to Thomas Stringer (

@trstringer

) for contributing source

material for this topic.

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/developing_python_3.html

Ansible and Python 3 ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Ansible and Python 3
Edit on GitHub
Ansible and Python 3
?
The
ansible-core
code runs Python 3 (for specific versions check
Control Node Requirements
Contributors to
ansible-core
and to Ansible Collections should be aware of the tips in this document so that they can write code
that will run on the same versions of Python as the rest of Ansible.
Minimum version of Python 3.x and Python 2.x
Developing Ansible code that supports Python 2 and Python 3
Understanding strings in Python 2 and Python 3
Control node string strategy: the Unicode Sandwich
Unicode Sandwich common borders: places to convert bytes to text in control node code
Reading and writing to files
Filesystem interaction
Interacting with other programs
Module string strategy: Native String
Module_utils string strategy: hybrid
Tips, tricks, and idioms for Python 2/Python 3 compatibility
Use forward-compatibility boilerplate
Prefix byte strings with
b_
Import Ansible's bundled Python
six
library
Handle exceptions with
as
Update octal numbers
String formatting for control node code
Use
str.format()
for Python 2.6 compatibility
Use percent format with byte strings
We do have some considerations depending on the types of Ansible code:

code on the control node - code that runs on the machine where you invoke

`/usr/bin/ansible`

, only needs to support the control node's Python versions.

modules - the code which Ansible transmits to and invokes on the managed machine. Modules need to support the managed node's Python versions, with some exceptions.

shared

`module_utils`

code - the common code that is used by modules to perform tasks and sometimes used by code on the control node.

Shared

`module_utils`

code needs to support the same range of Python as the modules.

However, the three types of code do not use the same string strategy. If you're developing a module or some

`module_utils`

code, be sure to read the section on string strategy carefully.

Minimum version of Python 3.x and Python 2.x

?

See

Control Node Requirements

and

Managed Node Requirements

for the

specific versions supported.

Your custom modules can support any version of Python (or other languages) you want, but the above are the requirements for the code contributed to the Ansible project.

Developing Ansible code that supports Python 2 and Python 3

?

The best place to start learning about writing code that supports both Python 2 and Python 3

is

Lennart Regebro's book: *Porting to Python 3*

.

The book describes several strategies for porting to Python 3. The one we're

using is

to support Python 2 and Python 3 from a single code base

Understanding strings in Python 2 and Python 3

?

Python 2 and Python 3 handle strings differently, so when you write code that supports Python 3

you must decide what string model to use. Strings can be an array of bytes (like in C) or

they can be an array of text. Text is what we think of as letters, digits,

numbers, other printable symbols, and a small number of unprintable symbols?

(control codes).

In Python 2, the two types for these (

`str`

for bytes and

`unicode`

for text) are often used interchangeably. When dealing only

with ASCII characters, the strings can be combined, compared, and converted

from one type to another automatically. When non-ASCII characters are

introduced, Python 2 starts throwing exceptions due to not knowing what encoding

the non-ASCII characters should be in.

Python 3 changes this behavior by making the separation between bytes (

`bytes`

)

and text (

`str`

) more strict. Python 3 will throw an exception when trying to combine and compare the two types. The programmer has to explicitly convert from one type to the other to mix values from each.

In Python 3 it is immediately apparent to the programmer when code is mixing the byte and text types inappropriately, whereas in Python 2, code that mixes those types may work until a user causes an exception by entering non-ASCII input.

Python 3 forces programmers to proactively define a strategy for working with strings in their program so that they don't mix text and byte strings unintentionally.

Ansible uses different strategies for working with strings in the code on the control node, in

:ref:

modules <module_string_strategy>

, and in

module_utils

code.

Control node string strategy: the Unicode Sandwich

?

Until recently

ansible-core

supported Python 2.x and followed this strategy, known as the Unicode Sandwich (named after Python 2's

unicode

text type). For Unicode Sandwich we know that

at the border of our code and the outside world (for example, file and network IO, environment variables, and some library calls) we are going to receive bytes.

We need to transform these bytes into text and use that throughout the internal portions of our code. When we have to send those strings back out to the outside world we first convert the text back into bytes.

To visualize this, imagine a ?sandwich? consisting of a top and bottom layer of bytes, a layer of conversion between, and all text type in the center.

For compatibility reasons you will see a bunch of custom functions we developed (

to_text

/

to_bytes

/

to_native

)

and while Python 2 is not a concern anymore we will continue to use them as they apply for other cases that make dealing with unicode problematic.

While we will not be using it most of it anymore, the documentation below is still useful for those developing modules that still need to support both Python 2 and 3 simultaneously.

Unicode Sandwich common borders: places to convert bytes to text in control node code

?

This is a partial list of places where we have to convert to and from bytes when using the Unicode Sandwich string strategy. It is not exhaustive but it gives you an idea of where to watch for problems.

Reading and writing to files

?

In Python 2, reading from files yields bytes. In Python 3, it can yield text.

To make code that's portable to both we don't make use of Python 3's ability to yield text but instead do the conversion explicitly ourselves. For example:

from

ansible.module_utils.common.text.converters

import

to_text

```

with
open
(
'filename-with-utf8-data.txt'
,
'rb'
)
as
my_file
:
b_data
=
my_file
.
read
()
try
:
data
=
to_text
(
b_data
,
errors
=
'surrogate_or_strict'
)
except
UnicodeError
:
# Handle the exception gracefully -- usually by displaying a good
# user-centric error message that can be traced back to this piece
# of code.
pass

```

Note

Much of Ansible assumes that all encoded text is UTF-8. At some point, if there is demand for other encodings we may change that, but for now it is safe to assume that bytes are UTF-8.

Writing to files is the opposite process:

```

from
ansible.module_utils.common.text.converters
import
to_bytes
with
open
(
'filename.txt'
,
'wb'
)
as
my_file
:

```

```
my_file
```

```
.
```

```
write
```

```
(
```

```
to_bytes
```

```
(
```

```
some_text_string
```

```
))
```

Note that we don't have to catch

UnicodeError

here because we're

transforming to UTF-8 and all text strings in Python can be transformed back to UTF-8.

Filesystem interaction

?

Dealing with file names often involves dropping back to bytes because on UNIX-like systems file names are bytes. On Python 2, if we pass a text string to these functions, the text string will be converted to a byte string inside of the function and a traceback will occur if non-ASCII characters are present. In Python 3, a traceback will only occur if the text string can't be decoded in the current locale, but it is still good to be explicit and have code which works on both versions:

```
import
```

```
os.path
```

```
from
```

```
ansible.module_utils.common.text.converters
```

```
import
```

```
to_bytes
```

```
filename
```

```
=
```

```
u
```

```
'/var/tmp/?????.txt'
```

```
f
```

```
=
```

```
open
```

```
(
```

```
to_bytes
```

```
(
```

```
filename
```

```
),
```

```
'wb'
```

```
)
```

```
mtime
```

```
=
```

```
os
```

```
.
```

```
path
```

```
.
```

```
getmtime
```

```
(
```

```
to_bytes
```

```
(
```

```
filename
```

```
))
```



```
b_filename
```

```
=
```

```
os
```

```
.
```

```
path
```

```
.
```

```
expandvars
```

```
(
```

```
to_bytes
```

```
(
```

```
filename
```

```
))
```

```
if
```

```
os
```

```
.
```

```
path
```

```
.
```

```
exists
```

```
(
```

```
to_bytes
```

```
(
```

```
filename
```

```
)):
```

```
pass
```

When you are only manipulating a filename as a string without talking to the filesystem (or a C library which talks to the filesystem) you can often get away without converting to bytes:

```
import
```

```
os.path
```

```
os
```

```
.
```

```
path
```

```
.
```

```
join
```

```
(
```

```
u
```

```
'/var/tmp/café'
```

```
,
```

```
u
```

```
'????'
```

```
)
```

```
os
```

```
.
```

```
path
```

```
.
```

```
split
```

```
(
```

```
u
```

```
'/var/tmp/café/????'
```

```
)
```

On the other hand, if the code needs to manipulate the file name and also talk to the filesystem, it can be more convenient to transform to bytes right away and manipulate in bytes.

Warning

Make sure all variables passed to a function are the same type.

If you're working with something like

```
os.path.join()
```

which takes

multiple strings and uses them in combination, you need to make sure that

all the types are the same (either all bytes or all text). Mixing

bytes and text will cause tracebacks.

Interacting with other programs

?

Interacting with other programs goes through the operating system and

C libraries and operates on things that the UNIX kernel defines. These

interfaces are all byte-oriented so the Python interface is byte oriented as

well. On both Python 2 and Python 3, byte strings should be given to Python's

subprocess library and byte strings should be expected back from it.

One of the main places in Ansible's control node code that we interact with

other programs is the connection plugins?

```
exec_command
```

methods. These

methods transform any text strings they receive in the command (and arguments

to the command) to execute into bytes and return stdout and stderr as byte strings

Higher level functions (like action plugins?

```
_low_level_execute_command
```

```
)
```

transform the output into text strings.

Module string strategy: Native String

?

In modules we use a strategy known as Native Strings. This makes things

easier on the community members who maintain so many of Ansible's

modules, by not breaking backwards compatibility by

mandating that all strings inside of modules are text and converting between

text and bytes at the borders.

Native strings refer to the type that Python uses when you specify a bare

string literal:

```
"This is a native string"
```

In Python 2, these are byte strings. In Python 3 these are text strings. Modules should be

coded to expect bytes on Python 2 and text on Python 3.

Module_utils string strategy: hybrid

?

In

```
module_utils
```

code we use a hybrid string strategy. Although Ansible's

```
module_utils
```

code is largely like module code, some pieces of it are

used by the control node as well. So it needs to be compatible with modules

and with the control node's assumptions, particularly the string strategy.

The module_utils code attempts to accept native strings as input

to its functions and emit native strings as their output.

In

```
module_utils
```

code:

Functions

must

accept string parameters as either text strings or byte strings.

Functions may return either the same type of string as they were given or the native string type for the Python version

they are run on.

Functions that return strings

must

document whether they return strings of the same type as they were given or native strings.

Module-utils functions are therefore often very defensive in nature.

They convert their string parameters into text (using

`ansible.module_utils.common.text.converters.to_text`

)

at the beginning of the function, do their work, and then convert

the return values into the native string type (using

`ansible.module_utils.common.text.converters.to_native`

)

or back to the string type that their parameters received.

Tips, tricks, and idioms for Python 2/Python 3 compatibility

?

Use forward-compatibility boilerplate

?

Use the following boilerplate code at the top of all python files

to make certain constructs act the same way on Python 2 and Python 3:

Make coding more python3-ish

from

`__future__`

import

(

`absolute_import`

,

`division`

,

`print_function`

)

`__metaclass__`

=

`type`

`__metaclass__`

=

`type`

makes all classes defined in the file into new-style

classes without explicitly inheriting from

`object`

.

The

`__future__`

imports do the following:

`absolute_import`

:

Makes imports look in

`sys.path`

for the modules being

imported, skipping the directory in which the module doing the importing

lives. If the code wants to use the directory in which the module doing

the importing, there's a new dot notation to do so.

`division`

:

Makes division of integers always return a float. If you need to

find the quotient use

x

//

y

instead of

x

/

y

.

print_function

:

Changes

print

from a keyword into a function.

See also

PEP 0328: Absolute Imports

PEP 0238: Division

PEP 3105: Print function

Prefix byte strings with

b_

?

Since mixing text and bytes types leads to tracebacks we want to be clear about what variables hold text and what variables hold bytes. We do this by prefixing any variable holding bytes with

b_

. For example:

filename

=

u

'/var/tmp/café.txt'

b_filename

=

to_bytes

(

filename

)

with

open

(

b_filename

)

as

f

:

data

=

f

.

read

()

We do not prefix the text strings instead because we only operate on byte strings at the borders, so there are fewer variables that need bytes than text.

Import Ansible's bundled Python

six
library
?
The third-party Python

six
library exists
to help projects create code that runs on both Python 2 and Python 3. Ansible
includes a version of the library in module_utils so that other modules can use it
without requiring that it is installed on the remote system. To make use of
it, import it like this:

```
from  
ansible.module_utils  
import
```

six
Note

Ansible can also use a system copy of six
Ansible will use a system copy of six if the system copy is a later
version than the one Ansible bundles.

Handle exceptions with

```
as  
?
```

In order for code to function on Python 2.6+ and Python 3, use the
new exception-catching syntax which uses the

```
as  
keyword:
```

```
try  
:  
a  
=  
2  
/  
0  
except  
ValueError
```

```
as  
e  
:  
module  
.  
fail_json
```

```
(  
msg  
=  
"Tried to divide by zero:
```

```
%s  
"
```

```
%  
e  
)
```

```
Do  
not
```

use the following syntax as it will fail on every version of Python 3:

```
try:  
    a = 2/0
```

except ValueError, e:

```
    module.fail_json(msg="Tried to divide by zero: %s" % e)
```

Update octal numbers

?

In Python 2.x, octal literals could be specified as

0755

. In Python 3,

octals must be specified as

0o755

.

String formatting for control node code

?

Use

```
str.format()
```

for Python 2.6 compatibility

?

Starting in Python 2.6, strings gained a method called

```
format()
```

to put

strings together. However, one commonly used feature of

```
format()
```

wasn't

added until Python 2.7, so you need to remember not to use it in Ansible code:

Does not work in Python 2.6!

```
new_string
```

```
=
```

```
"Dear
```

```
{}
```

```
, Welcome to
```

```
{}
```

```
"
```

.

```
format
```

```
(
```

```
    username
```

```
,
```

```
    location
```

```
)
```

Use this instead

```
new_string
```

```
=
```

```
"Dear
```

```
{0}
```

```
, Welcome to
```

```
{1}
```

```
"
```

.

```
format
```

```
(
```

```
    username
```

```
,
```

```
    location
```

```
)
```

Both of the format strings above map positional arguments of the

`format()`

method into the string. However, the first version doesn't work in

Python 2.6. Always remember to put numbers into the placeholders so the code is compatible with Python 2.6.

See also

Python documentation on format strings:

[format strings in 2.6](#)

[format strings in 3.x](#)

[Use percent format with byte strings](#)

?

In Python 3.5 and later, byte strings do not have a

`format()`

method. However, it

does have support for the older, percent-formatting.

`b_command_line`

=

`b`

`'ansible-playbook --become-user`

`%s`

`-K`

`%s`

`,`

`%`

`(`

`user`

`,`

`playbook_file`

`)`

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/module_lifecycle.html

The lifecycle of an Ansible module or plugin ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
The lifecycle of an Ansible module or plugin
Edit on GitHub
The lifecycle of an Ansible module or plugin
?

Modules and plugins in the main Ansible repository have a defined life cycle, from the first introduction to final removal. The module and plugin lifecycle is tied to the Ansible release cycle <release_cycle>

A module or plugin may move through these four stages:

When a module or plugin is first accepted into Ansible, we consider it in tech preview and will mark it as such in the documentation.

If a module or plugin matures, the ?preview? mark in the documentation is removed. Backward compatibility for these modules and plugins is maintained but not guaranteed, which means their parameters should be maintained with stable meanings.

If a module?s or plugin?s target API changes radically, or if someone creates a better implementation of its functionality, we may mark it deprecated. Modules and plugins that are deprecated are still available but they are reaching the end of their life cycle. We retain deprecated modules and plugins for 4 release cycles with deprecation warnings to help users update playbooks and roles that use them.

When a module or plugin has been deprecated for four release cycles, it is removed and replaced with a tombstone entry in the routing configuration. Modules and plugins that are removed are no longer shipped with Ansible. The tombstone entry helps users find alternative modules and plugins.

For modules and plugins in collections, the lifecycle is similar. Since ansible-base 2.10, it is no longer possible to mark modules as ?preview? or ?stable?.

Deprecating modules and plugins in the Ansible main repository

?

To deprecate a module in ansible-core, you must:

Rename the file so it starts with an

—

, for example, rename

old_cloud.py

to

_old_cloud.py

. This keeps the module available and marks it as deprecated on the module index pages.

Mention the deprecation in the relevant changelog (by creating a changelog fragment with a section

deprecated_features

).

Reference the deprecation in the relevant

porting_guide_core_x.y.rst

.

Add

deprecated:

to the documentation with the following sub-values:

removed_in

:

A

string

, such as

"2.10"

; the version of Ansible where the module will be replaced with a docs-only module stub. Usually current release +4.

Mutually exclusive with :removed_by_date:.

remove_by_date

:

(Added in ansible-base 2.10). An ISO 8601 formatted date when the module will be removed. Usually 2 years from the date the module is deprecated. Mutually exclusive with :removed_in:.

why

:

Optional string that used to detail why this has been removed.

alternatives

:

Inform users they should do instead, for example,

Use

M(whatmoduletouseinstead)

instead.

.

For an example of documenting deprecation, see this

PR that deprecates multiple modules

.

Some of the elements in the PR might now be out of date.

Deprecating modules and plugins in a collection

?

To deprecate a module in a collection, you must:

Add a

deprecation

entry to

plugin_routing

in

meta/runtime.yml

. For example, to deprecate the module

old_cloud

, add:

plugin_routing

:

modules

:

old_cloud

:

deprecation

:

removal_version

:

2.0.0

warning_text

:

Use `foo.bar.new_cloud` instead.

For other plugin types, you have to replace modules:

with

`<plugin_type>`:

, for example

lookup:

for lookup plugins. When deprecating action plugins, you need to add two entries: one for the action plugin and one for the module file that contains the documentation.

Instead of

`removal_version`

, you can also use

`removal_date`

with an ISO 8601 formatted date after which the module will be removed in a new major version of the collection.

Mention the deprecation in the relevant changelog. If the collection uses `antsibull-changelog`

, create a changelog fragment with a section `deprecated_features`

.

Add

`deprecated:`

to the documentation of the module or plugin with the following sub-values:

`removed_in`

:

A

string

, such as

`"2.10"`

; the version of Ansible where the module will be replaced with a docs-only module stub. Usually current release +4.

Mutually exclusive with `:removed_by_date:`.

`remove_by_date`

:

(Added in ansible-base 2.10). An ISO 8601 formatted date when the module will be removed. Usually 2 years from the date the module is deprecated. Mutually exclusive with `:removed_in:`.

`why`

:

String that used to detail why this has been removed.

`alternative`

:

Inform users they should do instead, for example,

Use

`M(whatmoduletouseinstead)`

instead.

. See

Linking within module documentation

for ways to reference entities other than modules.

Changing a module or plugin name in the Ansible main repository

?

You can also rename a module and keep a deprecated alias to the old name by using a symlink that starts with `_`.

This example allows the

`stat`

module to be called with

`fileinfo`

, making the following examples equivalent:

```
ln -s stat.py _fileinfo.py
```

```
ansible -m stat -a "path=/tmp" localhost
```

```
ansible -m fileinfo -a "path=/tmp" localhost
```

Renaming a module or plugin in a collection, or redirecting a module or plugin to another collection

?

To rename a module or plugin in a collection, or to redirect a module or plugin to another collection, you need to add a `redirect`

entry to

`plugin_routing`

in

`meta/runtime.yml`

. For example, to redirect the module

`old_cloud`

to

`foo.bar.new_cloud`

, add:

```
plugin_routing
```

```
:
```

```
modules
```

```
:
```

```
old_cloud
```

```
:
```

```
redirect
```

```
:
```

```
foo.bar.new_cloud
```

If you want to deprecate the old name, add a

deprecation:

entry (see above):

```
plugin_routing
```

```
:
```

```
modules
```

```
:
```

```
old_cloud
```

```
:
```

```
redirect
```

```
:
```

```
foo.bar.new_cloud
```

```
deprecation
```

```
:
```

```
removal_version
```

```
:
```

```
2.0.0
```

```
warning_text
```

```
:
```

Use `foo.bar.new_cloud` instead.

You need to use the Fully Qualified Collection Name (FQCN) of the new module/plugin name, even if it is located in the same collection as the redirect. By using a FQCN from another collection, you redirect the module/plugin to that collection.

If you need to support Ansible 2.9, please note that Ansible 2.9 does not know about

`meta/runtime.yml`

. With Ansible 2.9 you can still rename plugins and modules inside one collection by using symbolic links. Note that `ansible-base 2.10`, `ansible-core 2.11`, and newer will prefer

`meta/runtime.yml`

entries over symbolic links.

Tombstoning a module or plugin in a collection

?

To remove a deprecated module or plugin from a collection, you need to tombstone it:

Remove the module or plugin file with related files like tests, documentation references, and documentation.

Add a tombstone entry in

meta/runtime.yml

. For example, to tombstone the module

old_cloud

, add:

plugin_routing

:

modules

:

old_cloud

:

tombstone

:

removal_version

:

2.0.0

warning_text

:

Use foo.bar.new_cloud instead.

Instead of

removal_version

, you can also use

removal_date

with an ISO 8601 formatted date. The date should be the date of the next major release.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/overview_architecture.html

Ansible architecture ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Ansible architecture
Edit on GitHub
Ansible architecture
?

Ansible is a radically simple IT automation engine that automates cloud provisioning, configuration management, application deployment, intra-service orchestration, and many other IT needs.

Being designed for multi-tier deployments since day one, Ansible models your IT infrastructure by describing how all of your systems inter-relate, rather than just managing one system at a time.

It uses no agents and no additional custom security infrastructure, so it is easy to deploy - and most importantly, it uses a very simple language (YAML, in the form of Ansible Playbooks) that allows you to describe your automation jobs in a way that approaches plain English.

In this section, we'll give you a really quick overview of how Ansible works so you can see how the pieces fit together.

Modules

Module utilities

Plugins

Inventory

Playbooks

The Ansible search path

Modules

?

Ansible works by connecting to your nodes and pushing out scripts called "Ansible modules" to them. Most modules accept parameters that describe the desired state of the system.

Ansible then executes these modules (over SSH by default), and removes them when finished. Your library of modules can reside on any machine, and there are no servers, daemons, or databases required.

You can

write your own modules

, though you should first consider

whether you should

. Typically you'll work with your favorite terminal program, a text editor, and probably a version control system to keep track of changes to your content. You may write specialized modules in any language that can return JSON (Ruby, Python, bash, and so on).

Module utilities

?

When multiple modules use the same code, Ansible stores those functions as module utilities to minimize duplication and maintenance. For example, the code that parses URLs is

lib/ansible/module_utils/url.py

. You can

write your own module utilities

as well. Module utilities may only be written in Python or in PowerShell.

Plugins

?

Plugins

augment Ansible's core functionality. While modules execute on the target system in separate processes (usually that means on a remote system), plugins execute on the control node within the

/usr/bin/ansible

process. Plugins offer options and extensions for the core features of Ansible - transforming data, logging output, connecting to inventory, and more. Ansible ships with a number of handy plugins, and you can easily

write your own

. For example, you can write an

inventory plugin

to connect to any datasource that returns JSON. Plugins must be written in Python.

Inventory

?

By default, Ansible represents the machines it manages in a file (INI, YAML, and so on) that puts all of your managed machines in groups of your own choosing.

To add new machines, there is no additional SSL signing server involved, so there's never any hassle deciding why a particular machine didn't get linked up due to obscure NTP or DNS issues.

If there's another source of truth in your infrastructure, Ansible can also connect to that. Ansible can draw inventory, group, and variable information from sources like EC2, Rackspace, OpenStack, and more.

Here's what a plain text inventory file looks like:

```
[webservers]
```

```
www1.example.com
```

```
www2.example.com
```

```
[dbservers]
```

```
db0.example.com
```

```
db1.example.com
```

Once inventory hosts are listed, variables can be assigned to them in simple text files (in a subdirectory called ?group_vars/? or ?host_vars/?) or directly in the inventory file.

Or, as already mentioned, use a dynamic inventory to pull your inventory from data sources like EC2, Rackspace, or OpenStack.

Playbooks

?

Playbooks can finely orchestrate multiple slices of your infrastructure topology, with very detailed control over how many machines to tackle at a time. This is where Ansible starts to get most interesting.

Ansible's approach to orchestration is one of finely-tuned simplicity, as we believe your automation code should make perfect sense to you years down the road and there should be very little to remember about special syntax or features.

Here's what a simple playbook looks like:

```
---
```

```
-
```

```
hosts
```

```
:
```

```
webservers
```

```
serial
```

```
:
```

```
5
```

```
# update 5 machines at a time
```

```
roles
```

```
:
```

```
-
```

```
common
```

```
-
```


webapp

-

hosts

:

content_servers

roles

:

-

common

-

content

The Ansible search path

?

Modules, module utilities, plugins, playbooks, and roles can live in multiple locations. If you write your own code to extend Ansible's core features, you may have multiple files with similar or the same names in different locations on your Ansible control node. The search path determines which of these files Ansible will discover and use on any given playbook run.

Ansible's search path grows incrementally over a run. As

Ansible finds each playbook and role included in a given run, it appends any directories related to that playbook or role to the search path. Those

directories remain in scope for the duration of the run, even after the playbook or role

has finished executing. Ansible loads modules, module utilities, and plugins in this order:

Directories adjacent to a playbook specified on the command line. If you run Ansible with `ansible-playbook`

`/path/to/play.yml`

, Ansible appends these directories if they exist:

`/path/to/modules`

`/path/to/module_utils`

`/path/to/plugins`

Directories adjacent to a playbook that is statically imported by a `playbook` specified on the command line. If

`play.yml`

includes

-

`import_playbook:`

`/path/to/subdir/play1.yml`

, Ansible appends these directories if they exist:

`/path/to/subdir/modules`

`/path/to/subdir/module_utils`

`/path/to/subdir/plugins`

Subdirectories of a role directory referenced by a `playbook`. If

`play.yml`

runs

`myrole`

, Ansible appends these directories if they exist:

`/path/to/roles/myrole/modules`

`/path/to/roles/myrole/module_utils`

`/path/to/roles/myrole/plugins`

Directories specified as default paths in

`ansible.cfg`

or by the related

environment variables, including the paths for the various plugin types. See

Ansible Configuration Settings

for more information.

Sample
ansible.cfg

fields:

DEFAULT_MODULE_PATH
DEFAULT_MODULE_UTILS_PATH
DEFAULT_CACHE_PLUGIN_PATH
DEFAULT_FILTER_PLUGIN_PATH

Sample environment variables:

ANSIBLE_LIBRARY
ANSIBLE_MODULE_UTILS
ANSIBLE_CACHE_PLUGINS
ANSIBLE_FILTER_PLUGINS

The standard directories that ship as part of the Ansible distribution.

Caution

Modules, module utilities, and plugins in user-specified directories will
override the standard versions. This includes some files with generic names.

For example, if you have a file named

basic.py

in a user-specified

directory, it will override the standard

ansible.module_utils.basic

.

If you have more than one module, module utility, or plugin with the same name in different user-specified directories,
the order of commands at the command line and the order of includes and roles in each play will affect which one is
found and used on that particular play.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/testing.html

Testing Ansible ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Testing Ansible
Edit on GitHub
Testing Ansible
?
Why test your Ansible contributions?
Types of tests
Testing within GitHub & Azure Pipelines
Organization
Rerunning a failing CI job
How to test a PR
Setup: Installing Pytest and required Pytest libraries
Setup: Checking out a Pull Request
Testing the Pull Request
Run sanity tests
Run unit tests
Run integration tests
Code Coverage Online
Want to know more about testing?
Why test your Ansible contributions?
?

If you're a developer, one of the most valuable things you can do is to look at GitHub issues and help fix bugs, since bug-fixing is almost always prioritized over feature development. Even for non-developers, helping to test pull requests for bug fixes and features is still immensely valuable.

Ansible users who understand how to write playbooks and roles should be able to test their work. GitHub pull requests will automatically run a variety of tests (for example, Azure Pipelines) that show bugs in action. However, contributors must also test their work outside of the automated GitHub checks and show evidence of these tests in the pull request to ensure that their work will be more likely to be reviewed and merged.

Read on to learn how Ansible is tested, how to test your contributions locally, and how to extend testing capabilities.

If you want to learn about testing collections, read

Testing collections

Types of tests

?

At a high level, we have the following classifications of tests:

sanity

:

Sanity Tests

Sanity tests are made up of scripts and tools used to perform static code analysis.

The primary purpose of these tests is to enforce Ansible coding standards and requirements.

integration

:

Integration tests

Functional tests of modules and Ansible Core functionality.

units

:

Unit Tests

Tests directly against individual parts of the code base.

Testing within GitHub & Azure Pipelines

?

Organization

?

When Pull Requests (PRs) are created they are tested using Azure Pipelines, a Continuous Integration (CI) tool. Results are shown at the end of every PR.

When Azure Pipelines detects an error and it can be linked back to a file that has been modified in the PR then the relevant lines will be added as a GitHub comment. For example:

The test `ansible-test sanity --test pep8` failed with the following errors:

```
lib/ansible/modules/network/foo/bar.py:509:17: E265 block comment should start with '#'
```

The test `ansible-test sanity --test validate-modules` failed with the following error:

```
lib/ansible/modules/network/foo/bar.py:0:0: E307 version_added should be 2.4. Currently 2.3
```

From the above example we can see that

--test

pep8

and

--test

validate-modules

have identified an issue. The commands given allow you to run the same tests locally to ensure you've fixed all issues without having to push your changes to GitHub and wait for Azure Pipelines, for example:

If you haven't already got Ansible available, use the local checkout by running:

source hacking/env-setup

Then run the tests detailed in the GitHub comment:

ansible-test sanity --test pep8

ansible-test sanity --test validate-modules

If there isn't a GitHub comment stating what's failed you can inspect the results by clicking on the ?Details? button under the ?checks have failed? message at the end of the PR.

Rerunning a failing CI job

?

Occasionally you may find your PR fails due to a reason unrelated to your change. This could happen for several reasons, including:

a temporary issue accessing an external resource, such as a yum or Git repo

a timeout creating a virtual machine to run the tests on

If either issue appears to be the case, you can rerun the Azure Pipelines test by:

adding a comment with

/rebuild

(full rebuild) or

/rebuild_failed

(rebuild only failed CI nodes) to the pull request

closing and re-opening the pull request (full rebuild)

making another change to the branch and pushing to GitHub

If the issue persists, please contact the community. Visit the

Ansible communication guide

for details.

How to test a PR

?

Ideally, the code should add tests that prove that the code works. That's not always possible and tests are not always comprehensive, especially when a user doesn't have access to a wide variety of platforms, or is using an API or web service. In these cases, live testing against real equipment can be more valuable than automation that runs against simulated interfaces. In any case, things should always be tested manually the first time as well.

Thankfully, helping to test Ansible is pretty straightforward, assuming you are familiar with how Ansible works.

Setup: Installing Pytest and required Pytest libraries

?

Ansible's unit testing framework leverages the pytest library. Before diving into testing, ensure you have

pytest

installed alongside any additional pytest libraries such as

pytest-mock

and

pytest-xdist

.

Refer to the documentation for more information:

Unit Tests

.

Setup: Checking out a Pull Request

?

You can do this by:

checking out Ansible

fetching the proposed changes into a test branch

testing

commenting on that particular issue on GitHub

Here's how:

Warning

Testing source code from GitHub pull requests sent to us does have some inherent risk, as the source code sent may have mistakes or malicious code that could have a negative impact on your system. We recommend doing all testing on a virtual machine, whether a cloud instance, or locally. Some users like Vagrant or Docker for this, but they are optional. It is also useful to have virtual machines of different Linux or other flavors, since some features (for example, package managers such as apt or yum) are specific to those OS versions.

Create a fresh area to work:

```
git clone https://github.com/ansible/ansible.git ansible-pr-testing
```

```
cd ansible-pr-testing
```

Next, find the pull request you'd like to test and make a note of its number. It will look something like this:

Use `os.path.sep` instead of hardcoding `/` #65381

Note

Only test

```
ansible:devel
```

It is important that the PR request target be

```
ansible:devel
```

, as we do not accept pull requests into any other branch. Dot releases are cherry-picked manually by Ansible staff.

Use the pull request number when you fetch the proposed changes and create your branch for testing:

```
git fetch origin refs/pull/XXXX/head:testing_PRXXXX
```

```
git checkout testing_PRXXXX
```

The first command fetches the proposed changes from the pull request and creates a new branch named

```
testing_PRXXXX
```

, where the XXXX is the actual number associated with the pull request (for example, 65381). The second command checks out the newly created branch.

Note

If the GitHub user interface shows that the pull request will not merge cleanly, we do not recommend proceeding if you

are not somewhat familiar with Git and coding, as you will have to resolve a merge conflict. This is the responsibility of the original pull request contributor.

Note

Some users do not create feature branches, which can cause problems when they have multiple, unrelated commits in their version of devel

. If the source looks like

someuser:devel

, make sure there is only one commit listed on the pull request.

The Ansible source includes a script that allows you to use Ansible directly from source without requiring a full installation that is frequently used by developers on Ansible.

Simply source it (to use the Linux/Unix terminology) to begin using it immediately:

```
source ./hacking/env-setup
```

This script modifies the

PYTHONPATH

environment variables (along with a few other things), which will be temporarily set as long as your shell session is open.

Testing the Pull Request

?

At this point, you should be ready to begin testing!

Some ideas of what to test are:

Create a test Playbook with the examples in and check if they function correctly

Test to see if any Python backtraces returned (that's a bug)

Test on different operating systems, or against different library versions

Run sanity tests

?

```
ansible-test
```

```
sanity
```

More information:

Sanity Tests

Run unit tests

?

```
ansible-test
```

```
units
```

More information:

Unit Tests

Run integration tests

?

```
ansible-test
```

```
integration
```

```
-v
```

```
ping
```

More information:

Integration tests

Any potential issues should be added as comments on the pull request (and it is acceptable to comment if the feature works as well), remembering to include the output of

ansible

```
--version
```

Example:

Works for me! Tested on `Ansible 2.3.0`. I verified this on CentOS 6.5 and also Ubuntu 14.04.

If the PR does not resolve the issue, or if you see any failures from the unit/integration tests, just include that output instead:

This change causes errors for me.

When I ran this Ubuntu 16.04 it failed with the following:

...

some output

StackTrace

some other output

...

Code Coverage Online

?

The online code coverage reports

is a good way

to identify areas for testing improvement in Ansible. By following red colors you can

drill down through the reports to find files that have no tests at all. Adding both

integration and unit tests that show clearly how code should work, verify important

Ansible functions and increases testing coverage in areas where there is none is a valuable

way to help improve Ansible.

The code coverage reports only cover the

devel

branch of Ansible where new feature

development takes place. Pull requests and new code will be missing from the codecov.io

coverage reports so local reporting is needed. Most

ansible-test

commands allow you

to collect code coverage, this is particularly useful to indicate where to extend

testing. See

Testing Ansible and Collections

for more information.

Want to know more about testing?

?

If you'd like to know more about the plans for improving testing Ansible then why not join the

Ansible community forum

.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/testing_integration.html

Integration tests ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Integration tests
Edit on GitHub
Integration tests
?

Topics
Integration tests
Quick Start
Configuration
ansible-test command
integration_config.yml
Prerequisites
Non-destructive Tests
Destructive Tests
Windows Tests
Tests in containers
Running Integration Tests
Container Images
Other configuration for Cloud Tests
IAM policies for AWS
testing-policies
Other Definitions required
Network Tests
Where to find out more
The Ansible integration Test system.
Tests for playbooks, by playbooks.
Some tests may require credentials. Credentials may be specified with
credentials.yml

.
Some tests may require root.

Note
Every new module and plugin should have integration tests, even if the tests cannot be run on Ansible CI infrastructure.
In this case, the tests should be marked with the
unsupported
alias in
aliases file

.
Quick Start

?

It is highly recommended that you install and activate the

argcomplete

Python package.

It provides tab completion in

bash

for the

ansible-test

test runner.

Configuration

?

ansible-test command

?

The example below assumes

bin/

is in your

\$PATH

. An easy way to achieve that

is to initialize your environment with the

env-setup

command:

source hacking/env-setup

ansible-test --help

You can also call

ansible-test

with the full path:

bin/ansible-test --help

integration_config.yml

?

Making your own version of

integration_config.yml

can allow for setting some

tunable parameters to help run the tests better in your environment. Some

tests (for example, cloud tests) will only run when access credentials are provided. For more

information about supported credentials, refer to the various

cloud-config-*.template

files in the

tests/integration/

directory.

Prerequisites

?

Some tests assume things like

hg

,

svn

, and

git

are installed, and in path. Some tests

(such as those for Amazon Web Services) need separate definitions, which will be covered

later in this document.

(Complete list pending)

Non-destructive Tests

?

These tests will modify files in subdirectories, but will not do things that install or remove packages or things outside of those test subdirectories. They will also not reconfigure or bounce system services.

Note

Running integration tests within containers

To protect your system from any potential changes caused by integration tests, and to ensure a sensible set of dependencies are available we recommend that you always run integration tests with the

`--docker`

option, for example

`--docker`

`ubuntu2204`

. Get the list of supported container images by running

`ansible-test`

`integration`

`--help`

. You can find them in the

target docker images

section of the output. The

default

image is used for sanity and unit tests, as well as for platform independent integration tests such as those for cloud modules.

Run as follows for all POSIX platform tests executed by our CI system in a Fedora 34 container:

`ansible-test integration shippable/ --docker fedora34`

You can exclude a specific test as well, such as for individual modules:

`ansible-test integration --exclude git`

You can target a specific test as well, such as for individual modules:

`ansible-test integration ping`

You can use the

`-v`

option to make the output more verbose:

`ansible-test integration lineinfile -vvv`

Use the following command to list all the available targets:

`ansible-test integration --list-targets`

Note

Bash users

If you use

`bash`

with

`argcomplete`

, obtain a full list by doing:

`ansible-test`

`integration`

`<tab><tab>`

Destructive Tests

?

These tests are allowed to install and remove some trivial packages. You will likely want to devote these to a virtual environment, such as Docker. They won't reformat your filesystem:

`ansible-test integration destructive/ --docker fedora34`

Windows Tests

?

These tests exercise the

`winrm`

connection plugin and Windows modules. You'll

need to define an inventory with a remote Windows Server to use for testing,

and enable PowerShell Remoting to continue.

Running these tests may result in changes to your Windows host, so don't run them against a production/critical Windows environment.

Enable PowerShell Remoting (run on the Windows host by a Remote Desktop):

Enable-PSRemoting -Force

Define Windows inventory:

```
cp inventory.winrm.template inventory.winrm
```

```
$
```

```
{
```

```
EDITOR:-vi
```

```
}
```

```
inventory.winrm
```

Run the Windows tests executed by our CI system:

```
ansible-test windows-integration -v shippable/
```

Tests in containers

?

If you have a Linux system with Docker or Podman installed, running integration tests using the same containers used by

the Ansible continuous integration (CI) system is recommended.

Note

Podman

By default, Podman will only be used if the Docker CLI is not installed. If you have Docker installed but want to use Podman, you can change this behavior by setting the environment variable

```
ANSIBLE_TEST_PREFER_PODMAN
```

.

Note

Docker on non-Linux

Using Docker Engine to run Docker on a non-Linux host (such as macOS) is not recommended.

Some tests may fail, depending on the image used for testing.

Using the

```
--docker-privileged
```

option when running

integration

(not

```
network-integration
```

or

```
windows-integration
```

) may resolve the issue.

Running Integration Tests

?

To run all CI integration test targets for POSIX platforms in a Ubuntu 18.04 container:

```
ansible-test integration shippable/ --docker ubuntu1804
```

You can also run specific tests or select a different Linux distribution.

For example, to run tests for the

```
ping
```

module on a Ubuntu 18.04 container:

```
ansible-test integration ping --docker ubuntu1804
```

Container Images

?

Container images are updated regularly. To see the current list of container images:

```
ansible-test
```

```
integration
```

```
--help
```

The list is under the

target docker images and supported python version

heading.

Other configuration for Cloud Tests

?

To run some tests, you must provide access credentials in a file named `cloud-config-aws.yml`

or

`cloud-config-cs.ini`

in the `tests/integration` directory. Corresponding `.template`

files are available for syntax help. The newer AWS tests now use the file `tests/integration/cloud-config-aws.yml`

.

IAM policies for AWS

?

Ansible needs fairly wide ranging powers to run the tests in an AWS account. These rights can be provided to a dedicated user. These need to be configured before running the test.

`testing-policies`

?

The GitHub repository

`mattclay/aws-terminator`

contains two sets of policies used for all existing AWS module integration tests.

The

`hacking/aws_config/setup_iam.yml`

playbook can be used to setup two groups:

`ansible-integration-ci`

will have the policies applied necessary to run any integration tests not marked as

unsupported

and are designed to mirror those

used by Ansible's CI.

`ansible-integration-unsupported`

will have the additional policies applied

necessary to run the integration tests marked as

unsupported

including tests

for managing IAM roles, users and groups.

Once the groups have been created, you'll need to create a user and make the user a member of these

groups. The policies are designed to minimize the rights of that user. Please note that while this policy does limit

the user to one region, this does not fully restrict the user (primarily due to the limitations of the Amazon ARN

notation). The user will still have wide privileges for viewing account definitions, and will also be able to manage

some resources that are not related to testing (for example, AWS lambdas with different names). Tests should not be run in a primary production account in any case.

Other Definitions required

?

Apart from installing the policy and giving it to the user identity running the tests, a

lambda role

`ansible_integration_tests`

has to be created which has lambda basic execution

privileges.

Network Tests

?

For guidance on writing network test see

Resource module integration tests

.

Where to find out more

?

If you'd like to know more about the plans for improving testing Ansible, join the Ansible community forum

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/testing_sanity.html

Sanity Tests ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Sanity Tests
Edit on GitHub
Sanity Tests

?

Topics

Sanity Tests

Set up your environment

How to run

Available Tests

Sanity tests are made up of scripts and tools used to perform static code analysis.

The primary purpose of these tests is to enforce Ansible coding standards and requirements.

Tests are run with

ansible-test

sanity

.

All available tests are run unless the

--test

option is used.

Set up your environment

?

Install ansible-core

that provides the

ansible-test

tool.

If you want to run checks available in the development version of

ansible-core

,

install it from source code

.

Run

source

hacking/env-setup

from its source code directory in the same terminal session you run your tests.

Install

podman

or

docker

to avoid installing all the dependencies on your system.

If you test files in a collection:

Ensure you have your collection installed in the following path in your home directory:

```
~/ansible_collections/<NAMESPACE>/<COLLECTION_NAME>
```

. For instance, in case of the
community.general
collection, it will be

```
~/ansible_collections/community/general
```

If your collection is hosted on a remote server such as GitHub, clone it to that path as follows:

```
git
```

```
clone
```

```
<COLLECTION_REPO_URL>
```

```
~/ansible_collections/<NAMESPACE>/<COLLECTION_NAME>
```

How to run

?

Note

To run sanity tests using podman or docker, always use the default docker image

by passing the

```
--docker
```

argument without specifying the image name.

When testing files in a collection, change your location to your collection directory you created while
setting up your environment

:

```
cd
```

```
~/ansible_collections/<NAMESPACE>/<COLLECTION_NAME>
```

To run all sanity tests in a container:

```
ansible-test
```

```
sanity
```

```
--docker
```

To run a

specific test

, add the

```
--test
```

```
<NAME>
```

argument, for example,

```
--test
```

```
validate-modules
```

.

To list available tests, run:

```
ansible-test
```

```
sanity
```

```
--list-tests
```

To include disabled tests, add the

```
--allow-disabled
```

argument.

Available Tests

?

See the full list of

sanity tests

, which also details how to fix identified issues.

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/dev_guide/testing_units.html

Unit Tests ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Unit Tests
Edit on GitHub
Unit Tests
?

Unit tests are small isolated tests that target a specific library or module. Unit tests in Ansible are currently the only way of driving tests from python within Ansible's continuous integration process. This means that in some circumstances the tests may be a bit wider than just units.

Topics
Unit Tests
Available Tests
Running Tests
Installing dependencies
Extending unit tests
Structuring Unit Tests
Module test case common code
Fixtures files
Code Coverage For New or Updated Unit Tests
Available Tests
?

Unit tests can be found in
test/units
. Notice that the directory
structure of the tests matches that of
lib/ansible/

.
Running Tests
?

Note
To run unit tests using docker, always use the default docker image
by passing the
--docker
or
--docker
default
argument.

The Ansible unit tests can be run across the whole code base by doing:
cd
/path/to/ansible/source

```
source
hacking/env-setup
ansible-test
units
--docker
-v
```

Against a single file by doing:

```
ansible-test
units
--docker
-v
apt
```

Or against a specific Python version by doing:

```
ansible-test
units
--docker
-v
--python
2
.7
apt
```

If you are running unit tests against things other than modules, such as module utilities, specify the whole file path:

```
ansible-test
units
--docker
-v
test/units/module_utils/basic/test_imports.py
```

For advanced usage see the online help:

```
ansible-test
units
--help
```

You can also run tests in Ansible's continuous integration system by opening a pull request. This will automatically determine which tests to run based on the changes made in your pull request.

Installing dependencies
?

If you are running

```
ansible-test
with the
--docker
or
--venv
```

option you do not need to install dependencies manually.

Otherwise you can install dependencies using the

```
--requirements
option, which will
```

install all the required dependencies needed for unit tests. For example:

```
ansible-test
units
--python
2
.7
--requirements
apache2_module
```

The list of unit test requirements can be found at
test/units/requirements.txt

.

This does not include the list of unit test requirements for
ansible-test

itself,

which can be found at

test/lib/ansible_test/_data/requirements/units.txt

.

See also the

constraints

applicable to all test commands.

Extending unit tests

?

Warning

What a unit test isn't

If you start writing a test that requires external services then
you may be writing an integration test, rather than a unit test.

Structuring Unit Tests

?

Ansible drives unit tests through

pytest

. This

means that tests can either be written as simple functions which are included in any file
name like

test_<something>.py

or as classes.

Here is an example of a function:

#this function will be called simply because it is called test_*

def

test_add

():

a

=

10

b

=

23

c

=

33

assert

a

+

b

==

c

Here is an example of a class:

import

unittest

class

AddTester

(

unittest

```

.
TestCase
):
def
SetUp
():
self
.
a
=
10
self
.
b
=
23
# this function will
def
test_add
():
c
=
33
assert
self
.
a
+
self
.
b
==
c
# this function will
def
test_subtract
():
c
=
-
13
assert
self
.
a
-
self
.
b
==
c

```

Both methods work fine in most circumstances; the function-based interface is simpler and quicker and so that's probably where you should start when you are just trying to add a few basic tests for a module. The class-based test allows more tidy set up and tear down

of pre-requisites, so if you have many test cases for your module you may want to refactor to use that.

Assertions using the simple

`assert`

function inside the tests will give full

information on the cause of the failure with a trace-back of functions called during the assertion. This means that plain asserts are recommended over other external assertion libraries.

A number of the unit test suites include functions that are shared between several modules, especially in the networking arena. In these cases a file is created in the same directory, which is then included directly.

Module test case common code

?

Keep common code as specific as possible within the

`test/units/`

directory structure.

Don't import common unit test code from directories outside the current or parent directories.

Don't import other unit tests from a unit test. Any common code should be in dedicated files that aren't themselves tests.

Fixtures files

?

To mock out fetching results from devices, or provide other complex data structures that come from external libraries, you can use

`fixtures`

to read in pre-generated data.

You can check how

`fixtures`

are used in

`CPU_INFO_TEST_SCENARIOS`

If you are simulating APIs you may find that Python placebo is useful. See

Unit Testing Ansible Modules

for more information.

Code Coverage For New or Updated Unit Tests

?

New code will be missing from the codecov.io coverage reports (see

Testing Ansible

), so

local reporting is needed. Most

`ansible-test`

commands allow you to collect code

coverage; this is particularly useful when to indicate where to extend testing.

To collect coverage data add the

`--coverage`

argument to your

`ansible-test`

command line:

`ansible-test`

`units`

`--coverage`

`apt`

`ansible-test`

`coverage`

`html`

Results will be written to

<test/results/reports/coverage/index.html>

Reports can be generated in several different formats:

ansible-test

coverage

report

- Console report.

ansible-test

coverage

html

- HTML report.

ansible-test

coverage

xml

- XML report.

To clear data between test runs, use the

ansible-test

coverage

erase

command. See

Testing Ansible and Collections

for more information about generating coverage reports.

See also

Unit Testing Ansible Modules

Special considerations for unit testing modules

Testing Ansible and Collections

Running tests locally including gathering and reporting coverage data

Python 3 documentation - 26.4. unittest ? Unit testing framework

The documentation of the unittest framework in python 3

Python 2 documentation - 25.3. unittest ? Unit testing framework

The documentation of the earliest supported unittest framework - from Python 2.6

pytest: helps you write better programs

The documentation of pytest - the framework actually used to run Ansible unit tests

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/galaxy/user_guide.html

Galaxy User Guide ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Finding collections on Galaxy

Finding roles on Galaxy

Get more information about a role

Installing roles from Galaxy

Installing roles

Setting where to install roles

Installing a specific version of a role

Installing multiple roles from a file

Installing roles and collections from the same requirements.yml file

Installing multiple roles from multiple files

Dependencies
Using
meta/requirements.yml
Using
meta/main.yml
List installed roles
Remove an installed role
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Galaxy User Guide
Edit on GitHub
Galaxy User Guide
?

Ansible Galaxy
refers to the
Galaxy

website, a free site for finding, downloading, and sharing community developed collections and roles.

Use Galaxy to jump-start your automation project with great content from the Ansible community. Galaxy provides pre-packaged units of work such as

roles
, and
collections

.
The collection format provides a comprehensive package of automation that may include multiple playbooks, roles, modules, and plugins. See the
Galaxy documentation
for full details on Galaxy.

Finding collections on Galaxy

Finding roles on Galaxy

Get more information about a role

Installing roles from Galaxy

Installing roles

Installing a specific version of a role

Installing multiple roles from a file

Installing roles and collections from the same requirements.yml file

Installing multiple roles from multiple files

Dependencies

List installed roles

Remove an installed role

Finding collections on Galaxy

?

To find collections on Galaxy:

Click

Collections > Collections

in the left-hand navigation.

Type in your search term. You can filter by keyword, tags, and namespaces.

Galaxy presents a list of collections that match your search criteria.

See

Using Ansible collections

for complete details on installing and using collections.

Finding roles on Galaxy

?

To find standalone roles (that is roles that are not part of a collection):

Click

Roles > Roles

in the left-hand navigation.

Type in your search term. You can filter by keyword, tags, and namespaces.

Galaxy presents a list of roles that match your search criteria.

You can optionally search the Galaxy database by tags, platforms, author and multiple keywords using the ansible-galaxy

CLI command.

\$

ansible-galaxy

role

search

elasticsearch

--author

geerlingguy

The search command will return a list of the first 1000 results matching your search:

Found 6 roles matching your search:

Name	Description
----	-----
geerlingguy.elasticsearch	Elasticsearch for Linux.
geerlingguy.elasticsearch-curator	Elasticsearch curator for Linux.
geerlingguy.filebeat	Filebeat for Linux.
geerlingguy.fluentd	Fluentd for Linux.
geerlingguy.kibana	Kibana for Linux.

Get more information about a role

?

Use the

info

command to view more detail about a specific role:

\$

ansible-galaxy

role

info

username.role_name

This returns everything found in Galaxy for the role:

Role: username.role_name

description: Installs and configures a thing, a distributed, highly available NoSQL thing.

active: True

commit: c01947b7bc89ebc0b8a2e298b87ab416aed9dd57

commit_message: Adding travis

commit_url: https://github.com/username/repo_name/commit/c01947b7bc89ebc0b8a2e298b87ab

company: My Company, Inc.

created: 2015-12-08T14:17:52.773Z

download_count: 1

forks_count: 0

github_branch: main

github_repo: repo_name

github_user: username

id: 6381

is_valid: True

issue_tracker_url:

license: Apache

min_ansible_version: 2.15

modified: YYYY-MM-DDTHH:MM:SS.000Z

namespace: username

open_issues_count: 0

path: /Users/username/projects/roles

role_type: ANS

stargazers_count: 0

travis_status_url: https://travis-ci.org/username/repo_name.svg?branch=main

Installing roles from Galaxy

?

The

ansible-galaxy

command comes bundled with Ansible, and you can use it to install roles from Galaxy or directly from a Git based SCM.

You can

also use it to create a new role, remove roles, or perform tasks on the Galaxy website.

The command line tool by default communicates with the Galaxy website API using the server address

<https://galaxy.ansible.com>

. If you run your own internal Galaxy server

and want to use it instead of the default one, pass the

--server

option followed by the address of this galaxy server. You can set this option permanently by setting

the Galaxy server value in your

ansible.cfg

file. See

GALAXY_SERVER

for details on setting the value in

ansible.cfg

.

Installing roles

?

Use the

ansible-galaxy

command to download roles from the

Galaxy website

\$

ansible-galaxy

role

install

namespace.role_name

Setting where to install roles

?

By default, Ansible downloads roles to the first writable directory in the default list of paths

~/.ansible/roles:/usr/share/ansible/roles:/etc/ansible/roles

. This installs roles in the home directory of the user running

ansible-galaxy

.

You can override this with one of the following options:

Set the environment variable

ANSIBLE_ROLES_PATH

in your session.

Use the

--roles-path

option for the

ansible-galaxy

command.

Define

roles_path

in an

ansible.cfg

file.

The following provides an example of using

--roles-path

to install the role into the current working directory:

\$

ansible-galaxy

role

install

--roles-path

.

geerlingguy.apache

See also

Configuring Ansible

All about configuration files

Installing a specific version of a role

?

When the Galaxy server imports a role, it imports any Git tags matching the

Semantic Version

format as versions.

In turn, you can download a specific version of a role by specifying one of the imported tags.

To see the available versions for a role:

Locate the role on the Galaxy search page.

Click on the name to view more details, including the available versions.

To install a specific version of a role from Galaxy, append a comma and the value of a GitHub release tag. For example:

\$

ansible-galaxy

role

install

geerlingguy.apache,3.2.0

It is also possible to point directly to the Git repository and specify a branch name or commit hash as the version. For example, the following will

install a specific commit:

\$

ansible-galaxy

role

install

git+https://github.com/geerlingguy/ansible-role-apache.git,0b7cd353c0250e87a26e0499e59e7fd265cc2f25

Installing multiple roles from a file

?

You can install multiple roles by including the roles in a requirements.yml

file. The format of the file is YAML, and the

file extension must be either

.yml

or

.yaml

.

Use the following command to install roles included in requirements.yml:

\$

ansible-galaxy

install

-r

requirements.yml

Again, the extension is important. If the

.yml

extension is left off, the

ansible-galaxy

CLI assumes the file is in an older, now deprecated,

?basic? format.

Each role in the file will have one or more of the following attributes:

src

The source of the role. Use the format

namespace.role_name

, if downloading from Galaxy; otherwise, provide a URL pointing

to a repository within a Git based SCM. See the examples below. This is a required attribute.

scm

Specify the SCM. As of this writing only

git

or

hg

are allowed. See the examples below. Defaults to

git

.

version:

The version of the role to download. Provide a release tag value, commit hash, or branch name. Defaults to the branch set as a default in the repository, otherwise defaults to the

master

.

name:

Download the role to a specific name. Defaults to the Galaxy name when downloading from Galaxy, otherwise it defaults to the name of the repository.

Use the following example as a guide for specifying roles in requirements.yml

```
:
# from galaxy
-
name
:
yatesr.timezone
# from locally cloned Git repository (git+file:// requires full paths)
-
src
:
git+file:///home/bennojoy/nginx
# from GitHub
-
src
:
https://github.com/bennojoy/nginx
# from GitHub, overriding the name and specifying a specific tag
-
name
:
nginx_role
src
:
https://github.com/bennojoy/nginx
version
:
main
# from GitHub, specifying a specific commit hash
-
src
:
https://github.com/bennojoy/nginx
version
:
"ee8aa41"
# from a webserver, where the role is packaged in a tar.gz
-
name
:
http-role-gz
src
:
https://some.webserver.example.com/files/main.tar.gz
# from a webserver, where the role is packaged in a tar.bz2
-
name
:
http-role-bz2
src
:
https://some.webserver.example.com/files/main.tar.bz2
# from a webserver, where the role is packaged in a tar.xz (Python 3.x only)
```



```

-
name
:
http-role-xz
src
:
https://some.webserver.example.com/files/main.tar.xz
# from Bitbucket
-
src
:
git+https://bitbucket.org/willthames/git-ansible-galaxy
version
:
v1.4
# from Bitbucket, alternative syntax and caveats
-
src
:
https://bitbucket.org/willthames/hg-ansible-galaxy
scm
:
hg
# from GitLab or other git-based scm, using git+ssh
-
src
:
[email protected]
:mygroup/ansible-core.git
scm
:
git
version
:
"0.1"
# quoted, so YAML doesn't parse this as a floating-point value
Warning
Embedding credentials into a SCM URL is not secure. Make sure to use safe auth options for security reasons. For
example, use
SSH
,
netrc
or
http.extraHeader
/
url.<base>.pushInsteadOf
in Git config to prevent your credentials from being exposed in logs.
Installing roles and collections from the same requirements.yml file
?
You can install roles and collections from the same requirements files
---
roles
:
# Install a role from Ansible Galaxy.

```

```

-
name
:
geerlingguy.java
version
:
"1.9.6"
# note that ranges are not supported for roles
collections
:
# Install a collection from Ansible Galaxy.
-
name
:
community.general
version
:
">=7.0.0"
source
:
https://galaxy.ansible.com
Installing multiple roles from multiple files
?
For large projects, the
include
directive in a
requirements.yml
file provides the ability to split a large file into multiple smaller files.
For example, a project may have a
requirements.yml
file, and a
webserver.yml
file.
Below are the contents of the
webserver.yml
file:
# from github
-
src:
https://github.com/bennojoy/nginx
# from Bitbucket
-
src:
git+https://bitbucket.org/willthames/git-ansible-galaxy
version:
v1.4
The following shows the contents of the
requirements.yml
file that now includes the
webserver.yml
file:
# from galaxy
-
name:

```

yatesr.timezone

-

include:

<path_to_requirements>/webserver.yml

To install all the roles from both files, pass the root file, in this case requirements.yml

on the

command line, as follows:

\$

ansible-galaxy

role

install

-r

requirements.yml

Dependencies

?

Roles can also be dependent on other roles, and when you install a role that has dependencies, those dependencies will automatically be installed to the

roles_path

.

There are two ways to define the dependencies of a role:

using

meta/requirements.yml

using

meta/main.yml

Using

meta/requirements.yml

?

New in version 2.10.

You can create the file

meta/requirements.yml

and define dependencies in the same format used for

requirements.yml

described in the

Installing multiple roles from a file

section.

From there, you can import or include the specified roles in your tasks.

Using

meta/main.yml

?

Alternatively, you can specify role dependencies in the

meta/main.yml

file by providing a list of roles under the

dependencies

section. If the source of a role is Galaxy, you can simply specify the role in

the format

namespace.role_name

. You can also use the more complex format in

requirements.yml

, allowing you to provide

src

,

scm

,

version

, and

name

.

Dependencies installed that way, depending on other factors described below, will also be executed before

this role is executed during play execution.

To better understand how dependencies are handled during play execution, see

Roles

.

The following shows an example

meta/main.yml

file with dependent roles:

dependencies

:

-

geerlingguy.java

galaxy_info

:

author

:

geerlingguy

description

:

Elasticsearch for Linux.

company

:

"Midwestern

Mac,

LLC"

license

:

"license

(BSD,

MIT)"

min_ansible_version

:

2.4

galaxy_tags

:

-

web

-

system

-

monitoring

-

logging

-

lucene

-

elk

-

elasticsearch

Tags are inherited
down

the dependency chain. In order for tags to be applied to a role and all its dependencies, the tag should be applied to the role, not to all the tasks within a role.

Roles listed as dependencies are subject to conditionals and tag filtering, and may not execute fully depending on what tags and conditionals are applied.

If the source of a role is Galaxy, specify the role in the format
namespace.role_name

```
:
```

```
dependencies
```

```
:
```

```
-
```

```
geerlingguy.apache
```

```
-
```

```
geerlingguy.ansible
```

Alternately, you can specify the role dependencies in the complex form used in
requirements.yml

as follows:

```
dependencies
```

```
:
```

```
-
```

```
name
```

```
:
```

```
geerlingguy.ansible
```

```
-
```

```
name
```

```
:
```

```
composer
```

```
src
```

```
:
```

```
git+https://github.com/geerlingguy/ansible-role-composer.git
```

```
version
```

```
:
```

```
775396299f2da1f519f0d8885022ca2d6ee80ee8
```

Note

Galaxy expects all role dependencies to exist in Galaxy, and therefore dependencies to be specified in the
namespace.role_name

format. If you import a role with a dependency where the
src

value is a URL, the import process will fail.

List installed roles

?

Use

list

to show the name and version of each role installed in the
roles_path

```
.
```

```
$
```

```
ansible-galaxy
```

```
role
```

```
list
```

```
-
```

```
namespace-1.foo,
```

v2.7.2

-

namespace2.bar,

v2.6.2

Remove an installed role

?

Use

remove

to delete a role from

roles_path

:

\$

ansible-galaxy

role

remove

namespace.role_name

See also

Using Ansible collections

Shareable collections of modules, playbooks and roles

Roles

Reusable tasks, handlers, and other files in a known directory structure

Working with command line tools

Perform other related operations

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/getting_started/index.html

Getting started with Ansible ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Introduction to Ansible

Start automating with Ansible

Building an inventory

Creating a playbook

Ansible concepts

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

[Return Values](#)
[Ansible Configuration Settings](#)
[Controlling how Ansible behaves: precedence rules](#)
[YAML Syntax](#)
[Python 3 Support](#)
[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)

[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[ansible-core Roadmaps](#)

[Ansible](#)
[Getting started with Ansible](#)
[Edit on GitHub](#)
[Getting started with Ansible](#)
[?](#)

Ansible automates the management of remote systems and controls their desired state.

As shown in the preceding figure, most Ansible environments have three main components:

Control node

A system on which Ansible is installed.

You run Ansible commands such as

```
ansible
```

or

```
ansible-inventory
```

on a control node.

Inventory

A list of managed nodes that are logically organized.

You create an inventory on the control node to describe host deployments to Ansible.

Managed node

A remote system, or host, that Ansible controls.

Introduction to Ansible

[Start automating with Ansible](#)

[Building an inventory](#)

[Creating a playbook](#)

[Ansible concepts](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: <https://docs.ansible.com/ansible/latest/index.html>

[Ansible Documentation](#) ? [Ansible Community Documentation](#)

[Blog](#)

[Ansible community forum](#)

[Documentation](#)

[Ansible Community Documentation](#)

[Ansible](#)

Select version:

[latest](#)

[11](#)

[devel](#)

Search docs:

[Ansible getting started](#)

[Getting started with Ansible](#)

[Getting started with Execution Environments](#)

[Installation, Upgrade & Configuration](#)

[Installation Guide](#)

[Ansible Porting Guides](#)

[Using Ansible](#)

[Building Ansible inventories](#)

[Using Ansible command line tools](#)

[Using Ansible playbooks](#)

[Protecting sensitive data with Ansible vault](#)

[Using Ansible modules and plugins](#)

[Using Ansible collections](#)

[Using Ansible on Windows, BSD, and z/OS UNIX](#)

[Ansible tips and tricks](#)

[Contributing to Ansible](#)

[Ansible Community Guide](#)

[Ansible Collections Contributor Guide](#)

[ansible-core Contributors Guide](#)

[Advanced Contributor Guide](#)

[Ansible documentation style guide](#)

[Extending Ansible](#)

[Developer Guide](#)

[Common Ansible Scenarios](#)

[Legacy Public Cloud Guides](#)

[Network Automation](#)

[Network Getting Started](#)

[Network Advanced Topics](#)

[Network Developer Guide](#)

[Ansible Galaxy](#)

[Galaxy User Guide](#)

[Galaxy Developer Guide](#)

[Reference & Appendices](#)

[Collection Index](#)

[Indexes of all modules and plugins](#)

[Playbook Keywords](#)

[Return Values](#)

[Ansible Configuration Settings](#)

[Controlling how Ansible behaves: precedence rules](#)

[YAML Syntax](#)

[Python 3 Support](#)

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Ansible Documentation
Ansible Documentation
?

Welcome to Ansible community documentation!

This documentation covers the version of Ansible noted in the upper left corner of this page.

We maintain multiple versions of Ansible and of the documentation, so please be sure you are using the version of the documentation that covers the version of Ansible you're using.

For recent features, we note the version of Ansible where the feature was added.

Ansible releases a new major release approximately twice a year.

The core application evolves somewhat conservatively, valuing simplicity in language design and setup.

Contributors develop and change modules and plugins, hosted in collections, much more quickly.

Ansible getting started

Getting started with Ansible

Introduction to Ansible

Start automating with Ansible

Building an inventory

Creating a playbook

Ansible concepts

Getting started with Execution Environments

Introduction to Execution Environments

Setting up your environment

Building your first Execution Environment

Running your EE

Running Ansible with the community EE image

Installation, Upgrade & Configuration

Installation Guide

Installing Ansible

Installing Ansible on specific operating systems

Configuring Ansible

Ansible Porting Guides

Ansible 12 Porting Guide

Ansible 11 Porting Guide

Ansible 10 Porting Guide

Ansible 9 Porting Guide

Ansible 8 Porting Guide

Ansible 7 Porting Guide

Ansible 6 Porting Guide

Ansible 5 Porting Guide

Ansible 4 Porting Guide

- Ansible 3 Porting Guide
- Ansible 2.10 Porting Guide
- Ansible 2.9 Porting Guide
- Ansible 2.8 Porting Guide
- Ansible 2.7 Porting Guide
- Ansible 2.6 Porting Guide
- Ansible 2.5 Porting Guide
- Ansible 2.4 Porting Guide
- Ansible 2.3 Porting Guide
- Ansible 2.0 Porting Guide
- Using Ansible
 - Building Ansible inventories
 - How to build your inventory
 - Working with dynamic inventory
 - Patterns: targeting hosts and groups
 - Connection methods and details
 - Using Ansible command line tools
 - Introduction to ad hoc commands
 - Working with command line tools
 - Ansible CLI cheatsheet
 - Using Ansible playbooks
 - Ansible playbooks
 - Working with playbooks
 - Executing playbooks
 - Advanced playbook syntax
 - Manipulating data
 - Protecting sensitive data with Ansible vault
 - Ansible Vault
 - Managing vault passwords
 - Encrypting content with Ansible Vault
 - Using encrypted variables and files
 - Configuring defaults for using encrypted content
 - When are encrypted files made visible?
 - Format of files encrypted with Ansible Vault
 - Using Ansible modules and plugins
 - Introduction to modules
 - Boolean variables
 - Module maintenance and support
 - Rejecting modules
 - Working with plugins
 - Modules and plugins index
 - Using Ansible collections
 - Installing collections
 - Removing a collection
 - Downloading collections
 - Listing collections
 - Verifying collections
 - Using collections in a playbook
 - Collections index
 - Using Ansible on Windows, BSD, and z/OS UNIX
 - Managing BSD hosts with Ansible
 - Managing Windows hosts with Ansible
 - Managing z/OS UNIX hosts with Ansible
 - Ansible tips and tricks

General tips
Playbook tips
Inventory tips
Execution tricks
Sample Ansible setup
Contributing to Ansible
Ansible Community Guide
Getting started
Contributor path
Ansible Collections Contributor Guide
The Ansible Collections Development Cycle
Requesting changes to a collection
Creating your first collection pull request
Testing Collection Contributions
Review checklist for collection PRs
Ansible community package collections requirements
Guidelines for collection maintainers
Contributing to Ansible-maintained Collections
Ansible Community Steering Committee
Contributing to the Ansible Documentation
Other Tools and Programs
Working with the Ansible collection repositories
ansible-core Contributors Guide
Reporting bugs and requesting features
Contributing to the Ansible Documentation
The Ansible Development Cycle
Other Tools and Programs
Working with the Ansible repo
Advanced Contributor Guide
Committers Guidelines
Release Manager Guidelines
GitHub Admins
Ansible Ecosystem Project Development Resources
Ansible documentation style guide
Linguistic guidelines
reStructuredText guidelines
Markdown guidelines
Accessibility guidelines
More resources
Extending Ansible
Developer Guide
Adding modules and plugins locally
Should you develop a module?
Developing modules
Contributing your module to an existing Ansible collection
Conventions, tips, and pitfalls
Ansible and Python 3
Debugging modules
Module format and documentation
Ansible markup
Adjacent YAML documentation files
Windows module development walkthrough
Creating a new collection
Testing Ansible

The lifecycle of an Ansible module or plugin
Developing plugins
Developing dynamic inventory
Developing
ansible-core
Ansible module architecture
Python API
Rebasing a pull request
Using and developing module utilities
Ansible collection creator path
Developing collections
Migrating Roles to Roles in Collections on Galaxy
Collection Galaxy metadata structure
Ansible architecture
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Basic Concepts
How Network Automation is Different
Run Your First Command and Playbook
Build Your Inventory
Use Ansible network roles
Beyond the basics
Working with network connection options
Resources and next steps
Network Advanced Topics
Network Resource Modules
Ansible Network Examples
Parsing semi-structured text with Ansible
Validate data against set criteria with Ansible
Network Debug and Troubleshooting Guide
Working with command output and prompts in network modules
Ansible Network FAQ
Platform Options
Network Developer Guide
Developing network resource modules
Developing network plugins
Documenting new network platforms
Ansible Galaxy
Galaxy User Guide
Finding collections on Galaxy
Finding roles on Galaxy
Installing roles from Galaxy
Galaxy Developer Guide
Creating collections for Galaxy
Creating roles for Galaxy
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules

[YAML Syntax](#)
[Python 3 Support](#)
[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[Ansible project 12.0](#)
[Ansible project 11.0](#)
[Ansible project 10.0](#)
[Ansible project 9.0](#)
[Ansible project 8.0](#)
[Ansible project 7.0](#)
[Ansible project 6.0](#)
[Ansible project 5.0](#)
[Ansible project 4.0](#)
[Ansible project 3.0](#)
[Ansible project 2.10](#)
[Older Roadmaps](#)
[ansible-core Roadmaps](#)
[Ansible-core 2.19](#)
[Ansible-core 2.18](#)
[Ansible-core 2.17](#)
[Ansible-core 2.16](#)
[Ansible-core 2.15](#)
[Ansible-core 2.14](#)
[Ansible-core 2.13](#)
[Ansible-core 2.12](#)
[Ansible-core 2.11](#)
[Ansible-base 2.10](#)
[Next](#)
[© Copyright Ansible project contributors.](#)
[Last updated on Oct 08, 2025.](#)

Content from: https://docs.ansible.com/ansible/latest/installation_guide/intro_installation.html

Installing Ansible ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Installing Ansible

Control node requirements

Managed node requirements

Node requirement summary

Selecting an Ansible package and version to install

Installing and upgrading Ansible with pipx

Installing Ansible

Upgrading Ansible

Installing Extra Python Dependencies

Installing and upgrading Ansible with pip

Locating Python

Ensuring

pip

is available

Installing Ansible

Upgrading Ansible

Installing Ansible to containers

Installing for development

Installing

devel

from GitHub with

pip

Running the

devel

branch from a clone

Confirming your installation

Adding Ansible command shell completion

Installing

argcomplete

Configuring

argcomplete

Installing Ansible on specific operating systems

Configuring Ansible

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools
Using Ansible playbooks
Protecting sensitive data with Ansible vault
Using Ansible modules and plugins
Using Ansible collections
Using Ansible on Windows, BSD, and z/OS UNIX
Ansible tips and tricks
Contributing to Ansible
Ansible Community Guide
Ansible Collections Contributor Guide
ansible-core Contributors Guide
Advanced Contributor Guide
Ansible documentation style guide
Extending Ansible
Developer Guide
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Installation Guide
Installing Ansible
Edit on GitHub
Installing Ansible
?

Ansible is an agentless automation tool that you install on a single host (referred to as the control node).

From the control node, Ansible can manage an entire fleet of machines and other devices (referred to as managed nodes) remotely with SSH, Powershell remoting, and numerous other transports, all from a simple command-line interface with no databases or daemons required.

Control node requirements

Managed node requirements

Node requirement summary

Selecting an Ansible package and version to install

Installing and upgrading Ansible with pipx

Installing Ansible

Upgrading Ansible

Installing Extra Python Dependencies

Installing and upgrading Ansible with pip

Locating Python

Ensuring

pip

is available

Installing Ansible

Upgrading Ansible

Installing Ansible to containers

Installing for development

Installing

devel

from GitHub with

pip

Running the

devel

branch from a clone

Confirming your installation

Adding Ansible command shell completion

Installing

argcomplete

Configuring

argcomplete

Global configuration

Per command configuration

Using

argcomplete

with zsh or tcsh

Control node requirements

?

For your

control

node (the machine that runs Ansible), you can use nearly any UNIX-like machine with Python installed. This includes Red Hat, Debian, Ubuntu, macOS, BSDs, and Windows under a

Windows Subsystem for Linux (WSL) distribution

. Windows without WSL is not natively supported as a control node; see

Matt Davis' blog post

for more information.

Managed node requirements

?

The

managed

node (the machine that Ansible is managing) does not require Ansible to be installed, but requires Python to run Ansible-generated Python code.

The managed node also needs a user account that can connect through SSH to the node with an interactive POSIX shell.

Note

There can be exceptions in module requirements. For example, network modules do not require Python on the managed device. See documentation for the modules you use.

Node requirement summary

?

You can find details about control and managed node requirements, including Python versions, for each Ansible version in the

`ansible-core` control node Python support

and

`ansible-core` support matrix

sections.

Selecting an Ansible package and version to install

?

Ansible's community packages are distributed in two ways:

`ansible-core`

: a minimalist language and runtime package containing a set of built-in modules and plugins

.

`ansible`

: a much larger "batteries included" package, which adds a community-curated selection of Ansible Collections

for automating a wide variety of devices.

Choose the package that fits your needs.

The following instructions use

`ansible`

as a package name, but you can substitute

`ansible-core`

if you prefer to start with the minimal package and separately install only the Ansible Collections you require.

The

`ansible`

or

`ansible-core`

packages may be available in your operating systems package manager, and you are free to install these packages with your preferred method. For more information, see the

Installing Ansible on specific operating systems

guide. These installation instructions only cover the officially supported means of installing the python packages with `pip`

.

See the

Ansible package release status table

for the

`ansible-core`

version included in the package.

Installing and upgrading Ansible with `pipx`

?

On some systems, it may not be possible to install Ansible with

`pip`

, due to decisions made by the operating system developers. In such cases,

`pipx`

is a widely available alternative.

These instructions will not go over the steps to install

`pipx`

; if those instructions are needed, please continue to the
pipx installation instructions
for more information.

Installing Ansible

?

Use

`pipx`

in your environment to install the full Ansible package:

\$

`pipx`

`install`

`--include-deps`

`ansible`

You can install the minimal

`ansible-core`

package:

\$

`pipx`

`install`

`ansible-core`

Alternately, you can install a specific version of

`ansible-core`

:

\$

`pipx`

`install`

`ansible-core`

`==`

`2`

`.12.3`

Upgrading Ansible

?

To upgrade an existing Ansible installation to the latest released version:

\$

`pipx`

`upgrade`

`--include-injected`

`ansible`

Installing Extra Python Dependencies

?

To install additional python dependencies that may be needed, with the example of installing the
`argcomplete`

python package as described below:

\$

`pipx`

`inject`

`ansible`

`argcomplete`

Include the

`--include-apps`

option to make apps in the additional python dependency available on your PATH. This allows you to execute
commands for those apps from the shell.

\$

`pipx`

inject

--include-apps

ansible

argcomplete

If you need to install dependencies from a requirements file, for example when installing the Azure collection, you can use

runpip

.

\$

pipx

runpip

ansible

install

-r

~/ansible/collections/ansible_collections/azure/azcollection/requirements.txt

Installing and upgrading Ansible with pip

?

Locating Python

?

Locate and remember the path to the Python interpreter you wish to use to run Ansible. The following instructions refer to this Python as

python3

.

For example, if you have determined that you want the Python at

/usr/bin/python3.9

to be the one that you will install Ansible under, specify that instead of

python3

.

Ensuring

pip

is available

?

To verify whether

pip

is already installed for your preferred Python:

\$

python3

-m

pip

-V

If all is well, you should see something like the following:

\$

python3

-m

pip

-V

pip 21.0.1 from /usr/lib/python3.9/site-packages/pip (python 3.9)

If so,

pip

is available, and you can move on to the

next step

.

If you see an error like

No

module
named
pip
, you will need to install
pip
under your chosen Python interpreter before proceeding.
This may mean installing an additional OS package (for example,
python3-pip
) , or installing the latest
pip
directly from the Python Packaging Authority by running the following:

```
$  
curl  
https://bootstrap.pypa.io/get-pip.py  
-o  
get-pip.py  
$  
python3  
get-pip.py  
--user
```

You may need to perform some additional configuration before you are able to run Ansible. See the Python documentation on installing to the user site for more information.

Installing Ansible

?
Use
pip
in your selected Python environment to install the full Ansible package for the current user:

```
$  
python3  
-m  
pip  
install  
--user  
ansible
```

You can install the minimal
ansible-core
package for the current user:

```
$  
python3  
-m  
pip  
install  
--user  
ansible-core  
Alternately, you can install a specific version of  
ansible-core  
:  
$  
python3  
-m  
pip  
install
```

--user
ansible-core

==

2

.12.3

Upgrading Ansible

?

To upgrade an existing Ansible installation in this Python environment to the latest released version, simply add

--upgrade

to the command above:

\$

python3

-m

pip

install

--upgrade

--user

ansible

Installing Ansible to containers

?

Instead of installing Ansible content manually, you can simply build an execution environment container image or use one of the available community images as your control node.

See

Getting started with Execution Environments

for details.

Installing for development

?

If you are testing new features, fixing bugs, or otherwise working with the development team on changes to the core code, you can install and run the source from GitHub.

Note

You should only install and run the

devel

branch if you are modifying

ansible-core

or trying out features under development. This is a rapidly changing source of code and can become unstable at any point.

For more information on getting involved in the Ansible project, see the

Ansible Community Guide

.

For more information on creating Ansible modules and Collections, see the

Developer Guide

.

Installing

devel

from GitHub with

pip

?

You can install the

devel

branch of

ansible-core

directly from GitHub with

pip

:

```
$
python3
-m
pip
install
--user
https://github.com/ansible/ansible/archive/devel.tar.gz
You can replace
devel
in the URL mentioned above, with any other branch or tag on GitHub to install older versions of Ansible, tagged alpha or
beta versions, and release candidates.
Running the
devel
branch from a clone
?
ansible-core
is easy to run from source. You do not need
root
permissions to use it and there is no software to actually install. No daemons or database setup are required.
Clone the
ansible-core
repository
$
git
clone
https://github.com/ansible/ansible.git
$
cd
./ansible
Setup the Ansible environment
Using Bash
$
source
./hacking/env-setup
Using Fish
$
source
./hacking/env-setup.fish
To suppress spurious warnings/errors, use
-q
$
source
./hacking/env-setup
-q
Install Python dependencies
$
python3
-m
pip
install
--user
-r
./requirements.txt
Update the
```

devel
branch of
ansible-core
on your local machine
Use pull-with-rebase so any local changes are replayed.

```
$  
git  
pull  
--rebase  
Confirming your installation  
?
```

You can test that Ansible is installed correctly by checking the version:

```
$  
ansible  
--version  
The version displayed by this command is for the associated  
ansible-core  
package that has been installed.  
To check the version of the  
ansible  
package that has been installed:
```

```
$  
ansible-community  
--version  
Adding Ansible command shell completion  
?
```

You can add shell completion of the Ansible command line utilities by installing an optional dependency called `argcomplete`. It supports `bash`, and has limited support for `zsh` and `tcsh`. For more information about installation and configuration, see the `argcomplete` documentation.

```
.  
Installing  
argcomplete  
?
```

If you chose the `pipx` installation instructions:

```
$  
pipx  
inject  
--include-apps  
ansible  
argcomplete
```

If you chose the `pip` installation instructions:

```
$  
python3  
-m  
pip  
install  
--user  
argcomplete
```


Configuring
argcomplete

?

There are 2 ways to configure
argcomplete

to allow shell completion of the Ansible command line utilities: globally or per command.

Global configuration

?

Global completion requires bash 4.2.

\$

activate-global-python-argcomplete

--user

This will write a bash completion file to a user location. Use

--dest

to change the location or

sudo

to set up the completion globally.

Per command configuration

?

If you do not have bash 4.2, you must register each script independently.

\$

eval

\$(

register-python-argcomplete

ansible

)

\$

eval

\$(

register-python-argcomplete

ansible-config

)

\$

eval

\$(

register-python-argcomplete

ansible-console

)

\$

eval

\$(

register-python-argcomplete

ansible-doc

)

\$

eval

\$(

register-python-argcomplete

ansible-galaxy

)

\$

eval

\$(

register-python-argcomplete

```
ansible-inventory
)
$
eval
$(
register-python-argcomplete
ansible-playbook
)
$
eval
$(
register-python-argcomplete
ansible-pull
)
$
eval
$(
register-python-argcomplete
ansible-vault
)
```

You should place the above commands into your shell's profile file such as

`~/.profile`

or

`~/.bash_profile`

.

Using

`argcomplete`

with `zsh` or `tcsh`

?

See the

`argcomplete` documentation

.

See also

[Introduction to ad hoc commands](#)

[Examples of basic commands](#)

[Working with playbooks](#)

[Learning ansible's configuration management language](#)

[How do I handle the package dependencies required by Ansible package dependencies during Ansible installation ?](#)

[Ansible Installation related to FAQs](#)

[Forum](#)

[Join the Ansible community forum to get help and share insights](#)

[Real-time chat](#)

[How to join Ansible chat channels](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/inventory_guide/intro_dynamic_inventory.htm

Working with dynamic inventory ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

How to build your inventory

Working with dynamic inventory

Inventory script example: Cobbler

Other inventory scripts

Using inventory directories and multiple inventory sources

Static groups of dynamic groups

Patterns: targeting hosts and groups

Connection methods and details

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary

Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible

Building Ansible inventories
Working with dynamic inventory
Edit on GitHub
Working with dynamic inventory

?

Inventory script example: Cobbler
Other inventory scripts
Using inventory directories and multiple inventory sources

Static groups of dynamic groups

If your Ansible inventory fluctuates over time, with hosts spinning up and shutting down in response to business demands, the static inventory solutions described in

How to build your inventory
will not serve your needs.

You may need to track hosts from multiple sources: cloud providers, LDAP, Cobbler

, and/or enterprise CMDB systems.

Ansible integrates all of these options through a dynamic external inventory system.

Ansible supports two ways to connect with external inventory:

Inventory plugins

and

inventory scripts

.

Inventory plugins take advantage of the most recent updates to the Ansible Core code.

We recommend plugins over scripts for dynamic inventory.

You can

write your own plugin

to connect to additional dynamic inventory sources.

You can still use inventory scripts if you choose.

When we implemented inventory plugins, we ensured backwards compatibility through the script inventory plugin.

The examples below illustrate how to use inventory scripts.

If you prefer a GUI for handling dynamic inventory, the inventory database on AWX or

Red Hat Ansible Automation Platform

syncs with all your dynamic inventory sources, provides web and REST access to the results, and offers a graphical inventory editor.

With a database record of all of your hosts, you can correlate past event history and see which hosts have had failures on their last playbook runs.

Inventory script example: Cobbler

?

Ansible integrates seamlessly with

Cobbler

, a Linux installation server originally written by Michael DeHaan and now led by James Cammarata, who works for Ansible.

While primarily used to kickoff OS installations and manage DHCP and DNS, Cobbler has a generic layer that can represent data for multiple configuration management systems (even at the same time) and serve as a ?lightweight CMDB?.

To tie your Ansible inventory to Cobbler, copy

this script

to

/etc/ansible

and

chmod

+x

the file. Run

cobblerd

any time you use Ansible and use the

-i

command line option (for example,

-i

/etc/ansible/cobbler.py

) to communicate with Cobbler using Cobbler?s XMLRPC API.

Add a

cobbler.ini

file in

/etc/ansible

so Ansible knows where the Cobbler server is and some cache improvements can be used. For example:

[cobbler]

Set Cobbler's hostname or IP address

host = http://127.0.0.1/cobbler_api

API calls to Cobbler can be slow. For this reason, we cache the results of an API

call. Set this to the path you want cache files to be written to. Two files

will be written to this directory:

- ansible-cobbler.cache

- ansible-cobbler.index

cache_path = /tmp

The number of seconds a cache file is considered valid. After this many

seconds, a new API call will be made, and the cache file will be updated.

cache_max_age = 900

First test the script by running

/etc/ansible/cobbler.py

directly. You should see some JSON data output, but it may not have anything in it just yet.

Let's explore what this does. In Cobbler, assume a scenario somewhat like the following:

```
cobbler
profile
add
--name
=
webserver
--distro
=
CentOS6-x86_64
cobbler
profile
edit
--name
=
webserver
--mgmt-classes
=
"webserver"
--ksmeta
=
"a=2 b=3"
cobbler
system
edit
--name
=
foo
--dns-name
=
"foo.example.com"
--mgmt-classes
=
"atlanta"
--ksmeta
=
"c=4"
cobbler
system
edit
--name
=
bar
--dns-name
=
"bar.example.com"
--mgmt-classes
=
"atlanta"
--ksmeta
=
"c=5"
```

In the example above, the system `foo.example.com` is addressable by ansible directly, but is also addressable when using the group names `webserver` or `atlanta`. Since Ansible uses SSH, it contacts system foo over

?foo.example.com?, only, never just ?foo?. Similarly, if you tried ?ansible foo?, it would not find the system? but ?ansible ?foo*?? would do, because the system DNS name starts with ?foo?.

The script provides more than host and group info. In addition, as a bonus, when the ?setup? module is run (which happens automatically when using playbooks), the variables ?a?, ?b?, and ?c? will all be auto-populated in the templates:

```
# file: /srv/motd.j2
```

```
Welcome, I am templated with a value of a={{ a }}, b={{ b }}, and c={{ c }}
```

Which could be executed just like this:

```
ansible
```

```
webserver
```

```
-m
```

```
setup
```

```
ansible
```

```
webserver
```

```
-m
```

```
template
```

```
-a
```

```
"src=/tmp/motd.j2 dest=/etc/motd"
```

Note

The name ?webserver? came from Cobbler, as did the variables for the config file. You can still pass in your own variables like normal in Ansible, but variables from the external inventory script will override any that have the same name.

So, with the template above (

```
motd.j2
```

), this results in the following data being written to

```
/etc/motd
```

for system ?foo?:

```
Welcome, I am templated with a value of a=2, b=3, and c=4
```

And on system ?bar? (bar.example.com):

```
Welcome, I am templated with a value of a=2, b=3, and c=5
```

And technically, though there is no major good reason to do it, this also works:

```
ansible
```

```
webserver
```

```
-m
```

```
ansible.builtin.shell
```

```
-a
```

```
"echo {{ a }}"
```

So, in other words, you can use those variables in arguments/actions as well.

Other inventory scripts

?

In Ansible 2.10 and later, inventory scripts moved to their associated collections. Many are now in the `ansible-community/contrib-scripts` repository

. We recommend you use

Inventory plugins

instead.

Using inventory directories and multiple inventory sources

?

If the location given to

```
-i
```

in Ansible is a directory (or as so configured in

```
ansible.cfg
```

), Ansible can use multiple inventory sources

at the same time. When doing so, it is possible to mix both dynamic and statically managed inventory sources in the

same ansible run. Instant
hybrid cloud!

In an inventory directory, executable files are treated as dynamic inventory sources and most other files as static sources. Files which end with any of the following are ignored:

~, .orig, .bak, .ini, .cfg, .retry, .pyc, .pyo

You can replace this list with your own selection by configuring an
inventory_ignore_extensions
list in

ansible.cfg

, or setting the

ANSIBLE_INVENTORY_IGNORE

environment variable. The value in either case must be a comma-separated list of patterns, as shown above.

Any

group_vars

and

host_vars

subdirectories in an inventory directory are interpreted as expected, making inventory directories a powerful way to
organize different sets of configurations. See

Passing multiple inventory sources

for more information.

Static groups of dynamic groups

?

When defining groups of groups in the static inventory file, the child groups
must also be defined in the static inventory file, otherwise ansible returns an
error. If you want to define a static group of dynamic child groups, define
the dynamic groups as empty in the static inventory file. For example:

[tag_Name_staging_foo]

[tag_Name_staging_bar]

[staging:children]

tag_Name_staging_foo

tag_Name_staging_bar

See also

How to build your inventory

All about static inventory files

Communication

Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/inventory_guide/intro_inventory.html

How to build your inventory ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

How to build your inventory

Inventory basics: formats, hosts, and groups

Default groups

Hosts in multiple groups

Grouping groups: parent/child group relationships

Adding ranges of hosts

Passing multiple inventory sources

Organizing inventory in a directory

Managing inventory load order

Adding variables to inventory

Assigning a variable to one machine: host variables

Inventory aliases

Defining variables in INI format

Assigning a variable to many machines: group variables

Inheriting variable values: group variables for groups of groups

Organizing host and group variables

How variables are merged

Managing inventory variable load order

Connecting to hosts: behavioral inventory parameters

Non-SSH connection types

Inventory setup examples

Example: One inventory per environment

Example: Group by function

Example: Group by location

Working with dynamic inventory

Patterns: targeting hosts and groups

Connection methods and details

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks
Contributing to Ansible
Ansible Community Guide
Ansible Collections Contributor Guide
ansible-core Contributors Guide
Advanced Contributor Guide
Ansible documentation style guide
Extending Ansible
Developer Guide
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Building Ansible inventories
How to build your inventory
Edit on GitHub
How to build your inventory
?

Ansible automates tasks on managed nodes or `hosts` in your infrastructure by using a list or group of lists known as inventory. Ansible composes its inventory from one or more `inventory sources`. While one of these sources can be the list of host names you pass at the command line, most Ansible users create inventory files. Your inventory defines the managed nodes you automate and the variables associated with those hosts. You can also specify groups. Groups allow you to reference multiple associated hosts to target for your automation or to define variables in bulk.

Once you define your inventory, you use
patterns

to select the hosts or groups you want Ansible to run against.

The simplest inventory is a single file that contains a list of hosts and groups. The default location for this file is

/etc/ansible/hosts

. You can specify a different inventory source or sources at the command line by using the

-i

<path

or

expression>

option or by using the configuration system.

Ansible

Inventory plugins

supports a range of formats and sources, which makes your inventory flexible and customizable. As your inventory expands, you might need more than a single file to organize your hosts and groups. You have the following common options beyond the

/etc/ansible/hosts

file:

You can generate an inventory dynamically. For example, you can use an inventory plugin to list resources in one or more cloud providers or other sources. See

Working with dynamic inventory

.

You can use multiple sources for inventory, including both dynamic inventory and static files. See

Passing multiple inventory sources

.

You can create a directory with multiple inventory sources, static or dynamic. See

Organizing inventory in a directory

.

Inventory basics: formats, hosts, and groups

Default groups

Hosts in multiple groups

Grouping groups: parent/child group relationships

Adding ranges of hosts

Passing multiple inventory sources

Organizing inventory in a directory

Managing inventory load order

Adding variables to inventory

Assigning a variable to one machine: host variables

Inventory aliases

Defining variables in INI format

Assigning a variable to many machines: group variables

Inheriting variable values: group variables for groups of groups

Organizing host and group variables

How variables are merged

Managing inventory variable load order

Connecting to hosts: behavioral inventory parameters

Non-SSH connection types

Inventory setup examples

Example: One inventory per environment

Example: Group by function

Example: Group by location

The following YAML snippets include an ellipsis (?) to indicate that the snippets are part of a larger YAML file. You can find out more about YAML syntax at

YAML Basics

.

Inventory basics: formats, hosts, and groups

?

You can create your inventory file in one of many formats, depending on the inventory plugins you have.

The most common formats are INI and YAML because Ansible includes built-in support for them. This introduction focuses on these two formats, but many other formats and sources are possible.

A basic INI

/etc/ansible/hosts

might look like this:

mail.example.com

[webservers]

foo.example.com

bar.example.com

[dbservers]

one.example.com

two.example.com

three.example.com

The headings in brackets are group names. You can use group names to classify hosts and to decide which hosts you are controlling at what times and for what purpose.

Group names should follow the same guidelines as

Creating valid variable names

.

Here's the same basic inventory file in YAML format:

ungrouped

:

hosts

:

mail.example.com

:

webservers

:

hosts

:

foo.example.com

:

bar.example.com

:

dbservers

:

hosts

:

one.example.com

:

two.example.com

:

three.example.com

:

Default groups

?

Even if you do not define any groups in your inventory, Ansible creates two default groups:

all

and

ungrouped

. The

all
group contains every host. The
ungrouped
group contains all hosts that do not belong to any other group.
Every host always belongs to at least two groups (
all
and
ungrouped
, or
all
and another group). For example, in the basic inventory above, the host
mail.example.com
belongs to the
all
and
ungrouped
groups. The host
two.example.com
belongs to the
all
and
dbservers
groups. Although
all
and
ungrouped
are always present, they can be implicit and might not appear in group listings like
group_names

.
Hosts in multiple groups
?

You can put a host in more than one group. For example, you can include a production web server in a data center in
Atlanta in the
[prod]

,
[atlanta]
, and
[webservers]

groups. You can create groups that track the following criteria:

What

- An application, stack, or microservice (for example, database servers, web servers, and so on).

Where

- A datacenter or region, to talk to local DNS, storage, and so on (for example, east, west).

When

- The development stage, to avoid testing on production resources (for example, prod, test).

The following example extends the previous YAML inventory to include what, when, and where:

ungrouped

```
:
hosts
:
mail.example.com
:
webservers
:
```

```
hosts
:
foo.example.com
:
bar.example.com
:
dbservers
:
hosts
:
one.example.com
:
two.example.com
:
three.example.com
:
east
:
hosts
:
foo.example.com
:
one.example.com
:
two.example.com
:
west
:
hosts
:
bar.example.com
:
three.example.com
:
prod
:
hosts
:
foo.example.com
:
one.example.com
:
two.example.com
:
test
:
hosts
:
bar.example.com
:
three.example.com
:
As the example shows,
one.example.com
```

exists in the
dbservers

,
east
, and
prod
groups.

Grouping groups: parent/child group relationships

?

You can create parent/child relationships among groups. Parent groups are also known as nested groups or groups of groups. For example, if all your production hosts are already in groups such as

atlanta_prod

and

denver_prod

, you can create a

production

group that includes those smaller groups. This approach reduces maintenance because you add or remove hosts from the parent group by editing the child groups.

To create parent/child relationships for groups, use one of the following methods:

In INI format, use the

:children

suffix.

In YAML format, use the

children:

entry.

The following example shows the same inventory as above, simplified with parent groups for the

prod

and

test

groups:

ungrouped

:

hosts

:

mail.example.com

:

webservers

:

hosts

:

foo.example.com

:

bar.example.com

:

dbservers

:

hosts

:

one.example.com

:

two.example.com

:

three.example.com

:

```

east
:
hosts
:
foo.example.com
:
one.example.com
:
two.example.com
:
west
:
hosts
:
bar.example.com
:
three.example.com
:
prod
:
children
:
east
:
test
:
children
:
west
:

```

Note the following properties of child groups:

Any host that is a member of a child group is automatically a member of the parent group.

A group can have multiple parents and children, but not circular relationships.

A host can be in multiple groups, but Ansible processes only

one

instance of the host at runtime. Ansible merges the data from multiple groups.

Hosts and groups are always *global*. If you define a host or group more than once under different *branches* or *instances*, the host or group remains the same entity. Defining a host or group more than once either adds new information to it or overwrites any conflicting information with the latest definition.

Adding ranges of hosts

?

Some plugins, like YAML and INI, support adding ranges of hosts. If you have many hosts with a similar pattern, you can add the hosts as a range rather than listing each hostname separately:

In INI:

```

[webservers]
www[01:50].example.com

```

In YAML:

```

# ...
webservers
:
hosts
:
www[01:50].example.com
:

```


You can specify a stride (increments between sequence numbers) when you define a numeric range of hosts:

In INI:

```
[webservers]
```

```
www[01:50:2].example.com
```

In YAML:

```
# ...
```

```
webservers
```

```
:
```

```
hosts
```

```
:
```

```
www[01:50:2].example.com
```

```
:
```

The example above matches the subdomains `www01`, `www03`, `www05`, `?`, `www49`, but not `www00`, `www02`, `www50`, and so on, because the stride (increment) is 2 units for each step.

For numeric patterns, you can include or remove leading zeros as desired. Ranges are inclusive. You can also define alphabetic ranges:

```
[databases]
```

```
db-[a:f].example.com
```

Passing multiple inventory sources

?

You can target multiple inventory sources (static files, directories, dynamic inventory scripts or anything supported by inventory plugins) at the same time. To do this, specify multiple inventory sources from the command

line (see below) or by configuration, either by setting

```
ANSIBLE_INVENTORY
```

or in

```
ansible.cfg
```

```
(
```

```
DEFAULT_HOST_LIST
```

```
).
```

This capability can be useful when you want to target normally separate environments, like staging and production, at the same time for a specific action.

To target two inventory sources from the command line:

```
ansible-playbook
```

```
get_logs.yml
```

```
-i
```

```
staging
```

```
-i
```

```
production
```

Organizing inventory in a directory

?

You can consolidate multiple inventory sources in a single directory. The simplest version of this approach is a directory with multiple files instead of a single inventory file. Maintaining a single file becomes difficult when the file gets too long. If you have multiple teams and multiple automation projects, creating one inventory file per team or project lets everyone easily find the hosts and groups that matter to them. You can also still use the files individually or in subsets, depending on how you configure or call Ansible.

These files can use all formats or plugin configurations (for example, YAML or INI). In this case, your directory becomes your ?single? inventory source, and Ansible aggregates the multiple sources it finds in that directory. By default, Ansible ignores some directories and extensions, but you can change this behavior in the configuration (

```
INVENTORY_IGNORE_PATTERNS
```

and

```
INVENTORY_IGNORE_EXTS
```

```
).
```

You can also combine multiple inventory source types in an inventory directory. This method can be useful for

combining static and dynamic hosts and managing them as one inventory.

The following inventory directory combines an inventory plugin source, a dynamic inventory script, and a file with static hosts:

inventory/

```
openstack.yml      # configure inventory plugin to get hosts from OpenStack cloud
dynamic-inventory.py # add additional hosts with dynamic inventory script
on-prem           # add static hosts and groups
parent-groups     # add static hosts and groups
```

You can target this inventory directory as follows:

ansible-playbook

example.yml

-i

inventory

You can also configure the inventory directory in your

ansible.cfg

file. See

Configuring Ansible

for more details.

Ansible reads and loads files from the top directory down in alphabetically sorted order.

Managing inventory load order

?

Ansible loads inventory sources in the order you supply them. It defines hosts, groups, and variables as it encounters them in the source files, adding the

all

and

ungrouped

groups at the end if needed.

Depending on the inventory plugin or plugins you use, you might need to rearrange the order of sources to ensure that parent/child-defined groups or hosts exist as the plugins expect. Otherwise, you might encounter a parsing error. For example, the YAML and INI inventory plugins discard empty groups (groups with no associated hosts) when they finish processing each source.

If you define a variable multiple times, Ansible overwrites the previous value. The last definition wins.

Adding variables to inventory

?

You can define variables that relate to a specific host or group in your inventory. A simple way to start is by adding variables directly to the hosts and groups in a YAML or INI inventory source.

This guide documents how to add variables in the inventory source for simplicity. However, you can also use

Vars plugins

to add variables from many other sources. By default, Ansible ships with the

host_group_vars

plugin, which allows you to define variables in separate host and group variable files. Using separate files is a more robust approach to describing your system policy than defining variables in the inventory source. See

Organizing host and group variables

for guidelines on how to store variable values in individual files in the ?host_vars? and ?group_vars? directories.

Assigning a variable to one machine: host variables

?

You can easily assign a variable to a single host and then use that variable later in playbooks. You can do this directly in your inventory file.

In INI:

[atlanta]

host1 http_port=80 maxRequestsPerChild=808

host2 http_port=303 maxRequestsPerChild=909

In YAML:

atlanta

```

:
hosts
:
host1
:
http_port
:
80
maxRequestsPerChild
:
808
host2
:
http_port
:
303
maxRequestsPerChild
:
909

```

Unique values like non-standard SSH ports work well as host variables. You can add them to your Ansible inventory by adding the port number after the hostname with a colon:

```
badwolf.example.com:5309
```

You can use host variables to define ?Connection variables?. Connection variables configure connection

```

,
shell
, and
become

```

plugins to enable task execution on the host. For example:
[targets]

```

localhost      ansible_connection=local
other1.example.com  ansible_connection=ssh    ansible_user=myuser
other2.example.com  ansible_connection=ssh    ansible_user=myotheruser

```

Inventory aliases

?

The

inventory_hostname

is the unique identifier for a host in Ansible. This identifier can be an IP address or a hostname, but it can also be just an ?alias? or short name for the host.

In INI:

```
jumper ansible_port=5555 ansible_host=192.0.2.50
```

In YAML:

```

# ...
hosts
:
jumper
:
ansible_port
:
5555
ansible_host
:
192.0.2.50

```

In this example, running Ansible against the host alias `?jumper?` connects to 192.0.2.50 on port 5555. See [behavioral inventory parameters](#)

to further customize the connection to hosts.

This feature is also useful for targeting the same host more than once, but remember that tasks can run in parallel:

In INI:

```
jumper1 ansible_port=5555 ansible_host=192.0.2.50
```

```
jumper2 ansible_port=5555 ansible_host=192.0.2.50
```

In YAML:

```
# ...
```

```
hosts
```

```
:
```

```
jumper1
```

```
:
```

```
ansible_port
```

```
:
```

```
5555
```

```
ansible_host
```

```
:
```

```
192.0.2.50
```

```
jumper2
```

```
:
```

```
ansible_port
```

```
:
```

```
5555
```

```
ansible_host
```

```
:
```

```
192.0.2.50
```

Defining variables in INI format

?

Ansible interprets values that you pass in the INI format by using the

`key=value`

syntax differently depending on where you declare them:

When you declare a value inline with the host, Ansible interprets the INI value as a Python literal structure (for example, a string, number, tuple, list, dict, boolean, or None). Host lines accept multiple

`key=value`

parameters per line. Therefore, you need a way to indicate that a space is part of a value rather than a separator. You can quote values that contain whitespace (using single or double quotes). See the

[Python shlex parsing rules](#)

for details.

When you declare a value in a

`:vars`

section, Ansible interprets the INI value as a string. For example,

`var=FALSE`

creates a string with the value `?FALSE?`. Unlike host lines,

`:vars`

sections accept only a single entry per line, so everything after the

`=`

becomes the value for the entry.

If you need a variable from an INI inventory to have a certain type (for example, a string or a boolean), always specify the type with a filter in your task. Do not rely on types that you set in INI inventories when you consume variables.

Consider using the YAML format for inventory sources to avoid confusion about the actual type of a variable. The YAML inventory plugin processes variable values consistently and correctly.

Assigning a variable to many machines: group variables

?

If all hosts in a group share a variable value, you can apply that variable to an entire group at once.

In INI:

```
[atlanta]
```

```
host1
```

```
host2
```

```
[atlanta:vars]
```

```
ntp_server=ntp.atlanta.example.com
```

```
proxy=proxy.atlanta.example.com
```

In YAML:

```
atlanta
```

```
:
```

```
hosts
```

```
:
```

```
host1
```

```
:
```

```
host2
```

```
:
```

```
vars
```

```
:
```

```
ntp_server
```

```
:
```

```
ntp.atlanta.example.com
```

```
proxy
```

```
:
```

```
proxy.atlanta.example.com
```

Group variables are a convenient way to apply variables to multiple hosts at once. Before executing, however, Ansible always flattens variables, including inventory variables, to the host level. If a host is a member of multiple groups, Ansible reads variable values from all of those groups. If you assign different values to the same variable in different groups, Ansible chooses which value to use based on internal rules for merging

.

Inheriting variable values: group variables for groups of groups

?

You can apply variables to parent groups (nested groups or groups of groups) as well as to child groups. The syntax is the same:

```
:vars
```

for INI format and

```
vars:
```

for YAML format:

In INI:

```
[atlanta]
```

```
host1
```

```
host2
```

```
[raleigh]
```

```
host2
```

```
host3
```

```
[southeast:children]
```

```
atlanta
```

```
raleigh
```

```
[southeast:vars]
```

some_server=foo.southeast.example.com
halon_system_timeout=30
self_destruct_countdown=60
escape_pods=2

[usa:children]

southeast

northeast

southwest

northwest

In YAML:

usa

:

children

:

southeast

:

children

:

atlanta

:

hosts

:

host1

:

host2

:

raleigh

:

hosts

:

host2

:

host3

:

vars

:

some_server

:

foo.southeast.example.com

halon_system_timeout

:

30

self_destruct_countdown

:

60

escape_pods

:

2

northeast

:

northwest

:

southwest

:

A child group's variables have higher precedence (they override) than a parent group's variables.

Organizing host and group variables

?

Although you can define variables in the inventory source, you can also use

Vars plugins

to define alternate sources for your variables.

The default vars plugin that Ansible ships with,

host_group_vars

, lets you use separate host and group variable files. This method helps you organize your variable values more easily.

You can also use lists and hash data in these files, which you cannot do in your main inventory file.

For the

host_group_vars

plugin, your host and group variable files must use YAML syntax. Valid file extensions are ?.yaml?, ?.yml?, ?.json?, or no file extension. See

YAML Syntax

if you are new to YAML.

The

host_group_vars

plugin loads host and group variable files by searching paths relative to the inventory source or the playbook file. If your inventory file at

/etc/ansible/hosts

contains a host named ?foosball? that belongs to the

raleigh

and

webservers

groups, that host will use variables from the YAML files in the following locations:

/etc/ansible/group_vars/raleigh

can optionally end in '.yaml', '.yml', or '.json'

/etc/ansible/group_vars/webservers

/etc/ansible/host_vars/foosball

For example, if you group hosts in your inventory by datacenter, and each datacenter uses its own NTP server and database server, you can create a file named

/etc/ansible/group_vars/raleigh

to store the variables for the

raleigh

group:

ntp_server

:

acme.example.org

database_server

:

storage.example.org

You can also create

directories

named after your groups or hosts. Ansible reads all the files in these directories in lexicographical order. Here is an example with the ?raleigh? group:

/etc/ansible/group_vars/raleigh/db_settings

/etc/ansible/group_vars/raleigh/cluster_settings

All hosts in the ?raleigh? group have the variables that you define in these files

available to them. This method is very useful for keeping your variables organized when a single

file gets too big, or when you want to use

Ansible Vault

on some group variables.

Ansible?s

host_group_vars

vars plugin can also add

group_vars/

and

host_vars/

directories to your playbook directory when you use

ansible-playbook

. However, not all Ansible commands have a playbook (for example,

ansible

or

ansible-console

). For those commands, you can use the

--playbook-dir

option to provide the directory on the command line.

If you have sources for the vars plugins relative to both the playbook directory and the inventory directory, the variables that Ansible sources relative to the playbook override the variables that it sources relative to the inventory source.

To track changes to your inventory and variable definitions, keep your inventory sources and their relative variable directories and files in a Git repository or other version control system.

How variables are merged

?

Note

Ansible merges variables from different sources and applies precedence to some variables over others according to a set of rules. For example, variables that occur higher in an inventory can override variables that occur lower in the inventory. See

Variable precedence: where should I put a variable?

for more information.

Before it runs a play, Ansible merges and flattens variables to the specific host. This process keeps Ansible focused on the Host and Task, so groups do not survive outside of inventory and host matching. By default, Ansible overwrites variables, including the ones that you define for a group or host (see

DEFAULT_HASH_BEHAVIOUR

). The order/precedence for inventory entities is (from lowest to highest):

The following list shows the order of precedence for inventory entities, from lowest to highest:

all

group (because it is the ?parent? of all other groups)

parent group

child group

host

By default, Ansible merges groups at the same parent/child level in alphabetical order. Variables from the last group that Ansible loads overwrite variables from the previous groups. For example, Ansible merges an

a_group

with a

b_group

, and matching variables from

b_group

overwrite the variables in

a_group

.

You can fine-tune this merge behavior by setting the group variable

ansible_group_priority

. This variable overrides the alphabetical sorting for the merge order for groups of the same level (after Ansible resolves the parent/child order). The larger the number, the later Ansible merges the group, giving it higher priority. This variable defaults to


```

1
if you do not set it. For example:
a_group
:
vars
:
testvar
:
a
ansible_group_priority
:
10
b_group
:
vars
:
testvar
:
b

```

In this example, if both groups have the same priority, the result would normally be

```
testvar
```

```

==
b
. However, because we give
a_group
a higher priority, the result is
testvar
==
a
.

```

You can set

```
ansible_group_priority
```

only in an inventory source, not in

```
group_vars/
```

. Ansible uses this variable when it loads the

```
group_vars/
```

directory.

Managing inventory variable load order

?

This section describes how to control variable precedence by managing the load order of inventory sources. You can pass sources in a specific order at the command line or use prefixes in the filenames of sources within a directory. When you use multiple inventory sources, remember that Ansible resolves any variable conflicts according to the rules described in

How variables are merged

and

Variable precedence: where should I put a variable?

. You can control the merging order of variables in inventory sources to get the variable value you need.

When you pass multiple inventory sources at the command line, Ansible merges variables in the order you pass those parameters. If the

```
[all:vars]
```

section in the staging inventory defines

```
myvar
```

```

=
1

```

and the production inventory defines

myvar

=

2

, then the following outcomes are true:

If you pass

-i

staging

-i

production

, Ansible runs the playbook with

myvar

=

2

.

If you pass

-i

production

-i

staging

, Ansible runs the playbook with

myvar

=

1

.

When you put multiple inventory sources in a directory, Ansible merges the sources in alphabetical order according to their filenames. You can control the load order by adding prefixes to the files:

inventory/

01-openstack.yml # configure inventory plugin to get hosts from Openstack cloud

02-dynamic-inventory.py # add additional hosts with dynamic inventory script

03-static-inventory # add static hosts

group_vars/

all.yml # assign variables to all hosts

If

01-openstack.yml

defines

myvar

=

1

for the group

all

,

02-dynamic-inventory.py

defines

myvar

=

2

,

and

03-static-inventory

defines

myvar

=

3

, Ansible runs the playbook with
myvar
=
3

.
For more details on inventory plugins and dynamic inventory scripts see
Inventory plugins
and
Working with dynamic inventory

.
Connecting to hosts: behavioral inventory parameters
?

As described above, you can set the following variables to control how Ansible interacts with remote hosts.
Host connection:

Note
Ansible does not expose a channel to allow communication between the user and the ssh process to accept a password manually to decrypt an ssh key when using the ssh connection plugin (which is the default). The use of
ssh-agent
is highly recommended.

ansible_connection
Specifies the connection type to the host. This can be the name of any Ansible connection plugin. SSH protocol types
are
ssh
or
paramiko
. The default is
ssh

.
General for all connections:
ansible_host
Specifies the resolvable name or IP of the host to connect to, if it is different from the alias you wish to give to it. Never
set it to depend on
inventory_hostname
. If you really need something like that, use
inventory_hostname_short
so it can work with delegation.

ansible_port
The connection port number, if not the default (22 for ssh).
ansible_user
The username to use when connecting (logging in) to the host.
ansible_password
The password to use to authenticate to the host. (Never store this variable in plain text. Always use a vault. See
Keep vaulted variables safely visible

.)
Specific to the SSH connection plugin:
ansible_ssh_private_key_file
Private key file used by SSH. This is useful if you use multiple keys and you do not want to use SSH agent.
ansible_ssh_common_args
Ansible always appends this setting to the default command line for
sftp

,
scp
,
and

ssh

. This is useful for configuring a

ProxyCommand

for a certain host or

group.

ansible_sftp_extra_args

Ansible always appends this setting to the default

sftp

command line.

ansible_scp_extra_args

Ansible always appends this setting to the default

scp

command line.

ansible_ssh_extra_args

Ansible always appends this setting to the default

ssh

command line.

ansible_ssh_pipelining

Specifies whether to use SSH pipelining. This can override the

pipelining

setting in

ansible.cfg

.

ansible_ssh_executable (added in version 2.2)

This setting overrides the default behavior to use the system

ssh

. It can override the

ssh_executable

setting in the

ssh_connection

section of

ansible.cfg

.

Privilege escalation (see

Ansible Privilege Escalation

for further details):

ansible_become

Equivalent to

ansible_sudo

or

ansible_su

; allows you to force privilege escalation.

ansible_become_method

Allows you to set the privilege escalation method to a matching become plugin.

ansible_become_user

Equivalent to

ansible_sudo_user

or

ansible_su_user

; allows you to set the user you become through privilege escalation.

ansible_become_password

Equivalent to

ansible_sudo_password

or

`ansible_su_password`

; allows you to set the privilege escalation password. (Never store this variable in plain text. Always use a vault. See [Keep vaulted variables safely visible](#))

.)

`ansible_become_exe`

Equivalent to

`ansible_sudo_exe`

or

`ansible_su_exe`

; allows you to set the executable for the escalation method you selected.

`ansible_become_flags`

Equivalent to

`ansible_sudo_flags`

or

`ansible_su_flags`

; allows you to set the flags passed to the selected escalation method. You can also set this globally in

`ansible.cfg`

in the

`become_flags`

option under

`privilege_escalation`

.

Remote host environment parameters:

`ansible_shell_type`

Specifies the shell type of the target system. You should not use this setting unless you have set the

`ansible_shell_executable`

to a non-Bourne (sh) compatible shell. By default, Ansible

formats commands using

sh

-style syntax. If you set this to

csh

or

fish

, commands

that Ansible executes on target systems follow those shell's syntax instead.

`ansible_python_interpreter`

Specifies the target host Python path. This is useful for systems with more

than one Python or for systems where Python is not located at

`/usr/bin/python`

, such as *BSD, or where

`/usr/bin/python`

is not a 2.X series Python. We do not use the

`/usr/bin/env`

mechanism because that requires the remote user's

path to be set correctly and also assumes the

python

executable is named python, where the executable might

be named something like

`python2.6`

.

`ansible_*_interpreter`

Works for any language, such as Ruby or Perl, and works just like

`ansible_python_interpreter`

.

This variable replaces the shebang of modules that will run on that host.

New in version 2.1.

`ansible_shell_executable`

This setting sets the shell the Ansible control node will use on the target machine.

It overrides

`executable`

in

`ansible.cfg`

, which defaults to

`/bin/sh`

. You should only change this value if it is not possible to use

`/bin/sh`

(in other words, if

`/bin/sh`

is not installed on the target

machine or cannot be run from `sudo`).

Examples from an Ansible-INI host file:

`some_host ansible_port=2222 ansible_user=manager`

`aws_host ansible_ssh_private_key_file=/home/example/.ssh/aws.pem`

`freebsd_host ansible_python_interpreter=/usr/local/bin/python`

`ruby_module_host ansible_ruby_interpreter=/usr/bin/ruby.1.9.3`

Non-SSH connection types

?

As stated in the previous section, Ansible executes playbooks over SSH by default, but it is not limited to this connection type.

You can change the connection type with the host-specific parameter

`ansible_connection=<connection`

`plugin`

`name>`

.

For a full list of available plugins and examples, see

Plugin list

.

Inventory setup examples

?

See also

Sample Ansible setup

, which shows inventory along with playbooks and other Ansible artifacts.

Example: One inventory per environment

?

If you need to manage multiple environments, consider defining only the hosts of a single environment in each inventory. This way, it is harder to, for example, accidentally change the state of nodes inside the `?test?` environment when you wanted to update some `?staging?` servers.

For the example mentioned above, you could have an

`inventory_test`

file:

`[dbservers]`

`db01.test.example.com`

`db02.test.example.com`

`[appservers]`

`app01.test.example.com`

app02.test.example.com

app03.test.example.com

That file only includes hosts that are part of the ?test? environment. You can define the ?staging? machines in another file called

inventory_staging

:

[dbservers]

db01.staging.example.com

db02.staging.example.com

[appservers]

app01.staging.example.com

app02.staging.example.com

app03.staging.example.com

To apply a playbook called

site.yml

to all the app servers in the test environment, use the following command:

ansible-playbook

-i

inventory_test

-l

appservers

site.yml

Example: Group by function

?

In the previous section, you already saw an example of using groups to cluster hosts that have the same function. This approach allows you, for example, to define firewall rules inside a playbook or role that affect only database servers:

-

hosts

:

dbservers

tasks

:

-

name

:

Allow access from 10.0.0.1

ansible.builtin.iptables

:

chain

:

INPUT

jump

:

ACCEPT

source

:

10.0.0.1

Example: Group by location

?

Other tasks might focus on where a certain host is located. Let's

say that
db01.test.example.com
and
app01.test.example.com
are
located in DC1, while
db02.test.example.com
is in DC2:

```
[dc1]
db01.test.example.com
app01.test.example.com
[dc2]
db02.test.example.com
```

In practice, you might end up mixing all these setups. For example, you might need to update all nodes in a specific data center on one day, while on another day, you might need to update all the application servers no matter their location.

See also

[Inventory plugins](#)

[Pulling inventory from dynamic or static sources](#)

[Working with dynamic inventory](#)

[Pulling inventory from dynamic sources, such as cloud providers](#)

[Introduction to ad hoc commands](#)

[Examples of basic commands](#)

[Working with playbooks](#)

[Learning Ansible's configuration, deployment, and orchestration language.](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/module_plugin_guide/modules_intro.html

Introduction to modules ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Introduction to modules

Boolean variables

Module maintenance and support

Rejecting modules

Working with plugins

Modules and plugins index

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Using Ansible modules and plugins
Introduction to modules
Edit on GitHub
Introduction to modules
?

Modules (also referred to as ?task plugins? or ?library plugins?) are discrete units of code that can be used from the command line or in a playbook task. Ansible executes each module, usually on the remote managed node, and collects return values. In Ansible 2.10 and later, most modules are hosted in collections.

You can execute modules from the command line.

```
ansible webservers -m service -a "name=httpd state=started"
```

```
ansible webservers -m ping
```

```
ansible webservers -m command -a "/sbin/reboot -t now"
```

Each module supports arguments. Nearly all modules take

key=value

arguments, space delimited. Some modules take no arguments, and the command/shell modules simply take the string of the command you want to run.

From playbooks, Ansible modules are executed in a very similar way.

```
-  
name  
:  
reboot the servers  
command  
:  
/sbin/reboot -t now
```

Another way to pass arguments to a module is using YAML syntax, also called ?complex args?.

```
-  
name  
:  
restart webserver  
service  
:  
name  
:
```

httpd
state
:

restarted

All modules return JSON format data. This means modules can be written in any programming language. Modules should be idempotent, and should avoid making any changes if they detect that the current state matches the desired final state. When used in an Ansible playbook, modules can trigger `change` events in the form of notifying handlers

to run additional tasks.

You can access the documentation for each module from the command line with the `ansible-doc` tool.

`ansible-doc yum`

For a list of all available modules, see the

Collection docs

, or run the following at a command prompt.

`ansible-doc -l`

Boolean variables

?

Ansible accepts a broad range of values for

`bool`

in module arguments:

`true/false`

,
`1/0`

,
`yes/no`

,
`True/False`

and so on. The matching of valid strings is case insensitive.

While documentation examples focus on

`true/false`

to be compatible with

`ansible-lint`

default settings, you can use any of the following:

Valid values

Description

True

,
'true'

,
't'

,
'yes'

,
'y'

,
'on'

,
'1'

,
1

,
1.0

Truthy values

False

,
'false'

,
'f'

,
'no'

,
'n'

,
'off'

,
'0'

,
0

,
0.0

Falsy values

See also

[Introduction to ad hoc commands](#)

[Examples of using modules in /usr/bin/ansible](#)

[Working with playbooks](#)

[Examples of using modules with /usr/bin/ansible-playbook](#)

[Should you develop a module?](#)

[How to write your own modules](#)

[Python API](#)

[Examples of using modules with the Python API](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Indexes of all modules and plugins](#)

[All modules and plugins available](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/playbook_guide/playbooks_conditionals.html

Conditionals ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Ansible playbooks

Working with playbooks

Templating (Jinja2)

Using filters to manipulate data

Tests

Lookups

Python3 in templates

The now function: get the current time

The undef function: add hint for undefined variables

Loops

Controlling where tasks run: delegation and local actions

Conditionals

Basic conditionals with

when

Debugging conditionals

Commonly-used facts

Blocks

Handlers: running operations on change

Error handling in playbooks

Setting the remote environment

Working with language-specific version managers

Reusing Ansible artifacts

Roles

Module defaults

Interactive input: prompts

Using variables

Discovering variables: facts and magic variables

Playbook Example: Continuous Delivery and Rolling Upgrades

Executing playbooks

Advanced playbook syntax

Manipulating data

Protecting sensitive data with Ansible vault
Using Ansible modules and plugins
Using Ansible collections
Using Ansible on Windows, BSD, and z/OS UNIX
Ansible tips and tricks
Contributing to Ansible
Ansible Community Guide
Ansible Collections Contributor Guide
ansible-core Contributors Guide
Advanced Contributor Guide
Ansible documentation style guide
Extending Ansible
Developer Guide
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Using Ansible playbooks
Working with playbooks
Conditionals
Edit on GitHub
Conditionals
?

In a playbook, you may want to execute different tasks or have different goals, depending on the value of a fact (data about the remote system), a variable, or the result of a previous task. You may want the value of some variables to

depend on the value of other variables. Or you may want to create additional groups of hosts based on whether the hosts match other criteria. You can do all of these things with conditionals.

Ansible uses Jinja2

tests
and
filters

in conditionals. Ansible supports all the standard tests and filters and adds some unique ones as well.

Note

There are many options to control execution flow in Ansible. You can find more examples of supported conditionals at <https://jinja.palletsprojects.com/en/latest/templates/#comparisons>

.

Basic conditionals with

when

Conditionals based on `ansible_facts`

Conditions based on registered variables

Conditionals based on variables

Using conditionals in loops

Loading custom facts

Conditionals with reuse

Conditionals with imports

Conditionals with includes

Conditionals with roles

Selecting variables, files, or templates based on facts

Selecting variables files based on facts

Selecting files and templates based on facts

Debugging conditionals

Commonly-used facts

`ansible_facts[?distribution?]`

`ansible_facts[?distribution_major_version?]`

`ansible_facts[?os_family?]`

Basic conditionals with

when

?

The simplest conditional statement applies to a single task. Create the task, then add a

when

statement that applies a test. The

when

clause is a raw Jinja2 expression without double curly braces (see

Referencing simple variables

). When you run the task or playbook, Ansible evaluates the test for all hosts. On any host where the test passes (returns a value of True), Ansible runs that task. For example, if you are installing mysql on multiple machines, some of which have SELinux enabled, you might have a task to configure SELinux to allow mysql to run. You would only want that task to run on machines that have SELinux enabled:

tasks

:

-

name

:

Configure SELinux to start mysql on any port

`ansible.posix.seboolean`

:

name

:

`mysql_connect_any`

```
state
:
true
persistent
:
true
when
:
```

```
ansible_selinux.status == "enabled"
```

all variables can be used directly in conditionals without double curly braces

Conditionals based on ansible_facts

?

Often you want to execute or skip a task based on facts. Facts are attributes of individual hosts, including IP address, operating system, the status of a filesystem, and many more. With conditionals based on facts:

You can install a certain package only when the operating system is a particular version.

You can skip configuring a firewall on hosts with internal IP addresses.

You can perform cleanup tasks only when a filesystem is getting full.

See

Commonly-used facts

for a list of facts that frequently appear in conditional statements. Not all facts exist for all hosts. For example, the `lsb_major_release` fact used in the example below only exists when the

`lsb_release`

package

is installed on the target host. To see what facts are available on your systems, add a debug task to your playbook:

```
-
name
:
Show facts available on the system
ansible.builtin.debug
:
var
:
ansible_facts
```

Here is a sample conditional based on a fact:

```
tasks
:
-
name
:
Shut down Debian flavored systems
ansible.builtin.command
:
/sbin/shutdown -t now
when
:
ansible_facts['os_family'] == "Debian"
```

If you have multiple conditions, you can group them with parentheses:

```
tasks
:
-
name
:
Shut down CentOS 6 and Debian 7 systems
ansible.builtin.command
```



```

:
/sbin/shutdown -t now
when
:
(ansible_facts['distribution'] == "CentOS" and ansible_facts['distribution_major_version'] == "6") or
(ansible_facts['distribution'] == "Debian" and ansible_facts['distribution_major_version'] == "7")
You can use
logical operators
to combine conditions. When you have multiple conditions that all need to be true (that is, a logical
and
), you can specify them as a list:
tasks

```

```

:
-
name
:
Shut down CentOS 6 systems
ansible.builtin.command

```

```

:
/sbin/shutdown -t now
when

```

```

:
-
ansible_facts['distribution'] == "CentOS"
-
ansible_facts['distribution_major_version'] == "6"
If a fact or variable is a string, and you need to run a mathematical comparison on it, use a filter to ensure that Ansible
reads the value as an integer:
tasks

```

```

:
-
ansible.builtin.shell
:
echo "only on Red Hat 6, derivatives, and later"
when

```

```

:
ansible_facts['os_family'] == "RedHat" and ansible_facts['lsb']['major_release'] | int >= 6
You can store Ansible facts as variables to use for conditional logic, as in the following example:
tasks

```

```

:
-
name
:
Get the CPU temperature
set_fact
:
temperature
:
"{{
ansible_facts['cpu_temperature']
}}"
-
name
:

```

```
Restart the system if the temperature is too high
when
:
temperature | float > 90
shell
:
```

Conditions based on registered variables

?

Often in a playbook, you want to execute or skip a task based on the outcome of an earlier task. For example, you might want to configure a service after it is upgraded by an earlier task. To create a conditional based on a registered variable:

Register the outcome of the earlier task as a variable.

Create a conditional test based on the registered variable.

You create the name of the registered variable using the

register

keyword. A registered variable always contains the status of the task that created it as well as any output that the task generated. You can use registered variables in templates and action lines as well as in conditional

when

statements. You can access the string contents of the registered variable using

variable.stdout

. For example:

```
-
name
:
Test play
hosts
:
all
tasks
:
-
name
:
Register a variable
ansible.builtin.shell
:
cat /etc/motd
register
:
motd_contents
```

-

name

:

Use the variable in conditional statement

ansible.builtin.shell

:

echo "motd contains the word hi"

when

:

motd_contents.stdout.find('hi') != -1

You can use registered results in the loop of a task if the variable is a list. If the variable is not a list, you can convert it into a list, with either

stdout_lines

or with

```
variable.stdout.split()
```

. You can also split the lines by other fields:

```
-
```

```
name
```

```
:
```

Registered variable usage as a loop list

```
hosts
```

```
:
```

```
all
```

```
tasks
```

```
:
```

```
-
```

```
name
```

```
:
```

Retrieve the list of home directories

```
ansible.builtin.command
```

```
:
```

```
ls /home
```

```
register
```

```
:
```

```
home_dirs
```

```
-
```

```
name
```

```
:
```

Add home dirs to the backup spooler

```
ansible.builtin.file
```

```
:
```

```
path
```

```
:
```

```
/mnt/bkspool/{{ item }}
```

```
src
```

```
:
```

```
/home/{{ item }}
```

```
state
```

```
:
```

```
link
```

```
loop
```

```
:
```

```
"{{
```

```
home_dirs.stdout_lines
```

```
}}"
```

```
# same as loop: "{{ home_dirs.stdout.split() }}"
```

The string content of a registered variable can be empty. If you want to run another task only on hosts where the stdout of your registered variable is empty, check the registered variable's string contents for emptiness:

```
-
```

```
name
```

```
:
```

check registered variable for emptiness

```
hosts
```

```
:
```

```
all
```

```
tasks
```

```
:
```

```
-
```

```

name
:
List contents of directory
ansible.builtin.command
:
ls mydir
register
:
contents
-
name
:
Check contents for emptiness
ansible.builtin.debug
:
msg
:
"Directory
is
empty"
when
:
contents.stdout == ""

```

Ansible always registers something in a registered variable for every host, even on hosts where a task fails or Ansible skips a task because a condition is not met. To run a follow-up task on these hosts, query the registered variable for

```

is
skipped
(not for ?undefined? or ?default?). See
Registering variables
for more information. Here are sample conditionals based on the success or failure of a task. Remember to ignore errors
if you want Ansible to continue executing on a host when a failure occurs:

```

```

tasks
:
-
name
:
Register a variable, ignore errors and continue
ansible.builtin.command
:
/bin/false
register
:
result
ignore_errors
:
true
-
name
:
Run only if the task that registered the "result" variable fails
ansible.builtin.command
:
/bin/something
when

```

```

:
result is failed
-
name
:
Run only if the task that registered the "result" variable succeeds
ansible.builtin.command
:
/bin/something_else
when
:
result is succeeded
-
name
:
Run only if the task that registered the "result" variable is skipped
ansible.builtin.command
:
/bin/still/something_else
when
:
result is skipped
-
name
:
Run only if the task that registered the "result" variable changed something.
ansible.builtin.command
:
/bin/still/something_else
when
:
result is changed

```

Note
Older versions of Ansible used

```

success
and
fail
, but
succeeded
and
failed

```

use the correct tense. All of these options are now valid.

Conditionals based on variables

?

You can also create conditionals based on variables defined in the playbooks or inventory. Because conditionals require boolean input (a test must evaluate as True to trigger the condition), you must apply the

|

bool

filter to non-boolean variables, such as string variables with content like ?yes?, ?on?, ?1?, or ?true?. You can define variables like this:

vars

:

epic

:

```
true
monumental
```

```
:
```

```
"yes"
```

With the variables above, Ansible would run one of these tasks and skip the other:

```
tasks
```

```
:
```

```
-
```

```
name
```

```
:
```

Run the command if "epic" or "monumental" is true

```
ansible.builtin.shell
```

```
:
```

```
echo "This certainly is epic!"
```

```
when
```

```
:
```

```
epic or monumental | bool
```

```
-
```

```
name
```

```
:
```

Run the command if "epic" is false

```
ansible.builtin.shell
```

```
:
```

```
echo "This certainly isn't epic!"
```

```
when
```

```
:
```

```
not epic
```

If a required variable has not been set, you can skip or fail using Jinja2's

defined

test. For example:

```
tasks
```

```
:
```

```
-
```

```
name
```

```
:
```

Run the command if "foo" is defined

```
ansible.builtin.shell
```

```
:
```

```
echo "I've got '{{ foo }}' and am not afraid to use it!"
```

```
when
```

```
:
```

```
foo is defined
```

```
-
```

```
name
```

```
:
```

Fail if "bar" is undefined

```
ansible.builtin.fail
```

```
:
```

```
msg="Bailing out. This play requires 'bar'"
```

```
when
```

```
:
```

```
bar is undefined
```

This is especially useful in combination with the conditional import of

vars

files (see below).

As the examples show, you do not need to use

```
{{ }}
```

to use variables inside conditionals, as these are already implied.

Using conditionals in loops

?

If you combine a

when

statement with a

loop

, Ansible processes the condition separately for each item. This is by design, so you can execute the task on some items in the loop and skip it on other items. For example:

tasks

```
:
```

```
-
```

name

```
:
```

Run with items greater than 5

ansible.builtin.command

```
:
```

```
echo {{ item }}
```

loop

```
:
```

```
[
```

```
0
```

```
,
```

```
2
```

```
,
```

```
4
```

```
,
```

```
6
```

```
,
```

```
8
```

```
,
```

```
10
```

```
]
```

when

```
:
```

```
item > 5
```

If you need to skip the whole task when the loop variable is undefined, use the

|default

filter to provide an empty iterator. For example, when looping over a list:

```
-
```

name

```
:
```

Skip the whole task when a loop variable is undefined

ansible.builtin.command

```
:
```

```
echo {{ item }}
```

loop

```
:
```

```
"{{
```

```
mylist|default([])
```

```
}}"
```

when

:

item > 5

You can do the same thing when looping over a dict:

-

name

:

The same as above using a dict

ansible.builtin.command

:

echo {{ item.key }}

loop

:

```
"{{
query('dict',
mydict|default({}))
}}"
```

when

:

item.value > 5

Loading custom facts

?

You can provide your own facts, as described in

Should you develop a module?

. To run them, just make a call to your own custom fact gathering module at the top of your list of tasks, and the variables returned there will be accessible for future tasks:

tasks

:

-

name

:

Gather site specific fact data

action

:

site_facts

-

name

:

Use a custom fact

ansible.builtin.command

:

/usr/bin/thingy

when

:

my_custom_fact_just_retrieved_from_the_remote_system == '1234'

Conditionals with reuse

?

You can use conditionals with reusable tasks files, playbooks, or roles. Ansible executes these conditional statements differently for dynamic reuse (includes) and static reuse (imports). See

Reusing Ansible artifacts

for more information on reuse in Ansible.

Conditionals with imports

?

When you add a conditional to an import statement, Ansible applies the condition to all tasks within the imported file.

This behavior is the equivalent of

Tag inheritance: adding tags to multiple tasks

. Ansible applies the condition to every task and evaluates each task separately. For example, if you want to define and then display a variable that was not previously defined, you might have a playbook called `main.yml`

and a tasks file called

`other_tasks.yml`

```
:
```

```
# all tasks within an imported file inherit the condition from the import statement
```

```
# main.yml
```

```
-
```

```
hosts
```

```
:
```

```
all
```

```
tasks
```

```
:
```

```
-
```

```
import_tasks
```

```
:
```

```
other_tasks.yml
```

```
# note "import"
```

```
when
```

```
:
```

```
x is not defined
```

```
# other_tasks.yml
```

```
-
```

```
name
```

```
:
```

```
Set a variable
```

```
ansible.builtin.set_fact
```

```
:
```

```
x
```

```
:
```

```
foo
```

```
-
```

```
name
```

```
:
```

```
Print a variable
```

```
ansible.builtin.debug
```

```
:
```

```
var
```

```
:
```

```
x
```

Ansible expands this at execution time to the equivalent of:

```
-
```

```
name
```

```
:
```

```
Set a variable if not defined
```

```
ansible.builtin.set_fact
```

```
:
```

```
x
```

```
:
```

```
foo
```

```
when
```

```

:
x is not defined
# this task sets a value for x
-
name
:
Do the task if "x" is not defined
ansible.builtin.debug
:
var
:
x
when
:
x is not defined
# Ansible skips this task, because x is now defined
If
x
is initially defined, both tasks are skipped as intended. But if
x
is initially undefined, the debug task will be skipped since the conditional is evaluated for every imported task. The
conditional will evaluate to
true
for the
set_fact
task, which will define the variable and cause the
debug
conditional to evaluate to
false
.
If this is not the behavior you want, use an
include_*
statement to apply a condition only to that statement itself.
# using a conditional on include_* only applies to the include task itself
# main.yml
-
hosts
:
all
tasks
:
-
include_tasks
:
other_tasks.yml
# note "include"
when
:
x is not defined
Now if
x
is initially undefined, the debug task will not be skipped because the conditional is evaluated at the time of the include
and does not apply to the individual tasks.
You can apply conditions to

```

import_playbook
as well as to the other
import_*

statements. When you use this approach, Ansible returns a ?skipped? message for every task on every host that does not match the criteria, creating repetitive output. In many cases the

group_by module
can be a more streamlined way to accomplish the same objective; see
Handling OS and distro differences

.
Conditionals with includes
?

When you use a conditional on an
include_*

statement, the condition is applied only to the include task itself and not to any other tasks within the included file(s). To contrast with the example used for conditionals on imports above, look at the same playbook and tasks file, but using an include instead of an import:

Includes let you reuse a file to define a variable when it is not already defined
main.yml

```
-
include_tasks
:
other_tasks.yml
when
:
x is not defined
# other_tasks.yml
-
name
:
Set a variable
ansible.builtin.set_fact
:
x
:
foo
-
name
:
Print a variable
ansible.builtin.debug
:
var
:
x
```

Ansible expands this at execution time to the equivalent of:

```
# main.yml
-
include_tasks
:
other_tasks.yml
when
:
x is not defined
# if condition is met, Ansible includes other_tasks.yml
```

```
# other_tasks.yml
```

```
-
```

```
name
```

```
:
```

```
Set a variable
```

```
ansible.builtin.set_fact
```

```
:
```

```
x
```

```
:
```

```
foo
```

```
# no condition applied to this task, Ansible sets the value of x to foo
```

```
-
```

```
name
```

```
:
```

```
Print a variable
```

```
ansible.builtin.debug
```

```
:
```

```
var
```

```
:
```

```
x
```

```
# no condition applied to this task, Ansible prints the debug statement
```

```
By using
```

```
include_tasks
```

```
instead of
```

```
import_tasks
```

```
, both tasks from
```

```
other_tasks.yml
```

```
will be executed as expected. For more information on the differences between
```

```
include
```

```
v
```

```
import
```

```
see
```

```
Reusing Ansible artifacts
```

```
.
```

```
Conditionals with roles
```

```
?
```

```
There are three ways to apply conditions to roles:
```

```
Add the same condition or conditions to all tasks in the role by placing your
```

```
when
```

```
statement under the
```

```
roles
```

```
keyword. See the example in this section.
```

```
Add the same condition or conditions to all tasks in the role by placing your
```

```
when
```

```
statement on a static
```

```
import_role
```

```
in your playbook.
```

```
Add a condition or conditions to individual tasks or blocks within the role itself. This is the only approach that allows you to select or skip some tasks within the role based on your
```

```
when
```

```
statement. To select or skip tasks within the role, you must have conditions set on individual tasks or blocks, use the dynamic
```

```
include_role
```

```
in your playbook, and add the condition or conditions to the include. When you use this approach, Ansible applies the
```

condition to the include itself plus any tasks in the role that also have that

when

statement.

When you incorporate a role in your playbook statically with the

roles

keyword, Ansible adds the conditions you define to all the tasks in the role. For example:

-

hosts

:

webservers

roles

:

-

role

:

debian_stock_config

when

:

ansible_facts['os_family'] == 'Debian'

Selecting variables, files, or templates based on facts

?

Sometimes the facts about a host determine the values you want to use for certain variables or even the file or template you want to select for that host. For example, the names of packages are different on CentOS and Debian. The configuration files for common services are also different on different OS flavors and versions. To load different variables files, templates, or other files based on a fact about the hosts:

name your vars files, templates, or files to match the Ansible fact that differentiates them

select the correct vars file, template, or file for each host with a variable based on that Ansible fact

Ansible separates variables from tasks, keeping your playbooks from turning into arbitrary code with nested conditionals. This approach results in more streamlined and auditable configuration rules because there are fewer decision points to track.

Selecting variables files based on facts

?

You can create a playbook that works on multiple platforms and OS versions with a minimum of syntax by placing your variable values in vars files and conditionally importing them. If you want to install Apache on some CentOS and some Debian servers, create variables files with YAML keys and values. For example:

for vars/RedHat.yml

apache

:

httpd

somethingelse

:

42

Then import those variables files based on the facts you gather on the hosts in your playbook:

-

hosts

:

webservers

remote_user

:

root

vars_files

:

```

-
"vars/common.yml"
-
[
"vars/{{
ansible_facts['os_family']
}}.yml"
,
"vars/os_defaults.yml"
]
tasks
:
-
name
:
Make sure apache is started
ansible.builtin.service
:
name
:
'{{
apache
}}'
state
:
started

```

Ansible gathers facts on the hosts in the webservers group, then interpolates the variable `?ansible_facts[?os_family?]?` into a list of file names. If you have hosts with Red Hat operating systems (CentOS, for example), Ansible looks for `?vars/RedHat.yml?`. If that file does not exist, Ansible attempts to load `?vars/os_defaults.yml?`. For Debian hosts, Ansible first looks for `?vars/Debian.yml?`, before falling back on `?vars/os_defaults.yml?`. If no files in the list are found, Ansible raises an error.

Selecting files and templates based on facts

?

You can use the same approach when different OS flavors or versions require different configuration files or templates. Select the appropriate file or template based on the variables assigned to each host. This approach is often much cleaner than putting a lot of conditionals into a single template to cover multiple OS or package versions.

For example, you can template out a configuration file that is very different between, say, CentOS and Debian:

```

-
name
:
Template a file
ansible.builtin.template
:
src
:
"{{
item
}}"
dest
:
/etc/myapp/foo.conf
loop
:
"{{

```

```

query('first_found',
{
'files':
myfiles,
'paths':
mypaths}))
}}"
vars
:
myfiles
:
-
"{{
ansible_facts['distribution']
}}.conf"
-
default.conf
mypaths
:
[
'search_location_one/somedir/'
,
'/opt/other_location/somedir/'
]

```

Debugging conditionals

?

If your conditional

when

statement is not behaving as you intended, you can add a

debug

statement to determine if the condition evaluates to

true

or

false

. A common cause of unexpected behavior in conditionals is testing an integer as a string or a string as an integer. To debug a conditional statement, add the entire statement as the

var:

value in a

debug

task. Ansible then shows the test and how the statement evaluates. For example, here is a set of tasks and sample output:

```

-
name
:
check value of return code
ansible.builtin.debug
:
var
:
bar_status.rc
-
name
:
check test for rc value as string

```

```

ansible.builtin.debug
:
var
:
bar_status.rc == "127"
-
name
:
check test for rc value as integer
ansible.builtin.debug
:
var
:
bar_status.rc == 127
TASK
[
check value of return code
]
*****

ok
:
[
foo-1
]
=>
{
"bar_status.rc"
:
"127"
}
TASK
[
check test for rc value as string
]
*****

ok
:
[
foo-1
]
=>
{
"bar_status.rc == \"127\""
:
false
}
TASK
[
check test for rc value as integer
]
*****

ok
:
[

```



```
foo-1
]
=>
{
"bar_status.rc == 127"
:
true
}
```

Commonly-used facts

?

The following Ansible facts are frequently used in conditionals.

`ansible_facts[?distribution?]`

?

Possible values (sample, not complete list):

Alpine
Altlinux
Amazon
Archlinux
ClearLinux
Coreos
CentOS
Debian
Fedora
Gentoo
Mandriva
NA
OpenWrt
OracleLinux
RedHat
Slackware
SLES
SMGL
SUSE
Ubuntu
VMwareESX

`ansible_facts[?distribution_major_version?]`

?

The major version of the operating system. For example, the value is

16

for Ubuntu 16.04.

`ansible_facts[?os_family?]`

?

Possible values (sample, not complete list):

AIX
Alpine
Altlinux
Archlinux
Darwin
Debian
FreeBSD
Gentoo
HP-UX
Mandrake
RedHat

[SMGL](#)

[Slackware](#)

[Solaris](#)

[Suse](#)

[Windows](#)

[See also](#)

[Working with playbooks](#)

[An introduction to playbooks](#)

[Roles](#)

[Playbook organization by roles](#)

[General tips](#)

[Tips and tricks for playbooks](#)

[Using variables](#)

[All about variables](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/playbook_guide/playbooks_intro.html

Ansible playbooks ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Ansible playbooks

Playbook syntax

Playbook execution

Task execution

Desired state and idempotency

Running playbooks

Running playbooks in check mode

Ansible-Pull

Verifying playbooks

ansible-lint

Working with playbooks

Executing playbooks

Advanced playbook syntax

Manipulating data

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Using Ansible playbooks
Ansible playbooks
Edit on GitHub
Ansible playbooks
?

Ansible Playbooks provide a repeatable, reusable, simple configuration management and multimachine deployment system that is well suited to deploying complex applications. If you need to execute a task with Ansible more than once, you can write a playbook and put the playbook under source control. You can then use the playbook to push new configurations or confirm the configuration of remote systems.

Playbooks allow you to perform the following actions:

Declare configurations.

Orchestrate steps of any manual ordered process on multiple sets of machines in a defined order.

Launch tasks synchronously or
asynchronously

.

Playbook syntax

Playbook execution

Task execution

Desired state and idempotency

Running playbooks

Running playbooks in check mode

Ansible-Pull

Verifying playbooks

ansible-lint

Playbook syntax

?

You express playbooks in YAML format with a minimum of syntax. If you are not familiar with YAML, review the [YAML Syntax](#)

[overview](#) and consider installing an add-on for your text editor (see [Other Tools and Programs](#)

) to help you write clean YAML syntax in your playbooks.

A playbook consists of one or more `plays` in an ordered list. The terms `playbook` and `play` are sports analogies. Each play executes part of the overall goal of the playbook, running one or more tasks. Each task calls an Ansible module.

Playbook execution

?

A playbook runs in order from top to bottom. Within each play, tasks also run in order from top to bottom. Playbooks with multiple plays can orchestrate multimachine deployments, running one play on your web servers, another play on your database servers, and a third play on your network infrastructure. At a minimum, each play defines two things:

The managed nodes to target, using a `pattern`

.

At least one task to execute.

For Ansible 2.10 and later, you should use the fully-qualified collection name (FQCN) in your playbooks. Using the FQCN ensures that you have selected the correct module, because multiple collections can contain modules with the same name. For example,

`user`

. See

[Using collections in a playbook](#)

.

In the following example, the first play targets the web servers and the second play targets the database servers.

-

`name`

:

Update web servers

`hosts`

:

`web servers`

`remote_user`

:

`root`

`tasks`

:

-

`name`

:

Ensure apache is at the latest version

`ansible.builtin.yum`

:

`name`

:

`httpd`

`state`

:

`latest`

-

`name`

:

Write the apache config file

```
ansible.builtin.template
```

```
:
```

```
src
```

```
:
```

```
/srv/httpd.j2
```

```
dest
```

```
:
```

```
/etc/httpd.conf
```

```
-
```

```
name
```

```
:
```

```
Update db servers
```

```
hosts
```

```
:
```

```
databases
```

```
remote_user
```

```
:
```

```
root
```

```
tasks
```

```
:
```

```
-
```

```
name
```

```
:
```

```
Ensure postgresql is at the latest version
```

```
ansible.builtin.yum
```

```
:
```

```
name
```

```
:
```

```
postgresql
```

```
state
```

```
:
```

```
latest
```

```
-
```

```
name
```

```
:
```

```
Ensure that postgresql is started
```

```
ansible.builtin.service
```

```
:
```

```
name
```

```
:
```

```
postgresql
```

```
state
```

```
:
```

```
started
```

Your playbook can include more than just a hosts line and tasks. For example, the playbook above sets a

`remote_user`

for each play. The

`remote_user`

is the user account for the SSH connection. You can add other

Playbook Keywords

at the playbook, play, or task level to influence how Ansible behaves. Playbook keywords can control the connection plugin

, whether to use

privilege escalation

, how to handle errors, and more. To support a variety of environments, you can set many of these parameters as command-line flags in your Ansible configuration, or in your inventory. Learning the

precedence rules

for these sources of data helps you as you expand your Ansible ecosystem.

Task execution

?

By default, Ansible executes each task in order, one at a time, against all machines matched by the host pattern. Each task executes a module with specific arguments. After a task has executed on all target machines, Ansible moves to the next task. You can use

strategies

to change this default behavior. Within each play, Ansible applies the same task directives to all hosts. If a task fails on a host, Ansible removes that host from the rotation for the rest of the playbook.

When you run a playbook, Ansible returns information about connections, the name

lines of all your plays and tasks, whether each task has succeeded or failed on each machine, and whether each task has made a change on each machine. At the bottom of the playbook execution, Ansible provides a summary of the nodes that were targeted and how they performed. General failures and fatal ?unreachable? communication attempts are kept separate in the counts.

Desired state and idempotency

?

Most Ansible modules check whether the desired final state has already been achieved and exit without performing any actions if that state has been achieved. Repeating the task does not change the final state. Modules that behave this way are ?idempotent?. Whether you run a playbook once or multiple times, the outcome should be the same. However, not all playbooks and not all modules behave this way. If you are unsure, test your playbooks in a sandbox environment before running them multiple times in production.

Running playbooks

?

To run your playbook, use the

ansible-playbook

command.

ansible-playbook

playbook.yml

-f

10

Use the

--verbose

flag when running your playbook to see detailed output from successful and unsuccessful tasks.

Running playbooks in check mode

?

The Ansible check mode allows you to execute a playbook without applying any alterations to your systems. You can use check mode to test playbooks before you implement them in a production environment.

To run a playbook in check mode, pass the

-C

or

--check

flag to the

ansible-playbook

command:

ansible-playbook

--check

playbook.yaml

Executing this command runs the playbook normally. Instead of implementing any modifications, Ansible provides a report on the changes it would have made. This report includes details such as file modifications, command execution,

and module calls.

Check mode offers a safe and practical approach to examine the functionality of your playbooks without risking unintended changes to your systems. Check mode is also a valuable tool for troubleshooting playbooks that are not functioning as expected.

Ansible-Pull

?

You can invert the Ansible architecture so that nodes check in to a central location instead of you pushing configuration out to them.

The

ansible-pull

command is a small script that checks out a repo of configuration instructions from git and then runs

ansible-playbook

against that content.

If you load balance your checkout location,

ansible-pull

scales infinitely.

Run

ansible-pull

--help

for details.

Verifying playbooks

?

You may want to verify your playbooks to catch syntax errors and other problems before you run them. The

ansible-playbook

command offers several options for verification, including

--check

,

--diff

,

--list-hosts

,

--list-tasks

, and

--syntax-check

. The

Tools for validating playbooks

topic describes other tools for validating and testing playbooks.

ansible-lint

?

You can use

ansible-lint

for detailed, Ansible-specific feedback on your playbooks before you execute them. For example, if you run

ansible-lint

on the playbook called

verify-apache.yml

near the top of this page, you should get the following results:

\$

ansible-lint

verify-apache.yml

[

403

]

Package

installs

should

not

use

latest

verify-apache.yml:8

Task/Handler:

ensure

apache

is

at

the

latest

version

The

ansible-lint default rules

page describes each error.

See also

ansible-lint

Learn how to test Ansible Playbooks syntax

YAML Syntax

Learn about YAML syntax

General tips

Tips for managing playbooks in the real world

Collection Index

Browse existing collections, modules, and plugins

Should you develop a module?

Learn to extend Ansible by writing your own modules

Patterns: targeting hosts and groups

Learn about how to select hosts

Communication

Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/playbook_guide/playbooks_variables.html

Using variables ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Ansible playbooks

Working with playbooks

Templating (Jinja2)

Using filters to manipulate data

Tests

Lookups

Python3 in templates

The now function: get the current time

The undef function: add hint for undefined variables

Loops

Controlling where tasks run: delegation and local actions

Conditionals

Blocks

Handlers: running operations on change

Error handling in playbooks

Setting the remote environment

Working with language-specific version managers

Reusing Ansible artifacts

Roles

Module defaults

Interactive input: prompts

Using variables

Creating valid variable names

Simple variables

When to quote variables (a YAML gotcha)

List variables

Dictionary variables

Combining variables

Registering variables

Referencing nested variables

Transforming variables with Jinja2 filters

Where to set variables
Variable precedence: where should I put a variable?
Using advanced variable syntax
Discovering variables: facts and magic variables
Playbook Example: Continuous Delivery and Rolling Upgrades
Executing playbooks
Advanced playbook syntax
Manipulating data
Protecting sensitive data with Ansible vault
Using Ansible modules and plugins
Using Ansible collections
Using Ansible on Windows, BSD, and z/OS UNIX
Ansible tips and tricks
Contributing to Ansible
Ansible Community Guide
Ansible Collections Contributor Guide
ansible-core Contributors Guide
Advanced Contributor Guide
Ansible documentation style guide
Extending Ansible
Developer Guide
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible

Using Ansible playbooks

Working with playbooks

Using variables

Edit on GitHub

Using variables

?

Ansible uses variables to manage differences between systems. With Ansible, you can execute tasks and playbooks on multiple systems with a single command. To represent the variations among those different systems, you can create variables with standard YAML syntax, including lists and dictionaries. You can define these variables in your playbooks, in your

inventory

, in reusable

files

or

roles

, or at the command line. You can also create variables during a playbook run by registering the return value of a task as a new variable.

After you create a variable, you can use it in module arguments, in

conditional `when` statements

, in

templates

, and in

loops

.

After you understand the concepts and examples on this page, read about

Ansible facts

, which are variables you retrieve from remote systems.

Creating valid variable names

Simple variables

Defining simple variables

Referencing simple variables

When to quote variables (a YAML gotcha)

List variables

Defining variables as lists

Referencing list variables

Dictionary variables

Defining variables as key-value dictionaries

Referencing key-value dictionary variables

Combining variables

Combining list variables

Combining dictionary variables

Using the `merge_variables` lookup

Registering variables

Referencing nested variables

Transforming variables with Jinja2 filters

Where to set variables

Defining variables in inventory

Defining variables in a play

Defining variables in included files and roles

Defining variables at runtime

Key-value format

JSON string format

Vars from a JSON or YAML file

Variable precedence: where should I put a variable?

Understanding variable precedence

Scoping variables

Tips on where to set variables

Using advanced variable syntax

Creating valid variable names

?

Not all strings are valid Ansible variable names. A variable name can only include letters, numbers, and underscores.

Python keywords

or

playbook keywords

are not valid variable names. A variable name cannot begin with a number.

Variable names can begin with an underscore. In many programming languages, variables that begin with an underscore are private. This is not true in Ansible. Ansible treats variables that begin with an underscore the same as any other variable. Do not rely on this convention for privacy or security.

This table gives examples of valid and invalid variable names:

Valid variable names

Not valid

foo

*foo

,

Python keywords

such as

async

and

lambda

foo_env

playbook keywords

such as

environment

foo_port

foo-port

,

foo

port

,

foo.port

foo5

,

_foo

5foo

,

12

Ansible defines certain

variables

internally. You cannot define these variables.

Avoid variable names that overwrite Jinja2 global functions listed in

Working with playbooks

, such as

lookup

,

query

,

q

,

now
, and
undef

.

Simple variables

?

Simple variables combine a variable name with a single value. You can use this syntax, and the syntax for lists and dictionaries shown below, in a variety of places. For details about setting variables in inventory, in playbooks, in reusable files, in roles, or at the command line, see

Where to set variables

.

Defining simple variables

?

You can define a simple variable using standard YAML syntax. For example:

```
remote_install_path: /opt/my_app_config
```

Referencing simple variables

?

After you define a variable, use Jinja2 syntax to reference it. Jinja2 variables use double curly braces. For example, the expression

```
My
amp
goes
to
{{
max_amp_value
}}
```

demonstrates the most basic form of variable substitution. You can use Jinja2 syntax in playbooks. The following example shows a variable that defines the location of a file, which can vary from one system to another:

```
ansible.builtin.template
```

```
:
src
:
foo.cfg.j2
dest
:
'
{{
remote_install_path
}}
/foo.cfg'
```

Ansible allows Jinja2 loops and conditionals in templates

but not in playbooks. You cannot create a loop of tasks. Ansible playbooks are pure machine-parseable YAML.

When to quote variables (a YAML gotcha)

?

If you start a value with

```
{{
foo
}}
```

, you must quote the whole expression to create valid YAML syntax. If you do not quote the whole expression, the YAML parser cannot interpret the syntax. The parser cannot determine if it is a variable or the start of a YAML dictionary. For guidance on writing YAML, see the

YAML Syntax
documentation.

If you use a variable without quotes, like this:

```
- hosts: app_servers
  vars:
    app_path: {{ base_path }}/22
```

You will see:

ERROR!

Syntax

Error

while

loading

YAML.

If you add quotes, Ansible works correctly:

```
-
hosts
:
app_servers
vars
:
app_path
:
"
{{
base_path
}}
/22"
```

List variables

?

A list variable combines a variable name with multiple values. You can store the multiple values as an itemized list or in square brackets

```
[]
```

, separated with commas.

Defining variables as lists

?

You can define variables with multiple values using YAML lists. For example:

```
region
:
-
northeast
-
southeast
-
midwest
```

Referencing list variables

?

If you use a variable defined as a list (also called an array), you can use individual, specific items from that list. The first item in a list is item 0, the second item is item 1, and so on. For example:

```
region
:
"
{{
region
[
0
]
```

```
}}  
"
```

The value of this expression would be ?northeast?.

Dictionary variables

?

A dictionary stores data in key-value pairs. Usually, you use dictionaries to store related data, such as the information contained in an ID or a user profile.

Defining variables as key-value dictionaries

?

You can define more complex variables using YAML dictionaries. A YAML dictionary maps keys to values. For example:

```
foo
```

```
:
```

```
field1
```

```
:
```

```
one
```

```
field2
```

```
:
```

```
two
```

Referencing key-value dictionary variables

?

If you use a variable defined as a key-value dictionary (also called a hash), you can use individual, specific items from that dictionary using either bracket notation or dot notation:

```
foo['field1']
```

```
foo.field1
```

Both of these examples reference the same value (?one?). Bracket notation always works. Dot notation can cause problems because some keys collide with attributes and methods of python dictionaries. Use bracket notation if you use keys that start and end with two underscores, which are reserved for special meanings in python, or are any of the known public attributes:

```
add
```

```
,
```

```
append
```

```
,
```

```
as_integer_ratio
```

```
,
```

```
bit_length
```

```
,
```

```
capitalize
```

```
,
```

```
center
```

```
,
```

```
clear
```

```
,
```

```
conjugate
```

```
,
```

```
copy
```

```
,
```

```
count
```

```
,
```

```
decode
```

```
,
```

```
denominator
```

```
,
```

```
difference
```

```
,
```


difference_update
,
discard
,
encode
,
endswith
,
expandtabs
,
extend
,
find
,
format
,
fromhex
,
fromkeys
,
get
,
has_key
,
hex
,
imag
,
index
,
insert
,
intersection
,
intersection_update
,
isalnum
,
isalpha
,
isdecimal
,
isdigit
,
isdisjoint
,
is_integer
,
islower
,
isnumeric
,
isspace
,

issubset
,
issuperset
,
istitle
,
isupper
,
items
,
iteritems
,
iterkeys
,
itervalues
,
join
,
keys
,
ljust
,
lower
,
lstrip
,
numerator
,
partition
,
pop
,
popitem
,
real
,
remove
,
replace
,
reverse
,
rfind
,
rindex
,
rjust
,
rpartition
,
rsplit
,
rstrip
,

setdefault
,
sort
,
split
,
splitlines
,
startswith
,
strip
,
swapcase
,
symmetric_difference
,
symmetric_difference_update
,
title
,
translate
,
union
,
update
,
upper
,
values
,
viewitems
,
viewkeys
,
viewvalues
,
zfill

-
Combining variables

?

To merge variables that contain lists or dictionaries, you can use the following approaches.

Combining list variables

?

You can use the

set_fact

module to combine lists into a new

merged_list

variable as follows:

vars

:

list1

:

-

apple

```
-
banana
-
fig
list2
:
-
peach
-
plum
-
pear
tasks
:
-
name
:
Combine list1 and list2 into a merged_list var
ansible.builtin.set_fact
:
merged_list
:
"{{
list1
+
list2
}}"
Combining dictionary variables
?
To merge dictionaries, use the
combine
filter. For example:
vars
:
dict1
:
name
:
Leeroy Jenkins
age
:
25
occupation
:
Astronaut
dict2
:
location
:
Galway
country
:
Ireland
postcode
```

```
:
H71 1234
tasks
:
-
name
:
Combine dict1 and dict2 into a merged_dict var
ansible.builtin.set_fact
```

```
:
merged_dict
:
"{{
dict1
|
ansible.builtin.combine(dict2)
}}"
```

For more details, see
ansible.builtin.combine

.

Using the merge_variables lookup
?

To merge variables that match the given prefixes, suffixes, or regular expressions, you can use the
community.general.merge_variables
lookup. For example:
merged_variable

```
:
"{{
lookup('community.general.merge_variables',
'__my_pattern',
pattern_type='suffix')
}}"
```

For more details and example usage, refer to the
community.general.merge_variables lookup documentation

.

Registering variables
?

You can create a variable from the output of an Ansible task with the task keyword
register
. You can use the registered variable in any later task in your play. For example:

```
-
hosts
:
web_servers
tasks
:
-
name
:
Run a shell command and register its output as a variable
ansible.builtin.shell
:
/usr/bin/foo
register
```

```
:
foo_result
ignore_errors
:
```

```
true
```

```
-
```

```
name
```

```
:
```

Run a shell command using output of the previous task

```
ansible.builtin.shell
```

```
:
```

```
/usr/bin/bar
```

```
when
```

```
:
```

```
foo_result.rc == 5
```

For more examples of using registered variables in conditions on later tasks, see

Conditionals

. Registered variables may be simple variables, list variables, dictionary variables, or complex nested data structures.

The documentation for each module includes a

RETURN

section that describes the return values for that module. To see the values for a particular task, run your playbook with

```
-v
```

```
.
```

Registered variables are stored in memory. You cannot cache registered variables for use in future playbook runs. A registered variable is valid only on the host for the rest of the current playbook run, including subsequent plays within the same playbook run.

Registered variables are host-level variables. When you register a variable in a task with a loop, the registered variable contains a value for each item in the loop. The data structure placed in the variable during the loop contains a results

attribute, which is a list of all responses from the module. For a more in-depth example of how this works, see the

Loops

section on using register with a loop.

If a task fails or is skipped, Ansible still registers a variable with a failure or skipped status, unless the task is skipped based on tags. See

Tags

for information on adding and using tags.

Referencing nested variables

```
?
```

Many registered variables and

facts

are nested YAML or JSON data structures. You cannot access values from these nested data structures with the simple

```
{{
```

```
foo
```

```
}}
```

syntax. You must use either bracket notation or dot notation. For example, to reference an IP address from your facts using bracket notation:

```
,
```

```
{{
```

```
ansible_facts
```

```
[
```

```
"eth0"
```

```
][
```

```
"ipv4"
```

```
][
```

```
"address"
```

```
]
```

```
}}
```

```
,
```

To reference an IP address from your facts using dot notation:

```
{{
```

```
ansible_facts.eth0.ipv4.address
```

```
}}
```

Transforming variables with Jinja2 filters

?

Jinja2 filters let you transform the value of a variable within a template expression. For example, the `capitalize`

filter capitalizes any value passed to it; the

`to_yaml`

and

`to_json`

filters change the format of your variable values. Jinja2 includes many built-in filters

, and Ansible supplies many more filters. To find more examples of filters, see

[Using filters to manipulate data](#)

.

Where to set variables

?

You can define variables in a variety of places, such as in inventory, in playbooks, in reusable files, in roles, and at the command line. Ansible loads every possible variable it finds, then chooses the variable to apply based on variable precedence rules

.

Defining variables in inventory

?

You can define different variables for each host individually, or set shared variables for a group of hosts in your inventory. For example, if all machines in the

`[boston]`

group use `?boston.ntp.example.com?` as an NTP server, you can set a group variable. The

[How to build your inventory](#)

page has details on setting

host variables

and

group variables

in inventory.

Defining variables in a play

?

You can define variables directly in a playbook play:

-

hosts

:

webservers

vars

:

http_port

:

80

When you define variables in a play, they are visible only to tasks executed in that play.

Defining variables in included files and roles

?

You can define variables in reusable variables files or in reusable roles. If you define variables in reusable variable files, the sensitive variables are separated from playbooks. This separation enables you to store your playbooks in a source control software and even share the playbooks, without the risk of exposing passwords or other sensitive and personal data. For information about creating reusable files and roles, see [Reusing Ansible artifacts](#)

This example shows how you can include variables defined in an external file:

```
---
-
hosts
:
all
remote_user
:
root
vars
:
favcolor
:
blue
vars_files
:
-
/vars/external_vars.yml
tasks
:
-
name
:
This is just a placeholder
ansible.builtin.command
:
/bin/echo foo
```

The contents of each variables file is a simple YAML dictionary. For example:

```
---
# in the above example, this would be vars/external_vars.yml
somevar
:
somevalue
password
:
magic
```

You can keep per-host and per-group variables in similar files. To learn about organizing your variables, see [Organizing host and group variables](#)

Defining variables at runtime

?

You can define variables when you run your playbook by passing variables at the command line using the

`--extra-vars`

(or

`-e`

) argument. You can also request user input with a

`vars_prompt`

(see

Interactive input: prompts

). If you pass variables at the command line, use a single quoted string that contains one or more variables in one of the formats below.

Key-value format

?

Values passed in using the

key=value

syntax are interpreted as strings. Use the JSON format if you need to pass non-string values such as Booleans, integers, floats, and lists.

```
ansible-playbook release.yml --extra-vars "version=1.23.45 other_variable=foo"
```

JSON string format

?

```
ansible-playbook
```

```
release.yml
```

```
--extra-vars
```

```
'{"version": "1.23.45", "other_variable": "foo"}'
```

```
ansible-playbook
```

```
arcade.yml
```

```
--extra-vars
```

```
'{"pacman": "mrs", "ghosts": ["inky", "pinky", "clyde", "sue"]}'
```

When passing variables with

```
--extra-vars
```

, you must escape quotes and other special characters appropriately for both your markup (for example, JSON) and for your shell:

```
ansible-playbook
```

```
arcade.yml
```

```
--extra-vars
```

```
"{'name': 'Conan O'Brien'}"
```

```
ansible-playbook
```

```
arcade.yml
```

```
--extra-vars
```

```
'{"name": "Conan O'
```

```
\\'
```

```
'Brien"}'
```

```
ansible-playbook
```

```
script.yml
```

```
--extra-vars
```

```
"{'dialog': 'He said \\\"I just can't get enough of those single and double-quotes\"
```

```
\\'
```

```
\\\"\\\"}'"
```

Vars from a JSON or YAML file

?

If you have a lot of special characters, use a JSON or YAML file containing the variable definitions. Prepend both JSON and YAML file names with

@

.

```
ansible-playbook release.yml --extra-vars "@some_file.json"
```

```
ansible-playbook release.yml --extra-vars "@some_file.yaml"
```

Variable precedence: where should I put a variable?

?

You can set multiple variables with the same name in many different places. If you do this, Ansible loads every possible variable it finds, and then chooses the variable to apply based on variable precedence. In other words, the different variables will override each other in a certain order.

Teams and projects that agree on guidelines for defining variables (where to define certain types of variables) usually

avoid variable precedence concerns. You should define each variable in one place. Determine where to define a variable, and keep it simple. For examples, see

[Tips on where to set variables](#)

.

Some behavioral parameters that you can set in variables you can also set in Ansible configuration, as command-line options, and using playbook keywords. For example, you can define the user that Ansible uses to connect to remote devices as a variable with

`ansible_user`

, in a configuration file with

`DEFAULT_REMOTE_USER`

, as a command-line option with

`-u`

, and with the playbook keyword

`remote_user`

. If you define the same parameter in a variable and by another method, the variable overrides the other setting. This approach allows host-specific settings to override more general settings. For examples and more details on the precedence of these various settings, see

[Controlling how Ansible behaves: precedence rules](#)

.

Understanding variable precedence

?

Ansible does apply variable precedence, and you might have a use for it. Here is the order of precedence from least to greatest (the last listed variables override all other variables):

Command-line values (for example,

`-u`

`my_user`

, these are not variables)

Role defaults (as defined in

Role directory structure

)

[

1

]

Inventory file or script group vars

[

2

]

Inventory group_vars/all

[

3

]

Playbook group_vars/all

[

3

]

Inventory group_vars/*

[

3

]

Playbook group_vars/*

[

3

]

Inventory file or script host vars

```
[
2
]
Inventory host_vars/*
[
3
]
Playbook host_vars/*
[
3
]
Host facts and cached set_facts
[
4
]
Play vars
Play vars_prompt
Play vars_files
Role vars (as defined in
Role directory structure
)
Block vars (for tasks in block only)
Task vars (for the task only)
include_vars
Registered vars and set_facts
Role (and include_role) params
include params
Extra vars (for example,
-e
"user=my_user"
)(always win precedence)
```

In general, Ansible gives precedence to variables that were defined more recently, more actively, and with more explicit scope. Variables in the defaults folder inside a role are easily overridden. Anything in the vars directory of the role overrides previous versions of that variable in the namespace. Host or inventory variables override role defaults, but explicit includes such as the vars directory or an

```
include_vars
task override inventory variables.
```

Ansible merges different variables set in inventory so that more specific settings override more generic settings. For example,

```
ansible_ssh_user
specified as a group_var is overridden by
ansible_user
```

specified as a host_var. For details about the precedence of variables set in inventory, see [How variables are merged](#)

.

Footnotes

```
[
1
]
Tasks in each role see their own role?s defaults. Tasks defined outside of a role see the last role?s defaults.
[
2
]
(
```

```
1
,
2
)
```

Variables defined in inventory file or provided by dynamic inventory.

```
[
3
]
(
1
,
2
,
3
,
4
,
5
,
6
)
```

Includes vars added by ?vars plugins? as well as host_vars and group_vars which are added by the default vars plugin shipped with Ansible.

```
[
4
]
```

When created with set_facts?s cacheable option, variables have the high precedence in the play, but are the same as a host facts precedence when they come from the cache.

Note

Within any section, redefining a var overrides the previous instance. If multiple groups have the same variable, the last one loaded wins. If you define a variable twice in a play?s

vars:

section, the second one wins.

The previous text describes the default config

hash_behavior=replace

. Switch to

merge

to overwrite only partially.

Scoping variables

?

You can decide where to set a variable based on the scope you want that value to have. Ansible has three main scopes:

Global: this is set by config, environment variables and the command line

Play: each play and contained structures, vars entries (vars; vars_files; vars_prompt), role defaults and vars.

Host: variables directly associated to a host, like inventory, include_vars, facts or registered task outputs

Inside a template, you automatically have access to all variables that are in scope for a host, plus any registered variables, facts, and magic variables.

Tips on where to set variables

?

You should choose where to define a variable based on the kind of control you might want over values.

Set variables in inventory that deal with geography or behavior. Since groups are frequently the entity that maps roles to hosts, you can often set variables on the group instead of defining them on a role. Remember that child groups override parent groups, and host variables override group variables. See

Defining variables in inventory

for details on setting host and group variables.

Set common defaults in a

group_vars/all

file. See

Organizing host and group variables

for details on how to organize host and group variables in your inventory. You generally place group variables alongside your inventory file, but they can also be returned by dynamic inventory (see

Working with dynamic inventory

) or defined in AWX or on the

Red Hat Ansible Automation Platform

from the UI or API:

file: /etc/ansible/group_vars/all

this is the site wide default

ntp_server

:

default-time.example.com

Set location-specific variables in

group_vars/my_location

files. All groups are children of the

all

group, so variables set here override those set in

group_vars/all

:

file: /etc/ansible/group_vars/boston

ntp_server

:

boston-time.example.com

If one host used a different NTP server, you could set that in a host_vars file, which would override the group variable:

file: /etc/ansible/host_vars/xyz.boston.example.com

ntp_server

:

override.example.com

Set defaults in roles to avoid undefined-variable errors. If you share your roles, other users can rely on the reasonable defaults you added in the

roles/x/defaults/main.yml

file, or they can easily override those values in inventory or at the command line. See

Roles

for more info. For example:

file: roles/x/defaults/main.yml

if no other value is supplied in inventory or as a parameter, this value will be used

http_port

:

80

Set variables in roles to ensure a value is used in that role and is not overridden by inventory variables. If you are not sharing your role with others, you can define app-specific behaviors like ports this way, in

roles/x/vars/main.yml

. If you are sharing roles with others, putting variables here makes them harder to override, although they can still be overridden by passing a parameter to the role or setting a variable with

-e

:

```
# file: roles/x/vars/main.yml
# this will absolutely be used in this role
http_port
```

```
:
```

```
80
```

Pass variables as parameters when you call roles for maximum clarity, flexibility, and visibility. This approach overrides any defaults that exist for a role. For example:

```
roles
:
-
role
:
apache
vars
:
http_port
:
8080
```

When you read this playbook, it is clear that you have chosen to set a variable or override a default. You can also pass multiple values, which allows you to run the same role multiple times. See

Running a role multiple times in one play
for more details. For example:

```
roles
:
-
role
:
app_user
vars
:
myname
:
lan
-
role
:
app_user
vars
:
myname
:
Terry
-
role
:
app_user
vars
:
myname
:
Graham
-
role
:
```

app_user

vars

:

myname

:

John

Variables set in one role are available to later roles. You can set variables in the role's

vars

directory (as defined in

Role directory structure

) and use them in other roles and elsewhere in your playbook:

roles

:

-

role

:

common_settings

-

role

:

something

vars

:

foo

:

12

-

role

:

something_else

There are some protections in place to avoid the need to namespace variables. In this example, variables defined in `?common_settings?` are available to `?something?` and `?something_else?` tasks, but tasks in `?something?` have `foo` set at 12, even if `?common_settings?` sets `foo` to 20.

Instead of worrying about variable precedence, we encourage you to think about how easily or how often you want to override a variable when deciding where to set it. If you are not sure what other variables are defined and you need a particular value, use

`--extra-vars`

(

`-e`

) to override all other variables.

Using advanced variable syntax

?

For information about advanced YAML syntax used to declare variables and have more control over the data placed in YAML files used by Ansible, see

Advanced playbook syntax

.

See also

Ansible playbooks

An introduction to playbooks

Conditionals

Conditional statements in playbooks

Using filters to manipulate data

Jinja2 filters and their uses

Loops

[Looping in playbooks](#)

[Roles](#)

[Playbook organization by roles](#)

[General tips](#)

[Tips and tricks for playbooks](#)

[Special Variables](#)

[List of special variables](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/ansible/latest/reference_appendices/YAMLSyntax.html

YAML Syntax ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

YAML Basics

Gotchas
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
YAML Syntax
Edit on GitHub
YAML Syntax
?

This page provides a basic overview of correct YAML syntax, which is how Ansible playbooks (our configuration management language) are expressed.

We use YAML because it is easier for humans to read and write than other common data formats like XML or JSON. Further, there are libraries available in most programming languages for working with YAML.

You may also wish to read

Working with playbooks

at the same time to see how this is used in practice.

YAML Basics

?

For Ansible, nearly every YAML file starts with a list.

Each item in the list is a list of key/value pairs, commonly called a ?hash? or a ?dictionary?. So, we need to know how to write lists and dictionaries in YAML.

There?s another small quirk to YAML. All YAML files (regardless of their association with Ansible or not) can optionally begin with

and end with

...

. This is part of the YAML format and indicates the start and end of a document.

All members of a list are lines beginning at the same indentation level starting with a

"_

"

(a dash and a space):

A list of tasty fruits

-

Apple

-

Orange

-

Strawberry

-

Mango

...

A dictionary is represented in a simple

key:

value

form (the colon must be followed by a space):

An employee record

martin

:

name

:

Martin D'vloper

job

:

Developer

skill

:

Elite

More complicated data structures are possible, such as lists of dictionaries, dictionaries whose values are lists or a mix of both:

Employee records

-

martin

:

name

:

Martin D'vloper

job

:

Developer

skills

:

-

python

-

perl

-

pascal

-

tabitha

:

name

:

Tabitha Bitumen

job

:

Developer

skills

:

-

lisp

-

fortran

-

erlang

Dictionaries and lists can also be represented in an abbreviated form if you really want to:

martin

:

{

name

:

Martin D'vloper

,

job

:

Developer

,

skill

:

Elite

}

fruits

:

[

'Apple'

,

'Orange'

,

'Strawberry'

,

'Mango'

]

These are called ?Flow collections?.

Ansible doesn't really use these too much, but you can also specify a

boolean value

(true/false) in several forms:

create_key

:

true

needs_agent

:

false

knows_oop

:

True

likes_emacs

:

TRUE

uses_cvs

:

false

Use lowercase ?true? or ?false? for boolean values in dictionaries if you want to be compatible with default yamllint options.

Values can span multiple lines using

|

or

>

. Spanning multiple lines using a ?Literal Block Scalar?

|

will include the newlines and any trailing spaces.

Using a ?Folded Block Scalar?

>

will fold newlines to spaces; it is used to make what would otherwise be a very long line easier to read and edit.

In either case the indentation will be ignored.

Examples are:

include_newlines

:

|

exactly as you see

will appear these three

lines of poetry

fold_newlines

:

>

this is really a

single line of text

despite appearances

While in the above

>

example all newlines are folded into spaces, there are two ways to enforce a newline to be kept:

fold_some_newlines

:

>

a

b

c

d

e

f

Alternatively, it can be enforced by including newline

\n

characters:

fold_same_newlines

:

"a

b\n

c

d\n

e\nf\n"

Let?s combine what we learned so far in an arbitrary YAML example.

This really has nothing to do with Ansible, but will give you a feel for the format:

An employee record

name

:

Martin D'vloper

job

:

Developer

skill

:

```

Elite
employed
:
True
foods
:
-
Apple
-
Orange
-
Strawberry
-
Mango
languages
:
perl
:
Elite
python
:
Elite
pascal
:
Lame
education
:

```

```

|
4 GCSEs
3 A-Levels
BSc in the Internet of Things
That?s all you really need to know about YAML to start writing
Ansible
playbooks.
Gotchas
?

```

While you can put just about anything into an unquoted scalar, there are some exceptions.

A colon followed by a space (or newline)

```

":
"

```

is an indicator for a mapping.

A space followed by the pound sign

```

"
#"

```

starts a comment.

Because of this, the following is going to result in a YAML syntax error:

```
foo: somebody said I should put a colon here: so I did
```

```
windows_drive: c:
```

?but this will work:

```
windows_path
:
c:\windows
```

You will want to quote hash values using colons followed by a space or the end of the line:

```
foo
:
'somebody
said
I
should
put
a
colon
here:
so
I
did'
windows_drive
```

```
:
```

'c:'

?and then the colon will be preserved.
Alternatively, you can use double quotes:

```
foo
:
"somedbody
said
I
should
put
a
colon
here:
so
I
did"
windows_drive
```

```
:
```

"c:"

The difference between single quotes and double quotes is that in double quotes you can use escapes:

```
foo
:
"a
\t
TAB
and
a
\n
NEWLINE"
```

The list of allowed escapes can be found in the [YAML Specification](#) under ?Escape Sequences? (YAML 1.1) or ?Escape Characters? (YAML 1.2).

The following is invalid YAML:

```
foo: "an escaped \' single quote"
```

Further, Ansible uses ?{{ var }}? for variables. If a value after a colon starts with a ?{?, YAML will think it is a dictionary, so you must quote it, like so:

```
foo
:
"{{
```

variable

}}"

If your value starts with a quote the entire value must be quoted, not just part of it. Here are some additional examples of how to properly quote things:

foo

:

"{{

variable

}}/additional/string/literal"

foo2

:

"{{

variable

}}\\backslashes\\are\\also\\special\\characters"

foo3

:

"even

if

it

is

just

a

string

literal

it

must

all

be

quoted"

Not valid:

foo: "E:\\path\\"rest\\of\\path

In addition to

,

and

"

there are a number of characters that are special (or reserved) and cannot be used as the first character of an unquoted scalar:

[]

{}

>

|

*

&

!

%

#

`

@

,

.

You should also be aware of

?

:

-

. In YAML, they are allowed at the beginning of a string if a non-space character follows, but YAML processor implementations differ, so it is better to use quotes.

In Flow Collections, the rules are a bit more strict:

a scalar in block mapping: this } is [all , valid

flow mapping: { key: "you { should [use , quotes here" }

Boolean conversion is helpful, but this can be a problem when you want a literal

yes

or other boolean values as a string.

In these cases just use quotes:

non_boolean

:

"yes"

other_string

:

"False"

YAML converts certain strings into floating-point values, such as the string

1.0

. If you need to specify a version number (in a requirements.yml file, for example), you will need to quote the value if it looks like a floating-point value:

version

:

"1.0"

See also

[Working with playbooks](#)

Learn what playbooks can do and how to write/run them.

[YAMLLint](#)

YAML Lint (online) helps you debug YAML syntax if you are having problems

[Wikipedia YAML syntax reference](#)

[A good guide to YAML syntax](#)

[YAML 1.1 Specification](#)

The Specification for YAML 1.1, which PyYAML and libyaml are currently implementing

[YAML 1.2 Specification](#)

For completeness, YAML 1.2 is the successor of 1.1

[Communication](#)

Got questions? Need help? Want to share your ideas? Visit the [Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: <https://docs.ansible.com/archive/ansible/2.3/>

Ansible Documentation ? Ansible Documentation

AnsibleFest

Products

Community

Webinars & Training

Blog

Documentation

Ansible Core

v2.3

For previous versions, see the
documentation archive.

Introduction

Installation

Getting Started

Inventory

Dynamic Inventory

Patterns

Introduction To Ad-Hoc Commands

Configuration file

BSD Support

Windows Support

Networking Support

Quickstart Video

Playbooks

Intro to Playbooks

Playbook Roles and Include Statements

Variables

Templating (Jinja2)

Conditionals

Loops

Blocks

Strategies

Best Practices

Playbooks: Special Topics

Become (Privilege Escalation)

Accelerated Mode

Asynchronous Actions and Polling

Check Mode (?Dry Run?)

Playbook Debugger

Delegation, Rolling Updates, and Local Actions

Setting the Environment (and Working With Proxies)

Working With Language-Specific Version Managers

Error Handling In Playbooks

Advanced Syntax

Lookups

Examples

Prompts

Tags

Vault

Start and Step

Directives Glossary

About Modules

Introduction
Return Values
Module Support
Module Index
All Modules
Cloud Modules
Clustering Modules
Commands Modules
Crypto Modules
Database Modules
Files Modules
Identity Modules
Inventory Modules
Messaging Modules
Monitoring Modules
Network Modules
Notification Modules
Packaging Modules
Remote Management Modules
Source Control Modules
Storage Modules
System Modules
Utilities Modules
Web Infrastructure Modules
Windows Modules
Detailed Guides
Amazon Web Services Guide
Getting Started with Azure
Rackspace Cloud Guide
Google Cloud Platform Guide
CloudStack Cloud Guide
Using Vagrant and Ansible
Continuous Delivery and Rolling Upgrades
Getting Started with Docker
Using Ansible with the Packet host
Developer Information
Ansible Developer Guide
Community Information
Ansible Community Guide
Ansible Tower
Ansible Galaxy
The Website
The command line tool
Testing Strategies
Integrating Testing With Ansible Playbooks
The Right Level of Testing
Check Mode As A Drift Test
Modules That Are Useful for Testing
Testing Lifecycle
Integrating Testing With Rolling Updates
Achieving Continuous Deployment
Conclusion
Frequently Asked Questions
How can I set the PATH or any other environment variable for a task or entire playbook?

How do I handle different machines needing different user accounts or ports to log in with?
How do I get ansible to reuse connections, enable Kerberized SSH, or have Ansible pay attention to my local SSH config file?
How do I configure a jump host to access servers that I have no direct access to?
How do I speed up management inside EC2?
How do I handle python pathing not having a Python 2.X in /usr/bin/python on a remote machine?
What is the best way to make content reusable/redistributable?
Where does the configuration file live and what can I configure in it?
How do I disable cowsay?
How do I see a list of all of the ansible_ variables?
How do I see all the inventory vars defined for my host?
How do I loop over a list of hosts in a group, inside of a template?
How do I access a variable name programmatically?
How do I access a variable of the first host in a group?
How do I copy files recursively onto a target host?
How do I access shell environment variables?
How do I generate crypted passwords for the user module?
Can I get training on Ansible?
Is there a web interface / REST API / etc?
How do I submit a change to the documentation?
How do I keep secret data in my playbook?
When should I use {{ }}? Also, how to interpolate variables or dynamic variable names
Why don't you ship in X format?
I don't see my question here

Glossary

YAML Syntax

YAML Basics

Gotchas

Ansible Porting Guides

Python 3 Support

Testing Python 3 with commands and playbooks

Testing Python 3 module support

Release and maintenance

AnsibleFest

Products

Community

Webinars & Training

Blog

Ansible Documentation

Docs

»

Ansible Documentation

You are reading an unmaintained version of the Ansible documentation. Unmaintained Ansible versions can contain unfixed security vulnerabilities (CVE). Please upgrade to a maintained version. See the latest Ansible documentation

.

Ansible Documentation

¶

About Ansible

¶

Welcome to the Ansible documentation!

Ansible is an IT automation tool. It can configure systems, deploy software, and orchestrate more advanced IT tasks such as continuous deployments or zero downtime rolling updates.

Ansible's main goals are simplicity and ease-of-use. It also has a strong focus on security and reliability, featuring a

minimum of moving parts, usage of OpenSSH for transport (with an accelerated socket mode and pull modes as alternatives), and a language that is designed around auditability by humans?even those not familiar with the program. We believe simplicity is relevant to all sizes of environments, so we design for busy users of all types: developers, sysadmins, release engineers, IT managers, and everyone in between. Ansible is appropriate for managing all environments, from small setups with a handful of instances to enterprise environments with many thousands of instances.

Ansible manages machines in an agent-less manner. There is never a question of how to upgrade remote daemons or the problem of not being able to manage systems because daemons are uninstalled. Because OpenSSH is one of the most peer-reviewed open source components, security exposure is greatly reduced. Ansible is decentralized?it relies on your existing OS credentials to control access to remote machines. If needed, Ansible can easily connect with Kerberos, LDAP, and other centralized authentication management systems.

This documentation covers the current released version of Ansible (2.2) and also some development version features (2.3). For recent features, we note in each section the version of Ansible where the feature was added.

Ansible, Inc. releases a new major release of Ansible approximately every two months. The core application evolves somewhat conservatively, valuing simplicity in language design and setup. However, the community around new modules and plugins being developed and contributed moves very quickly, typically adding 20 or so new modules in each release.

[Introduction](#)

[Quickstart Video](#)

[Playbooks](#)

[Playbooks: Special Topics](#)

[About Modules](#)

[Module Index](#)

[Detailed Guides](#)

[Developer Information](#)

[Community Information](#)

[Ansible Tower](#)

[Ansible Galaxy](#)

[Testing Strategies](#)

[Frequently Asked Questions](#)

[Glossary](#)

[YAML Syntax](#)

[Ansible Porting Guides](#)

[Python 3 Support](#)

[Release and maintenance](#)

[Next](#)

Copyright © 2017 Red Hat, Inc.

Last updated on Sep 01, 2022.

Ansible docs are generated from

GitHub sources

using

Sphinx

using a theme provided by

Read the Docs

.

Content from: <https://docs.ansible.com/archive/ansible/2.5/>

Ansible Documentation ? Ansible Documentation

AnsibleFest

Products

Community

Webinars & Training

Blog

Documentation

Ansible

2.5

For previous versions, see the
documentation archive.

Installation, Upgrade & Configuration

Installation Guide

Basics / What Will Be Installed

What Version To Pick?

Control Machine Requirements

Managed Node Requirements

Installing the Control Machine

Latest Release via DNF or Yum

Latest Releases Via Apt (Ubuntu)

Latest Releases Via Apt (Debian)

Latest Releases Via Portage (Gentoo)

Latest Releases Via pkg (FreeBSD)

Latest Releases on Mac OSX

Latest Releases Via OpenCSW (Solaris)

Latest Releases Via Pacman (Arch Linux)

Latest Releases Via Pip

Tarballs of Tagged Releases

Running From Source

Ansible on GitHub

Configuring Ansible

Configuration file

Getting the latest configuration

Environmental configuration

Command line options

Ansible Porting Guides

Using Ansible

User Guide

Ansible Quickstart

Getting Started

Foreword

Remote Connection Information

Your first commands

Host Key Checking

Working with Command Line Tools

ansible

Synopsis

Description

Common Options

Environment

Files

Author

Copyright
See also
ansible-config
Synopsis
Description
Common Options
Actions
Environment
Files
Author
Copyright
See also
ansible-console
Synopsis
Description
Common Options
Environment
Files
Author
Copyright
See also
ansible-doc
Synopsis
Description
Common Options
Environment
Files
Author
Copyright
See also
ansible-galaxy
Synopsis
Description
Common Options
Actions
Environment
Files
Author
Copyright
See also
ansible-inventory
Synopsis
Description
Common Options
Environment
Files
Author
Copyright
See also
ansible-playbook
Synopsis
Description
Common Options
Environment

Files
Author
Copyright
See also
ansible-pull
Synopsis
Description
Common Options
Environment
Files
Author
Copyright
See also
ansible-vault
Synopsis
Description
Common Options
Actions
Environment
Files
Author
Copyright
See also
Introduction To Ad-Hoc Commands
Parallelism and Shell Commands
File Transfer
Managing Packages
Users and Groups
Deploying From Source Control
Managing Services
Time Limited Background Operations
Gathering Facts
Working with Inventory
Hosts and Groups
Host Variables
Group Variables
Groups of Groups, and Group Variables
Default groups
Splitting Out Host and Group Specific Data
How Variables Are Merged
List of Behavioral Inventory Parameters
Non-SSH connection types
Working With Dynamic Inventory
Example: The Cobbler External Inventory Script
Example: AWS EC2 External Inventory Script
Example: OpenStack External Inventory Script
Explicit use of inventory script
Implicit use of inventory script
Refresh the cache
Other inventory scripts
Using Inventory Directories and Multiple Inventory Sources
Static Groups of Dynamic Groups
Working With Playbooks
Intro to Playbooks

- About Playbooks
- Playbook Language Example
- Basics
- Action Shorthand
- Handlers: Running Operations On Change
- Executing A Playbook
- Ansible-Pull
- Tips and Tricks
- Creating Reusable Playbooks
- Including and Importing
- Roles
- Dynamic vs. Static
- Differences Between Static and Dynamic
- Tradeoffs and Pitfalls Between Includes and Imports
- Variables
- What Makes A Valid Variable Name
- Variables Defined in Inventory
- Variables Defined in a Playbook
- Variables defined from included files and roles
- Using Variables: About Jinja2
- Jinja2 Filters
- Hey Wait, A YAML Gotcha
- Information discovered from systems: Facts
- Turning Off Facts
- Local Facts (Facts.d)
- Ansible version
- Fact Caching
- Registered Variables
- Accessing Complex Variable Data
- Magic Variables, and How To Access Information About Other Hosts
- Variable File Separation
- Passing Variables On The Command Line
- Variable Precedence: Where Should I Put A Variable?
- Variable Scopes
- Variable Examples
- Advanced Syntax
- Templating (Jinja2)
- Filters
- Tests
- Lookups
- Python Version and Templating
- Conditionals
- The When Statement
- Loops and Conditionals
- Loading in Custom Facts
- Applying ?when? to roles, imports, and includes
- Conditional Imports
- Selecting Files And Templates Based On Variables
- Register Variables
- Commonly Used Facts
- Loops
- Standard Loops
- Complex loops
- Using lookup vs query with loop

Do-Until Loops
Using register with a loop
Looping over the inventory
Loop Control
Loops and Includes in 2.0
Blocks
Error Handling
Advanced Playbooks Features
Understanding Privilege Escalation
Asynchronous Actions and Polling
Check Mode (?Dry Run?)
Playbook Debugger
Delegation, Rolling Updates, and Local Actions
Setting the Environment (and Working With Proxies)
Working With Language-Specific Version Managers
Error Handling In Playbooks
Advanced Syntax
Working With Plugins
Prompts
Tags
Using Vault in playbooks
Start and Step
Playbook Keywords
Lookups
Strategies
Strategy Plugins
Best Practices
Content Organization
Staging vs Production
Rolling Updates
Always Mention The State
Group By Roles
Operating System and Distribution Variance
Bundling Ansible Modules With Playbooks
Whitespace and Comments
Always Name Tasks
Keep It Simple
Version Control
Variables and Vaults
Understanding Privilege Escalation
Become
Directives
Connection variables
Command line options
For those from Pre 1.9 , sudo and su still work!
Limitations
Become and Networks
network_cli and become
Setting enable mode for all tasks
authorize and auth_pass
Become and Windows
Administrative Rights
Local Service Accounts
Accounts without a Password

Become Flags
Limitations
Ansible Vault
What Can Be Encrypted With Vault
Creating Encrypted Files
Editing Encrypted Files
Rekeying Encrypted Files
Encrypting Unencrypted Files
Decrypting Encrypted Files
Viewing Encrypted Files
Use encrypt_string to create encrypted variables to embed in yaml
Vault Ids and Multiple Vault Passwords
Providing Vault Passwords
Multiple vault passwords
Speeding Up Vault Operations
Vault Format
Vault Payload Format 1.1
Working with Patterns
Working With Modules
Introduction
Return Values
Common
Internal use
Module Maintenance & Support
Core
Network
Certified
Community
Issue Reporting
Support
Module Index
All modules
Cloud modules
Clustering modules
Commands modules
Crypto modules
Database modules
Files modules
Identity modules
Inventory modules
Messaging modules
Monitoring modules
Net Tools modules
Network modules
Notification modules
Packaging modules
Remote Management modules
Source Control modules
Storage modules
System modules
Utilities modules
Web Infrastructure modules
Windows modules
Working With Plugins

Action Plugins
Enabling Action Plugins
Using Action Plugins
Plugin List
Cache Plugins
Enabling Cache Plugins
Using Cache Plugins
Plugin List
Callback Plugins
Example Callback Plugins
Enabling Callback Plugins
Plugin List
Connection Plugins
ssh Plugins
Enabling Connection Plugins
Using Connection Plugins
Plugin List
Inventory Plugins
Enabling Inventory Plugins
Using Inventory Plugins
Plugin List
Lookup Plugins
Enabling Lookup Plugins
Using Lookup Plugins
query
Plugin List
Shell Plugins
Enabling Shell Plugins
Using Shell Plugins
Strategy Plugins
Enabling Strategy Plugins
Using Strategy Plugins
Plugin List
Vars Plugins
Enabling Vars Plugins
Using Vars Plugins
Plugin Lists
Filters
Filters For Formatting Data
Forcing Variables To Be Defined
Defaulting Undefined Variables
Omitting Parameters
List Filters
Set Theory Filters
Random Number Filter
Shuffle Filter
Math
JSON Query Filter
IP address filter
Network CLI filters
Network XML filters
Hashing filters
Combining hashes/dictionaries
Extracting values from containers

- Comment Filter
- URL Split Filter
- Regular Expression Filters
- Other Useful Filters
- Combination Filters
- Debugging Filters
- Tests
- Test syntax
- Testing strings
- Version Comparison
- Group theory tests
- Testing paths
- Task results
- Plugin Filter Configuration
- BSD Support
- Working with BSD
- Bootstrapping BSD
- Setting the Python interpreter
- Which modules are available?
- Using BSD as the control machine
- BSD Facts
- BSD Efforts and Contributions
- Windows Guides
- Setting up a Windows Host
- Host Requirements
- WinRM Setup
- Windows Remote Management
- What is WinRM?
- Authentication Options
- Non-Administrator Accounts
- WinRM Encryption
- Inventory Options
- IPv6 Addresses
- HTTPS Certificate Validation
- Limitations
- Using Ansible and Windows
- Use Cases
- Path Formatting for Windows
- Limitations
- Developing Windows Modules
- Desired State Configuration
- What is Desired State Configuration?
- Host Requirements
- Why Use DSC?
- How to Use DSC?
- Custom DSC Resources
- Examples
- Windows Frequently Asked Questions
- Does Ansible work with Windows XP or Server 2003?
- Can I Manage Windows Nano Server?
- Can Ansible run on Windows?
- Can I use SSH keys to authenticate?
- Why can I run a command locally that does not work under Ansible?
- This program won't install with Ansible

What modules are available?
Can I run Python modules?
Can I connect over SSH?
Why is connecting to the host via ssh failing?
Why are my credentials are being rejected?
Why am I getting an error SSL CERTIFICATE_VERIFY_FAILED?
Contributing to Ansible
Ansible Community Guide
Community Code of Conduct
Anti-harassment policy
Policy violations
The Ansible Development Process
Road Maps
Pull Requests
Backport Pull Request Process
Ansibullbot
Overview
Community Maintainers
Workflow
PR Labels
Workflow Labels
Informational Labels
Special Labels
Reporting Bugs And Requesting Features
Reporting A Bug
Requesting a feature
How To Help
Become a power user
Ask and answer questions online
Participate in your local meetup
File and verify issues
Review and submit pull requests
Become a module maintainer
Join a working group
Teach Ansible to others
Social media
Module Maintainer Guidelines
Maintainer Responsibilities
Pull Requests, Issues, and Workflow
Pull Requests
Issues
PR Workflow
Maintainers (BOTMETA.yml)
Changing Maintainership
Tools and other Resources
Release Managers
Pre-releases: What and Why
What is Beta?
What is a Release Candidate?
Release Process
Communicating
Code of Conduct
Mailing List Information
IRC Channel

General Channels
Working Group
Language specific channels
IRC Meetings
Tower Support Questions
Contributors License Agreement
Other Tools And Programs
Extending Ansible
Developer Guide
Ansible Architecture
Modules
Plugins
Inventory
Playbooks
Extending Ansible with Plug-ins and the API
Developing Modules
Welcome
Should You Develop A Module?
How To Develop A Module
Ansible Module Architecture
Types of Modules
Action Plugins
New-style Modules
Non-native want JSON modules
Binary Modules
Old-style Modules
How modules are executed
executor/task_executor
Normal Action Plugin
executor/module_common.py
Special Considerations
Appendix: Module Utilities
Developing Plugins
General Guidelines
Raising Errors
String Encoding
Plugin Configuration
Callback Plugins
Connection Plugins
Inventory Plugins
Lookup Plugins
Vars Plugins
Filter Plugins
Test Plugins
Distributing Plugins
Developing Dynamic Inventory Sources
Script Conventions
Tuning the External Inventory Script
Developing the Ansible Core Engine
Ansible Module Architecture
Types of Modules
How modules are executed
Ansible and Python 3
Minimum Version of Python-3.x and Python-2.x

Porting Controller Code to Python 3
Controller String Strategy
Tips, tricks, and idioms to adopt
Porting Modules to Python 3
Module String Strategy
Tips, tricks, and idioms to adopt
Porting module_utils code to Python 3
Module_utils String Strategy
Python API
Python API example
Rebasing a Pull Request
Configuring Your Remotes
Rebasing Your Branch
Updating Your Pull Request
Getting Help Rebasing
Testing Ansible
Introduction
Types of tests
Testing within GitHub & Shippable
Organization
Rerunning a failing CI job
How to test a PR
Setup: Checking out a Pull Request
Testing the Pull Request
Want to know more about testing?
Repo Merge
Background
Why Are We Doing This (Again?)?
Metadata - Support/Ownership and Module Status
Release and maintenance
Committers Guidelines (for people with commit rights to Ansible on GitHub)
Features, High Level Design, and Roadmap
Our Workflow on GitHub
Addendum to workflow for Committers:
Roles on Core
General Rules
People
Ansible Style Guide
Ansible Style Guide
1. Why Use a Style Guide?
2. Resources
3. Basic Rules
3.1. Use Standard American English
3.2. Write for a Global Audience
3.3. Follow Naming Conventions
3.4. Important Information First
3.5. Sentence Structure
3.6. Avoid padding
3.7. Avoid redundant prepositional phrases
3.8. Avoid verbosity
3.9. Avoid pomposity
3.10. Action verbs, menus, and commands
4. Voice Style
4.1. Active Voice

5. Trademark Usage

5.1. General Rules:

5.2. Guidelines for the proper use of trademarks:

5.3. The importance of Ansible trademarks

5.4. Common Ansible Trademarks

5.5. Other Common Trademarks and Resource Sites:

6. Grammar and Punctuation

6.1. Common Styles and Usage, and Common Mistakes

7. Spelling - Word Usage - Common Words and Phrases to Use and Avoid

7.1. Acronyms

7.2. Applications

7.3. As

7.4. Asks for

7.5. Assure/Ensure/Insure

7.6. Back up

7.7. Backup

7.8. Backward

7.9. Backwards compatibility

7.10. By way of

7.11. Can/May

7.12. CD or cd

7.13. CD-ROM

7.14. Command line

7.15. Daylight saving time (DST)

7.16. Download

7.17. e.g.

7.18. Failover

7.19. Fail over

7.20. Fewer

7.21. File name

7.22. File system

7.23. For instance

7.24. For further/additional/whatever information

7.25. For this reason

7.26. Forward

7.27. Gigabyte (GB)

7.28. Got

7.29. High-availability

7.30. Highly available

7.31. Hostname

7.32. i.e.

7.33. It?s and its

7.34. Less

7.35. Linux

7.36. Login

7.37. Log in

7.38. Log on

7.39. Lots of

7.40. Make sure

7.41. Manual/man page

7.42. MB

7.43. MBps

7.44. MySQL

7.45. Need to

7.46. Read-only
7.47. Real time/real-time
7.48. Refer to
7.49. See
7.50. Since
7.51. Tells
7.52. That/which
7.53. Then/than
7.54. Third-party
7.55. Troubleshoot
7.56. UK
7.57. UNIX®
7.58. Unset
7.59. US
7.60. User
7.61. Username
7.62. View
7.63. Within
7.64. World Wide Web
7.65. Webpage
7.66. Web server
7.67. Website
7.68. Who/whom
7.69. Will
7.70. Wish
7.71. x86
7.72. x86_64
7.73. You
7.74. You may

Scenario Guides

Cisco ACI Guide

What is Cisco ACI ?

Application Centric Infrastructure (ACI)

Application Policy Infrastructure Controller (APIC)

ACI Fabric

More information

Using the ACI modules

Querying ACI configuration

Common parameters

Proxy support

Return values

More information

ACI authentication

Password-based authentication

Signature-based authentication using certificates

Generate certificate and private key

Configure your local user

Use signature-based authentication with Ansible

More information

Using ACI REST with Ansible

Built-in idempotency

Using the aci_rest module

More information

Operational examples

Waiting for all controllers to be ready
Waiting for cluster to be fully-fit
APIC error messages
Known issues
ACI Ansible community
Amazon Web Services Guide
Introduction
Authentication
Provisioning
Host Inventory
Tags And Groups And Variables
Autoscaling with Ansible Pull
Autoscaling with Ansible Tower
Ansible With (And Versus) CloudFormation
AWS Image Building With Ansible
Next Steps: Explore Modules
Microsoft Azure Guide
Requirements
Authenticating with Azure
Using Service Principal
Using Active Directory Username/Password
Providing Credentials to Azure Modules
Using Environment Variables
Storing in a File
Passing as Parameters
Other Cloud Environments
Creating Virtual Machines
Creating Individual Components
Creating a Virtual Machine with Default Options
Dynamic Inventory Script
Host Groups
Examples
Disabling certificate validation on Azure endpoints
CloudStack Cloud Guide
Introduction
Prerequisites
Limitations and Known Issues
Credentials File
Regions
Environment Variables
Use Cases
Use Case: Provisioning in a Advanced Networking CloudStack setup
Use Case: Provisioning on a Basic Networking CloudStack setup
Getting Started with Docker
Requirements
Connecting to the Docker API
Parameters
Environment Variables
Dynamic Inventory Script
Groups
Examples
Configuration
Environment Variables
Configuration File

Google Cloud Platform Guide

Introduction

Credentials

Calling Modules By Passing Credentials

Configuring Modules with secrets.py

Configuring Modules with Environment Variables

GCE Dynamic Inventory

Use Cases

Create an instance

Configuring instances in a group

Using Ansible with the Packet host

Introduction

Requirements

Device Creation

Updating Devices

More Complex Playbooks

Dynamic Inventory Script

Rackspace Cloud Guide

Introduction

Credentials File

Running from a Python Virtual Environment (Optional)

Provisioning

Host Inventory

rax.py

Standard Inventory

Use Cases

Network and Server

Complete Environment

RackConnect and Managed Cloud

Using a Control Machine

Using Ansible Pull

Using Ansible Pull with XenStore

Advanced Usage

Autoscaling with Tower

Orchestration in the Rackspace Cloud

Continuous Delivery and Rolling Upgrades

Introduction

Site Deployment

Reusable Content: Roles

Configuration: Group Variables

The Rolling Upgrade

Managing Other Load Balancers

Continuous Delivery End-To-End

Using Vagrant and Ansible

Introduction

Vagrant Setup

Running Ansible Manually

Advanced Usages

Getting Started with VMware

Introduction

Requirements

vmware_guest module

Debugging

Ansible for Network Automation

Ansible for Network Automation

Introduction

An Introduction to Network Automation with Ansible

Who should use this guide?

Basic Concepts

Control Node

Managed Nodes

Inventory

Modules

Tasks

Playbooks

How Network Automation is Different

Execution on the Control Node

Multiple Communication Protocols

Modules Organized by Network Platform

Privilege Escalation:

authorize

and

become

Run Your First Command and Playbook

Prerequisites

Install Ansible

Establish a Manual Connection to a Managed Node

Run Your First Network Ansible Command

Create and Run Your First Network Ansible Playbook

Build Your Inventory

Basic Inventory

Add Variables to Inventory

Group Variables within Inventory

Variable Syntax

Group Inventory by Platform

Protecting Sensitive Variables with

ansible-vault

Beyond the Basics

Beyond Playbooks: Moving Tasks and Variables into Roles

A Typical Ansible Filetree

Tracking Changes to Inventory and Playbooks: Source Control with Git

Resources and Next Steps

User Guide to to Network Automation with Ansible

Who should use this guide?

Ansible Network FAQ

How can I improve performance for network playbooks?

Why is my output sometimes replaced with

?

Why do the

*_config

modules always return

changed=true

with abbreviated commands?

Network Best Practices for Ansible 2.5

Overview

Concepts

Example

Implementation Notes

Troubleshooting

Network Debug and Troubleshooting Guide

Introduction

How to troubleshoot

Category ?socket_path issue?

Category ?Unable to open shell?

Timeout issues

Playbook issues

Proxy Issues

Working with Command Output in Network Modules

Conditionals in Networking Modules

Platform Options

EOS Platform Options

IOS Platform Options

Junos OS Platform Options

NXOS Platform Options

Settings by Platform

An Introduction to Network Automation with Ansible

Who should use this guide?

Reference & Appendices

Module Index

All modules

Cloud modules

Amazon

Atomic

Azure

Centurylink

Cloudscale

Cloudstack

Digital_Ocean

Dimensiondata

Docker

Google

Linode

Lxc

Lxd

Misc

Oneandone

Openstack

Ovh

Ovirt

Packet

Profitbricks

Pubnub

Rackspace

Smartos

Softlayer

Spotinst

Univention

Vmware

Vultr

Webfaction

Clustering modules

K8S
Openshift
Commands modules
Crypto modules
Database modules
Influxdb
Misc
Mongodb
Mssql
Mysql
Postgresql
Proxysql
Vertica
Files modules
Identity modules
Cyberark
Ipa
Keycloak
Opendj
Inventory modules
Messaging modules
Monitoring modules
Zabbix
Net Tools modules
Basics
Exoscale
Infinity
Ldap
Nios
Network modules
A10
Aci
Aireos
Aos
Aruba
Asa
Avi
Bigswitch
Citrix
Cloudengine
Cloudvision
Cnos
Cumulus
Dellos10
Dellos6
Dellos9
Edgeos
Enos
Eos
F5
Fortimanager
Fortios
Illumos
Interface

ios
iosxr
Ironware
Junos
Layer2
Layer3
Netact
Netconf
Netscaler
Netvisor
Nso
Nuage
Nxos
Onyx
Ordnance
Ovs
Panos
Protocol
Radware
Routing
Sros
System
Vynos
Notification modules
Packaging modules
Language
Os
Remote Management modules
Foreman
Hpilo
Imc
Ipmi
Manageiq
Oneview
Stacki
Ucs
Source Control modules
Storage modules
Infinidat
Netapp
Purestorage
Zfs
System modules
Utilities modules
Helper
Logic
Web Infrastructure modules
Ansible_Tower
Windows modules
Playbook Keywords
Play
Role
Block
Task

Ansible Galaxy
The Website
The command line tool
Installing Roles
roles_path
version
Installing multiple roles from a file
Installing multiple roles from multiple files
Dependencies
Create roles
Force
Container Enabled
Using a Custom Role Skeleton
Search for Roles
Get more information about a role
List installed roles
Remove an installed role
Authenticate with Galaxy
Import a role
Branch
Role name
No wait
Delete a role
Travis integrations
List Travis integrations
Remove Travis integrations
Return Values
Common
backup_file
changed
failed
invocation
msg
rc
results
skipped
stderr
stderr_lines
stdout
stdout_lines
Internal use
ansible_facts
exception
warnings
deprecations
Ansible Configuration Settings
The configuration file
Avoiding security risks with
ansible.cfg
in the current directory
Common Options
ACTION_WARNINGS
AGNOSTIC_BECOME_PROMPT
ALLOW_WORLD_READABLE_TMPFILES

ANSIBLE_COW_PATH
ANSIBLE_COW_SELECTION
ANSIBLE_COW_WHITELIST
ANSIBLE_FORCE_COLOR
ANSIBLE_NOCOLOR
ANSIBLE_NOCOWS
ANSIBLE_PIPELINING
ANSIBLE_SSH_ARGS
ANSIBLE_SSH_CONTROL_PATH
ANSIBLE_SSH_CONTROL_PATH_DIR
ANSIBLE_SSH_EXECUTABLE
ANSIBLE_SSH_RETRIES
ANY_ERRORS_FATAL
BECOME_ALLOW_SAME_USER
CACHE_PLUGIN
CACHE_PLUGIN_CONNECTION
CACHE_PLUGIN_PREFIX
CACHE_PLUGIN_TIMEOUT
COLOR_CHANGED
COLOR_DEBUG
COLOR_DEPRECATE
COLOR_DIFF_ADD
COLOR_DIFF_LINES
COLOR_DIFF_REMOVE
COLOR_ERROR
COLOR_HIGHLIGHT
COLOR_OK
COLOR_SKIP
COLOR_UNREACHABLE
COLOR_VERBOSE
COLOR_WARN
COMMAND_WARNINGS
DEFAULT_ACTION_PLUGIN_PATH
DEFAULT_ALLOW_UNSAFE_LOOKUPS
DEFAULT_ASK_PASS
DEFAULT_ASK_SU_PASS
DEFAULT_ASK_SUDO_PASS
DEFAULT_ASK_VAULT_PASS
DEFAULT_BECOME
DEFAULT_BECOME_ASK_PASS
DEFAULT_BECOME_EXE
DEFAULT_BECOME_FLAGS
DEFAULT_BECOME_METHOD
DEFAULT_BECOME_USER
DEFAULT_CACHE_PLUGIN_PATH
DEFAULT_CALLABLE_WHITELIST
DEFAULT_CALLBACK_PLUGIN_PATH
DEFAULT_CALLBACK_WHITELIST
DEFAULT_CONNECTION_PLUGIN_PATH
DEFAULT_DEBUG
DEFAULT_EXECUTABLE
DEFAULT_FACT_PATH
DEFAULT_FILTER_PLUGIN_PATH
DEFAULT_FORCE_HANDLERS

DEFAULT_FORKS
DEFAULT_GATHER_SUBSET
DEFAULT_GATHER_TIMEOUT
DEFAULT_GATHERING
DEFAULT_HANDLER_INCLUDES_STATIC
DEFAULT_HASH_BEHAVIOUR
DEFAULT_HOST_LIST
DEFAULT_INTERNAL_POLL_INTERVAL
DEFAULT_INVENTORY_PLUGIN_PATH
DEFAULT_JINJA2_EXTENSIONS
DEFAULT_KEEP_REMOTE_FILES
DEFAULT_LIBVIRT_LXC_NOSECLABEL
DEFAULT_LOAD_CALLBACK_PLUGINS
DEFAULT_LOCAL_TMP
DEFAULT_LOG_FILTER
DEFAULT_LOG_PATH
DEFAULT_LOOKUP_PLUGIN_PATH
DEFAULT_MANAGED_STR
DEFAULT_MODULE_ARGS
DEFAULT_MODULE_COMPRESSION
DEFAULT_MODULE_LANG
DEFAULT_MODULE_NAME
DEFAULT_MODULE_PATH
DEFAULT_MODULE_SET_LOCALE
DEFAULT_MODULE_UTILS_PATH
DEFAULT_NO_LOG
DEFAULT_NO_TARGET_SYSLOG
DEFAULT_NULL_REPRESENTATION
DEFAULT_POLL_INTERVAL
DEFAULT_PRIVATE_KEY_FILE
DEFAULT_PRIVATE_ROLE_VARS
DEFAULT_REMOTE_PORT
DEFAULT_REMOTE_USER
DEFAULT_ROLES_PATH
DEFAULT_SCP_IF_SSH
DEFAULT_SELINUX_SPECIAL_FS
DEFAULT_SFTP_BATCH_MODE
DEFAULT_SQUASH_ACTIONS
DEFAULT_SSH_TRANSFER_METHOD
DEFAULT_STDOUT_CALLBACK
DEFAULT_STRATEGY
DEFAULT_STRATEGY_PLUGIN_PATH
DEFAULT_SU
DEFAULT_SU_EXE
DEFAULT_SU_FLAGS
DEFAULT_SU_USER
DEFAULT_SUDO
DEFAULT_SUDO_EXE
DEFAULT_SUDO_FLAGS
DEFAULT_SUDO_USER
DEFAULT_SYSLOG_FACILITY
DEFAULT_TASK_INCLUDES_STATIC
DEFAULT_TEST_PLUGIN_PATH
DEFAULT_TIMEOUT

DEFAULT_TRANSPORT
DEFAULT_UNDEFINED_VAR_BEHAVIOR
DEFAULT_VARS_PLUGIN_PATH
DEFAULT_VAULT_ENCRYPT_IDENTITY
DEFAULT_VAULT_ID_MATCH
DEFAULT_VAULT_IDENTITY
DEFAULT_VAULT_IDENTITY_LIST
DEFAULT_VAULT_PASSWORD_FILE
DEFAULT_VERBOSITY
DEPRECATION_WARNINGS
DIFF_ALWAYS
DIFF_CONTEXT
DISPLAY_ARGS_TO_STDOUT
DISPLAY_SKIPPED_HOSTS
ENABLE_TASK_DEBUGGER
ERROR_ON_MISSING_HANDLER
GALAXY_IGNORE_CERTS
GALAXY_ROLE_SKELETON
GALAXY_ROLE_SKELETON_IGNORE
GALAXY_SERVER
GALAXY_TOKEN
HOST_KEY_CHECKING
INJECT_FACTS_AS_VARS
INVENTORY_ENABLED
INVENTORY_EXPORT
INVENTORY_IGNORE_EXTS
INVENTORY_IGNORE_PATTERNS
INVENTORY_UNPARSED_IS_FAILED
MAX_FILE_SIZE_FOR_DIFF
MERGE_MULTIPLE_CLI_TAGS
NETWORK_GROUP_MODULES
PARAMIKO_HOST_KEY_AUTO_ADD
PARAMIKO_LOOK_FOR_KEYS
PERSISTENT_COMMAND_TIMEOUT
PERSISTENT_CONNECT_RETRY_TIMEOUT
PERSISTENT_CONNECT_TIMEOUT
PERSISTENT_CONTROL_PATH_DIR
PLAYBOOK_VARS_ROOT
PLUGIN_FILTERS_CFG
RETRY_FILES_ENABLED
RETRY_FILES_SAVE_PATH
SHOW_CUSTOM_STATS
STRING_TYPE_FILTERS
SYSTEM_WARNINGS
TAGS_RUN
TAGS_SKIP
USE_PERSISTENT_CONNECTIONS
VARIABLE_PRECEDENCE
YAML_FILENAME_EXTENSIONS
Environment Variables
YAML Syntax
YAML Basics
Gotchas
Python 3 Support

On the controller side
Using Python 3 on the managed machines with commands and playbooks
What to do if an incompatibility is found
Release and maintenance
Testing Strategies
Integrating Testing With Ansible Playbooks
The Right Level of Testing
Check Mode As A Drift Test
Modules That Are Useful for Testing
Testing Lifecycle
Integrating Testing With Rolling Updates
Achieving Continuous Deployment
Conclusion
Sanity Tests
Sanity Tests » ansible-doc
Sanity Tests » azure-requirements
Sanity Tests » boilerplate
Sanity Tests » compile
Sanity Tests » configure-remoting-ps1
Sanity Tests » empty-init
Sanity Tests » import
Sanity Tests » integration-aliases
Groups
Requirements
Skipping
Unstable
Disabled
Unsupported
Cloud
Untested
Issues
Questions
Sanity Tests » line-endings
Sanity Tests » no-assert
Sanity Tests » no-basestring
Sanity Tests » no-dict-iteritems
Sanity Tests » no-dict-iterkeys
Sanity Tests » no-dict-itervalues
Sanity Tests » no-get-exception
Sanity Tests » no-illegal-filenames
Illegal Characters
Illegal Names
Sanity Tests » no-smart-quotes
Sanity Tests » no-tests-as-filters
Sanity Tests » no-underscore-variable
Sanity Tests » no-unicode_literals
Sanity Tests » pep8
Sanity Tests » ps lint
Sanity Tests » pylint
Sanity Tests » replace-urlopen
Sanity Tests » required-and-default-attributes
Sanity Tests » rstcheck
Sanity Tests » sanity-docs
Sanity Tests » shebang

[Sanity Tests » shellcheck](#)

[Sanity Tests » symlinks](#)

[Sanity Tests » test-constraints](#)

[Sanity Tests » use-argspec-type-path](#)

[Sanity Tests » use-compatible-six](#)

[Sanity Tests » validate-modules](#)

[Sanity Tests » yamllint](#)

[Frequently Asked Questions](#)

[How can I set the PATH or any other environment variable for a task or entire playbook?](#)

[How do I handle different machines needing different user accounts or ports to log in with?](#)

[How do I get ansible to reuse connections, enable Kerberized SSH, or have Ansible pay attention to my local SSH config file?](#)

[How do I configure a jump host to access servers that I have no direct access to?](#)

[How do I speed up management inside EC2?](#)

[How do I handle python not having a Python interpreter at /usr/bin/python on a remote machine?](#)

[Common Platform Issues](#)

[Running in a virtualenv](#)

[Running on BSD](#)

[Running on Solaris](#)

[What is the best way to make content reusable/redistributable?](#)

[Where does the configuration file live and what can I configure in it?](#)

[How do I disable cowsay?](#)

[How do I see a list of all of the ansible_ variables?](#)

[How do I see all the inventory vars defined for my host?](#)

[How do I loop over a list of hosts in a group, inside of a template?](#)

[How do I access a variable name programmatically?](#)

[How do I access a variable of the first host in a group?](#)

[How do I copy files recursively onto a target host?](#)

[How do I access shell environment variables?](#)

[How do I generate crypted passwords for the user module?](#)

[Ansible supports dot notation and array notation for variables. Which notation should I use?](#)

[Can I get training on Ansible?](#)

[Is there a web interface / REST API / etc?](#)

[How do I submit a change to the documentation?](#)

[How do I keep secret data in my playbook?](#)

[When should I use {{ }}? Also, how to interpolate variables or dynamic variable names](#)

[Why don't you ship in X format?](#)

[I don't see my question here](#)

[Glossary](#)

[Ansible Tower](#)

[Roadmaps](#)

[Ansible Roadmap](#)

[Ansible 2.1](#)

[Windows](#)

[Network](#)

[VMware](#)

[Azure](#)

[Docker](#)

[Cloud](#)

[Ansiballz](#)

[Diff-support](#)

[Other](#)

[Community](#)

[Ansible 2.2](#)

Docker
Extras split from Core
Tweaks/Fixes
AWS
Google
OpenStack
Azure load balancer
VMware
Windows
Network
Role revamp
Vault
Python3
Infrastructure Buildout and Changes
Ansible 2.3
General Comments from the Core Team
Repo Merge
Metadata
Documentation
Windows
Azure
Networking
Python3
Testing and CI
Amazon
Plugin Loader
ansible-ssh
Ansible 2.4
Administrivia and Process
Python 2.4 and 2.5 support discontinuation
Python 3
Ansible-Config
Inventory
Facts
PluginLoader
Static Loop Keyword
Vault
Globalize Callbacks
Plugins
Group Priorities
Runtime Check on Modules for Blacklisting
Disambiguate Includes
Windows
AWS
Azure
Google Cloud Platform
Network Roadmap
Contributor Quality of Life
Ansible 2.5
Engine improvements
Ansible Content Management
Ansible-Config
Inventory
Facts

PluginLoader
Static Loop Keyword
Vault
Role Versioning
Globalize Callbacks
Runtime Check on Modules for Blacklisting
Windows
General Cloud
AWS
Azure
Network Roadmap
Documentation
Contributor Quality of Life
Ansible Documentation
Docs
»

Ansible Documentation

You are reading an unmaintained version of the Ansible documentation. Unmaintained Ansible versions can contain unfixed security vulnerabilities (CVE). Please upgrade to a maintained version. See [the latest Ansible documentation](#)

.

Ansible Documentation



About Ansible



Ansible is an IT automation tool. It can configure systems, deploy software, and orchestrate more advanced IT tasks such as continuous deployments or zero downtime rolling updates.

Ansible's main goals are simplicity and ease-of-use. It also has a strong focus on security and reliability, featuring a minimum of moving parts, usage of OpenSSH for transport (with other transports and pull modes as alternatives), and a language that is designed around auditability by humans—even those not familiar with the program.

We believe simplicity is relevant to all sizes of environments, so we design for busy users of all types: developers, sysadmins, release engineers, IT managers, and everyone in between. Ansible is appropriate for managing all environments, from small setups with a handful of instances to enterprise environments with many thousands of instances.

Ansible manages machines in an agent-less manner. There is never a question of how to upgrade remote daemons or the problem of not being able to manage systems because daemons are uninstalled. Because OpenSSH is one of the most peer-reviewed open source components, security exposure is greatly reduced. Ansible is decentralized—it relies on your existing OS credentials to control access to remote machines. If needed, Ansible can easily connect with Kerberos, LDAP, and other centralized authentication management systems.

This documentation covers the current released version of Ansible (2.5) and also some development version features. For recent features, we note in each section the version of Ansible where the feature was added.

Ansible releases a new major release of Ansible approximately every two months. The core application evolves somewhat conservatively, valuing simplicity in language design and setup. However, the community around new modules and plugins being developed and contributed moves very quickly, adding many new modules in each release.

Installation, Upgrade & Configuration

Installation Guide

Basics / What Will Be Installed

What Version To Pick?

Control Machine Requirements

Managed Node Requirements

Installing the Control Machine

Ansible on GitHub

Configuring Ansible

Configuration file

Environmental configuration
Command line options
Ansible Porting Guides
Using Ansible
User Guide
Ansible Quickstart
Getting Started
Working with Command Line Tools
Introduction To Ad-Hoc Commands
Working with Inventory
Working With Dynamic Inventory
Working With Playbooks
Understanding Privilege Escalation
Ansible Vault
Working with Patterns
Working With Modules
Working With Plugins
BSD Support
Windows Guides
Contributing to Ansible
Ansible Community Guide
Community Code of Conduct
The Ansible Development Process
Road Maps
Pull Requests
Ansibullbot
Reporting Bugs And Requesting Features
How To Help
Module Maintainer Guidelines
Release Managers
Release Process
Communicating
Contributors License Agreement
Other Tools And Programs
Extending Ansible
Developer Guide
Ansible Architecture
Developing Modules
Ansible Module Architecture
Appendix: Module Utilities
Developing Plugins
Developing Dynamic Inventory Sources
Developing the Ansible Core Engine
Ansible and Python 3
Python API
Rebasing a Pull Request
Testing Ansible
Repo Merge
Release and maintenance
Committers Guidelines (for people with commit rights to Ansible on GitHub)
Ansible Style Guide
Scenario Guides
Cisco ACI Guide
What is Cisco ACI ?

Using the ACI modules
ACI authentication
Using ACI REST with Ansible
Operational examples
APIC error messages
Known issues
ACI Ansible community
Amazon Web Services Guide
Introduction
Authentication
Provisioning
Host Inventory
Tags And Groups And Variables
Autoscaling with Ansible Pull
Autoscaling with Ansible Tower
Ansible With (And Versus) CloudFormation
AWS Image Building With Ansible
Next Steps: Explore Modules
Microsoft Azure Guide
Requirements
Authenticating with Azure
Other Cloud Environments
Creating Virtual Machines
Dynamic Inventory Script
CloudStack Cloud Guide
Introduction
Prerequisites
Limitations and Known Issues
Credentials File
Regions
Environment Variables
Use Cases
Getting Started with Docker
Requirements
Connecting to the Docker API
Dynamic Inventory Script
Google Cloud Platform Guide
Introduction
Credentials
GCE Dynamic Inventory
Use Cases
Using Ansible with the Packet host
Introduction
Requirements
Device Creation
Updating Devices
More Complex Playbooks
Dynamic Inventory Script
Rackspace Cloud Guide
Introduction
Credentials File
Provisioning
Host Inventory
Use Cases

Advanced Usage
Continuous Delivery and Rolling Upgrades
Introduction
Site Deployment
Reusable Content: Roles
Configuration: Group Variables
The Rolling Upgrade
Managing Other Load Balancers
Continuous Delivery End-To-End
Using Vagrant and Ansible
Introduction
Vagrant Setup
Running Ansible Manually
Advanced Usages
Getting Started with VMware
Introduction
Requirements
vmware_guest module
Debugging
Ansible for Network Automation
Ansible for Network Automation
Introduction
An Introduction to Network Automation with Ansible
Who should use this guide?
Reference & Appendices
Module Index
All modules
Cloud modules
Clustering modules
Commands modules
Crypto modules
Database modules
Files modules
Identity modules
Inventory modules
Messaging modules
Monitoring modules
Net Tools modules
Network modules
Notification modules
Packaging modules
Remote Management modules
Source Control modules
Storage modules
System modules
Utilities modules
Web Infrastructure modules
Windows modules
Playbook Keywords
Play
Role
Block
Task
Ansible Galaxy

The Website
The command line tool
Return Values
Common
Internal use
Ansible Configuration Settings
The configuration file
Common Options
Environment Variables
YAML Syntax
YAML Basics
Gotchas
Python 3 Support
On the controller side
Using Python 3 on the managed machines with commands and playbooks
What to do if an incompatibility is found
Release and maintenance
Testing Strategies
Integrating Testing With Ansible Playbooks
The Right Level of Testing
Check Mode As A Drift Test
Modules That Are Useful for Testing
Testing Lifecycle
Integrating Testing With Rolling Updates
Achieving Continuous Deployment
Conclusion
Sanity Tests
Sanity Tests » ansible-doc
Sanity Tests » azure-requirements
Sanity Tests » boilerplate
Sanity Tests » compile
Sanity Tests » configure-remoting-ps1
Sanity Tests » empty-init
Sanity Tests » import
Sanity Tests » integration-aliases
Sanity Tests » line-endings
Sanity Tests » no-assert
Sanity Tests » no-basestring
Sanity Tests » no-dict-iteritems
Sanity Tests » no-dict-iterkeys
Sanity Tests » no-dict-itervalues
Sanity Tests » no-get-exception
Sanity Tests » no-illegal-filenames
Sanity Tests » no-smart-quotes
Sanity Tests » no-tests-as-filters
Sanity Tests » no-underscore-variable
Sanity Tests » no-unicode_literals
Sanity Tests » pep8
Sanity Tests » ps lint
Sanity Tests » pylint
Sanity Tests » replace-urlopen
Sanity Tests » required-and-default-attributes
Sanity Tests » rstcheck
Sanity Tests » sanity-docs

[Sanity Tests » shebang](#)

[Sanity Tests » shellcheck](#)

[Sanity Tests » symlinks](#)

[Sanity Tests » test-constraints](#)

[Sanity Tests » use-argspec-type-path](#)

[Sanity Tests » use-compat-six](#)

[Sanity Tests » validate-modules](#)

[Sanity Tests » yamllint](#)

[Frequently Asked Questions](#)

[How can I set the PATH or any other environment variable for a task or entire playbook?](#)

[How do I handle different machines needing different user accounts or ports to log in with?](#)

[How do I get ansible to reuse connections, enable Kerberized SSH, or have Ansible pay attention to my local SSH config file?](#)

[How do I configure a jump host to access servers that I have no direct access to?](#)

[How do I speed up management inside EC2?](#)

[How do I handle python not having a Python interpreter at /usr/bin/python on a remote machine?](#)

[Common Platform Issues](#)

[What is the best way to make content reusable/redistributable?](#)

[Where does the configuration file live and what can I configure in it?](#)

[How do I disable cowsay?](#)

[How do I see a list of all of the ansible_ variables?](#)

[How do I see all the inventory vars defined for my host?](#)

[How do I loop over a list of hosts in a group, inside of a template?](#)

[How do I access a variable name programmatically?](#)

[How do I access a variable of the first host in a group?](#)

[How do I copy files recursively onto a target host?](#)

[How do I access shell environment variables?](#)

[How do I generate crypted passwords for the user module?](#)

[Ansible supports dot notation and array notation for variables. Which notation should I use?](#)

[Can I get training on Ansible?](#)

[Is there a web interface / REST API / etc?](#)

[How do I submit a change to the documentation?](#)

[How do I keep secret data in my playbook?](#)

[When should I use {{ }}? Also, how to interpolate variables or dynamic variable names](#)

[Why don't you ship in X format?](#)

[I don't see my question here](#)

[Glossary](#)

[Ansible Tower](#)

[Roadmaps](#)

[Ansible Roadmap](#)

[Ansible 2.1](#)

[Ansible 2.2](#)

[Ansible 2.3](#)

[Ansible 2.4](#)

[Ansible 2.5](#)

[Next](#)

[Last updated on Nov 29, 2022.](#)

[Copyright © 2018 Red Hat, Inc.](#)

Content from: <https://docs.ansible.com/archive/ansible/2.6/>

Ansible Documentation ? Ansible Documentation

AnsibleFest

Products

Community

Webinars & Training

Blog

Documentation

Ansible

2.6

Installation, Upgrade & Configuration

Installation Guide

Configuring Ansible

Ansible Porting Guides

Using Ansible

User Guide

Contributing to Ansible

Ansible Community Guide

Extending Ansible

Developer Guide

Scenario Guides

Cisco ACI Guide

Amazon Web Services Guide

Microsoft Azure Guide

CloudStack Cloud Guide

Getting Started with Docker

Google Cloud Platform Guide

Using Ansible with the Packet host

Rackspace Cloud Guide

Continuous Delivery and Rolling Upgrades

Using Vagrant and Ansible

Ansible for VMWare

Ansible for VMware

Ansible for Network Automation

Ansible for Network Automation

Reference & Appendices

Module Index

Playbook Keywords

Ansible Galaxy

Return Values

Ansible Configuration Settings

YAML Syntax

Python 3 Support

Release and maintenance

Testing Strategies

Sanity Tests

Frequently Asked Questions

Glossary

Ansible Tower

Roadmaps

Ansible Roadmap

Ansible Documentation

Docs

»

Ansible Documentation

You are reading an unmaintained version of the Ansible documentation. Unmaintained Ansible versions can contain unfixed security vulnerabilities (CVE). Please upgrade to a maintained version. See the latest Ansible documentation

.

Ansible Documentation



About Ansible



Ansible is an IT automation tool. It can configure systems, deploy software, and orchestrate more advanced IT tasks such as continuous deployments or zero downtime rolling updates.

Ansible's main goals are simplicity and ease-of-use. It also has a strong focus on security and reliability, featuring a minimum of moving parts, usage of OpenSSH for transport (with other transports and pull modes as alternatives), and a language that is designed around auditability by humans—even those not familiar with the program.

We believe simplicity is relevant to all sizes of environments, so we design for busy users of all types: developers, sysadmins, release engineers, IT managers, and everyone in between. Ansible is appropriate for managing all environments, from small setups with a handful of instances to enterprise environments with many thousands of instances.

Ansible manages machines in an agent-less manner. There is never a question of how to upgrade remote daemons or the problem of not being able to manage systems because daemons are uninstalled. Because OpenSSH is one of the most peer-reviewed open source components, security exposure is greatly reduced. Ansible is decentralized—it relies on your existing OS credentials to control access to remote machines. If needed, Ansible can easily connect with Kerberos, LDAP, and other centralized authentication management systems.

This documentation covers the current released version of Ansible (2.5) and also some development version features.

For recent features, we note in each section the version of Ansible where the feature was added.

Ansible releases a new major release of Ansible approximately every two months. The core application evolves somewhat conservatively, valuing simplicity in language design and setup. However, the community around new modules and plugins being developed and contributed moves very quickly, adding many new modules in each release.

Installation, Upgrade & Configuration

Installation Guide

Basics / What Will Be Installed

What Version To Pick?

Control Machine Requirements

Managed Node Requirements

Installing the Control Machine

Ansible on GitHub

Configuring Ansible

Configuration file

Environmental configuration

Command line options

Ansible Porting Guides

Using Ansible

User Guide

Ansible Quickstart

Getting Started

Working with Command Line Tools

Introduction To Ad-Hoc Commands

Working with Inventory

Working With Dynamic Inventory

Working With Playbooks

Understanding Privilege Escalation

Ansible Vault

Working with Patterns

Working With Modules
Working With Plugins
BSD Support
Windows Guides
Contributing to Ansible
Ansible Community Guide
Community Code of Conduct
The Ansible Development Process
Road Maps
Pull Requests
Ansibullbot
Reporting Bugs And Requesting Features
How To Help
Module Maintainer Guidelines
Release Managers
Release Process
Communicating
Contributors License Agreement
Other Tools And Programs
Extending Ansible
Developer Guide
Ansible Architecture
Developing Modules
Ansible Module Architecture
Appendix: Module Utilities
Developing Plugins
Developing Dynamic Inventory Sources
Developing the Ansible Core Engine
Ansible and Python 3
Python API
Rebasing a Pull Request
Testing Ansible
Repo Merge
Release and maintenance
Committers Guidelines (for people with commit rights to Ansible on GitHub)
Ansible Style Guide
Scenario Guides
Cisco ACI Guide
What is Cisco ACI ?
Using the ACI modules
ACI authentication
Using ACI REST with Ansible
Operational examples
APIC error messages
Known issues
ACI Ansible community
Amazon Web Services Guide
Introduction
Authentication
Provisioning
Host Inventory
Tags And Groups And Variables
Autoscaling with Ansible Pull
Autoscaling with Ansible Tower

Ansible With (And Versus) CloudFormation
AWS Image Building With Ansible
Next Steps: Explore Modules
Microsoft Azure Guide
Requirements
Authenticating with Azure
Other Cloud Environments
Creating Virtual Machines
Dynamic Inventory Script
CloudStack Cloud Guide
Introduction
Prerequisites
Limitations and Known Issues
Credentials File
Regions
Environment Variables
Use Cases
Getting Started with Docker
Requirements
Connecting to the Docker API
Dynamic Inventory Script
Google Cloud Platform Guide
Introduction
Introduction
Credentials
GCE Dynamic Inventory
Using Ansible with the Packet host
Introduction
Requirements
Device Creation
Updating Devices
More Complex Playbooks
Dynamic Inventory Script
Rackspace Cloud Guide
Introduction
Credentials File
Provisioning
Host Inventory
Use Cases
Advanced Usage
Continuous Delivery and Rolling Upgrades
Introduction
Site Deployment
Reusable Content: Roles
Configuration: Group Variables
The Rolling Upgrade
Managing Other Load Balancers
Continuous Delivery End-To-End
Using Vagrant and Ansible
Introduction
Vagrant Setup
Running Ansible Manually
Advanced Usages
Ansible for VMWare

Ansible for VMware
Introduction to Ansible for VMware
Ansible for VMware Concepts
VMware Prerequisites
Getting Started with Ansible for VMware
Ansible for VMware Scenarios
Ansible VMware Module Guide
Troubleshooting Ansible for VMware
List of useful links to VMware
Ansible VMware FAQ
Ansible for Network Automation
Ansible for Network Automation
Getting Started with Ansible for Network Automation
Advanced Topics with Ansible for Network Automation
Reference & Appendices
Module Index
All modules
Cloud modules
Clustering modules
Commands modules
Crypto modules
Database modules
Files modules
Identity modules
Inventory modules
Messaging modules
Monitoring modules
Net Tools modules
Network modules
Notification modules
Packaging modules
Remote Management modules
Source Control modules
Storage modules
System modules
Utilities modules
Web Infrastructure modules
Windows modules
Playbook Keywords
Play
Role
Block
Task
Ansible Galaxy
The Website
The command line tool
Return Values
Common
Internal use
Ansible Configuration Settings
The configuration file
Common Options
Environment Variables
YAML Syntax

YAML Basics
Gotchas
Python 3 Support
On the controller side
Using Python 3 on the managed machines with commands and playbooks
What to do if an incompatibility is found
Release and maintenance
Testing Strategies
Integrating Testing With Ansible Playbooks
The Right Level of Testing
Check Mode As A Drift Test
Modules That Are Useful for Testing
Testing Lifecycle
Integrating Testing With Rolling Updates
Achieving Continuous Deployment
Conclusion
Sanity Tests
ansible-doc
azure-requirements
boilerplate
changelog
compile
configure-remoting-ps1
empty-init
import
integration-aliases
line-endings
no-assert
no-basestring
no-dict-iteritems
no-dict-iterkeys
no-dict-itervalues
no-get-exception
no-illegal-filenames
no-smart-quotes
no-tests-as-filters
no-underscore-variable
no-unicode_literals
pep8
pslint
pylint
replace-urlopen
required-and-default-attributes
rstcheck
sanity-docs
shebang
shellcheck
symlinks
test-constraints
use-argspec-type-path
use-compatible-six
validate-modules
yamllint
Frequently Asked Questions

How can I set the PATH or any other environment variable for a task or entire playbook?

How do I handle different machines needing different user accounts or ports to log in with?

How do I get ansible to reuse connections, enable Kerberized SSH, or have Ansible pay attention to my local SSH config file?

How do I configure a jump host to access servers that I have no direct access to?

How do I speed up management inside EC2?

How do I handle python not having a Python interpreter at /usr/bin/python on a remote machine?

Common Platform Issues

What is the best way to make content reusable/redistributable?

Where does the configuration file live and what can I configure in it?

How do I disable cowsay?

How do I see a list of all of the ansible_ variables?

How do I see all the inventory vars defined for my host?

How do I loop over a list of hosts in a group, inside of a template?

How do I access a variable name programmatically?

How do I access a variable of the first host in a group?

How do I copy files recursively onto a target host?

How do I access shell environment variables?

How do I generate crypted passwords for the user module?

Ansible supports dot notation and array notation for variables. Which notation should I use?

Can I get training on Ansible?

Is there a web interface / REST API / etc?

How do I submit a change to the documentation?

How do I keep secret data in my playbook?

When should I use {{ }}? Also, how to interpolate variables or dynamic variable names

Why don't you ship in X format?

I don't see my question here

Glossary

Ansible Tower

Roadmaps

Ansible Roadmap

Ansible 2.6

Ansible 2.5

Ansible 2.4

Ansible 2.3

Ansible 2.2

Ansible 2.1

Next

Last updated on Jan 30, 2023.

Copyright © 2018 Red Hat, Inc.

Content from: <https://docs.ansible.com/archive/ansible/2.7/>

Ansible Documentation ? Ansible Documentation

AnsibleFest

Products

Community

Webinars & Training

Blog

Documentation

Ansible

2.7

Installation, Upgrade & Configuration

Installation Guide

Configuring Ansible

Ansible Porting Guides

Using Ansible

User Guide

Contributing to Ansible

Ansible Community Guide

Extending Ansible

Developer Guide

Scenario Guides

Cisco ACI Guide

Amazon Web Services Guide

Microsoft Azure Guide

CloudStack Cloud Guide

Getting Started with Docker

Google Cloud Platform Guide

Infoblox Guide

Getting Started with Kubernetes and OpenShift

Cisco Meraki Guide

Using Ansible with the Packet host

Rackspace Cloud Guide

Continuous Delivery and Rolling Upgrades

Using Vagrant and Ansible

Vultr Guide

Ansible for VMWare

Ansible for VMware

Ansible for Network Automation

Ansible for Network Automation

Reference & Appendices

Module Index

Playbook Keywords

Ansible Galaxy

Return Values

Ansible Configuration Settings

YAML Syntax

Python 3 Support

Release and maintenance

Testing Strategies

Sanity Tests

Frequently Asked Questions

Glossary

Ansible Reference: Module Utilities

Special Variables
Ansible Tower
Roadmaps
Ansible Roadmap
Ansible
Docs
»

Ansible Documentation

You are reading an unmaintained version of the Ansible documentation. Unmaintained Ansible versions can contain unfixed security vulnerabilities (CVE). Please upgrade to a maintained version. See [the latest Ansible documentation](#)

.

Ansible Documentation



About Ansible



Ansible is an IT automation tool. It can configure systems, deploy software, and orchestrate more advanced IT tasks such as continuous deployments or zero downtime rolling updates.

Ansible's main goals are simplicity and ease-of-use. It also has a strong focus on security and reliability, featuring a minimum of moving parts, usage of OpenSSH for transport (with other transports and pull modes as alternatives), and a language that is designed around auditability by humans—even those not familiar with the program.

We believe simplicity is relevant to all sizes of environments, so we design for busy users of all types: developers, sysadmins, release engineers, IT managers, and everyone in between. Ansible is appropriate for managing all environments, from small setups with a handful of instances to enterprise environments with many thousands of instances.

Ansible manages machines in an agent-less manner. There is never a question of how to upgrade remote daemons or the problem of not being able to manage systems because daemons are uninstalled. Because OpenSSH is one of the most peer-reviewed open source components, security exposure is greatly reduced. Ansible is decentralized—it relies on your existing OS credentials to control access to remote machines. If needed, Ansible can easily connect with Kerberos, LDAP, and other centralized authentication management systems.

This documentation covers the version of Ansible noted in the upper left corner of this page. We maintain multiple versions of Ansible and of the documentation, so please be sure you are using the version of the documentation that covers the version of Ansible you're using. For recent features, we note the version of Ansible where the feature was added.

Ansible releases a new major release of Ansible approximately three to four times per year. The core application evolves somewhat conservatively, valuing simplicity in language design and setup. However, the community around new modules and plugins being developed and contributed moves very quickly, adding many new modules in each release.

Installation, Upgrade & Configuration

Installation Guide

Basics / What Will Be Installed

What Version To Pick?

Control Machine Requirements

Managed Node Requirements

Installing the Control Machine

Ansible on GitHub

Configuring Ansible

Configuration file

Environmental configuration

Command line options

Ansible Porting Guides

Using Ansible

User Guide

Ansible Quickstart

Getting Started
Working with Command Line Tools
Introduction To Ad-Hoc Commands
Working with Inventory
Working With Dynamic Inventory
Working With Playbooks
Understanding Privilege Escalation
Ansible Vault
Working with Patterns
Working With Modules
Working With Plugins
BSD Support
Windows Guides
Contributing to Ansible
Ansible Community Guide
Community Code of Conduct
How can I help?
Reporting Bugs And Requesting Features
Contributing to the Ansible Documentation
Communicating
The Ansible Development Process
Contributors License Agreement
Triage Process
Other Tools And Programs
Ansible Style Guide
Committers Guidelines
Module Maintainer Guidelines
Release manager guidelines
GitHub Admins
Extending Ansible
Developer Guide
Adding modules and plugins locally
Should you develop a module?
Ansible module development: getting started
Contributing your module to Ansible
Conventions, tips, and pitfalls
Ansible and Python 3
Debugging modules
Module format and documentation
Windows module development walkthrough
Developing Cisco ACI modules
Information for submitting a group of modules
Testing Ansible
The lifecycle of an Ansible module
Developing plugins
Developing dynamic inventory
Developing the Ansible Core Engine
Ansible module architecture
Python API
Rebasing a pull request
Appendix: Module Utilities
Ansible Architecture
Scenario Guides
Cisco ACI Guide

What is Cisco ACI ?
Using the ACI modules
ACI authentication
Using ACI REST with Ansible
Operational examples
APIC error messages
Known issues
ACI Ansible community
Amazon Web Services Guide
Introduction
Authentication
Provisioning
Host Inventory
Tags And Groups And Variables
Autoscaling with Ansible Pull
Autoscaling with Ansible Tower
Ansible With (And Versus) CloudFormation
AWS Image Building With Ansible
Next Steps: Explore Modules
Microsoft Azure Guide
Requirements
Authenticating with Azure
Other Cloud Environments
Creating Virtual Machines
Dynamic Inventory Script
CloudStack Cloud Guide
Introduction
Prerequisites
Limitations and Known Issues
Credentials File
Regions
Environment Variables
Use Cases
Getting Started with Docker
Requirements
Connecting to the Docker API
Dynamic Inventory Script
Google Cloud Platform Guide
Introduction
Introduction
Credentials
GCE Dynamic Inventory
Infoblox Guide
Prerequisites
Credentials and authenticating
NIOS lookup plugins
Use cases with modules
Dynamic inventory script
Getting Started with Kubernetes and OpenShift
Requirements
Installation and use
Authenticating with the API
Cisco Meraki Guide
What is Cisco Meraki?

Using the Meraki modules
Common Parameters
Meraki Authentication
Returned Data Structures
Handling Returned Data
Error Handling
Using Ansible with the Packet host
Introduction
Requirements
Device Creation
Updating Devices
More Complex Playbooks
Dynamic Inventory Script
Rackspace Cloud Guide
Introduction
Credentials File
Provisioning
Host Inventory
Use Cases
Advanced Usage
Continuous Delivery and Rolling Upgrades
Introduction
Site Deployment
Reusable Content: Roles
Configuration: Group Variables
The Rolling Upgrade
Managing Other Load Balancers
Continuous Delivery End-To-End
Using Vagrant and Ansible
Introduction
Vagrant Setup
Running Ansible Manually
Advanced Usages
Vultr Guide
Requirements
Configuration
Authentication
Usage
Dynamic Inventory
Ansible for VMWare
Ansible for VMware
Introduction to Ansible for VMware
Ansible for VMware Concepts
VMware Prerequisites
Getting Started with Ansible for VMware
Using VMware dynamic inventory plugin
Ansible for VMware Scenarios
Ansible VMware Module Guide
Troubleshooting Ansible for VMware
List of useful links to VMware
Ansible VMware FAQ
Ansible for Network Automation
Ansible for Network Automation
Getting Started with Ansible for Network Automation

[Advanced Topics with Ansible for Network Automation](#)

[Reference & Appendices](#)

[Module Index](#)

[Playbook Keywords](#)

[Ansible Galaxy](#)

[Return Values](#)

[Ansible Configuration Settings](#)

[YAML Syntax](#)

[Python 3 Support](#)

[Release and maintenance](#)

[Testing Strategies](#)

[Sanity Tests](#)

[Frequently Asked Questions](#)

[Glossary](#)

[Ansible Reference: Module Utilities](#)

[Special Variables](#)

[Ansible Tower](#)

[Roadmaps](#)

[Ansible Roadmap](#)

[Ansible 2.7](#)

[Ansible 2.6](#)

[Ansible 2.5](#)

[Ansible 2.4](#)

[Ansible 2.3](#)

[Ansible 2.2](#)

[Ansible 2.1](#)

[Next](#)

© Copyright 2019 Red Hat, Inc.

Last updated on Mar 31, 2023.

Content from: <https://docs.ansible.com/automation-controller/>

Index of /automation-controller
Index of /automation-controller
Name
Last modified
Size
Description
Parent Directory

-
- 4.0.0/
2022-12-09 03:20
-
- 4.0.1/
2022-12-09 04:04
-
- 4.1.0/
2025-10-18 22:24
-
- 4.1.1/
2025-10-18 22:55
-
- 4.1.2/
2025-10-19 22:03
-
- 4.1.3/
2025-10-18 22:13
-
- 4.1.4/
2025-10-18 22:35
-
- 4.2.0/
2025-10-18 22:39
-
- 4.2.1/
2025-10-18 22:52
-
- 4.2.2/
2025-10-18 22:18
-
- 4.3.0_old/
2022-12-16 03:15
-
- 4.3/
2025-10-18 22:33
-
- 4.4.0/
2022-11-30 18:41
-
- 4.4/
2025-10-18 22:49
-
- 4.5/
2025-10-19 22:01

-
4.6/
2025-10-18 22:05

-
4.7/
2025-09-10 05:28

-
latest/
2025-10-18 22:05

-

Content from: <https://docs.ansible.com/automation-controller/4.1.2/>

Index of /automation-controller/4.1.2

Index of /automation-controller/4.1.2

Name

Last modified

Size

Description

Parent Directory

-

html/

2025-10-18 22:02

-

html_ja/

2022-04-21 21:33

-

html_zh/

2022-04-21 21:46

-

pdf/

2023-02-10 05:45

-

Content from: <https://docs.ansible.com/automation-controller/4.1.2/html/>

Index of /automation-controller/4.1.2/html

Index of /automation-controller/4.1.2/html

Name
Last modified
Size
Description
Parent Directory
-
administration/
2025-10-18 22:02
-
controllerapi/
2025-10-18 22:03
-
controllercli/
2022-05-16 16:17
-
quickstart/
2025-10-18 22:03
-
release-notes/
2025-10-18 22:03
-
towerapi/
2025-10-18 22:03
-
towercli/
2022-05-16 16:17
-
upgrade-migration-gu..>
2025-10-18 22:03
-
userguide/
2025-10-18 22:03
-

Content from: <https://docs.ansible.com/automation-controller/4.1.3/html/>

Index of /automation-controller/4.1.3/html

Index of /automation-controller/4.1.3/html

Name
Last modified
Size
Description
Parent Directory

-
- administration/
2025-10-18 22:13
-
- controllerapi/
2025-10-18 22:13
-
- quickstart/
2025-10-18 22:13
-
- release-notes/
2025-10-18 22:13
-
- towerapi/
2025-10-18 22:13
-
- upgrade-migration-gu..>
2025-10-18 22:13
-
- userguide/
2025-10-18 22:13
-

Content from: <https://docs.ansible.com/automation-controller/4.2.0/>

Index of /automation-controller/4.2.0

Index of /automation-controller/4.2.0

Name
Last modified
Size
Description
Parent Directory

-
access/
2022-09-01 02:11
-
html/
2025-10-18 22:39
-
html_ja/
2022-05-26 06:24
-
html_ko/
2022-05-26 06:10
-
html_zh/
2022-05-26 06:39
-
pdf/
2023-02-10 05:34
-

Content from: <https://docs.ansible.com/automation-controller/4.2.0/html/>

Index of /automation-controller/4.2.0/html

Index of /automation-controller/4.2.0/html

Name
Last modified
Size
Description
Parent Directory
-
administration/
2025-10-18 22:39
-
controllerapi/
2025-10-18 22:39
-
controllercli/
2022-06-02 18:55
-
quickstart/
2025-10-18 22:39
-
release-notes/
2025-10-18 22:39
-
towerapi/
2025-10-18 22:39
-
towercli/
2022-06-02 18:55
-
upgrade-migration-gu..>
2025-10-18 22:39
-
userguide/
2025-10-18 22:39
-

Content from: <https://docs.ansible.com/automation-controller/4.2.1/>

Index of /automation-controller/4.2.1

Index of /automation-controller/4.2.1

Name
Last modified
Size
Description
Parent Directory

-
html/
2025-10-18 22:52
-
html_ja/
2022-12-07 19:28
-
html_ko/
2022-12-07 17:58
-
html_zh/
2022-12-07 18:43
-
pdf/
2023-02-11 03:08
-

Content from: <https://docs.ansible.com/automation-controller/4.2.1/html/>

Index of /automation-controller/4.2.1/html

Index of /automation-controller/4.2.1/html

Name
Last modified
Size
Description
Parent Directory
-
administration/
2025-10-18 22:52
-
controllerapi/
2025-10-18 22:52
-
controllercli/
2022-11-30 19:10
-
quickstart/
2025-10-18 22:52
-
release-notes/
2025-10-18 22:52
-
towerapi/
2025-10-18 22:52
-
towercli/
2022-11-30 19:10
-
upgrade-migration-gu..>
2025-10-18 22:52
-
userguide/
2025-10-18 22:52
-

Content from: <https://docs.ansible.com/automation-controller/4.2.2/html/>

Index of /automation-controller/4.2.2/html

Index of /automation-controller/4.2.2/html

Name
Last modified
Size
Description
Parent Directory

-
- administration/
2025-10-18 22:19
-
- controllerapi/
2025-10-18 22:19
-
- quickstart/
2025-10-18 22:19
-
- release-notes/
2025-10-18 22:19
-
- towerapi/
2025-10-18 22:19
-
- upgrade-migration-gu..>
2025-10-18 22:19
-
- userguide/
2025-10-18 22:19
-

Content from: <https://docs.ansible.com/automation-controller/4.4/html/>

Index of /automation-controller/4.4/html

Index of /automation-controller/4.4/html

Name
Last modified
Size
Description
Parent Directory
-
administration/
2025-10-18 22:49
-
controllerapi/
2025-10-18 22:49
-
controllercli/
2023-07-13 20:43
-
quickstart/
2025-10-18 22:49
-
release-notes/
2025-10-18 22:49
-
towerapi/
2025-10-18 22:49
-
towercli/
2023-07-13 20:43
-
upgrade-migration-gu..>
2025-10-18 22:49
-
userguide/
2025-10-18 22:49
-

Content from: <https://docs.ansible.com/automation-controller/latest/html/controllercli/index.html>

AWX Command Line Interface ? AWX CLI Automation Controller 4.6 documentation

AnsibleFest

Products

Community

Webinars & Training

Blog

Documentation

AUTOMATION

CONTROLLER

v4.6

Basic Usage

Installation

Synopsis

Getting Started

Resources and Actions

Global Options

Authentication

Generating a Personal Access Token

Working with OAuth2.0 Applications

OAuth2.0 Token Scoping

Session Authentication

Output Formatting

YAML Formatting

Human-Readable (Tabular) Formatting

Custom Formatting with jq

Colorized Output

Usage Examples

Verifying CLI Configuration

Printing the History of a Particular Job

Creating and Launching a Job Template

Updating a Job Template with Extra Vars

Importing an SSH Key

Import/Export

Reference Guide

awx activity_stream

awx activity_stream get

awx activity_stream list

awx ad_hoc_commands

awx ad_hoc_commands create

awx ad_hoc_commands delete

awx ad_hoc_commands get

awx ad_hoc_commands list

awx ad_hoc_commands stdout

awx analytics

awx analytics get

awx analytics list

awx applications

awx applications create

awx applications delete

awx applications get

awx applications list

awx applications modify

awx bulk
awx bulk get
awx bulk host_create
awx bulk host_delete
awx bulk job_launch
awx bulk list
awx config
awx constructed_inventory
awx constructed_inventory create
awx constructed_inventory delete
awx constructed_inventory get
awx constructed_inventory list
awx constructed_inventory modify
awx credential_input_sources
awx credential_input_sources create
awx credential_input_sources delete
awx credential_input_sources get
awx credential_input_sources list
awx credential_input_sources modify
awx credential_types
awx credential_types create
awx credential_types delete
awx credential_types get
awx credential_types list
awx credential_types modify
awx credentials
awx credentials create
awx credentials delete
awx credentials get
awx credentials list
awx credentials modify
awx execution_environments
awx execution_environments create
awx execution_environments delete
awx execution_environments get
awx execution_environments list
awx execution_environments modify
awx export
awx groups
awx groups create
awx groups delete
awx groups get
awx groups list
awx groups modify
awx host_metric_summary_monthly
awx host_metric_summary_monthly get
awx host_metric_summary_monthly list
awx host_metrics
awx host_metrics delete
awx host_metrics get
awx host_metrics list
awx hosts
awx hosts create
awx hosts delete

awx hosts get
awx hosts list
awx hosts modify
awx import
awx instance_groups
awx instance_groups create
awx instance_groups delete
awx instance_groups get
awx instance_groups list
awx instance_groups modify
awx instances
awx instances create
awx instances get
awx instances list
awx instances modify
awx inventory
awx inventory create
awx inventory delete
awx inventory get
awx inventory list
awx inventory modify
awx inventory_sources
awx inventory_sources associate
awx inventory_sources create
awx inventory_sources delete
awx inventory_sources disassociate
awx inventory_sources get
awx inventory_sources list
awx inventory_sources modify
awx inventory_sources update
awx inventory_updates
awx inventory_updates delete
awx inventory_updates get
awx inventory_updates list
awx inventory_updates stdout
awx job_templates
awx job_templates associate
awx job_templates create
awx job_templates delete
awx job_templates disassociate
awx job_templates get
awx job_templates launch
awx job_templates list
awx job_templates modify
awx jobs
awx jobs delete
awx jobs get
awx jobs list
awx jobs monitor
awx jobs stdout
awx labels
awx labels create
awx labels get
awx labels list

awx labels modify
awx login
awx me
awx mesh_visualizer
awx metrics
awx notification_templates
awx notification_templates create
awx notification_templates delete
awx notification_templates get
awx notification_templates list
awx notification_templates modify
awx notifications
awx notifications get
awx notifications list
awx organizations
awx organizations associate
awx organizations create
awx organizations delete
awx organizations disassociate
awx organizations get
awx organizations list
awx organizations modify
awx ping
awx project_updates
awx project_updates delete
awx project_updates get
awx project_updates list
awx project_updates stdout
awx projects
awx projects associate
awx projects create
awx projects delete
awx projects disassociate
awx projects get
awx projects list
awx projects modify
awx projects update
awx receptor_addresses
awx receptor_addresses get
awx receptor_addresses list
awx role_definitions
awx role_definitions create
awx role_definitions delete
awx role_definitions get
awx role_definitions list
awx role_definitions modify
awx role_team_assignments
awx role_team_assignments create
awx role_team_assignments delete
awx role_team_assignments get
awx role_team_assignments list
awx role_user_assignments
awx role_user_assignments create
awx role_user_assignments delete

awx role_user_assignments get
awx role_user_assignments list
awx roles
awx roles get
awx roles list
awx schedules
awx schedules create
awx schedules delete
awx schedules get
awx schedules list
awx schedules modify
awx service_index
awx service_index get
awx service_index list
awx settings
awx settings list
awx settings modify
awx system_job_templates
awx system_job_templates get
awx system_job_templates list
awx system_jobs
awx system_jobs delete
awx system_jobs get
awx system_jobs list
awx teams
awx teams create
awx teams delete
awx teams get
awx teams grant
awx teams list
awx teams modify
awx teams revoke
awx tokens
awx tokens create
awx tokens delete
awx tokens get
awx tokens list
awx tokens modify
awx unified_job_templates
awx unified_job_templates get
awx unified_job_templates list
awx unified_jobs
awx unified_jobs get
awx unified_jobs list
awx users
awx users create
awx users delete
awx users get
awx users grant
awx users list
awx users modify
awx users revoke
awx workflow_approvals
awx workflow_approvals delete

awx workflow_approvals get
awx workflow_approvals list
awx workflow_job_nodes
awx workflow_job_nodes get
awx workflow_job_nodes list
awx workflow_job_template_nodes
awx workflow_job_template_nodes create
awx workflow_job_template_nodes delete
awx workflow_job_template_nodes get
awx workflow_job_template_nodes list
awx workflow_job_template_nodes modify
awx workflow_job_templates
awx workflow_job_templates associate
awx workflow_job_templates create
awx workflow_job_templates delete
awx workflow_job_templates disassociate
awx workflow_job_templates get
awx workflow_job_templates launch
awx workflow_job_templates list
awx workflow_job_templates modify
awx workflow_jobs
awx workflow_jobs delete
awx workflow_jobs get
awx workflow_jobs list
awx workflow_jobs monitor

AnsibleFest

Products

Community

Webinars & Training

Blog

AWX CLI

Ansible Docs Home

»

Ansible Automation Platform Docs

»

? AWX CLI Automation Controller 4.6 documentation

»

AWX Command Line Interface

¶

awx is the official command-line client for AWX and Red Hat Ansible Automation Platform. It:

Uses naming and structure consistent with the AWX HTTP API

Provides consistent output formats with optional machine-parsable formats

To the extent possible, auto-detects API versions, available endpoints, and

feature support across multiple versions of AWX and Red Hat Ansible Automation Platform.

Potential uses include:

Configuring and launching jobs/playbooks

Checking on the status and output of job runs

Managing objects like organizations, users, teams, etc?

Basic Usage

Installation

Synopsis

Getting Started

Resources and Actions

Global Options

Authentication

Generating a Personal Access Token

Working with OAuth2.0 Applications

OAuth2.0 Token Scoping

Session Authentication

Output Formatting

YAML Formatting

Human-Readable (Tabular) Formatting

Custom Formatting with jq

Colorized Output

Usage Examples

Verifying CLI Configuration

Printing the History of a Particular Job

Creating and Launching a Job Template

Updating a Job Template with Extra Vars

Importing an SSH Key

Import/Export

Reference Guide

awx activity_stream

awx activity_stream get

awx activity_stream list

awx ad_hoc_commands

awx ad_hoc_commands create

awx ad_hoc_commands delete

awx ad_hoc_commands get

awx ad_hoc_commands list

awx ad_hoc_commands stdout

awx analytics

awx analytics get

awx analytics list

awx applications

awx applications create

awx applications delete

awx applications get

awx applications list

awx applications modify

awx bulk

awx bulk get

awx bulk host_create

awx bulk host_delete

awx bulk job_launch

awx bulk list

awx config

awx constructed_inventory

awx constructed_inventory create

awx constructed_inventory delete

awx constructed_inventory get

awx constructed_inventory list

awx constructed_inventory modify

awx credential_input_sources

awx credential_input_sources create

awx credential_input_sources delete

awx credential_input_sources get

awx credential_input_sources list

awx credential_input_sources modify
awx credential_types
awx credential_types create
awx credential_types delete
awx credential_types get
awx credential_types list
awx credential_types modify
awx credentials
awx credentials create
awx credentials delete
awx credentials get
awx credentials list
awx credentials modify
awx execution_environments
awx execution_environments create
awx execution_environments delete
awx execution_environments get
awx execution_environments list
awx execution_environments modify
awx export
awx groups
awx groups create
awx groups delete
awx groups get
awx groups list
awx groups modify
awx host_metric_summary_monthly
awx host_metric_summary_monthly get
awx host_metric_summary_monthly list
awx host_metrics
awx host_metrics delete
awx host_metrics get
awx host_metrics list
awx hosts
awx hosts create
awx hosts delete
awx hosts get
awx hosts list
awx hosts modify
awx import
awx instance_groups
awx instance_groups create
awx instance_groups delete
awx instance_groups get
awx instance_groups list
awx instance_groups modify
awx instances
awx instances create
awx instances get
awx instances list
awx instances modify
awx inventory
awx inventory create
awx inventory delete

awx inventory get
awx inventory list
awx inventory modify
awx inventory_sources
awx inventory_sources associate
awx inventory_sources create
awx inventory_sources delete
awx inventory_sources disassociate
awx inventory_sources get
awx inventory_sources list
awx inventory_sources modify
awx inventory_sources update
awx inventory_updates
awx inventory_updates delete
awx inventory_updates get
awx inventory_updates list
awx inventory_updates stdout
awx job_templates
awx job_templates associate
awx job_templates create
awx job_templates delete
awx job_templates disassociate
awx job_templates get
awx job_templates launch
awx job_templates list
awx job_templates modify
awx jobs
awx jobs delete
awx jobs get
awx jobs list
awx jobs monitor
awx jobs stdout
awx labels
awx labels create
awx labels get
awx labels list
awx labels modify
awx login
awx me
awx mesh_visualizer
awx metrics
awx notification_templates
awx notification_templates create
awx notification_templates delete
awx notification_templates get
awx notification_templates list
awx notification_templates modify
awx notifications
awx notifications get
awx notifications list
awx organizations
awx organizations associate
awx organizations create
awx organizations delete

awx organizations disassociate
awx organizations get
awx organizations list
awx organizations modify
awx ping
awx project_updates
awx project_updates delete
awx project_updates get
awx project_updates list
awx project_updates stdout
awx projects
awx projects associate
awx projects create
awx projects delete
awx projects disassociate
awx projects get
awx projects list
awx projects modify
awx projects update
awx receptor_addresses
awx receptor_addresses get
awx receptor_addresses list
awx role_definitions
awx role_definitions create
awx role_definitions delete
awx role_definitions get
awx role_definitions list
awx role_definitions modify
awx role_team_assignments
awx role_team_assignments create
awx role_team_assignments delete
awx role_team_assignments get
awx role_team_assignments list
awx role_user_assignments
awx role_user_assignments create
awx role_user_assignments delete
awx role_user_assignments get
awx role_user_assignments list
awx roles
awx roles get
awx roles list
awx schedules
awx schedules create
awx schedules delete
awx schedules get
awx schedules list
awx schedules modify
awx service_index
awx service_index get
awx service_index list
awx settings
awx settings list
awx settings modify
awx system_job_templates

awx system_job_templates get
awx system_job_templates list
awx system_jobs
awx system_jobs delete
awx system_jobs get
awx system_jobs list
awx teams
awx teams create
awx teams delete
awx teams get
awx teams grant
awx teams list
awx teams modify
awx teams revoke
awx tokens
awx tokens create
awx tokens delete
awx tokens get
awx tokens list
awx tokens modify
awx unified_job_templates
awx unified_job_templates get
awx unified_job_templates list
awx unified_jobs
awx unified_jobs get
awx unified_jobs list
awx users
awx users create
awx users delete
awx users get
awx users grant
awx users list
awx users modify
awx users revoke
awx workflow_approvals
awx workflow_approvals delete
awx workflow_approvals get
awx workflow_approvals list
awx workflow_job_nodes
awx workflow_job_nodes get
awx workflow_job_nodes list
awx workflow_job_template_nodes
awx workflow_job_template_nodes create
awx workflow_job_template_nodes delete
awx workflow_job_template_nodes get
awx workflow_job_template_nodes list
awx workflow_job_template_nodes modify
awx workflow_job_templates
awx workflow_job_templates associate
awx workflow_job_templates create
awx workflow_job_templates delete
awx workflow_job_templates disassociate
awx workflow_job_templates get
awx workflow_job_templates launch

[awx workflow_job_templates list](#)
[awx workflow_job_templates modify](#)
[awx workflow_jobs](#)
[awx workflow_jobs delete](#)
[awx workflow_jobs get](#)
[awx workflow_jobs list](#)
[awx workflow_jobs monitor](#)

[Indices and tables](#)

[¶](#)

[Index](#)

[Module Index](#)

[Search Page](#)

[Next](#)

Copyright © Red Hat, Inc.

Red Hat Ansible Automation Platform docs are generated using

Sphinx

using a theme provided by

[Read the Docs](#)

.

Content from: <https://docs.ansible.com/automation-controller/latest/html/release-notes/index.html>

Automation Controller Release Notes v4.6 ? Automation Controller Release Notes v4.6

AnsibleFest

Products

Community

Webinars & Training

Blog

Documentation

AUTOMATION

CONTROLLER

v4.6

Are you using the latest and greatest version of Automation Controller? Find the Automation Controller documentation

set which best matches your version of the controller.

Automation Controller Release Notes v4.6

Release Notes

Known Issues

The

CSRF_TRUSTED_ORIGIN

setting

Launching the ansible-runner component not working as expected

Deleted default orgs produces duplicate Ansible-Galaxy credentials

Isolated nodes unsupported in an OpenShift deployment

Browsers ignoring the

autocomplete=off

setting

Login via HTTP requires workaround

Job slicing and limit interactions

Misuse of job slicing can cause errors in job scheduling

Default LDAP directory must be configured to use LDAP Authentication

Potential security issue using

X_FORWARDED_FOR

in

REMOTE_HOST_HEADERS

Server error when accessing SAML metadata via hostname

SAML authentication revokes admin role upon login

Live events status indicators

VMWare Self-Signed Certs

awx-manage inventory_import user

Database on Disk Becomes Corrupted

Safari unable to establish connection to web socket

Local management not functioning as expected

Problems when using SSH customization

Database server installed on nodes

Reactivating OAuth authentication accounts which have been deleted

Using vaulted variables in inventory sourced from a project

Saved scheduled and workflow configurations and surveys

Supported Locales

AnsibleFest

Products

Community

Webinars & Training

Blog

Automation Controller Release Notes

Ansible Docs Home

»

Ansible Automation Platform Docs

»

? Automation Controller Release Notes v4.6

»

Automation Controller Release Notes v4.6

¶

Thank you for your interest in Red Hat Ansible Automation Platform controller. Automation controller is a commercial offering that helps teams manage complex multi-tier deployments by adding control, knowledge, and delegation to Ansible-powered environments.

The

Automation Controller Release Notes

provides release notes, known issues, and related reference materials. This document has been updated to include information for the latest release of Automation Controller v4.6.

We Need Feedback!

If you spot a typo in this documentation, or if you have thought of a way to make this manual better, we would love to hear from you! Please send an email to:

docs

@

ansible

.

com

If you have a suggestion, try to be as specific as possible when describing it. If you have found an error, please include the manual's title, chapter number/section number, and some of the surrounding text so we can find it easily. We may not be able to respond to every message sent to us, but you can be sure that we will be reading them all!

Automation Controller Version 4.6; September 30, 2024;

<https://access.redhat.com/>

Automation Controller Release Notes v4.6

Release Notes

Known Issues

The

CSRF_TRUSTED_ORIGIN

setting

Launching the ansible-runner component not working as expected

Deleted default orgs produces duplicate Ansible-Galaxy credentials

Isolated nodes unsupported in an OpenShift deployment

Browsers ignoring the

autocomplete=off

setting

Login via HTTP requires workaround

Job slicing and limit interactions

Misuse of job slicing can cause errors in job scheduling

Default LDAP directory must be configured to use LDAP Authentication

Potential security issue using

X_FORWARDED_FOR

in

REMOTE_HOST_HEADERS

Server error when accessing SAML metadata via hostname

SAML authentication revokes admin role upon login

Live events status indicators

VMWare Self-Signed Certs

awx-manage inventory_import user

Database on Disk Becomes Corrupted
Safari unable to establish connection to web socket
Local management not functioning as expected
Problems when using SSH customization
Database server installed on nodes
Reactivating OAuth authentication accounts which have been deleted
Using vaulted variables in inventory sourced from a project
Saved scheduled and workflow configurations and surveys
Supported Locales

Index



Index

Copyright © Red Hat, Inc.



Ansible, Ansible Automation Platform, Red Hat, and Red Hat Enterprise Linux are trademarks of Red Hat, Inc., registered in the United States and other countries.

If you distribute this document, or a modified version of it, you must provide attribution to Red Hat, Inc. and provide a link to the original version.

Third Party Rights

Ubuntu and Canonical are registered trademarks of Canonical Ltd.

The CentOS Project is copyright protected. The CentOS Marks are trademarks of Red Hat, Inc. (?Red Hat?).

Microsoft, Windows, Windows Azure, and Internet Explore are trademarks of Microsoft, Inc.

VMware is a registered trademark or trademark of VMware, Inc.

Amazon Web Services?, ?AWS?, ?Amazon EC2?, and ?EC2?, are trademarks of Amazon Web Services, Inc. or its affiliates.

OpenStack? and OpenStack logo are trademarks of OpenStack, LLC.

Chrome? and Google Compute Engine? service registered trademarks of Google Inc.

Safari® is a registered trademark of Apple, Inc.

Firefox® is a registered trademark of the Mozilla Foundation.

All other trademarks are the property of their respective owners.

Next

Copyright © Red Hat, Inc.

Red Hat Ansible Automation Platform docs are generated using

Sphinx

using a theme provided by

Read the Docs

.

Content from: <https://docs.ansible.com/automation-tower-prior-versions.html>

[Automation Controller documentation archive](#) | [Ansible Documentation](#)

[Ansible documentation](#)

[Toggle navigation](#)

[Products](#)

[Blog](#)

[Community](#)

[Webinars and training](#)

[Try it now](#)

[Toggle navigation](#)

[Join the community](#)

[Users](#)

[Developers](#)

[Maintainers](#)

[Ansible core](#)

[Ansible ecosystem](#)

[Red Hat Ansible Automation Platform](#)

[Automation Controller documentation archive](#)

Archive page for the Automation Controller documentation that provides older versions of content.

Got thoughts or feedback on this site? We want to hear from you!

[Join us in the](#)

[Ansible Forum](#)

[or open a](#)

[GitHub issue](#)

[in the docsite repository.](#)

[Automation Controller version 4.5 \(Supported\)](#)

[Automation Controller version 4.4 \(Supported\)](#)

[Automation Controller version 4.3 \(Supported\)](#)

[Automation Controller version 4.2.2 \(Supported\)](#)

[Automation Controller version 4.2.1 \(Supported\)](#)

[Automation Controller version 4.2.0 \(Supported\)](#)

[Automation Controller version 4.1.4 \(Supported\)](#)

[Automation Controller version 4.1.3 \(Supported\)](#)

[Automation Controller version 4.1.2 \(Supported\)](#)

[Automation Controller version 4.1.1 \(Supported\)](#)

[Automation Controller version 4.1.0 \(Supported\)](#)

[Automation Tower version 3.8.6 \(Supported\)](#)

[Automation Tower version 3.8.5 \(Supported\)](#)

[Automation Tower version 3.8.4 \(Supported\)](#)

[Automation Tower version 3.8.3 \(Supported\)](#)

[Automation Tower version 3.8.2 \(Supported\)](#)

[Automation Tower version 3.8.1 \(Supported\)](#)

[Automation Tower version 3.8.0 \(Supported\)](#)

[CC BY-SA 4.0](#) |

[Privacy policy](#) |

[Sponsored by](#)

Content from: <https://docs.ansible.com/community.html>

[Join the Ansible community | Ansible Documentation](#)

[Ansible documentation](#)

[Toggle navigation](#)

[Products](#)

[Blog](#)

[Community](#)

[Webinars and training](#)

[Try it now](#)

[Toggle navigation](#)

[Join the community](#)

[Users](#)

[Developers](#)

[Maintainers](#)

[Ansible core](#)

[Ansible ecosystem](#)

[Red Hat Ansible Automation Platform](#)

[Join the Ansible community](#)

Open-source and collaboration are at the heart of the Ansible community, helping more people experience the power of automation and work better and faster together.

Got thoughts or feedback on this site? We want to hear from you!

[Join us in the](#)

[Ansible Forum](#)

[or open a](#)

[GitHub issue](#)

[in the docsite repository.](#)

[Collaborate with us](#)

[Read the code of conduct](#)

[Explore ways to contribute](#)

[Find an Ansible project](#)

[Review the contributor path](#)

[Join the conversation](#)

[Join the Ansible Forum](#)

[Find public community meetings on the Forum Events calendar](#)

[Check out the guide to community communication](#)

[Participate in community discussions](#)

[Share your ideas and vote in the Community topics](#)

[Take part in a working group](#)

[Follow Ansible community news and announcements](#)

[Subscribe to the Bullhorn](#)

[CC BY-SA 4.0 |](#)

[Privacy policy |](#)

[Sponsored by](#)

Content from: <https://docs.ansible.com/core-translated-ja.html>

????????? (Japanese) | Ansible Documentation

Ansible documentation

Toggle navigation

Products

Blog

Community

Webinars and training

Try it now

Toggle navigation

Join the community

Users

Developers

Maintainers

Ansible core

Ansible ecosystem

Red Hat Ansible Automation Platform

????????? (Japanese)

Access Japanese translations for Ansible Core documentation.

Got thoughts or feedback on this site? We want to hear from you!

Join us in the

Ansible Forum

or open a

GitHub issue

in the docsite repository.

ansible-core ?????? 2.15

Visit the documentation

ansible-core ?????? 2.14

Visit the documentation

ansible-core ?????? 2.13

Visit the documentation

ansible-core ?????? 2.12

Visit the documentation

ansible-core ?????? 2.11

Visit the documentation

Ansible ?????? 2.9

Visit the documentation

CC BY-SA 4.0 |

Privacy policy |

Sponsored by

Content from: <https://docs.ansible.com/core.html>

[Ansible Core | Ansible Documentation](#)

[Ansible documentation](#)

[Toggle navigation](#)

[Products](#)

[Blog](#)

[Community](#)

[Webinars and training](#)

[Try it now](#)

[Toggle navigation](#)

[Join the community](#)

[Users](#)

[Developers](#)

[Maintainers](#)

[Ansible core](#)

[Ansible ecosystem](#)

[Red Hat Ansible Automation Platform](#)

[Ansible Core](#)

Core includes the Ansible language and runtime, a set of built-in modules and command-line tools, and a framework for extending automation with collections.

Got thoughts or feedback on this site? We want to hear from you!

[Join us in the](#)

[Ansible Forum](#)

[or open a](#)

[GitHub issue](#)

[in the docsite repository.](#)

[Ansible Core](#)

Ansible Core is the language and runtime that powers automation.

[Ansible Core documentation](#)

[Ansible Test](#)

Ansible Test is a command-line tool for performing sanity, unit, and integration testing with Ansible Core and collections.

[Ansible Test documentation](#)

????????? (Japanese)

[Access Japanese translations for Ansible Core documentation.](#)

[Learn more](#)

[Ansible community documentation archive](#)

Archive page for the Ansible community documentation that provides older versions of content.

[Visit the archive](#)

[CC BY-SA 4.0 |](#)

[Privacy policy |](#)

[Sponsored by](#)

Content from: https://docs.ansible.com/dev_guide.html

Developer Guide ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Adding modules and plugins locally

Should you develop a module?

Developing modules

Contributing your module to an existing Ansible collection

Conventions, tips, and pitfalls

Ansible and Python 3

Debugging modules

Module format and documentation

Ansible markup

Adjacent YAML documentation files

Windows module development walkthrough

Creating a new collection

Testing Ansible

The lifecycle of an Ansible module or plugin

Developing plugins

Developing dynamic inventory

Developing

ansible-core

Ansible module architecture
Python API
Rebasing a pull request
Using and developing module utilities
Ansible collection creator path
Developing collections
Migrating Roles to Roles in Collections on Galaxy
Collection Galaxy metadata structure
Ansible architecture
Common Ansible Scenarios
Legacy Public Cloud Guides
Network Automation
Network Getting Started
Network Advanced Topics
Network Developer Guide
Ansible Galaxy
Galaxy User Guide
Galaxy Developer Guide
Reference & Appendices
Collection Index
Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Edit on GitHub
Developer Guide
?

Note

Making Open Source More Inclusive

Red Hat is committed to replacing problematic language in our code, documentation, and web properties. We are beginning with these four terms: master, slave, blacklist, and whitelist. We ask that you open an issue or pull request if you come upon a term that we have missed. For more details, see our CTO Chris Wright's message

.

Welcome to the Ansible Developer Guide!

Who should use this guide?

If you want to extend Ansible by using a custom module or plugin locally, creating a module or plugin, adding functionality to an existing module, or expanding test coverage, this guide is for you. We've included detailed information for developers on how to test and document modules, as well as the prerequisites for getting your module or plugin accepted into the main Ansible repository.

Find the task that best describes what you want to do:

I'm looking for a way to address a use case:

I want to

add a custom plugin or module locally

.

I want to figure out if

developing a module is the right approach

for my use case.

I want to understand

what a successful collection creator path looks like

.

I want to

develop a collection

.

I want to

contribute to an Ansible-maintained collection

.

I want to

contribute to a community-maintained collection

.

I want to

migrate a role to a collection

.

I've read the info above, and I'm sure I want to develop a module:

What do I need to know before I start coding?

I want to

set up my Python development environment

.

I want to

get started writing a module

.

I want to write a specific kind of module:

a

network module

a

Windows module

.

an

Amazon module

.

an

oVirt/RHV module

.

a

VMware module

.

I want to

write a series of related modules

that integrate Ansible with a new product (for example, a database, cloud provider, network platform, and so on).

I want to refine my code:

I want to

debug my module code

.

I want to

add tests

.

I want to

document my module

.

I want to

improve documentation by using Ansible markup

.

I want to

document my set of modules for a network platform

.

I want to follow

conventions and tips for clean, usable module code

.

I want to

make sure my code runs on Python 2 and Python 3

.

I want to work on other development projects:

I want to

write a plugin

.

I want to

connect Ansible to a new source of inventory

.

I want to

deprecate an outdated module

.

I want to contribute back to the Ansible project:

I want to

understand how to contribute to Ansible

.

I want to

contribute my module or plugin

.

I want to

understand the DCO agreement

for contributions to the

Ansible Core

and

Ansible Documentation

repositories.

If you prefer to read the entire guide, here's a list of the pages in order.

Adding modules and plugins locally

Modules and plugins: what is the difference?

Adding modules and plugins in collections

Adding a module or plugin outside of a collection

Adding a non-module plugin locally outside of a collection

Using

ansible.legacy

- to access custom versions of an `ansible.builtin` module
- Should you develop a module?
- Developing modules
 - Preparing an environment for developing Ansible modules
 - Creating a standalone module
 - Creating a module in a collection
 - Creating an info or a facts module
 - Verifying your module code
 - Testing your newly-created module
 - Contributing back to Ansible
 - Communication and development support
- Credit
 - Contributing your module to an existing Ansible collection
- Contributing modules: objective requirements
- Contributing to Ansible: subjective requirements
- Other checklists
 - Conventions, tips, and pitfalls
- Scoping your module(s)
- Designing module interfaces
- General guidelines & tips
- Functions and Methods
- Python tips
 - Importing and using shared code
- Handling module failures
 - Handling exceptions (bugs) gracefully
- Creating correct and informative module output
- Following Ansible conventions
- Module Security
 - Ansible and Python 3
 - Minimum version of Python 3.x and Python 2.x
 - Developing Ansible code that supports Python 2 and Python 3
- Debugging modules
 - Detailed debugging steps
 - Simple debugging
- Module format and documentation
 - Non-Python modules documentation
 - Python shebang & UTF-8 coding
 - Copyright and license
- DOCUMENTATION block
- EXAMPLES block
- RETURN block
- Python imports
- Testing module documentation
- Ansible markup
 - Semantic markup within module documentation
 - Linking within module documentation
 - Format macros within module documentation
- Adjacent YAML documentation files
- YAML documentation for plugins
- YAML format
- Supported plugin types
- Windows module development walkthrough

Windows environment setup
Create a Windows server in a VM
Create an Ansible inventory
Provisioning the environment
Windows new module development
Windows module utilities
Windows playbook module testing
Windows debugging
Windows unit testing
Windows integration testing
Windows communication and development support
Creating a new collection
Before you start coding
Naming conventions
Speak to us
Where to get support
Required files
New to Git or GitHub
Testing Ansible
Why test your Ansible contributions?
Types of tests
Testing within GitHub & Azure Pipelines
How to test a PR
Want to know more about testing?
The lifecycle of an Ansible module or plugin
Deprecating modules and plugins in the Ansible main repository
Deprecating modules and plugins in a collection
Changing a module or plugin name in the Ansible main repository
Renaming a module or plugin in a collection, or redirecting a module or plugin to another collection
Tombstoning a module or plugin in a collection
Developing plugins
Writing plugins in Python
Raising errors
String encoding
Plugin configuration & documentation standards
Developing particular plugin types
Developing dynamic inventory
Inventory sources
Inventory plugins
Inventory scripts
Developing
ansible-core
ansible-core
project branches and tags
Ansible module architecture
Ansible module architecture
Types of modules
How modules are executed
Python API
Rebasing a pull request
Configuring your remotes
Rebasing your branch
Updating your pull request
Getting help rebasing

Using and developing module utilities
Naming and finding module utilities
Standard module utilities
Ansible collection creator path
Examine currently available solutions
Create your content
Put your content in a collection
Write good user collection documentation
Publish your collection source code
Follow a versioning convention
Understand and implement testing and CI
Provide good contributor & maintainer documentation
Publish your collection on distribution servers
Make your collection a part of Ansible community package
Maintain
Communicate
Developing collections
Creating collections
Using shared resources in collections
Testing collections
Distributing collections
Documenting collections
Migrating Ansible content to a different collection
Contributing to collections
Generating changelogs and porting guide entries in a collection
Collection structure
Collection Galaxy metadata structure
Migrating Roles to Roles in Collections on Galaxy
Comparing standalone roles to collection roles
Migrating a role to a collection
Migrating a role that contains plugins to a collection
Using
ansible.legacy
to access local custom modules from collections-based roles
Collection Galaxy metadata structure
Structure
Examples
Ansible architecture
Modules
Module utilities
Plugins
Inventory
Playbooks
The Ansible search path
Previous
Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: <https://docs.ansible.com/developers.html>

Developers | Ansible Documentation

Ansible documentation

Toggle navigation

Products

Blog

Community

Webinars and training

Try it now

Toggle navigation

Join the community

Users

Developers

Maintainers

Ansible core

Ansible ecosystem

Red Hat Ansible Automation Platform

Developers

Extend automation with custom Ansible modules, add functionality to existing modules, or fix bugs to improve existing code.

Got thoughts or feedback on this site? We want to hear from you!

Join us in the

Ansible Forum

or open a

GitHub issue

in the docsite repository.

Quicklinks

Developing modules

Ansible and Python 3

Python API

Start writing code

Set up your development environment

Learn how Ansible works

Write custom modules or plugins

Contribute code to a collection

Make your first contribution

Explore the Collection contributor guide

Contribute your module to an existing collection

Review the lifecycle of an Ansible module or plugin

Test plugins and modules

Understand testing in Ansible

Run sanity tests

Write integration tests

Write unit tests

Create new collections

Set things up with the collection skeleton

Test your collection for code quality

Publish your collection on a distribution server

Request a collection be added to the Ansible package

CC BY-SA 4.0 |

Privacy policy |

Sponsored by

Content from: https://docs.ansible.com/developing_api.html

Python API ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[ansible-core Roadmaps](#)
[Ansible](#)
[Developer Guide](#)
[Python API](#)
[Edit on GitHub](#)
[Python API](#)
[?](#)

[Topics](#)

[Python API](#)

[Attention](#)

The Ansible API is intended for internal Ansible use. Ansible may make changes to this API at any time that could break backward compatibility with older versions of the API. Because of this, external use is not supported by Ansible. If you want to use Python API only for executing playbooks or modules, consider [ansible-runner](#)

first.

If you would like to use Ansible programmatically from a language other than Python, trigger events asynchronously, or have access control and logging demands, please see the [AWX project](#)

.

[See also](#)

[Developing dynamic inventory](#)

[Developing dynamic inventory integrations](#)

[Developing modules](#)

[Getting started on developing a module](#)

[Developing plugins](#)

[How to develop plugins](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_collections.html

Developing collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

- Interpreter Discovery
- Releases and maintenance
- Testing Strategies
- Sanity Tests
- Frequently Asked Questions
- Glossary
- Ansible Reference: Module Utilities
- Special Variables
- Red Hat Ansible Automation Platform
- Ansible Automation Hub
- Logging Ansible output
- Roadmaps
- Ansible Roadmap
- ansible-core Roadmaps
- Ansible
- Developer Guide
- Developing collections
- Edit on GitHub
- Developing collections
- ?

Collections are a distribution format for Ansible content. You can package and distribute playbooks, roles, modules, and plugins using collections. A typical collection addresses a set of related use cases. For example, the `cisco.ios`

collection automates management of Cisco IOS devices.

You can create a collection and publish it to

Ansible Galaxy

or to a private Automation Hub instance. You can publish certified collections to the Red Hat Automation Hub, part of the Red Hat Ansible Automation Platform.

Examine the

Ansible collection creator path

to understand how to go from creating a collection to having it included in the Ansible package distribution.

Developing new collections

Creating collections

Naming your collection

Creating a new collection

Creating a collection from a custom template

Creating collections with `ansible-creator`

Using shared resources in collections

Using documentation fragments in collections

Leveraging optional module utilities in collections

Listing collection dependencies

Testing collections

Testing tools

Distributing collections

Initial configuration of your distribution server or servers

Building your collection tarball

Preparing to publish your collection

Publishing your collection

Documenting collections

Documenting modules and plugins

Documenting roles

Verifying your collection documentation

Build a docsite with `antsibull-docs`

Working with existing collections

[Migrating Ansible content to a different collection](#)

[Migrating content](#)

[Contributing to collections](#)

[Contributing to a collection: community.general](#)

[Generating changelogs and porting guide entries in a collection](#)

[Understanding ansible-changelog](#)

[Including collection changelogs into Ansible](#)

[Collections references](#)

[Collection structure](#)

[Collection directories and files](#)

[Collection Galaxy metadata structure](#)

[Structure](#)

[Examples](#)

[For instructions on developing modules, see](#)

[Developing modules](#)

.

[See also](#)

[Using Ansible collections](#)

[Learn how to install and use collections in playbooks and roles](#)

[Contributing to Ansible-maintained Collections](#)

[Guidelines for contributing to selected collections](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_collections_changelogs.html

Generating changelogs and porting guide entries in a collection ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing collections
Generating changelogs and porting guide entries in a collection
Edit on GitHub
Generating changelogs and porting guide entries in a collection
?

You can create and share changelog and porting guide entries for your collection. If your collection is part of the Ansible Community package, we recommend that you use the `antsibull-changelog` tool to generate Ansible-compatible changelogs. The Ansible changelog uses the output of this tool to collate all the collections included in an Ansible release into one combined changelog for the release.

Note

Ansible here refers to the Ansible 2.10 or later release that includes a curated set of collections.

Understanding `antsibull-changelog`

Generating changelogs

Porting Guide entries from changelog fragments

Including collection changelogs into Ansible

Understanding `antsibull-changelog`

?

The

`antsibull-changelog`

tool allows you to create and update changelogs for Ansible collections that are compatible with the combined Ansible changelogs. This is an update to the changelog generator used in prior Ansible releases. The tool adds three new changelog fragment categories:

`breaking_changes`

,

`security_fixes`

and

`trivial`

. The tool also generates the

`changelog.yaml`

file that Ansible uses to create the combined

`CHANGELOG.rst`

file and Porting Guide for the release.

See

Creating a changelog fragment

and the

`antsibull-changelog` documentation

for complete details.

Note

The collection maintainers set the changelog policy for their collections. See the individual collection contributing guidelines for complete details.

Generating changelogs

?

To initialize changelog generation:

Install

antsibull-changelog

:

pip

install

antsibull-changelog

.

Initialize changelogs for your repository:

antsibull-changelog

init

<path/to/your/collection>

.

Optionally, edit the

changelogs/config.yaml

file to customize the location of the generated changelog

.rst

file or other options. See

Bootstrapping changelogs for collections

for details.

To generate changelogs from the changelog fragments you created:

Optionally, validate your changelog fragments:

antsibull-changelog

lint

.

Generate the changelog for your release:

antsibull-changelog

release

[--version

version_number]

.

Note

Add the

--reload-plugins

option if you ran the

antsibull-changelog

release

command previously and the version of the collection has not changed.

antsibull-changelog

caches the information on all plugins and does not update its cache until the collection version changes.

Porting Guide entries from changelog fragments

?

The Ansible changelog generator automatically adds several changelog fragment categories to the Ansible Porting Guide:

major_changes

breaking_changes

deprecated_features

removed_features

Including collection changelogs into Ansible

?

If your collection is part of Ansible, use one of the following three options to include your changelog into the Ansible release changelog:

Use the

`antsibull-changelog`

tool.

If are not using this tool, include the properly formatted

`changelog.yaml`

file into your collection. See the

`changelog.yaml` format

for details.

Add a link to own changelogs or release notes in any format by opening an issue at

<https://github.com/ansible-community/ansible-build-data/>

with the HTML link to that information.

Note

For the first two options, Ansible pulls the changelog details from Galaxy so your changelogs must be included in the collection version on Galaxy that is included in the upcoming Ansible release.

See also

Generating changelogs and porting guide entries in a collection

Learn how to create good changelog fragments.

Using Ansible collections

Learn how to install and use collections.

Contributing to Ansible-maintained Collections

Guidelines for contributing to selected collections

Communication

Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_collections_contributing.html

Contributing to collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing collections
Contributing to collections
Edit on GitHub
Contributing to collections

?

If you want to add functionality to an existing collection, modify a collection you are using to fix a bug, or change the behavior of a module in a collection, clone the Git repository for that collection and make changes on a branch. You can combine changes to a collection with a local checkout of Ansible (

source

hacking/env-setup

).

You should first check the collection repository to see if it has specific contribution guidelines. These are typically listed in the README.md or CONTRIBUTING.md files within the repository.

See

Creating your first collection pull request

for more general guidelines and

Testing Ansible and Collections

for testing guidelines.

Contributing to a collection: community.general

?

These instructions apply to collections hosted in the

ansible_collections GitHub organization

. For other collections, especially for collections not hosted on GitHub, check the

README.md

of the collection for information on contributing to it.

This example uses the

community.general collection

. To contribute to other collections in the same GitHub org, replace the folder names

community

and

general

with the namespace and collection name of a different collection.

Prerequisites

?

Include

~/dev/ansible/collections/

in

COLLECTIONS_PATHS

If that path mentions multiple directories, make sure that no other directory earlier in the search path contains a copy of `community.general`

.

Creating a PR

?

Create the directory

```
~/dev/ansible/collections/ansible_collections/community
```

:

```
mkdir
```

```
-p
```

```
~/dev/ansible/collections/ansible_collections/community
```

Clone

the `community.general` Git repository

or a fork of it into the directory

`general`

:

```
cd
```

```
~/dev/ansible/collections/ansible_collections/community
```

```
git
```

```
clone
```

```
[email protected]
```

```
:ansible-collections/community.general.git
```

`general`

If you clone from a fork, add the original repository as a remote

`upstream`

:

```
cd
```

```
~/dev/ansible/collections/ansible_collections/community/general
```

```
git
```

```
remote
```

```
add
```

```
upstream
```

```
[email protected]
```

```
:ansible-collections/community.general.git
```

Create a branch and commit your changes on the branch.

Remember to add tests for your changes, see

[Testing collections](#)

.

Push your changes to your fork of the collection and create a Pull Request.

You can test your changes by using this checkout of

`community.general`

in playbooks and roles with whichever version of Ansible you have installed locally, including a local checkout of

`ansible/ansible`

?s

`devel`

branch.

See also

[Using Ansible collections](#)

[Learn how to install and use collections.](#)

[Contributing to Ansible-maintained Collections](#)

[Guidelines for contributing to selected collections](#)

[Communication](#)

Got questions? Need help? Want to share your ideas? Visit the [Ansible communication guide](#)

[Previous](#)

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_collections_creating.html

Creating collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing collections
Creating collections
Edit on GitHub
Creating collections
?

To create a collection:

Create a
new collection
, optionally using a custom
collection template
, with the
ansible-galaxy
collection
init
command.

Add modules and other content to the collection.

Build the collection into a collection artifact with
ansible-galaxy collection build

.
Publish the collection artifact to Galaxy with
ansible-galaxy collection publish

.
A user can then install your collection on their systems.

Naming your collection

Creating a new collection

Creating a collection from a custom template

Creating collections with ansible-creator

Naming your collection

?
Collection names consist of a namespace and a name, separated by a period (

.
) . Both namespace and name should be valid Python identifiers. This means that they should consist of ASCII letters, digits, and underscores.

Note

Usually namespaces and names use lower-case letters, digits, and underscores, but no upper-case letters.

You should make sure that the namespace you use is not registered by someone else by checking on

Ansible Galaxy's namespace list

. If you chose a namespace or even a full collection name that collides with another collection on Galaxy, it can happen

that if you or someone else runs

ansible-galaxy

collection

install

with your collection name, you end up with another collection. Even if the namespace currently does not exist, it could be created later by someone else.

If you want to request a new namespace on Ansible Galaxy,

create an issue on github.com/ansible/galaxy

.

There are a few special namespaces:

ansible

:

The

ansible namespace

is owned by Red Hat and reserved for official Ansible collections. Two special members are the synthetic

ansible.builtin

and

ansible.legacy

collections. These cannot be found on Ansible Galaxy, but are built-in into ansible-core.

community

:

The

community namespace

is owned by the Ansible community. Collections from this namespace generally live in the

GitHub ansible-collection organization

. If you want to create a collection in this namespace,

request

it on the forum.

local

:

The

local namespace

does not contain any collection on Ansible Galaxy, and the intention is that this will never change. You can use the

local

namespace for collections that are locally on your machine or locally in your Git repositories, without having to fear collisions with actually existing collections on Ansible Galaxy.

Creating a new collection

?

Create your collection skeleton in a path that includes

ansible_collections

, for example

collections/ansible_collections/

.

To start a new collection, run the following command in your collections directory:

ansible_collections#>

ansible-galaxy

collection

init

my_namespace.my_collection

Note

Both the namespace and collection names use the same strict set of requirements. Both are limited to alphanumeric characters and underscores, must have a minimum length of two characters, and cannot start with an underscore.

It will create the structure

[my_namespace]/[my_collection]/[collection

skeleton]

.

Hint

If Git is used for version control, the corresponding repository should be initialized in the collection directory.

Once the collection exists, you can populate the directories with the content you want inside the collection. See

ansible-collections

GitHub Org to get a better idea of what you can place inside a collection.

Reference: the

ansible-galaxy

collection

command

Currently the

ansible-galaxy

collection

command implements the following sub commands:

init

: Create a basic collection based on the default template included with Ansible or your own template.

build

: Create a collection artifact that can be uploaded to Galaxy or your own repository.

publish

: Publish a built collection artifact to Galaxy.

install

: Install one or more collections.

To learn more about the

ansible-galaxy

command-line tool, see the

ansible-galaxy

man page.

Creating a collection from a custom template

?

The built-in collection template is a simple example of a collection that works with

ansible-core

, but if you want to simplify your development process you may want to create a custom collection template to pass to

ansible-galaxy

collection

init

.

A collection skeleton is a directory that looks like a collection directory but any

.j2

files (excluding those in

templates/

and

roles/*/templates/

) will be templated by

ansible-galaxy

collection

init

. The skeleton's

galaxy.yml.j2

file should use the variables

namespace

and

collection_name

which are derived from

ansible-galaxy
init
namespace.collection_name
, and will populate the metadata in the initialized collection?
galaxy.yml
file. There are a few additional variables available by default (for example,
version
is
1.0.0
) , and these can be supplemented/overridden using
--extra-vars

.
An example
galaxy.yml.j2
file that accepts an optional dictionary variable
dependencies
could look like this:

```
namespace:
{{
namespace
}}
name:
{{
collection_name
}}
version:
{{
(
version
|
quote
)
is
version
(
'0.0.0'
,
operator
=
'gt'
,
version_type
=
'semver'
)|
ternary
(
version
,
undef
(
'version must be a valid semantic version greater than 0.0.0'
))
}}
```

dependencies:

```
{{
dependencies
|
default
({},
true
)
}}
```

To initialize a collection using the new template, pass the path to the skeleton with

ansible-galaxy

collection

init

:

ansible_collections#>

ansible-galaxy

collection

init

--collection-skeleton

/path/to/my/namespace/skeleton

--extra-vars

"@my_vars_file.json"

my_namespace.my_collection

Note

Before

ansible-core

2.17, collection skeleton templating is limited to the few hardcoded variables including

namespace

,

collection_name

, and

version

.

Note

The default collection skeleton uses an internal filter

comment_ify

that isn't accessibly to

--collection-skeleton

. Use

ansible-doc

-t

filter|test

--list

to see available plugins.

Creating collections with ansible-creator

?

ansible-creator

is designed to quickly scaffold an Ansible collection project.

Note

The

Ansible Development Tools

package offers a convenient way to install

ansible-creator

along with a curated set of tools for developing automation content.

After

installing

ansible-creator

you can initialize a project in one of the following ways:

Use the

init

subcommand.

Use

ansible-creator

with the

Ansible extension

in Visual Studio Code.

See also

[Using Ansible collections](#)

[Learn how to install and use collections.](#)

[Collection structure](#)

[Directories and files included in the collection skeleton](#)

[Ansible Development Tools \(ADT\)](#)

[Python package of tools to create and test Ansible content.](#)

[Communication](#)

Got questions? Need help? Want to share your ideas? Visit the [Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_collections_documenting.html

Documenting collections ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing collections
Documenting collections
Edit on GitHub
Documenting collections
?
Documenting modules and plugins
?
Documenting modules is thoroughly documented in
Module format and documentation
. Plugins can be documented the same way as modules, that is with
DOCUMENTATION
,
EXAMPLES
, and
RETURN
blocks.
Documenting roles
?
To document a role, you have to add a role argument spec by creating a file
meta/argument_specs.yml
in your role. See
Role argument validation
for details. As an example, you can look at
the argument specs file
of the
telekom_mms.icinga_director.ansible_icinga role
on GitHub.
Verifying your collection documentation
?
You can use
antsibull-docs
to lint your collection documentation.
See
Linting collection documentation
.
for details.
Build a docsite with antsibull-docs
?

You can use

`antsibull-docs`

to build a Sphinx-based docsite for your collection:

Create your collection and make sure you can use it with `ansible-core` by adding it to your

`COLLECTIONS_PATHS`

.

Create a directory

`dest`

and run

`antsibull-docs`

`sphinx-init`

`--use-current`

`--dest-dir`

`dest`

`namespace.name`

, where

`namespace.name`

is the name of your collection.

Go into

`dest`

and run

`pip`

`install`

`-r`

`requirements.txt`

. You might want to create a venv and activate it first to avoid installing this globally.

Then run

`./build.sh`

.

Open

`build/html/index.html`

in a browser of your choice.

See

`antsibull-docs` documentation

for complete details.

If you want to add additional documentation to your collection next to the plugin, module, and role documentation, see `docs` directory

.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_collections_migrating.html

Migrating Ansible content to a different collection ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing collections
Migrating Ansible content to a different collection
Edit on GitHub
Migrating Ansible content to a different collection
?

You might decide to move content from one collection to another; for example, to extract a set of related modules out of `community.general` or `community.network` to create a more focused collection.

When you migrate content between collections, you must take certain steps to ensure users can follow the transition.

Migrating content

Adding the content to the new collection

Removing the content from the old collection

Updating `BOTMETA.yml`

Migrating content

?

If the collection from which you are going to migrate content is included in the Ansible community package

, ensure the target collection satisfies the

Ansible community package collections requirements

. After you satisfy the requirements, you can migrate the content as follows:

Copy content from the source (old) collection to the target (new) collection.

Change

`M()`

, examples,

seealso

,

`extended_documentation_fragments`

to use actual FQCNs in moved content, old collection, and in other collections that have references to the content.

Move all related issues, pull requests, and wiki pages.

Look through the

`docs/docsite`

directory of the

ansible-documentation GitHub repository

(for example, using the

`grep`

command-line utility) to check if there are examples using the moved modules and plugins so that you can update those

FQCNs.

Deprecate the module/plugin with

removal_version

scheduled for the next major version in

meta/runtime.yml

of the old collection. The deprecation must be released after the copied content has been included in a release of the new collection.

When the next major release of the old collection is prepared:

remove the module/plugin from the old collection

remove related unit and integration tests

remove specific module utils (if they are NOT used by other modules/plugins or

module_utils

)

remove specific documentation fragments if there are any in the old collection

add a changelog fragment containing entries for

removed_features

and

breaking_changes

; you can see an example of a changelog fragment in this

pull request

change

meta/runtime.yml

in the old collection:

add

redirect

to the corresponding module/plugin's entry

in particular, add

redirect

for the removed module utils and documentation fragments if applicable

remove

removal_version

from there

remove related entries from

tests/sanity/ignore.txt

files if exist

remove changelog fragments for removed content that are not yet part of the changelog (in other words, do not modify changelogs/changelog.yml

and do not delete files mentioned in it)

remove requirements that are no longer required in

tests/unit/requirements.txt

,

tests/requirements.yml

and

galaxy.yml

To implement these changes, you need to create at least three PRs:

Create a PR against the new collection to copy the content.

Deprecate the module/plugin in the old collection.

Later create a PR against the old collection to remove the content according to the schedule.

Adding the content to the new collection

?

Create a PR in the new collection to:

Copy ALL the related files from the old collection.

If it is an action plugin, include the corresponding module with documentation.

If it is a module, check if it has a corresponding action plugin that should move with it.

Check
meta/
for relevant updates to
runtime.yml
if it exists.
Carefully check the moved
tests/integration
and
tests/units
and update for FQCN.
Review
tests/sanity/ignore-*.txt
entries in the old collection.

Update
meta/runtime.yml
in the old collection.
Removing the content from the old collection
?

Create a PR against the source collection repository to remove the modules, module_utils, plugins, and docs_fragments related to this migration:

If you are removing an action plugin, remove the corresponding module that contains the documentation.

If you are removing a module, remove any corresponding action plugin that should stay with it.

Remove any entries about removed plugins from
meta/runtime.yml

. Ensure they are added into the new repo.

Remove sanity ignore lines from
tests/sanity/ignore*.txt

Remove associated integration tests from
tests/integrations/targets/
and unit tests from
tests/units/plugins/
.

if you are removing from content from
community.general

or
community.network
, remove entries from
.github/BOTMETA.yml
.

Carefully review
meta/runtime.yml
for any entries you may need to remove or update, in particular deprecated entries.

Update
meta/runtime.yml
to contain redirects for EVERY PLUGIN, pointing to the new collection name.

Warning

Maintainers for the old collection have to make sure that the PR is merged in a way that it does not break user experience and semantic versioning:

A new version containing the merged PR must not be released before the collection the content has been moved to has been released again, with that content contained in it. Otherwise the redirects cannot work and users relying on that content will experience breakage.

Once 1.0.0 of the collection from which the content has been removed has been released, such PRs can only be merged for a new

major

version (in other words, 2.0.0, 3.0.0, and so on).

Updating BOTMETA.yml

?

The

BOTMETA.yml

, for example in

community.general collection repository

, is the source of truth for:

ansibullbot

If the old and/or new collection has

ansibullbot

, its

BOTMETA.yml

must be updated correspondingly.

Ansibullbot will know how to redirect existing issues and PRs to the new repo. The build process for docs.ansible.com will know where to find the module docs.

\$modules/monitoring/grafana/grafana_plugin.py

:

migrated_to

:

community.grafana

\$modules/monitoring/grafana/grafana_dashboard.py

:

migrated_to

:

community.grafana

\$modules/monitoring/grafana/grafana_datasource.py

:

migrated_to

:

community.grafana

\$plugins/callback/grafana_annotations.py

:

maintainers

:

\$team_grafana

labels

:

monitoring grafana

migrated_to

:

community.grafana

\$plugins/doc_fragments/grafana.py

:

maintainers

:

\$team_grafana

labels

:

monitoring grafana

migrated_to

:

community.grafana

Example PR

The
migrated_to:
key must be added explicitly for every
file
. You cannot add
migrated_to
at the directory level. This is to allow module and plugin webdocs to be redirected to the new collection docs.
migrated_to:
MUST be added for every:
module
plugin
module_utils
contrib/inventory script
You do NOT need to add
migrated_to
for:
Unit tests
Integration tests
ReStructured Text docs (anything under
docs/docsite/rst/
)
Files that never existed in
ansible/ansible:devel
See also
Using Ansible collections
Learn how to install and use collections.
Contributing to Ansible-maintained Collections
Guidelines for contributing to selected collections
Communication
Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide
Previous
Next
© Copyright Ansible project contributors.
Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_collections_path.html

Ansible collection creator path ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Ansible collection creator path
Edit on GitHub
Ansible collection creator path
?

Note

If you are unfamiliar with Ansible collections, first take a look at the
Using Ansible collections guide

Ansible collections are a distribution format for Ansible content that can include playbooks, roles, modules, and plugins. A typical collection addresses a set of related use cases. For example, the `community.dns` collection includes modules and plugins to work with DNS.

You can install collections made by others or share yours with the community through a distribution server such as Ansible Galaxy

. Certified collections can be published to the Red Hat Automation Hub, a part of the Red Hat Ansible Automation Platform.

Creating and sharing collections is a great way of contributing to the Ansible project.

The Ansible community package consists of
`ansible-core`

, which, among other core components, includes the
`ansible.builtin`

collection maintained by the Core team, and a set of collections maintained by the community.

The purpose of this guide is to give you as a (potential) content creator a consistent overview of the Ansible collection creator journey from an idea for the first module/role to having your collection included in the Ansible community package. The

Collection development guidelines section

provides references to more detailed aspects of this journey.

The overall journey consists of the following milestones:

Examine currently available solutions

Create your content

Put your content in a collection

Write good user collection documentation

Publish your collection source code

Follow a versioning convention

Understand and implement testing and CI

Add tests

Implement continuous integration

Provide good contributor & maintainer documentation

Publish your collection on distribution servers

Make your collection a part of Ansible community package

Maintain

Communicate

Examine currently available solutions

?

If you have an idea for a new role or module/plugin, there is no need to reinvent the wheel if there is already a sufficient solution that solves your automation issue.

Therefore, first examine the currently available content including:

Ansible builtin modules and plugins

Ansible package collection index

Ansible Galaxy

Ansible Automation Hub

if you have the Ansible Automation Platform subscription

In case the solutions you found are not fully sufficient or have flaws, consider improving them rather than creating your own. Each collection includes information on where to create issues for that collection to propose your enhancement ideas.

If you already have your content written and used in your workflows, you can still consider integrating it to the existing solutions.

However, if these options do not apply to your collection ideas, we encourage you to create and share your own.

Create your content

?

You

tried

but have not found any sufficient solution for your automation issue.

Use one of the following guides:

Roles guide

: if you want to create a role.

Developer guide

: if you want to create a new Ansible module or plugin for your personal use.

Put your content in a collection

?

You

created

new content.

Now it is time to create a reusable and sharable collection.

Use the

Developing collections guide

to learn how.

We recommend you to use the

collection_template repository

as a basis for your collection.

Write good user collection documentation

?

Your collection

README.md

file should contain a quick-start installation and usage guides.

You can use the

community.general collection README file

as an example.

If your collection contains modules or plugins, make sure their documentation is comprehensive.

Use the

Module format and documentation guide

and

Ansible documentation style guide
to learn more.

Publish your collection source code

?

Publish your collection on a platform for software development and version control such as
GitHub

.

It can be your personal repository or your organization's one.

You can also

request

a repository under the

ansible-collections

organization.

Make sure your collection contains exhaustive license information.

Ansible is an open source project, so we encourage you to license it under one of open source licenses.

If you plan to submit your collection for inclusion in the Ansible community package, your collection must satisfy the
licensing requirements

.

If you have used the

collection_template repository

we recommended earlier as a skeleton for your collection, it already contains the

GNU

GPL

v3

license.

Follow a versioning convention

?

When releasing new versions of your collections, take the following recommended practices into consideration:

Follow a versioning convention. Using

SemVer

is highly recommended.

Base your releases on

Git tags

.

Understand and implement testing and CI

?

This section is applicable to collections containing modules and plugins.

For role testing, see the

Ansible Molecule

project.

Add tests

?

Testing your collection ensures that your code works well and integrates with other components such as
ansible-core

.

Take a look at the following documents:

Testing Ansible guide

: provides general information about testing.

Testing collections guide

: contains collection-specific testing information.

Implement continuous integration

?

Now make sure when pull requests are created in your collection repository they are automatically tested using a CI tool
such as GitHub Actions or Azure Pipelines.

The
collection_template repository
contains GitHub Actions
templates
you can adjust and use to enable the workflows in your repository.
Provide good contributor & maintainer documentation
?

See the
collection_template/README.md
as an example.

Publish your collection on distribution servers
?

To distribute your collection and allow others to conveniently use it, publish your collection on one or more distribution servers.

See the
Distributing collections guide
to learn how.

Make your collection a part of Ansible community package
?

Make your collection satisfy the
Ansible community package collections requirements
and submit it for inclusion.

See the
inclusion process description
to learn how.

Maintain
?

Maintain your collection.

See the
Ansible collection maintainer guidelines
for details.

Communicate
?

Engage with the community.

Take a look at the
Ansible communication guide
to see available communication options.

See also

Developing collections

A set of guidelines about collection development aspects

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_core.html

Developing ansible-core ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

[Interpreter Discovery](#)
[Releases and maintenance](#)
[Testing Strategies](#)
[Sanity Tests](#)
[Frequently Asked Questions](#)
[Glossary](#)
[Ansible Reference: Module Utilities](#)
[Special Variables](#)
[Red Hat Ansible Automation Platform](#)
[Ansible Automation Hub](#)
[Logging Ansible output](#)
[Roadmaps](#)
[Ansible Roadmap](#)
[ansible-core Roadmaps](#)
[Ansible](#)

[Developer Guide](#)

[Developing](#)

[ansible-core](#)

[Edit on GitHub](#)

[Developing](#)

[ansible-core](#)

[?](#)

[Although](#)

[ansible-core](#)

[\(the code hosted in the](#)

[ansible/ansible repository](#)

[on GitHub\)](#) includes a few plugins that can be swapped out by the playbook directives or configuration, much of the code there is not modular. The documents here give insight into how the parts of

[ansible-core](#)

[work together.](#)

[ansible-core](#)

[project branches and tags](#)

[Ansible module architecture](#)

[See also](#)

[Python API](#)

[Learn about the Python API for task execution](#)

[Developing plugins](#)

[Learn about developing plugins](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_inventory.html

Developing dynamic inventory ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing dynamic inventory
Edit on GitHub
Developing dynamic inventory
?

Ansible can pull inventory information from dynamic sources, including cloud sources, by using the supplied inventory plugins

. For details about how to pull inventory information, see

Working with dynamic inventory

. If the source you want is not currently covered by existing plugins, you can create your own inventory plugin as with any other plugin type.

In previous versions, you had to create a script or program that could output JSON in the correct format when invoked with the proper arguments.

You can still use and write inventory scripts, as we ensured backwards compatibility through the script inventory plugin

and there is no restriction on the programming language used.

If you choose to write a script, however, you will need to implement some features yourself such as caching, configuration management, dynamic variable and group composition, and so on.

If you use

inventory plugins

instead, you can use the Ansible codebase and add these common features automatically.

Topics

Inventory sources

Inventory plugins

Developing an inventory plugin

verify_file method

parse method

inventory object

inventory cache

constructed features

Common format for inventory sources

The ?auto? plugin

Inventory scripts

Inventory script conventions

Tuning the external inventory script

Inventory sources

?

Inventory sources are the input strings that inventory plugins work with.

An inventory source can be a path to a file or to a script, or it can be raw data that the plugin can interpret.

The table below shows some examples of inventory plugins and the source types that you can pass to them with -i

on the command line.

Plugin

Source

host list

A comma-separated list of hosts

yaml

Path to a YAML format data file

constructed

Path to a YAML configuration file

ini

Path to an INI formatted data file

virtualbox

Path to a YAML configuration file

script plugin

Path to an executable that outputs JSON

Inventory plugins

?

Like most plugin types (except modules), inventory plugins must be developed in Python. They execute on the control node and should therefore adhere to the

Control node requirements

.

Most of the documentation in

Developing plugins

also applies here. You should read that document first for a general understanding and then come back to this document for specifics on inventory plugins.

Normally, inventory plugins are executed at the start of a run, and before the playbooks, plays, or roles are loaded.

However, you can use the

meta:

refresh_inventory

task to clear the current inventory and execute the inventory plugins again, and this task will generate a new inventory.

If you use the persistent cache, inventory plugins can also use the configured cache plugin to store and retrieve data.

Caching inventory avoids making repeated and costly external calls.

Developing an inventory plugin

?

The first thing you want to do is use the base class:

from

ansible.plugins.inventory

import

BaseInventoryPlugin

class

InventoryModule

(

BaseInventoryPlugin

):

NAME

=

'myplugin'

used internally by Ansible, it should match the file name but not required

If the inventory plugin is in a collection, the NAME should be in the ?namespace.collection_name.myplugin? format. The base class has a couple of methods that each plugin should implement and a few helpers for parsing the inventory source and updating the inventory.

After you have the basic plugin working, you can incorporate other features by adding more base classes:

```
from
ansible.plugins.inventory
import
BaseInventoryPlugin
```

```
,
```

```
Constructable
```

```
,
```

```
Cacheable
```

```
class
InventoryModule
```

```
(
BaseInventoryPlugin
```

```
,
```

```
Constructable
```

```
,
```

```
Cacheable
```

```
):
NAME
```

```
=
'myplugin'
```

For the bulk of the work in a plugin, we mostly want to deal with 2 methods

```
verify_file
```

```
and
```

```
parse
```

```
.
```

```
verify_file method
```

```
?
```

Ansible uses this method to quickly determine if the inventory source is usable by the plugin. The determination does not need to be 100% accurate, as there might be an overlap in what plugins can handle and by default Ansible will try the enabled plugins as per their sequence.

```
def
verify_file
```

```
(
self
```

```
,
```

```
path
```

```
):
''' return true/false if this is possibly a valid file for this plugin to consume '''
```

```
valid
```

```
=
False
```

```
if
super
```

```
(
InventoryModule
```

```
,
```

```
self
```

```
)
```

```
.
```

```
verify_file
```

```
(
path
```

```
):
# base class verifies that file exists and is readable by current user
```



```

if
path
.
endswith
((
'virtualbox.yaml'
,
'virtualbox.yml'
,
'vbox.yaml'
,
'vbox.yml'
)):
valid
=
True
return
valid

```

In the above example, from the
virtualbox inventory plugin

, we screen for specific file name patterns to avoid attempting to consume any valid YAML file. You can add any type of condition here, but the most common one is ?extension matching?. If you implement extension matching for YAML configuration files, the path suffix <plugin_name>.<yml|yaml> should be accepted. All valid extensions should be documented in the plugin description.

The following is another example that does not use a ?file? but the inventory source string itself,

from the

host list

plugin:

```
def
```

```
verify_file
```

```
(
```

```
self
```

```
,
```

```
path
```

```
):
```

```
""" don't call base class as we don't expect a path, but a host list """
```

```
host_list
```

```
=
```

```
path
```

```
valid
```

```
=
```

```
False
```

```
b_path
```

```
=
```

```
to_bytes
```

```
(
```

```
host_list
```

```
,
```

```
errors
```

```
=
```

```
'surrogate_or_strict'
```

```
)
```

```
if
```

```
not
```

```

os
.
path
.
exists
(
b_path
)
and
','
in
host_list
:
# the path does NOT exist and there is a comma to indicate this is a 'host list'
valid
=
True
return
valid

```

This method is just to expedite the inventory process and avoid unnecessary parsing of sources that are easy to filter out before causing a parse error.

parse method
?

This method does the bulk of the work in the plugin.

It takes the following parameters:

inventory: inventory object with existing data and the methods to add hosts/groups/variables to inventory

loader: Ansible's DataLoader. The DataLoader can read files, auto load JSON/YAML and decrypt vaulted data, and cache read files.

path: string with inventory source (this is usually a path, but is not required)

cache: indicates whether the plugin should use or avoid caches (cache plugin and/or loader)

The base class does some minimal assignment for reuse in other methods.

```

def
parse
(
self
,
inventory
,
loader
,
path
,
cache
=
True
):
self
.
loader
=
loader
self
.
inventory

```

```

=
inventory
self
.
templar
=
Templar
(
loader
=
loader
)

```

It is up to the plugin now to parse the provided inventory source and translate it into Ansible inventory. To facilitate this, the example below uses a few helper functions:

NAME

```

=
'myplugin'
def
parse
(
self
,
inventory
,
loader
,
path
,
cache
=
True
):
# call base method to ensure properties are available for use with other helper methods
super
(
InventoryModule
,
self
)
.
parse
(
inventory
,
loader
,
path
,
cache
)
# this method will parse 'common format' inventory sources and
# update any options declared in DOCUMENTATION as needed
config
=

```

```

self
.
_read_config_data
(
path
)
# if NOT using _read_config_data you should call set_options directly,
# to process any defined configuration for this plugin,
# if you don't define any options you can skip
#self.set_options()
# example consuming options from inventory source
mysession
=
apilib
.
session
(
user
=
self
.
get_option
(
'api_user'
),
password
=
self
.
get_option
(
'api_pass'
),
server
=
self
.
get_option
(
'api_server'
)
)
# make requests to get data to feed into inventory
mydata
=
mysession
.
getitall
()
#parse data and create inventory objects:
for
colo
in
mydata

```

```

:
for
server
in
mydata
[
colo
][
'servers'
]:
self
.
inventory
.
add_host
(
server
[
'name'
])
self
.
inventory
.
set_variable
(
server
[
'name'
],
'ansible_host'
,
server
[
'external_ip'
])

```

The specifics will vary depending on API and structure returned. Remember that if you get an inventory source error or any other issue, you should

raise

`AnsibleParserError`

to let Ansible know that the source was invalid or the process failed.

For examples on how to implement an inventory plugin, see the source code here:

`lib/ansible/plugins/inventory`

```

.
inventory object
?

```

The
inventory
object passed to
`parse`
has helpful methods for populating inventory.
`add_group`
adds a group to inventory if it doesn't already exist. It takes the group name as the only positional argument.
`add_child`

adds a group or host that exists in inventory to a parent group in inventory. It takes two positional arguments, the name of the parent group and the name of the child group or host.

`add_host`

adds a host to inventory if it doesn't already exist, optionally to a specific group. It takes the host name as the first argument and accepts two optional keyword arguments,

`group`

and

`port`

.

`group`

is the name of a group in inventory, and

`port`

is an integer.

`set_variable`

adds a variable to a group or host in inventory. It takes three positional arguments: the name of the group or host, the name of the variable, and the value of the variable.

To create groups and variables using Jinja2 expressions, see the section on implementing constructed features below.

To see other inventory object methods, see the source code here:

`lib/ansible/inventory/data.py`

.

inventory cache

?

To cache the inventory, extend the inventory plugin documentation with the `inventory_cache` documentation fragment and use the `Cacheable` base class.

`extends_documentation_fragment`

:

-

`inventory_cache`

class

`InventoryModule`

(

`BaseInventoryPlugin`

,

`Constructable`

,

`Cacheable`

):

`NAME`

=

`'myplugin'`

Next, load the cache plugin specified by the user to read from and update the cache. If your inventory plugin uses YAML-based configuration files and the

`_read_config_data`

method, the cache plugin is loaded within that method. If your inventory plugin does not use

`_read_config_data`

, you must load the cache explicitly with

`load_cache_plugin`

.

`NAME`

=

`'myplugin'`

def

```

parse
(
self
,
inventory
,
loader
,
path
,
cache
=
True
):
super
(
InventoryModule

```

```

,
self
)
.
parse
(
inventory
,
loader
,
path
)
self
.

```

```

load_cache_plugin
()

```

Before using the cache plugin, you must retrieve a unique cache key by using the `get_cache_key`

method. This task needs to be done by all inventory modules using the cache, so that you don't use/overwrite other parts of the cache.

```

def
parse
(
self
,
inventory
,
loader
,
path
,
cache
=
True
):
super
(

```

```
InventoryModule
```

```
,
self
)
.
parse
(
inventory
,
loader
,
path
)
self
.
load_cache_plugin
()
cache_key
=
self
.
get_cache_key
(
path
)
```

Now that you've enabled caching, loaded the correct plugin, and retrieved a unique cache key, you can set up the flow of data between the cache and your inventory using the

cache

parameter of the

parse

method. This value comes from the inventory manager and indicates whether the inventory is being refreshed (such as by the

--flush-cache

or the meta task

refresh_inventory

). Although the cache shouldn't be used to populate the inventory when being refreshed, the cache should be updated with the new inventory if the user has enabled caching. You can use

self._cache

like a dictionary. The following pattern allows refreshing the inventory to work in conjunction with caching.

```
def
parse
(
self
,
inventory
,
loader
,
path
,
cache
=
True
):
```



```

super
(
InventoryModule
,
self
)
.
parse
(
inventory
,
loader
,
path
)
self
.
load_cache_plugin
()
cache_key
=
self
.
get_cache_key
(
path
)
# cache may be True or False at this point to indicate if the inventory is being refreshed
# get the user's cache option too to see if we should save the cache if it is changing
user_cache_setting
=
self
.
get_option
(
'cache'
)
# read if the user has caching enabled and the cache isn't being refreshed
attempt_to_read_cache
=
user_cache_setting
and
cache
# update if the user has caching enabled and the cache is being refreshed; update this value to True if the cache has
expired below
cache_needs_update
=
user_cache_setting
and
not
cache
# attempt to read the cache if inventory isn't being refreshed and the user has caching enabled
if
attempt_to_read_cache

```

```

:
try
:
results
=
self
.
_cache
[
cache_key
]
except
KeyError
:
# This occurs if the cache_key is not in the cache or if the cache_key expired, so the cache needs to be updated
cache_needs_update
=
True
if
not
attempt_to_read_cache
or
cache_needs_update
:
# parse the provided inventory source
results
=
self
.
get_inventory
()
if
cache_needs_update
:
self
.
_cache
[
cache_key
]
=
results
# submit the parsed data to the inventory object (add_host, set_variable, etc)
self
.
populate
(
results
)
After the
parse
method is complete, the contents of
self._cache
is used to set the cache plugin if the contents of the cache have changed.

```

You have three other cache methods available:

`set_cache_plugin`

forces the cache plugin to be set with the contents of

`self._cache`

, before the

`parse`

method completes

`update_cache_if_changed`

sets the cache plugin only if

`self._cache`

has been modified, before the

`parse`

method completes

`clear_cache`

flushes the cache, ultimately by calling the cache plugin's

`flush()`

method, whose implementation is dependent upon the particular cache plugin in use. Note that if the user is using the same cache backend for facts and inventory, both will get flushed. To avoid this, the user can specify a distinct cache backend in their inventory plugin configuration.

constructed features

?

Inventory plugins can create host variables and groups from Jinja2 expressions and variables by using features from the constructed

inventory plugin. To do this, use the

`Constructable`

base class and extend the inventory plugin's documentation with the

constructed

documentation fragment.

`extends_documentation_fragment`

:

-

constructed

class

`InventoryModule`

(

`BaseInventoryPlugin`

,

`Constructable`

):

`NAME`

=

`'ns.coll.myplugin'`

There are three main options in the

constructed

documentation fragment:

`compose`

creates variables using Jinja2 expressions. This is implemented by calling the

`_set_composite_vars`

method.

`keyed_groups`

creates groups of hosts based on variable values. This is implemented by calling the

`_add_host_to_keyed_groups`

method.

groups

creates groups based on Jinja2 conditionals. This is implemented by calling the `_add_host_to_composed_groups` method.

Each method should be called for every host added to inventory. Three positional arguments are required: the constructed option, a dictionary of variables, and a host name. Calling the method

`_set_composite_vars`

first will allow

keyed_groups

and

groups

to use the composed variables.

By default, undefined variables are ignored. This is permitted by default for

compose

so you can make the variable definitions depend on variables that will be populated later in a play from other sources.

For groups, it allows using variables that are not always present without having to use the

default

filter. To support configuring undefined variables to be an error, pass the constructed option

strict

to each of the methods as a keyword argument.

keyed_groups

and

groups

use any variables already associated with the host (for example, from an earlier inventory source).

`_add_host_to_keyed_groups`

and

`add_host_to_composed_groups`

can turn this off by passing the keyword argument

`fetch_hostvars`

.

Here is an example using all three methods:

def

add_host

(

self

,

hostname

,

host_vars

):

self

.

inventory

.

add_host

(

hostname

,

group

=

'all'

)

for

var_name

,

```

var_value
in
host_vars
.
items
():
self
.
inventory
.
set_variable
(
hostname
,
var_name
,
var_value
)
strict
=
self
.
get_option
(
'strict'
)
# Add variables created by the user's Jinja2 expressions to the host
self
.
_set_composite_vars
(
self
.
get_option
(
'compose'
),
host_vars
,
hostname
,
strict
=
True
)
# Create user-defined groups using variables and Jinja2 conditionals
self
.
_add_host_to_composed_groups
(
self
.
get_option
(

```

```

'groups'
),
host_vars
,
hostname
,
strict
=
strict
)
self
.
_add_host_to_keyed_groups
(
self
.
get_option
(
'keyed_groups'
),
host_vars
,
hostname
,
strict
=
strict
)

```

By default, group names created with `_add_host_to_composed_groups()` and `_add_host_to_keyed_groups()` are valid Python identifiers. Invalid characters are replaced with an underscore

. A plugin can change the sanitization used for the constructed features by setting `self._sanitize_group_name` to a new function. The core engine also does sanitization, so if the custom function is less strict it should be used in conjunction with the configuration setting `TRANSFORM_INVALID_GROUP_CHARS`

```

.
from
ansible.inventory.group
import
to_safe_group_name
class
InventoryModule
(
BaseInventoryPlugin
,
Constructable
):
NAME
=
'ns.coll.myplugin'

```

```

@staticmethod
def
custom_sanitizer
(
name
):
return
to_safe_group_name
(
name
,
replacer
=
"
)
def
parse
(
self
,
inventory
,
loader
,
path
,
cache
=
True
):
super
(
InventoryModule
,
self
)
.
parse
(
inventory
,
loader
,
path
)
self
.
__sanitize_group_name
=
custom_sanitizer
Common format for inventory sources
?

```

To simplify development, most plugins use a standard YAML-based configuration file as the inventory source. The file has only one required field

plugin
, which should contain the name of the plugin that is expected to consume the file.
Depending on other common features used, you might need other fields, and you can add custom options in each plugin as required.
For example, if you use the integrated caching,
cache_plugin
,
cache_timeout
and other cache-related fields could be present.
The ?auto? plugin
?
From Ansible 2.5 onwards, we include the
auto inventory plugin
and enable it by default. If the
plugin
field in your standard configuration file matches the name of your inventory plugin, the
auto
inventory plugin will load your plugin. The ?auto? plugin makes it easier to use your plugin without having to update configurations.
Inventory scripts
?
Even though we now have inventory plugins, we still support inventory scripts, not only for backwards compatibility but also to allow users to use other programming languages.
Inventory script conventions
?
Inventory scripts must accept the
--list
and
--host
<hostname>
arguments. Although other arguments are allowed, Ansible will not use them.
Such arguments might still be useful for executing the scripts directly.
When the script is called with the single argument
--list
, the script must output to stdout a JSON object that contains all the groups to be managed. Each group?s value should be either an object containing a list of each host, any child groups, and potential group variables, or simply a list of hosts:

```
{
"group001"
:
{
"hosts"
:
[
"host001"
,
"host002"
],
"vars"
:
{
"var1"
:
true
```



```

},
"children"
:
[
"group002"
],
},
"group002"
:
{
"hosts"
:
[
"host003"
,
"host004"
],
"vars"
:
{
"var2"
:
500
},
"children"
:[]
}
}

```

If any of the elements of a group are empty, they may be omitted from the output.

When called with the argument

```

--host
<hostname>

```

(where <hostname> is a host from above), the script must print a JSON object, either empty or containing variables to make them available to templates and playbooks. For example:

```

{
"VAR001"
:
"VALUE"
,
"VAR002"
:
"VALUE"
}

```

Printing variables is optional. If the script does not print variables, it should print an empty JSON object.

Tuning the external inventory script

?

New in version 1.3.

The stock inventory script system mentioned above works for all versions of Ansible, but calling

```

--host

```

for every host can be rather inefficient, especially if it involves API calls to a remote subsystem.

To avoid this inefficiency, if the inventory script returns a top-level element called ?_meta?, it is possible to return all the host variables in a single script execution. When this meta element contains a value for ?hostvars?, the inventory script will not be invoked with

```

--host

```

for each host. This behavior results in a significant performance increase for large numbers of hosts. The data to be added to the top-level JSON object looks like this:

```
{  
  
    # results of inventory script as above go here  
    # ...  
  
    "_meta": {  
        "hostvars": {  
            "host001": {  
                "var001": "value"  
            },  
            "host002": {  
                "var002": "value"  
            }  
        }  
    }  
}
```

To satisfy the requirements of using
_meta
, to prevent ansible from calling your inventory with
--host
you must at least populate
_meta
with an empty
hostvars
object.
For example:

```
{  
  
    # results of inventory script as above go here  
    # ...  
  
    "_meta": {  
        "hostvars": {}  
    }  
}
```

If you intend to replace an existing static inventory file with an inventory script, it must return a JSON object which contains an ?all? group that includes every host in the inventory as a member and every group in the inventory as a child. It should also include an ?ungrouped? group which contains all hosts which are not members of any other group. A skeleton example of this JSON object is:

```
{  
    "_meta"  
    :  
    {  
        "hostvars"  
        :  
        {}  
    },  
    "all"  
    :  
    {  
        "children"
```

```
[  
  "ungrouped"  
]  
},  
"ungrouped"  
:  
{  
  "children"  
:  
  [  
  ]  
}  
}
```

An easy way to see how this should look is using
`ansible-inventory`

, which also supports

`--list`

and

`--host`

parameters like an inventory script would.

See also

[Python API](#)

[Python API to Playbooks and Ad Hoc Task Execution](#)

[Developing modules](#)

[Get started with developing a module](#)

[Developing plugins](#)

[How to develop plugins](#)

[AWX](#)

[REST API endpoint and GUI for Ansible, syncs with dynamic inventory](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_locally.html

Adding modules and plugins locally ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Adding modules and plugins locally
Edit on GitHub
Adding modules and plugins locally

?

You can extend Ansible by adding custom modules or plugins. You can create them from scratch or copy existing ones for local use. You can store a local module or plugin on your Ansible control node and share it with your team or organization. You can also share plugins and modules by including them in a collection, then publishing the collection on Ansible Galaxy.

If you are using a local module or plugin but Ansible cannot find it, this page is all you need.

If you want to create a plugin or a module, see

Developing plugins

,

Developing modules

and

Developing collections

.

Extending Ansible with local modules and plugins offers shortcuts such as:

You can copy other people's modules and plugins.

When writing a new module, you can choose any programming language you like.

You do not have to clone any repositories.

You do not have to open a pull request.

You do not have to add tests (though we recommend that you do!).

Modules and plugins: what is the difference?

Adding modules and plugins in collections

Adding a module or plugin outside of a collection

Adding standalone local modules for all playbooks and roles

Adding standalone local modules for selected playbooks or a single role

Adding a non-module plugin locally outside of a collection

Adding local non-module plugins for all playbooks and roles

Adding standalone local plugins for selected playbooks or a single role

Using

ansible.legacy

to access custom versions of an

ansible.builtin

module

Modules and plugins: what is the difference?

?

If you are looking to add functionality to Ansible, you might wonder whether you need a module or a plugin. Here is a

quick overview to help you understand what you need:

Plugins

extend Ansible's core functionality. Most plugin types execute on the control node within the

/usr/bin/ansible

process. Plugins offer options and extensions for the core features of Ansible: transforming data, logging output, connecting to inventory, and more.

Modules are a type of plugin that execute automation tasks on a ?target? (usually a remote system). Modules work as standalone scripts that Ansible executes in their own process outside of the control node. Modules interface with Ansible mostly with JSON, accepting arguments and returning information by printing a JSON string to stdout before exiting. Unlike the other plugins (which must be written in Python), modules can be written in any language; although Ansible provides modules in Python and Powershell only.

Adding modules and plugins in collections

?

You can add modules and plugins by

creating a collection

. With a collection, you can use custom modules and plugins in any playbook or role. You can share your collection easily at any time through Ansible Galaxy.

The rest of this page describes other methods of using local, standalone modules or plugins.

Adding a module or plugin outside of a collection

?

You can configure Ansible to load standalone local modules or plugins in specific locations and make them available to all playbooks and roles (using configured paths). Alternatively, you can make a non-collection local module or plugin available only to certain playbooks or roles (with adjacent paths).

Adding standalone local modules for all playbooks and roles

?

To load standalone local modules automatically and make them available to all playbooks and roles, use the

DEFAULT_MODULE_PATH

configuration setting or the

ANSIBLE_LIBRARY

environment variable. The configuration setting and environment variable take a colon-separated list, similar to

\$PATH

. You have two options:

Add your standalone local module to one of the default configured locations. See the

DEFAULT_MODULE_PATH

configuration setting for details. Default locations may change without notice.

Add the location of your standalone local module to an environment variable or configuration:

the

ANSIBLE_LIBRARY

environment variable

the

DEFAULT_MODULE_PATH

configuration setting

To view your current configuration settings for modules:

ansible-config dump |grep DEFAULT_MODULE_PATH

After you save your module file in one of these locations, Ansible loads it and you can use it in any local task, playbook, or role.

To confirm that

my_local_module

is available:

type

ansible

localhost

-m

my_local_module

to see the output for that module, or

type

ansible-doc

-t

module

my_local_module

to see the documentation for that module

Note

This applies to all plugin types but requires specific configuration and/or adjacent directories for each plugin type, see below.

Note

The

ansible-doc

command can parse module documentation from modules written in Python or an adjacent YAML file. If you have a module written in a programming language other than Python, you should write the documentation in a Python or YAML file adjacent to the module file.

Adjacent YAML documentation files

Adding standalone local modules for selected playbooks or a single role

?

Ansible automatically loads all executable files from certain directories adjacent to your playbook or role as modules. Standalone modules in these locations are available only to the specific playbook, playbooks, or role in the parent directory.

To use a standalone module only in a selected playbook or playbooks, store the module in a subdirectory called library

in the directory that contains the playbook or playbooks.

To use a standalone module only in a single role, store the module in a subdirectory called

library

within that role.

Note

This applies to all plugin types but requires specific configuration and/or adjacent directories for each plugin type, see below.

Warning

Roles contained in collections cannot contain any modules or other plugins. All plugins in a collection must live in the collection

plugins

directory tree. All plugins in that tree are accessible to all roles in the collection. If you are developing new modules, we recommend distributing them in

collections

, not in roles.

Adding a non-module plugin locally outside of a collection

?

You can configure Ansible to load standalone local plugins in a specified location or locations and make them available to all playbooks and roles. Alternatively, you can make a standalone local plugin available only to specific playbooks or roles.

Note

Although modules are plugins, the naming patterns for directory names and environment variables that apply to other plugin types do not apply to modules. See

Adding a module or plugin outside of a collection

.

Adding local non-module plugins for all playbooks and roles

?

To load standalone local plugins automatically and make them available to all playbooks and roles, use the configuration setting or environment variable for the type of plugin you are adding. These configuration settings and environment variables take a colon-separated list, similar to

\$PATH

. You have two options:

Add your local plugin to one of the default configured locations. See

configuration settings

for details on the correct configuration setting for the plugin type. Default locations may change without notice.

Add the location of your local plugin to an environment variable or configuration:

the relevant

ANSIBLE_plugin_type_PLUGINS

environment variable - for example,

\$ANSIBLE_INVENTORY_PLUGINS

or

\$ANSIBLE_VARS_PLUGINS

the relevant

plugin_type_PATH

configuration setting, most of which begin with

DEFAULT_

- for example,

DEFAULT_CALLBACK_PLUGIN_PATH

or

DEFAULT_FILTER_PLUGIN_PATH

or

BECOME_PLUGIN_PATH

To view your current configuration settings for non-module plugins:

ansible-config dump |grep plugin_type_PATH

After your plugin file is added to one of these locations, Ansible loads it and you can use it in any local module, task, playbook, or role. For more information on environment variables and configuration settings, see

Ansible Configuration Settings

.

To confirm that

plugins/plugin_type/my_local_plugin

is available:

type

ansible-doc

-t

<plugin_type>

my_local_lookup_plugin

to see the documentation for that plugin - for example,

ansible-doc

-t

lookup

my_local_lookup_plugin

The

ansible-doc

command works for most plugin types, but not for action, filter, or test plugins. See

ansible-doc

for more details.

Adding standalone local plugins for selected playbooks or a single role

?

Ansible automatically loads all plugins from certain directories adjacent to your playbook or role, loading each type of plugin separately from a directory named for the type of plugin. Standalone plugins in these locations are available only to the specific playbook, playbooks, or role in the parent directory.

To use a standalone plugin only in a selected playbook or playbooks, store the plugin in a subdirectory for the correct plugin_type

(for example,

callback_plugins

or

inventory_plugins

) in the directory that contains the playbooks. These directories must use the

_plugins

suffix. For a full list of plugin types, see

Working with plugins

.

To use a standalone plugin only in a single role, store the plugin in a subdirectory for the correct

plugin_type

(for example,

cache_plugins

or

strategy_plugins

) within that role. When shipped as part of a role, the plugin is available as soon as the role is executed. These

directories must use the

_plugins

suffix. For a full list of plugin types, see

Working with plugins

.

Warning

Roles contained in collections cannot contain any plugins. All plugins in a collection must live in the collection

plugins

directory tree. All plugins in that tree are accessible to all roles in the collection. If you are developing new plugins, we

recommend distributing them in

collections

, not in roles.

Warning

Some plugin types are needed early during Ansible execution, such as callbacks, inventory, and cache. These plugin types cannot be loaded dynamically and must exist in configured paths or be referenced by FQCN in configuration.

Using

ansible.legacy

to access custom versions of an

ansible.builtin

module

?

If you need to override one of the

ansible.builtin

modules and are using FQCN, you need to use

ansible.legacy

as part of the fully-qualified collection name (FQCN). For example, if you had your own

copy

module, you would access it as

ansible.legacy.copy

. See

Using ansible.legacy to access local custom modules from collections-based roles

for details on how to use custom modules with roles within a collection.

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_modules_best_practices.html

Conventions, tips, and pitfalls ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Conventions, tips, and pitfalls
Edit on GitHub
Conventions, tips, and pitfalls

?

Topics

Scoping your module(s)

Designing module interfaces

General guidelines & tips

Functions and Methods

Python tips

Importing and using shared code

Handling module failures

Handling exceptions (bugs) gracefully

Creating correct and informative module output

Following Ansible conventions

Module Security

As you design and develop modules, follow these basic conventions and tips for clean, usable code:

Scoping your module(s)

?

Especially if you want to contribute your module(s) to an existing Ansible Collection, make sure each module includes enough logic and functionality, but not too much. If these guidelines seem confusing, consider whether you really need to write a module at all.

Each module should have a concise and well-defined functionality. Basically, follow the UNIX philosophy of doing one thing well.

Do not add

get

,

list

or

info

state options to an existing module - create a new

_info

or

_facts

module.

Modules should not require that a user know all the underlying options of an API/tool to be used. For example, if the legal values for a required module option cannot be documented, the module does not belong in Ansible Core.

Modules should encompass much of the logic for interacting with a resource. A lightweight wrapper around a complex API forces users to offload too much logic into their playbooks. If you want to connect Ansible to a complex API, create multiple modules that interact with smaller individual pieces of the API.

Avoid creating a module that does the work of other modules; this leads to code duplication and divergence, and makes things less uniform, unpredictable and harder to maintain. Modules should be the building blocks. If you are asking ?how can I have a module execute other modules? ? you want to write a role.

Designing module interfaces

?

If your module is addressing an object, the option for that object should be called `name`

whenever possible, or accept

`name`

as an alias.

Modules accepting boolean status should accept

`yes`

,

`no`

,

`true`

,

`false`

, or anything else a user may likely throw at them. The AnsibleModule common code supports this with `type='bool'`

.

Avoid

`action`

/

`command`

, they are imperative and not declarative, there are other ways to express the same thing.

General guidelines & tips

?

Each module should be self-contained in one file, so it can be auto-transferred by `ansible-core`

.

Module name MUST use underscores instead of hyphens or spaces as a word separator. Using hyphens and spaces will prevent

`ansible-core`

from importing your module.

Always use the

`hacking/test-module.py`

script when developing modules - it will warn you about common pitfalls.

If you have a local module that returns information specific to your installations, a good name for this module is `site_info`

.

Eliminate or minimize dependencies. If your module has dependencies, document them at the top of the module file and raise JSON error messages when dependency import fails.

Don't write to files directly; use a temporary file and then use the

`atomic_move`

function from

`ansible.module_utils.basic`

to move the updated temporary file into place. This prevents data corruption and ensures that the correct context for the file is kept.

Avoid creating caches. Ansible is designed without a central server or authority, so you cannot guarantee it will not run

with different permissions, options or locations. If you need a central authority, have it on top of Ansible (for example, using bastion/cm/ci server, AWX, or the Red Hat Ansible Automation Platform); do not try to build it into modules.

If you package your module(s) in an RPM, install the modules on the control machine in

/usr/share/ansible

. Packaging modules in RPMs is optional.

Functions and Methods

?

Each function should be concise and should describe a meaningful amount of work.

?Don?t repeat yourself? is generally a good philosophy.

Function names should use underscores:

my_function_name

.

The name of each function should describe what the function does.

Each function should have a docstring.

If your code is too nested, that?s usually a sign the loop body could benefit from being a function. Parts of our existing code are not the best examples of this at times.

Python tips

?

Include a

main

function that wraps the normal execution.

Call your

main

function from a conditional so you can import it into unit tests - for example:

if

__name__

==

'__main__'

:

main

()

Importing and using shared code

?

Use shared code whenever possible - don?t reinvent the wheel. Ansible offers the

AnsibleModule

common Python code, plus

utilities

for many common use cases and patterns. You can also create documentation fragments for docs that apply to multiple modules.

Import

ansible.module_utils

code in the same place as you import other libraries.

Do NOT use wildcards (*) for importing other python modules; instead, list the function(s) you are importing (for example,

from

some.other_python_module.basic

import

otherFunction

).

Import custom packages in

try

/

except

, capture any import errors, and handle them with

```

fail_json()
in
main()
. For example:
import
traceback
from
ansible.module_utils.basic
import
missing_required_lib
LIB_IMP_ERR
=
None
try
:
import
foo
HAS_LIB
=
True
except
:
HAS_LIB
=
False
LIB_IMP_ERR
=
traceback
.
format_exc
()
Then in
main()
, just after the argspec, do
if
not
HAS_LIB
:
module
.
fail_json
(
msg
=
missing_required_lib
(
"foo"
),
exception
=
LIB_IMP_ERR
)

```

And document the dependency in the requirements

section of your module?
DOCUMENTATION block

.

Handling module failures

?

When your module fails, help users understand what went wrong. If you are using the `AnsibleModule`

common Python code, the

`failed`

element will be included for you automatically when you call

`fail_json`

. For polite module failure behavior:

Include a key of

`failed`

along with a string explanation in

`msg`

. If you don't do this, Ansible will use standard return codes: 0=success and non-zero=failure.

Don't raise a `traceback` (`stacktrace`). Ansible can deal with `stacktraces` and automatically converts anything unparsable into a failed result, but raising a `stacktrace` on module failure is not user-friendly.

Do not use

`sys.exit()`

. Use

`fail_json()`

from the module object.

Handling exceptions (bugs) gracefully

?

Validate upfront?fail fast and return useful and clear error messages.

Use defensive programming?use a simple design for your module, handle errors gracefully, and avoid direct `stacktraces`.

Fail predictably?if we must fail, do it in a way that is the most expected. Either mimic the underlying tool or the general way the system works.

Give out a useful message on what you were doing and add exception messages to that.

Avoid `catchall` exceptions, they are not very useful unless the underlying API gives very good error messages pertaining the attempted action.

Creating correct and informative module output

?

Modules must output valid JSON only. Follow these guidelines for creating correct, useful module output:

Module return data must be encoded as strict UTF-8. Modules that cannot return UTF-8 encoded data should return the data encoded by something such as `base64`. Optionally modules can make the determination if they can encode as UTF-8 and utilize

`errors='replace'`

to replace non UTF-8 characters making the return values lossy.

Make your top-level return type a hash (dictionary).

Nest complex return values within the top-level hash.

Incorporate any lists or simple scalar values within the top-level return hash.

Do not send module output to standard error, because the system will merge standard out with standard error and prevent the JSON from parsing.

Capture standard error and return it as a variable in the JSON on standard out. This is how the `command` module is implemented.

Never do

`print("some`

`status`

`message")`

in a module, because it will not produce valid JSON output.

Always return useful data, even when there is no change.

Be consistent about returns (some modules are too random), unless it is detrimental to the state/action.

Make returns reusable?most of the time you don't want to read it, but you do want to process it and re-purpose it.

Return diff if in diff mode. This is not required for all modules, as it won't make sense for certain ones, but please include it when applicable.

Enable your return values to be serialized as JSON with Python's standard JSON encoder and decoder

library. Basic python types (strings, int, dicts, lists, and so on) are serializable.

Do not return an object using `exit_json()`. Instead, convert the fields you need from the object into the fields of a dictionary and return the dictionary.

Results from many hosts will be aggregated at once, so your module should return only relevant output. Returning the entire contents of a log file is generally bad form.

If a module returns `stderr` or otherwise fails to produce valid JSON, the actual output will still be shown in Ansible, but the command will not succeed.

Following Ansible conventions

?

Ansible conventions offer a predictable user interface across all modules, playbooks, and roles. To follow Ansible conventions in your module development:

Use consistent names across modules (yes, we have many legacy deviations - don't make the problem worse!).

Use consistent options (arguments) within your module(s).

Do not use `?message?` or `?syslog_facility?` as an option name, because this is used internally by Ansible.

Normalize options with other modules - if Ansible and the API your module connects to use different names for the same option, add aliases to your options so the user can choose which names to use in tasks and playbooks.

Return facts from

`*_facts`

modules in the

`ansible_facts`

field of the

result dictionary

so other modules can access them.

Implement

`check_mode`

in all

`*_info`

and

`*_facts`

modules. Playbooks which conditionalize based on fact information will only conditionalize correctly in

`check_mode`

if the facts are returned in

`check_mode`

. Usually you can add

`supports_check_mode=True`

when instantiating

`AnsibleModule`

.

Use module-specific environment variables. For example, if you use the helpers in

`module_utils.api`

for basic authentication with

`module_utils.urls.fetch_url()`

and you fall back on environment variables for default values, use a module-specific environment variable like

`API_<MODULENAME>_USERNAME`

to avoid conflicts between modules.

Keep module options simple and focused - if you're loading a lot of choices/states on an existing option, consider adding a new, simple option instead.

Keep options small when possible. Passing a large data structure to an option might save us a few tasks, but it adds a complex requirement that we cannot easily validate before passing on to the module.

If you want to pass complex data to an option, write an expert module that allows this, along with several smaller modules that provide a more ?atomic? operation against the underlying APIs and services. Complex operations require complex data. Let the user choose whether to reflect that complexity in tasks and plays or in vars files.

Implement declarative operations (not CRUD) so the user can ignore existing state and focus on final state. For example, use

started/stopped

,

present/absent

.

Strive for a consistent final state (aka idempotency). If running your module twice in a row against the same system would result in two different states, see if you can redesign or rewrite to achieve consistent final state. If you can?t, document the behavior and the reasons for it.

Provide consistent return values within the standard Ansible return structure, even if NA/None are used for keys normally returned under other options.

Module Security

?

Avoid passing user input from the shell.

Always check return codes.

You must always use

module.run_command

, not

subprocess

or

Popen

or

os.system

.

Avoid using the shell unless absolutely necessary.

If you must use the shell, you must pass

use_unsafe_shell=True

to

module.run_command

.

If any variables in your module can come from user input with

use_unsafe_shell=True

, you must wrap them with

pipes.quote(x)

.

When fetching URLs, use

fetch_url

or

open_url

from

ansible.module_utils.urls

. Do not use

urllib2

, which does not natively verify TLS certificates and so is insecure for https.

Sensitive values marked with

no_log=True

will automatically have that value stripped from module return values. If your module could return these sensitive values as part of a dictionary key name, you should call the

ansible.module_utils.basic.sanitize_keys()

function to strip the values from the keys. See the [uri](#) module for an example.

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_modules_general_windows.html

Windows module development walkthrough ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Windows module development walkthrough
Edit on GitHub
Windows module development walkthrough
?

In this section, we will walk through developing, testing, and debugging an Ansible Windows module.

Because Windows modules are written in Powershell and need to be run on a Windows host, this guide differs from the usual development walkthrough guide.

What's covered in this section:

Windows environment setup
Create a Windows server in a VM
Create an Ansible inventory
Provisioning the environment
Windows new module development
Windows module utilities
Exposing shared module options
Windows playbook module testing
Windows debugging
Windows unit testing
Windows integration testing
Windows communication and development support
Windows environment setup
?

Unlike Python module development which can be run on the host that runs Ansible, Windows modules need to be written and tested for Windows hosts.

While evaluation editions of Windows can be downloaded from Microsoft, these images are usually not ready to be used by Ansible without further modification. The easiest way to set up a Windows host so that it is ready to be used by Ansible is to set up a virtual machine using Vagrant.

Vagrant can be used to download existing OS images called boxes

that are then

deployed to a hypervisor like VirtualBox. These boxes can either be created and stored offline or they can be downloaded from a central repository called Vagrant Cloud.

This guide will use the Vagrant boxes created by the

packer-windoze

repository which have also been uploaded to

Vagrant Cloud

.

To find out more info on how these images are created, please go to the GitHub repo and look at the

README

file.

Before you can get started, the following programs must be installed (please consult the Vagrant and VirtualBox documentation for installation instructions):

Vagrant

VirtualBox

Create a Windows server in a VM

?

To create a single Windows Server 2016 instance, run the following:

vagrant

init

jborean93/WindowsServer2016

vagrant

up

This will download the Vagrant box from Vagrant Cloud and add it to the local boxes on your host and then start up that instance in VirtualBox. When starting for the first time, the Windows VM will run through the sysprep process and then create a HTTP and HTTPS WinRM listener automatically. Vagrant will finish its process once the listeners are online, after which the VM can be used by Ansible.

Create an Ansible inventory

?

The following Ansible inventory file can be used to connect to the newly created Windows VM:

[windows]

WindowsServer ansible_host

=

127.0.0.1

[windows:vars]

ansible_user

=

vagrant

ansible_password

=

vagrant

ansible_port

=

55986

ansible_connection

=

winrm

ansible_winrm_transport

=

ntlm

ansible_winrm_server_cert_validation

=

ignore

Note

The port

55986

is automatically forwarded by Vagrant to the

Windows host that was created, if this conflicts with an existing local port then Vagrant will automatically use another one at random and display show that in the output.

The OS that is created is based on the image set. The following images can be used:

jborean93/WindowsServer2012

jborean93/WindowsServer2012R2

jborean93/WindowsServer2016

jborean93/WindowsServer2019

jborean93/WindowsServer2022

When the host is online, it can be accessed by RDP on 127.0.0.1:3389

but the

port may differ depending if there was a conflict. To get rid of the host, run `vagrant`

`destroy`

`--force`

and Vagrant will automatically remove the VM and any other files associated with that VM.

While this is useful when testing modules on a single Windows instance, these hosts won't work without modification with domain-based modules. The Vagrantfile at

`ansible-windows`

can be used to create a test domain environment to be used in Ansible. This repo contains three files which are used by both Ansible and Vagrant to create multiple Windows hosts in a domain environment. These files are:

`Vagrantfile`

: The Vagrant file that reads the inventory setup of `inventory.yml`

and provisions the hosts that are required in `inventory.yml`

: Contains the hosts that are required and other connection information such as IP addresses and forwarded ports in `main.yml`

: Ansible playbook called by Vagrant to provision the domain control node and join the child hosts to the domain

By default, these files will create the following environment:

A single AD domain controller running on Windows Server 2016

Five child hosts for each major Windows Server version joined to that domain

A domain with the DNS name

`domain.local`

A local administrator account on each host with the username

`vagrant`

and password

`vagrant`

A domain admin account

`vagrant-domain@domain.local`

with the password

`VagrantPass1`

The domain name and accounts can be modified by changing the variables

`domain_*`

in the

`inventory.yml`

file if it is required. The inventory

file can also be modified to provision more or less servers by changing the hosts that are defined under the

domain_children

key. The host variable

ansible_host

is the private IP that will be assigned to the VirtualBox host

only network adapter while

vagrant_box

is the box that will be used to

create the VM.

Provisioning the environment

?

To provision the environment as is, run the following:

git

clone

<https://github.com/jborean93/ansible-windows.git>

cd

vagrant

vagrant

up

Note

Vagrant provisions each host sequentially so this can take some time

to complete. If any errors occur during the Ansible phase of setting up the

domain, run

vagrant

provision

to rerun just that step.

Unlike setting up a single Windows instance with Vagrant, these hosts can also

be accessed using the IP address directly as well as through the forwarded

ports. It is easier to access it over the host only network adapter as the

normal protocol ports are used, for example RDP is still over

3389

. In cases where

the host cannot be resolved using the host only network IP, the following

protocols can be access over

127.0.0.1

using these forwarded ports:

RDP

: 295xx

SSH

: 296xx

WinRM

HTTP

: 297xx

WinRM

HTTPS

: 298xx

SMB

: 299xx

Replace

xx

with the entry number in the inventory file where the domain

controller started with

00

and is incremented from there. For example, in

the default

inventory.yml
file, WinRM over HTTPS for
SERVER2012R2
is
forwarded over port
29804
as it is the fourth entry in
domain_children

.
Windows new module development
?

When creating a new module there are a few things to keep in mind:

Module code is in Powershell (.ps1) files while the documentation is contained in Python (.py) files of the same name

Avoid using

Write-Host/Debug/Verbose/Error

in the module and add what needs to be returned to the

\$module.Result

variable

To fail a module, call

\$module.FailJson("failure

message

here")

, an Exception or ErrorRecord can be set to the second argument for a more descriptive error message

You can pass in the exception or ErrorRecord as a second argument to

FailJson("failure",

\$_)

to get a more detailed output

Most new modules require check mode and integration tests before they are merged into the main Ansible codebase

Avoid using try/catch statements over a large code block, rather use them for individual calls so the error message can be more descriptive

Try and catch specific exceptions when using try/catch statements

Avoid using PSCustomObjects unless necessary

Look for common functions in

./lib/ansible/module_utils/powershell/

and use the code there instead of duplicating work. These can be imported by adding the line

#Requires

-Module

*

where * is the file name to import, and will be automatically included with the module code sent to the Windows target when run through Ansible

As well as PowerShell module utils, C# module utils are stored in

./lib/ansible/module_utils/csharp/

and are automatically imported in a module execution if the line

#AnsibleRequires

-CSharpUtil

*

is present

C# and PowerShell module utils achieve the same goal but C# allows a developer to implement low level tasks, such as calling the Win32 API, and can be faster in some cases

Ensure the code runs under Powershell v5.1 and higher on Windows Server 2016 and higher; if higher minimum Powershell or OS versions are required, ensure the documentation reflects this clearly

Ansible runs modules under strictmode version 2.0. Be sure to test with that enabled by putting

Set-StrictMode

-Version

2.0

at the top of your dev script

Favor native Powershell cmdlets over executable calls if possible

Use the full cmdlet name instead of aliases, for example

Remove-Item

over

rm

Use named parameters with cmdlets, for example

Remove-Item

-Path

C:\temp

over

Remove-Item

C:\temp

A very basic Powershell module

win_environment

incorporates best practices for Powershell modules. It demonstrates how to implement check-mode and diff-support, and also shows a warning to the user when a specific condition is met.

A slightly more advanced module is

win_uri

which additionally shows how to use different parameter types (bool, str, int, list, dict, path) and a selection of choices for parameters, how to fail a module and how to handle exceptions.

As part of the new

AnsibleModule

wrapper, the input parameters are defined and validated based on an argument

spec. The following options can be set at the root level of the argument spec:

mutually_exclusive

: A list of lists, where the inner list contains module options that cannot be set together

no_log

: Stops the module from emitting any logs to the Windows Event log

options

: A dictionary where the key is the module option and the value is the spec for that option

required_by

: A dictionary where the option(s) specified by the value must be set if the option specified by the key is also set

required_if

: A list of lists where the inner list contains 3 or 4 elements;

The first element is the module option to check the value against

The second element is the value of the option specified by the first element, if matched then the required if check is run

The third element is a list of required module options when the above is matched

An optional fourth element is a boolean that states whether all module options in the third elements are required (default:

\$false

) or only one (

\$true

)

required_one_of

: A list of lists, where the inner list contains module options where at least one must be set

required_together

: A list of lists, where the inner list contains module options that must be set together

supports_check_mode

: Whether the module supports check mode, by default this is

\$false

The actual input options for a module are set within the

options

value as a dictionary. The keys of this dictionary

are the module option names while the values are the spec of that module option. Each spec can have the following options set:

aliases

: A list of aliases for the module option

choices

: A list of valid values for the module option, if

type=list

then each list value is validated against the choices and not the list itself

default

: The default value for the module option if not set

deprecated_aliases

: A list of hashtables that define aliases that are deprecated and the versions they will be removed in. Each entry must contain the keys

name

and

collection_name

with either

version

or

date

elements

: When

type=list

, this sets the type of each list value, the values are the same as

type

no_log

: Will sanitize the input value before being returned in the

module_invocation

return value

removed_in_version

: States when a deprecated module option is to be removed, a warning is displayed to the end user if set

removed_at_date

: States the date (YYYY-MM-DD) when a deprecated module option will be removed, a warning is displayed to the end user if set

removed_from_collection

: States from which collection the deprecated module option will be removed; must be specified if one of

removed_in_version

and

removed_at_date

is specified

required

: Will fail when the module option is not set

type

: The type of the module option, if not set then it defaults to

str

. The valid types are;

bool

: A boolean value

dict

: A dictionary value, if the input is a JSON or key=value string then it is converted to dictionary

float

: A float or

Single

value

int
: An Int32 value

json
: A string where the value is converted to a JSON string if the input is a dictionary

list
: A list of values,
elements=<type>
can convert the individual list value types if set. If
elements=dict
then
options
is defined, the values will be validated against the argument spec. When the input is a string then the string is split by
,
and any whitespace is trimmed

path
: A string where values likes
%TEMP%
are expanded based on environment values. If the input value starts with
\\?\
then no expansion is run

raw
: No conversions occur on the value passed in by Ansible

sid
: Will convert Windows security identifier values or Windows account names to a
SecurityIdentifier

value

str
: The value is converted to a string

When
type=dict
, or
type=list
and
elements=dict
, the following keys can also be set for that module option:
apply_defaults
: The value is based on the
options
spec defaults for that key if
True
and null if
False
. Only valid when the module option is not defined by the user and
type=dict
.
mutually_exclusive
: Same as the root level
mutually_exclusive
but validated against the values in the sub dict
options
: Same as the root level
options
but contains the valid options for the sub option
required_if

: Same as the root level

required_if

but validated against the values in the sub dict

required_by

: Same as the root level

required_by

but validated against the values in the sub dict

required_together

: Same as the root level

required_together

but validated against the values in the sub dict

required_one_of

: Same as the root level

required_one_of

but validated against the values in the sub dict

A module type can also be a delegate function that converts the value to whatever is required by the module option. For example the following snippet shows how to create a custom type that creates a

UInt64

value:

\$spec

=

@{

uint64_type

=

@{

type

=

[Func[[Object], [UInt64]]]

{

[System.UInt64]

::

Parse

(

\$args

[

0

])

}

}

}

\$uint64_type

=

\$module

.

Params

.

uint64_type

When in doubt, look at some of the other core modules and see how things have been implemented there.

Sometimes there are multiple ways that Windows offers to complete a task; this is the order to favor when writing modules:

Native Powershell cmdlets like

Remove-Item

-Path

C:\temp
 -Recurse
 .NET classes like
 [System.IO.Path]::GetRandomFileName()
 WMI objects through the
 New-CimInstance
 cmdlet
 COM objects through
 New-Object
 -ComObject
 cmdlet
 Calls to native executables like
 Secedit.exe
 PowerShell modules support a small subset of the
 #Requires
 options built
 into PowerShell as well as some Ansible-specific requirements specified by
 #AnsibleRequires
 . These statements can be placed at any point in the script,
 but are most commonly near the top. They are used to make it easier to state the
 requirements of the module without writing any of the checks. Each
 requires
 statement must be on its own line, but there can be multiple requires statements
 in one script.
 These are the checks that can be used within Ansible modules:
 #Requires
 -Module
 Ansible.ModuleUtils.<module_util>
 : Added in Ansible 2.4, specifies a module_util to load in for the module execution.
 #Requires
 -Version
 x.y
 : Added in Ansible 2.5, specifies the version of PowerShell that is required by the module. The module will fail if this
 requirement is not met.
 #AnsibleRequires
 -PowerShell
 <module_util>
 : Added in Ansible 2.8, like
 #Requires
 -Module
 , this specifies a module_util to load in for module execution.
 #AnsibleRequires
 -CSharpUtil
 <module_util>
 : Added in Ansible 2.8, specifies a C# module_util to load in for the module execution.
 #AnsibleRequires
 -OSVersion
 x.y
 : Added in Ansible 2.5, specifies the OS build version that is required by the module and will fail if this requirement is not
 met. The actual OS version is derived from
 [Environment]::OSVersion.Version
 .
 #AnsibleRequires
 -Become

: Added in Ansible 2.5, forces the exec runner to run the module with become

, which is primarily used to bypass WinRM restrictions. If

ansible_become_user

is not specified then the

SYSTEM

account is used instead.

The

#AnsibleRequires

-PowerShell

and

#AnsibleRequires

-CSharpUtil

support further features such as:

Importing a util contained in a collection (added in Ansible 2.9)

Importing a util by relative names (added in Ansible 2.10)

Specifying the util is optional by adding

-Optional

to the import

declaration (added in Ansible 2.12).

See the below examples for more details:

Imports the PowerShell Ansible.ModuleUtils.Legacy provided by Ansible itself

#AnsibleRequires -PowerShell Ansible.ModuleUtils.Legacy

Imports the PowerShell my_util in the my_namespace.my_name collection

#AnsibleRequires -PowerShell ansible_collections.my_namespace.my_name.plugins.module_utils.my_util

Imports the PowerShell my_util that exists in the same collection as the current module

#AnsibleRequires -PowerShell ../module_utils.my_util

Imports the PowerShell Ansible.ModuleUtils.Optional provided by Ansible if it exists.

If it does not exist then it will do nothing.

#AnsibleRequires -PowerShell Ansible.ModuleUtils.Optional -Optional

Imports the C# Ansible.Process provided by Ansible itself

#AnsibleRequires -CSharpUtil Ansible.Process

Imports the C# my_util in the my_namespace.my_name collection

#AnsibleRequires -CSharpUtil ansible_collections.my_namespace.my_name.plugins.module_utils.my_util

Imports the C# my_util that exists in the same collection as the current module

#AnsibleRequires -CSharpUtil ../module_utils.my_util

Imports the C# Ansible.Optional provided by Ansible if it exists.

If it does not exist then it will do nothing.

#AnsibleRequires -CSharpUtil Ansible.Optional -Optional

For optional require statements, it is up to the module code to then verify

whether the util has been imported before trying to use it. This can be done by

checking if a function or type provided by the util exists or not.

While both

#Requires

-Module

and

#AnsibleRequires

-PowerShell

can be

used to load a PowerShell module it is recommended to use

#AnsibleRequires

.

This is because

#AnsibleRequires

supports collection module utils, imports
by relative util names, and optional util imports.
C# module utils can reference other C# utils by adding the line
using
Ansible.<module_util>;
to the top of the script with all the other
using statements.
Windows module utilities
?

Like Python modules, PowerShell modules also provide a number of module
utilities that provide helper functions within PowerShell. These module_utils
can be imported by adding the following line to a PowerShell module:

```
#Requires -Module Ansible.ModuleUtils.Legacy
```

This will import the module_util at

```
./lib/ansible/module_utils/powershell/Ansible.ModuleUtils.Legacy.psm1
```

and enable calling all of its functions. As of Ansible 2.8, Windows module
utils can also be written in C# and stored at
lib/ansible/module_utils/csharp

These module_utils can be imported by adding the following line to a PowerShell
module:

```
#AnsibleRequires -CSharpUtil Ansible.Basic
```

This will import the module_util at

```
./lib/ansible/module_utils/csharp/Ansible.Basic.cs
```

and automatically load the types in the executing process. C# module utils can
reference each other and be loaded together by adding the following line to the
using statements at the top of the util:

```
using
```

```
Ansible.Become
```

```
;
```

There are special comments that can be set in a C# file for controlling the
compilation parameters. The following comments can be added to the script;

```
//AssemblyReference
```

```
-Name
```

```
<assembly
```

```
dll>
```

```
[-CLR
```

```
[Core|Framework]]
```

```
: The assembly DLL to reference during compilation, the optional
```

```
-CLR
```

flag can also be used to state whether to reference when running under .NET Core, Framework, or both (if omitted)

```
//NoWarn
```

```
-Name
```

```
<error
```

```
id>
```

```
[-CLR
```

```
[Core|Framework]]
```

```
: A compiler warning ID to ignore when compiling the code, the optional
```

```
-CLR
```

works the same as above. A list of warnings can be found at

Compiler errors

As well as this, the following pre-processor symbols are defined;

```
CORECLR
```

```
: This symbol is present when PowerShell is running through .NET Core
```

WINDOWS

: This symbol is present when PowerShell is running on Windows

UNIX

: This symbol is present when PowerShell is running on Unix

A combination of these flags help to make a module util interoperable on both

.NET Framework and .NET Core, here is an example of them in action:

```
#if CORECLR
using
Newtonsoft.Json
;
#else
using
System.Web.Script.Serialization
;
#endif
//AssemblyReference -Name Newtonsoft.Json.dll -CLR Core
//AssemblyReference -Name System.Web.Extensions.dll -CLR Framework
// Ignore error CS1702 for all .NET types
//NoWarn -Name CS1702
// Ignore error CS1956 only for .NET Framework
//NoWarn -Name CS1956 -CLR Framework
```

The following is a list of module_utils that are packaged with Ansible and a general description of what they do:

ArgvParser: Utility used to convert a list of arguments to an escaped string compliant with the Windows argument parsing rules.

CamelConversion: Utility used to convert camelCase strings/lists/dicts to snake_case.

CommandUtil: Utility used to execute a Windows process and return the stdout/stderr and rc as separate objects.

FileUtil: Utility that expands on the

Get-ChildItem

and

Test-Path

to work with special files like

C:\pagefile.sys

.

Legacy: General definitions and helper utilities for Ansible module.

LinkUtil: Utility to create, remove, and get information about symbolic links, junction points and hard inks.

SID: Utilities used to convert a user or group to a Windows SID and vice versa.

For more details on any specific module utility and their requirements, please see the

Ansible

module utilities source code

.

PowerShell module utilities can be stored outside of the standard Ansible distribution for use with custom modules. Custom module_utils are placed in a folder called

module_utils

located in the root folder of the playbook or role directory.

C# module utilities can also be stored outside of the standard Ansible distribution for use with custom modules. Like PowerShell utils, these are stored in a folder called

module_utils

and the file name must end in the extension

.cs

, start with

Ansible.

and be named after the namespace defined in the util.

The below example is a role structure that contains two PowerShell custom module_utils called Ansible.ModuleUtils.ModuleUtil1

```
,
Ansible.ModuleUtils.ModuleUtil2
, and a C# util containing the namespace
Ansible.CustomUtil
```

```
:
meta/
main.yml
defaults/
main.yml
module_utils/
Ansible.ModuleUtils.ModuleUtil1.psm1
Ansible.ModuleUtils.ModuleUtil2.psm1
Ansible.CustomUtil.cs
tasks/
main.yml
```

Each PowerShell module_util must contain at least one function that has been exported with

Export-ModuleMember

at the end of the file. For example

Export-ModuleMember

-Function

Invoke-CustomUtil

```
,
Get-CustomInfo
```

Exposing shared module options

?

PowerShell module utils can easily expose common module options that a module can use when building its argument spec.

This allows common features to be stored and maintained in one location and have those features used by multiple modules with minimal effort. Any new features or bugfixes added to one of these utils are then automatically used by the various modules that call that util.

An example of this would be to have a module util that handles authentication and communication against an API. This util can be used by multiple modules to expose a common set of module options like the API endpoint, username, password, timeout, cert validation, and so on without having to add those options to each module spec.

The standard convention for a module util that has a shared argument spec would have

A

Get-<namespace.name.util

name>Spec

function that outputs the common spec for a module

It is highly recommended to make this function name be unique to the module to avoid any conflicts with other utils that can be loaded

The format of the output spec is a Hashtable in the same format as the

\$spec

used for normal modules

A function that takes in an

AnsibleModule

object called under the

-Module

parameter which it can use to get the shared options

Because these options can be shared across various module it is highly recommended to keep the module option names and

aliases in the shared spec as specific as they can be. For example do not have a util option called

password

,

rather you should prefix it with a unique name like

acme_password

.

Warning

Failure to have a unique option name or alias can prevent the util being used by module that also use those names or aliases for its own options.

The following is an example module util called

ServiceAuth.psm1

in a collection that implements a common way for modules to authentication with a service.

Invoke-MyServiceResource

```
{
[
CmdletBinding
()]
param
(
[
Parameter
(
Mandatory
=
$true
)]
[
ValidateScript
({
$_
.
GetType
().
FullName
-eq
'Ansible.Basic.AnsibleModule'
})]
$Module
,
[
Parameter
(
Mandatory
=
$true
)]
[String]
$ResourceId
,
[String]
$State
=
'present'
)
```

```
# Process the common module options known to the util
```

```
$params
```

```
=
```

```
@{
```

```
ServerUri
```

```
=
```

```
$Module
```

```
.
```

```
Params
```

```
.
```

```
my_service_url
```

```
}
```

```
if
```

```
(
```

```
$Module
```

```
.
```

```
Params
```

```
.
```

```
my_service_username
```

```
)
```

```
{
```

```
$params
```

```
.
```

```
Credential
```

```
=
```

```
Get-MyServiceCredential
```

```
}
```

```
if
```

```
(
```

```
$State
```

```
-eq
```

```
'absent'
```

```
)
```

```
{
```

```
Remove-MyService
```

```
@params
```

```
-ResourceId
```

```
$ResourceId
```

```
}
```

```
else
```

```
{
```

```
New-MyService
```

```
@params
```

```
-ResourceId
```

```
$ResourceId
```

```
}
```

```
}
```

```
Get-MyNamespaceMyCollectionServiceAuthSpec
```

```
{
```

```
# Output the util spec
```

```
@{
```

```
options
```

```
=
```

```
@{
```

```

my_service_url
=
@{
type
=
'str'
;
required
=
$true
}
my_service_username
=
@{
type
=
'str'
}
my_service_password
=
@{
type
=
'str'
;
no_log
=
$true
}
}
required_together
=
@(
,@(
'my_service_username'
,
'my_service_password'
)
)
}
$exportMembers
=
@{
Function
=
'Get-MyNamespaceMyCollectionServiceAuthSpec'
,
'Invoke-MyServiceResource'
}
Export-ModuleMember
@exportMembers

```

For a module to take advantage of this common argument spec it can be set out like
 #!powershell

```

# Include the module util ServiceAuth.psm1 from the my_namespace.my_collection collection
#AnsibleRequires -PowerShell ansible_collections.my_namespace.my_collection.plugins.module_utils.ServiceAuth
# Create the module spec like normal
$spec
=
@{
options
=
@{
resource_id
=
@{
type
=
'str'
;
required
=
$true
}
state
=
@{
type
=
'str'
;
choices
=
'absent'
,
'present'
}
}
}
# Create the module from the module spec but also include the util spec to merge into our own.
$module
=
[Ansible.Basic.AnsibleModule]
::
Create
(
$args
,
$spec
,
@(
Get-MyNamespaceMyCollectionServiceAuthSpec
))
# Call the ServiceAuth module util and pass in the module object so it can access the module options.
Invoke-MyServiceResource
-Module
$module
-ResourceId

```

\$module

.

Params

.

resource_id

-State

\$module

.

params

.

state

\$module

.

ExitJson

()

Note

Options defined in the module spec will always have precedence over a util spec. Any list values under the same key in a util spec will be appended to the module spec for that same key. Dictionary values will add any keys that are missing from the module spec and merge any values that are lists or dictionaries. This is similar to how the doc fragment plugins work when extending module documentation.

To document these shared util options for a module, create a doc fragment plugin that documents the options implemented

by the module util and extend the module docs for every module that implements the util to include that fragment in its docs.

Windows playbook module testing

?

You can test a module with an Ansible playbook. For example:

Create a playbook in any directory

touch

testmodule.yml

.

Create an inventory file in the same directory

touch

hosts

.

Populate the inventory file with the variables required to connect to a Windows host(s).

Add the following to the new playbook file:

-

name

:

test out windows module

hosts

:

windows

tasks

:

-

name

:

test out module

win_module

:

name

```

:
test name
Run the playbook
ansible-playbook
-i
hosts
testmodule.yml
This can be useful for seeing how Ansible runs with
the new module end to end. Other possible ways to test the module are
shown below.
Windows debugging
?
Debugging a module currently can only be done on a Windows host. This can be
useful when developing a new module or implementing bug fixes. These
are some steps that need to be followed to set this up:
Copy the module script to the Windows server
Copy the folders
./lib/ansible/module_utils/powershell
and
./lib/ansible/module_utils/csharp
to the same directory as the script above
Add an extra
#
to the start of any
#Requires
-Module
lines in the module code, this is only required for any lines starting with
#Requires
-Module
Add the following to the start of the module script that was copied to the server:
# Set $ErrorActionPreference to what's set during Ansible execution
$ErrorActionPreference
=
"Stop"
# Set the first argument as the path to a JSON file that contains the module args
$args
=
@(
"
$(
$pwd
.
Path
)
\args.json"
)
# Or instead of an args file, set $complex_args to the pre-processed module args
$complex_args
=
@{
  _ansible_check_mode
=
>false
  _ansible_diff

```

```

=
$false
path
=
"C:\temp"
state
=
"present"
}
# Import any C# utils referenced with '#AnsibleRequires -CSharpUtil' or 'using Ansible.;
# The $_csharp_utils entries should be the context of the C# util files and not the path
Import-Module
-Name
"
$(
$pwd
.
Path
)
\powershell\Ansible.ModuleUtils.AddType.psm1"
$_csharp_utils
=
@(
[System.IO.File]
::
ReadAllText
(
"
$(
$pwd
.
Path
)
\csharp\Ansible.Basic.cs"
)
)
Add-CSharpType
-References
$_csharp_utils
-IncludeDebugInfo
# Import any PowerShell modules referenced with '#Requires -Module`
Import-Module
-Name
"
$(
$pwd
.
Path
)
\powershell\Ansible.ModuleUtils.Legacy.psm1"
# End of the setup code and start of the module code
#!powershell
You can add more args to
$complex_args

```


as required by the module or define the module options through a JSON file with the structure:

```
{
  "ANSIBLE_MODULE_ARGS"
  :
  {
    "_ansible_check_mode"
    :
    false

    ,
    "_ansible_diff"
    :
    false

    ,
    "path"
    :
    "C:\\temp"

    ,
    "state"
    :
    "present"
  }
}
```

There are multiple IDEs that can be used to debug a Powershell script, two of the most popular ones are

Powershell ISE

Visual Studio Code

To be able to view the arguments as passed by Ansible to the module follow these steps.

Prefix the Ansible command with

ANSIBLE_KEEP_REMOTE_FILES=1

to specify that Ansible should keep the exec files on the server.

Log onto the Windows server using the same user account that Ansible used to execute the module.

Navigate to

%TEMP%\..

. It should contain a folder starting with
ansible-tmp-

.

Inside this folder, open the PowerShell script for the module.

In this script is a raw JSON script under

\$json_raw

which contains the module arguments under
module_args

. These args can be assigned manually to the

\$complex_args

variable that is defined on your debug script or put in the
args.json

file.

Windows unit testing

?

Currently there is no mechanism to run unit tests for Powershell modules under Ansible CI.

Windows integration testing

?

Integration tests for Ansible modules are typically written as Ansible roles. These test

roles are located in
./test/integration/targets
. You must first set up your testing
environment, and configure a test inventory for Ansible to connect to.
In this example we will set up a test inventory to connect to two hosts and run the integration
tests for win_stat:
Run the command
source
./hacking/env-setup
to prepare environment.
Create a copy of
./test/integration/inventory.winrm.template
and name it
inventory.winrm
.
Fill in entries under
[windows]
and set the required variables that are needed to connect to the host.
Install the required Python modules
to support WinRM and a configured authentication method.
To execute the integration tests, run
ansible-test
windows-integration
win_stat
; you can replace
win_stat
with the role you want to test.
This will execute all the tests currently defined for that role. You can set
the verbosity level using the
-v
argument just as you would with
ansible-playbook.
When developing tests for a new module, it is recommended to test a scenario once in
check mode and twice not in check mode. This ensures that check mode
does not make any changes but reports a change, as well as that the second run is
idempotent and does not report changes. For example:

```
-  
name  
:  
  remove a file (check mode)  
win_file  
:  
  path  
:  
  C:\temp  
state  
:  
  absent  
register  
:  
  remove_file_check  
check_mode  
:  
  true
```

```
-
name
:
get result of remove a file (check mode)
win_command
:
powershell.exe "if (Test-Path -Path 'C:\temp') { 'true' } else { 'false' }"
register
:
remove_file_actual_check
-
name
:
assert remove a file (check mode)
assert
:
that
:
-
remove_file_check is changed
-
remove_file_actual_check.stdout == 'true\r\n'
-
name
:
remove a file
win_file
:
path
:
C:\temp
state
:
absent
register
:
remove_file
-
name
:
get result of remove a file
win_command
:
powershell.exe "if (Test-Path -Path 'C:\temp') { 'true' } else { 'false' }"
register
:
remove_file_actual
-
name
:
assert remove a file
assert
:
that
```

```
:
```

-

remove_file is changed

-

```
remove_file_actual.stdout == 'false\r\n'
```

-

name

```
:
```

remove a file (idempotent)

win_file

```
:
```

path

```
:
```

C:\temp

state

```
:
```

absent

register

```
:
```

remove_file_again

-

name

```
:
```

assert remove a file (idempotent)

assert

```
:
```

that

```
:
```

-

not remove_file_again is changed

Windows communication and development support

?

Join the

Ansible Forum

and use the

windows

tag for discussions about Ansible development for Windows.

Previous

Next

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_plugins.html

Developing plugins ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing plugins
Edit on GitHub
Developing plugins

?

Writing plugins in Python
Raising errors
String encoding
Plugin configuration & documentation standards
General precedence rules
Accessing configuration settings
Developing particular plugin types
Action plugins
Cache plugins
Callback plugins
Connection plugins
Filter plugins
Inventory plugins
Lookup plugins
Test plugins
Vars plugins

Plugins augment Ansible's core functionality with logic and features that are accessible to all modules. Ansible collections include a number of handy plugins, and you can easily write your own. All plugins must:

be written in Python
raise errors
return strings in unicode
conform to Ansible's configuration and documentation standards

Once you've reviewed these general guidelines, you can skip to the particular type of plugin you want to develop.

Writing plugins in Python

?

You must write your plugin in Python so it can be loaded by the
PluginLoader

and returned as a Python object that any module can use. Since your plugin will execute on the control node, you must write it in a
compatible version of Python

.

Raising errors

?

You should return errors encountered during plugin execution by raising

`AnsibleError()`

or a similar class with a message describing the error. When wrapping other exceptions into error messages, you should always use the

`to_native`

Ansible function to ensure proper string compatibility across Python versions:

from

`ansible.module_utils.common.text.converters`

import

`to_native`

try

:

`cause_an_exception`

()

except

`Exception`

as

`e`

:

raise

`AnsibleError`

(

'Something happened, this was original exception:

`%s`

,

`%`

`to_native`

(

`e`

))

Since Ansible evaluates variables only when they are needed, filter and test plugins should propagate the exceptions

`jinja2.exceptions.UndefinedError`

and

`AnsibleUndefinedVariable`

to ensure undefined variables are only fatal when necessary.

Check the different

`AnsibleError` objects

and see which one applies best to your situation.

Check the section on the specific plugin type you're developing for type-specific error handling details.

String encoding

?

You must convert any strings returned by your plugin into Python's unicode type. Converting to unicode ensures that these strings can run through Jinja2. To convert strings:

from

`ansible.module_utils.common.text.converters`

import

`to_text`

`result_string`

=

`to_text`

(

`result_string`

)

Plugin configuration & documentation standards

?

To define configurable options for your plugin, describe them in the DOCUMENTATION

section of the python file. Callback and connection plugins have declared configuration requirements this way since Ansible version 2.4; most plugin types now do the same. This approach ensures that the documentation of your plugin's options will always be correct and up-to-date. To add a configurable option to your plugin, define it in this format:

options

```
:
option_name
:
description
:
describe this config option
default
:
default value for this config option
env
:
-
name
:
MYCOLLECTION_NAME_ENV_VAR_NAME
ini
:
-
section
:
mycollection_section_of_ansible.cfg_where_this_config_option_is_defined
key
:
key_used_in_ansible.cfg
vars
:
-
name
:
mycollection_name_of_ansible_var
-
name
:
mycollection_name_of_second_var
version_added
:
X.x
required
:
True/False
type
:
boolean/float/integer/list/none/path/pathlist/pathspec/string/tmpfilepath
version_added
:
X.x
```

The supported configuration fields are:

env

List of environment variables that can be used to set this option.

Each entry includes a

name

field specifying the environment variable name.

The name should be in uppercase and should be prefixed with the collection name.

Multiple environment variables can be listed for the same option.

The last set environment variable in the list takes precedence if multiple are set.

This is commonly used for plugins (especially inventory plugins) to allow configuration through environment variables.

Examples:

VMWARE_PORT

,

GRAFANA_PASSWORD

ini

List of configuration file settings that can be used to set this option.

Each entry includes a

section

field for the configuration file section and a

key

field for the configuration key. Both should be in lowercase and should be prefixed with the collection name.

Multiple configuration settings can be listed for the same option.

The last set configuration setting in the list takes precedence if multiple are set.

This allows plugins to be configured with ansible.cfg.

Example:

grafana_password

vars

List of Ansible variables that can be used to set this option.

Each entry includes a

name

field specifying the variable name.

The name should be in lowercase and should be prefixed with the collection name.

Multiple variables can be listed for the same option.

The last set variable in the list takes precedence if multiple are set.

Variables follow Ansible's variable precedence rules.

This allows plugins to be configured with Ansible variables.

Example:

ansible_vmware_port

General precedence rules

?

The precedence rules for configuration sources are listed below, starting with the highest precedence values:

Keywords

CLI settings

Environment variables (

env

)

Values defined in

ansible.cfg

Default value for the option, if present.

Accessing configuration settings

?

To access the configuration settings in your plugin, use

self.get_option(<option_name>)

.

Some plugin types handle this differently:

Become, callback, connection and shell plugins are guaranteed to have the engine call

set_options()

.

Lookup plugins always require you to handle it in the
run()
method.

Inventory plugins are done automatically if you use the
base

_read_config_file()

method. If not, you must use

self.get_option(<option_name>)

.

Cache plugins do it on load.

Cliconf, httpapi and netconf plugins indirectly piggy back on connection plugins.

Vars plugin settings are populated when first accessed (using the

self.get_option()

or

self.get_options()

method.

If you need to populate settings explicitly, use a

self.set_options()

call.

Configuration sources follow the precedence rules for values in Ansible. When there are multiple values from the same
category, the value defined last takes precedence. For example, in the above configuration block, if both

name_of_ansible_var

and

name_of_second_var

are defined, the value of the

option_name

option will be the value of

name_of_second_var

. Refer to

Controlling how Ansible behaves: precedence rules
for further information.

Plugins that support embedded documentation (see
ansible-doc

for the list) should include well-formed doc strings. If you inherit from a plugin, you must document the options it takes,
either through a documentation fragment or as a copy. See

Module format and documentation

for more information on correct documentation. Thorough documentation is a good idea even if you're developing a
plugin for local use.

In ansible-core 2.14 we added support for documenting filter and test plugins. You have two options for providing
documentation:

Define a Python file that includes inline documentation for each plugin.

Define a Python file for multiple plugins and create adjacent documentation files in YAML format.

Developing particular plugin types

?

Action plugins

?

Action plugins let you integrate local processing and local data with module functionality.

To create an action plugin, create a new class with the Base(ActionBase) class as the parent:

from

ansible.plugins.action

import

ActionBase

```
class
ActionModule
(
    ActionBase
):
    pass
```

From there, execute the module using the `_execute_module` method to call the original module.

After successful execution of the module, you can modify the module return data.

```
module_return
=
self
.
_execute_module
(
    module_name
    =
    '<NAME_OF_MODULE>'
    ,
    module_args
    =
    module_args
    ,
    task_vars
    =
    task_vars
    ,
    tmp
    =
    tmp
)

```

For example, if you wanted to check the time difference between your Ansible control node and your target machine(s), you could write an action plugin to check the local time and compare it to the return data from Ansible's

```
setup
module:
#!/usr/bin/python
# Make coding more python3-ish, this is required for contributions to Ansible
from
__future__
import
(
    absolute_import
    ,
    division
    ,
    print_function
)
__metaclass__
=
type
from
ansible.plugins.action
import
```

```
ActionBase
from
datetime
import
datetime
class
ActionModule
(
    ActionBase
):
    def
    run
    (
        self
        ,
        tmp
        =
        None
        ,
        task_vars
        =
        None
    ):
        super
        (
            ActionModule
            ,
            self
        )
        .
        run
        (
            tmp
            ,
            task_vars
        )
        module_args
        =
        self
        .
        _task
        .
        args
        .
        copy
        ()
        module_return
        =
        self
        .
        _execute_module
        (
            module_name
            =
```

```
'setup'
,
module_args
=
module_args
,
task_vars
=
task_vars
,
tmp
=
tmp
)
ret
=
dict
()
remote_date
=
None
if
not
module_return
.
get
(
'failed'
):
for
key
,
value
in
module_return
[
'ansible_facts'
]
.
items
():
if
key
==
'ansible_date_time'
:
remote_date
=
value
[
'iso8601'
]
if
remote_date
```

```
:
remote_date_obj
=
datetime
.
strptime
(
remote_date
,
'%Y-%m-
%d
T%H:%M:%SZ'
)
time_delta
=
datetime
.
utcnow
()
-
remote_date_obj
ret
[
'delta_seconds'
]
=
time_delta
.
seconds
ret
[
'delta_days'
]
=
time_delta
.
days
ret
[
'delta_microseconds'
]
=
time_delta
.
microseconds
return
dict
(
ansible_facts
=
dict
(
ret
))
```

This code checks the time on the control node, captures the date and time for the remote machine using the setup

module, and calculates the difference between the captured time and the local time, returning the time delta in days, seconds and microseconds.

For practical examples of action plugins, see the source code for the action plugins included with Ansible Core
Cache plugins

?

Cache plugins store gathered facts and data retrieved by inventory plugins.

Import cache plugins using the cache_loader so you can use

```
self.set_options()
```

and

```
self.get_option(<option_name>)
```

. If you import a cache plugin directly in the code base, you can only access options by the `ansible.constants`

, and you break the cache plugin's ability to be used by an inventory plugin.

from

```
ansible.plugins.loader
```

```
import
```

```
cache_loader
```

```
[
```

```
...
```

```
]
```

```
plugin
```

```
=
```

```
cache_loader
```

```
.
```

```
get
```

```
(
```

```
'custom_cache'
```

```
,
```

```
**
```

```
cache_kwargs
```

```
)
```

There are two base classes for cache plugins,

`BaseCacheModule`

for database-backed caches, and

`BaseCacheFileModule`

for file-backed caches.

To create a cache plugin, start by creating a new

`CacheModule`

class with the appropriate base class. If you're creating a plugin using an

`__init__`

method you should initialize the base class with any provided args and kwargs to be compatible with inventory plugin cache options. The base class calls

```
self.set_options(direct=kwargs)
```

. After the base class

`__init__`

method is called

```
self.get_option(<option_name>)
```

should be used to access cache options.

New cache plugins should take the options

`_uri`

```
,
_prefix
, and
_timeout
to be consistent with existing cache plugins.
from
ansible.plugins.cache
import
BaseCacheModule
class
CacheModule
(
BaseCacheModule
):
def
__init__(
self
,
*
args
,
**
kwargs
):
super(
CacheModule
,
self
)
.
__init__(
self
,
*
args
,
**
kwargs
)
self
.
_connection
=
self
.
_get_option(
'_uri'
)
self
.
_prefix
=
```



```
self
.  
get_option  
(  
    '_prefix'  
)
```

```
self
.  
_timeout  
=  
self
.  
get_option  
(  
    '_timeout'  
)
```

If you use the
BaseCacheModule
, you must implement the methods
get

,

contains

,

keys

,

set

,

delete

,

flush

, and
copy

. The
contains

method should return a boolean that indicates if the key exists and has not expired. Unlike file-based caches, the
get

method does not raise a KeyError if the cache has expired.

If you use the
BaseFileCacheModule
, you must implement

_load
and
_dump

methods that will be called from the base class methods

get
and
set

.

If your cache plugin stores JSON, use
AnsibleJSONEncoder
in the
_dump

or
set

method and
AnsibleJSONDecoder
in the
_load
or
get
method.

For example cache plugins, see the source code for the
cache plugins included with Ansible Core

.

Callback plugins

?

Callback plugins add new behaviors to Ansible when responding to events. By default, callback plugins control most of the output you see when running the command line programs.

To create a callback plugin, create a new class with the Base(Callbacks) class as the parent:

from
ansible.plugins.callback
import
CallbackBase
class
CallbackModule
(
CallbackBase
):
pass

From there, override the specific methods from the CallbackBase that you want to provide a callback for.

For plugins intended for use with Ansible version 2.0 and later, you should only override methods that start with v2

.

For a complete list of methods that you can override, please see

__init__.py
in the
lib/ansible/plugins/callback
directory.

The following is a modified example of how Ansible's timer plugin is implemented,
but with an extra option so you can see how configuration works in Ansible version 2.4 and later:

Make coding more python3-ish, this is required for contributions to Ansible

from
__future__
import
(
absolute_import
,
division
,
print_function
)
__metaclass__
=
type
not only visible to ansible-doc, it also 'declares' the options the plugin requires and how to configure them.
DOCUMENTATION
=
'''

```

name: timer
callback_type: aggregate
requirements:
- enable in configuration
short_description: Adds time to play stats
version_added: "2.0" # for collections, use the collection version, not the Ansible version
description:
- This callback just adds total play duration to the play stats.
options:
format_string:
description: format of the string shown to user at play end
ini:
- section: callback_timer
key: format_string
env:
- name: ANSIBLE_CALLBACK_TIMER_FORMAT
default: "Playbook run took
%s
days,
%s
hours,
%s
minutes,
%s
seconds"
"""

from
datetime
import
datetime
from
ansible.plugins.callback
import
CallbackBase
class
CallbackModule
(
CallbackBase
):
"""

This callback module tells you how long your plays ran for.
"""

CALLBACK_VERSION
=
2.0
CALLBACK_TYPE
=
'aggregate'
CALLBACK_NAME
=
'namespace.collection_name.timer'
# only needed if you ship it and don't want to enable by default
CALLBACK_NEEDS_ENABLED
=

```

```

True
def
__init__
(
self
):
# make sure the expected objects are present, calling the base's __init__
super
(
CallbackModule
,
self
)
.
__init__
()
# start the timer when the plugin is loaded, the first play should start a few milliseconds after.
self
.
start_time
=
datetime
.
now
()
def
_days_hours_minutes_seconds
(
self
,
runtime
):
""" internal helper method for this callback """
minutes
=
(
runtime
.
seconds
//
60
)
%
60
r_seconds
=
runtime
.
seconds
-
(
minutes
*
60

```

```

)
return
runtime
.
days
,
runtime
.
seconds
//
3600
,
minutes
,
r_seconds
# this is only event we care about for display, when the play shows its summary stats; the rest are ignored by the base
class
def
v2_playbook_on_stats
(
self
,
stats
):
end_time
=
datetime
.
now
()
runtime
=
end_time
-
self
.
start_time
# Shows the usage of a config option declared in the DOCUMENTATION variable. Ansible will have set it when it loads
the plugin.
# Also note the use of the display object to print to screen. This is available to all callbacks, and you should use this over
printing yourself
self
.
_display
.
display
(
self
.
_plugin_options
[
'format_string'
]
%

```

```
(
self
.
_days_hours_minutes_seconds
(
runtime
)))
```

Note that the

`CALLBACK_VERSION`

and

`CALLBACK_NAME`

definitions are required for properly functioning plugins for Ansible version 2.0 and later.

`CALLBACK_TYPE`

is mostly needed to distinguish `stdout` plugins from the rest, since you can only load one plugin that writes to `stdout`.

For example callback plugins, see the source code for the

callback plugins included with Ansible Core

New in ansible-core 2.11, callback plugins are notified (by the

`v2_playbook_on_task_start`

) of

meta

tasks. By default, only explicit

meta

tasks that users list in their plays are sent to callbacks.

There are also some tasks which are generated internally and implicitly at various points in execution. Callback plugins

can opt-in to receiving these implicit tasks as well, by setting

`self.wants_implicit_tasks`

=

`True`

. Any

Task

object received by a callback hook will have an

`.implicit`

attribute, which can be consulted to determine whether the

Task

originated from within Ansible, or explicitly by the user.

Connection plugins

?

Connection plugins allow Ansible to connect to target hosts so it can execute tasks on them. Ansible ships with many

connection plugins, but only one can be used per host at a time. The most commonly used connection plugins are native

`ssh`

,

`paramiko`

, and

`local`

. All of these can be used with ad-hoc tasks and in playbooks.

To create a new connection plugin (for example, to support SNMP, Message bus, or other transports), copy the format

of one of the existing connection plugins and drop it into

connection

directory on your

local plugin path

.

Connection plugins can support common options (such as the

`--timeout`

flag) by defining an entry in the documentation for the attribute name (in this case

timeout

). If the common option has a non-null default, the plugin should define the same default since a different default would be ignored.

For example connection plugins, see the source code for the connection plugins included with Ansible Core

.

Filter plugins

?

Filter plugins manipulate data. They are a feature of Jinja2 and are also available in Jinja2 templates used by the template

module. As with all plugins, they can be easily extended, but instead of having a file for each one you can have several per file. Most of the filter plugins shipped with Ansible reside in a `core.py`

.

Filter plugins do not use the standard configuration system described above, but since `ansible-core 2.14` can use it as plain documentation.

Since Ansible evaluates variables only when they are needed, filter plugins should propagate the exceptions `jinja2.exceptions.UndefinedError`

and

`AnsibleUndefinedVariable`

to ensure undefined variables are only fatal when necessary.

try

:

`cause_an_exception`

 (

`with_undefined_variable`

)

except

`jinja2`

.

 exceptions

.

`UndefinedError`

as

`e`

:

 raise

`AnsibleUndefinedVariable`

 (

 "Something happened, this was the original exception:

 %s

 "

 %

`to_native`

 (

`e`

))

except

`Exception`

as

`e`

:

 raise

`AnsibleFilterError`

```
(
"Something happened, this was the original exception:
%s
"
%
to_native
(
e
))
```

For example filter plugins, see the source code for the filter plugins included with Ansible Core

.

Inventory plugins

?

Inventory plugins parse inventory sources and form an in-memory representation of the inventory. Inventory plugins were added in Ansible version 2.4.

You can see the details for inventory plugins in the [Developing dynamic inventory](#) page.

Lookup plugins

?

Lookup plugins pull in data from external data stores. Lookup plugins can be used within playbooks both for looping ?

playbook language constructs like

`with_fileglob`

and

`with_items`

are implemented through lookup plugins ? and to return values into a variable or parameter.

Lookup plugins are expected to return lists, even if just a single element.

Ansible includes many

filters which can be used to manipulate the data returned by a lookup plugin. Sometimes it makes sense to do the filtering inside the lookup plugin, other times it is better to return results that can be filtered in the playbook. Keep in mind how the data will be referenced when determining the appropriate level of filtering to be done inside the lookup plugin.

Here's a simple lookup plugin implementation ? this lookup returns the contents of a text file as a variable:

python 3 headers, required if submitting to Ansible

```
from
__future__
import
(
absolute_import
,
division
,
print_function
)
__metaclass__
=
type
DOCUMENTATION
=
r
"""
name: file
author: Daniel Hokka Zakrisson (@dhozac) <
```


>

version_added: "0.9" # for collections, use the collection version, not the Ansible version

short_description: read file contents

description:

- This lookup returns the contents from a file on the Ansible control node's file system.

options:

_terms:

description: path(s) of files to read

required: True

option1:

description:

- Sample option that could modify plugin behavior.

- This one can be set directly ``option1='x'`` or in ansible.cfg, but can also use vars or environment.

type: string

ini:

- section: file_lookup

key: option1

notes:

- if read in variable context, the file can be interpreted as YAML if the content is valid to the parser.

- this lookup does not understand globbing --- use the fileglob lookup instead.

"""

from

ansible.errors

import

AnsibleError

,

AnsibleParserError

from

ansible.plugins.lookup

import

LookupBase

from

ansible.utils.display

import

Display

display

=

Display

()

class

LookupModule

(

LookupBase

):

def

run

(

self

,

terms

,

variables

=

None

```
,
**

kwargs
):
# First of all populate options,
# this will already take into account env vars and ini config
self
.
set_options
(
var_options
=
variables
,
direct
=
kwargs
)
# lookups in general are expected to both take a list as input and output a list
# this is done so they work with the looping construct 'with_'.
ret
=
[]
for
term
in
terms
:
display
.
debug
(
"File lookup term:
%s
"
%
term
)
# Find the file in the expected search path, using a class method
# that implements the 'expected' search path for Ansible plugins.
lookupfile
=
self
.
find_file_in_search_path
(
variables
,
'files'
,
term
)
)
# Don't use print or your own logging, the display class
```

```

# takes care of it in a unified way.
display

.
vvvv
(
u
"File lookup using
%s
as file"
%
lookupfile
)
try
:
if
lookupfile
:
contents
,
show_data
=
self
.
_loader
.
_get_file_contents
(
lookupfile
)
ret
.
append
(
contents
.
rstrip
())
else
:
# Always use ansible error classes to throw 'final' exceptions,
# so the Ansible engine will know how to deal with them.
# The Parser error indicates invalid options passed
raise
AnsibleParserError
()
except
AnsibleParserError
:
raise
AnsibleError
(
"could not locate file in lookup:
%s
"

```

```

%
term
)
# consume an option: if this did something useful, you can retrieve the option value here
if
self
.
get_option
(
'option1'
)
==
'do something'
:
pass
return
ret
The following is an example of how this lookup is called:
---
-
hosts
:
all
vars
:
contents
:
"{{
lookup('namespace.collection_name.file',
/etc/foo.txt')
}}"
contents_with_option
:
"{{
lookup('namespace.collection_name.file',
/etc/foo.txt',
option1='donothing')
}}"
tasks
:
-
debug
:
msg
:
the value of foo.txt is {{ contents }} as seen today {{ lookup('pipe', 'date +"%Y-%m-%d") }}
For example lookup plugins, see the source code for the
lookup plugins included with Ansible Core
.
For more usage examples of lookup plugins, see
Using Lookups
.
Test plugins
?
```

Test plugins verify data. They are a feature of Jinja2 and are also available in Jinja2 templates used by the template module. As with all plugins, they can be easily extended, but instead of having a file for each one you can have several per file. Most of the test plugins shipped with Ansible reside in a `core.py`

. These are specially useful in conjunction with some filter plugins like `map` and `select`; they are also available for conditional directives like `when`:

. Test plugins do not use the standard configuration system described above. Since `ansible-core 2.14` test plugins can use plain documentation.

Since Ansible evaluates variables only when they are needed, test plugins should propagate the exceptions `jinja2.exceptions.UndefinedError`

and `AnsibleUndefinedVariable` to ensure undefined variables are only fatal when necessary.

```
try
:
cause_an_exception
(
with_undefined_variable
)
except
jinja2
.
exceptions
.
UndefinedError
as
e
:
raise
AnsibleUndefinedVariable
(
"Something happened, this was the original exception:
%s
"
%
to_native
(
e
))
except
Exception
as
e
:
raise
AnsibleFilterError
(
"Something happened, this was the original exception:
```

```
%s
"
%
to_native
(
e
))
```

For example test plugins, see the source code for the test plugins included with Ansible Core

```
.
```

Vars plugins
?

Vars plugins inject additional variable data into Ansible runs that did not come from an inventory source, playbook, or command line. Playbook constructs like `?host_vars?` and `?group_vars?` work using vars plugins.

Vars plugins were partially implemented in Ansible 2.0 and rewritten to be fully implemented starting with Ansible 2.4.

Vars plugins are supported by collections starting with Ansible 2.10.

Older plugins used a `run` method as their main body/work:

```
def
run
(
self
,
name
,
vault_password
=
None
):
pass
# your code goes here
```

Ansible 2.0 did not pass passwords to older plugins, so vaults were unavailable.

Most of the work now happens in the

`get_vars` method which is called from the `VariableManager` when needed.

```
def
get_vars
(
self
,
loader
,
path
,
entities
):
pass
# your code goes here
```

The parameters are:

loader: Ansible's `DataLoader`. The `DataLoader` can read files, auto-load JSON/YAML and decrypt vaulted data, and cache read files.

path: this is `?directory data?` for every inventory source and the current play's playbook directory, so they can search for data in reference to them.

get_vars

will be called at least once per available path.

entities: these are host or group names that are pertinent to the variables needed. The plugin will get called once for hosts and again for groups.

This

get_vars

method just needs to return a dictionary structure with the variables.

Since Ansible version 2.4, vars plugins only execute as needed when preparing to execute a task. This avoids the costly ?always execute? behavior that occurred during inventory construction in older versions of Ansible. Since Ansible version 2.10, vars plugin execution can be toggled by the user to run when preparing to execute a task or after importing an inventory source.

The user must explicitly enable vars plugins that reside in a collection. See

Enabling vars plugins

for details.

Legacy vars plugins are always loaded and run by default. You can prevent them from automatically running by setting

REQUIRES_ENABLED

to True.

class

VarsModule

(

BaseVarsPlugin

):

REQUIRES_ENABLED

=

True

Include the

vars_plugin_staging

documentation fragment to allow users to determine when vars plugins run.

DOCUMENTATION

=

"""

name: custom_hostvars

version_added: "2.10" # for collections, use the collection version, not the Ansible version

short_description: Load custom host vars

description: Load custom host vars

options:

stage:

ini:

- key: stage

section: vars_custom_hostvars

env:

- name: ANSIBLE_VARS_PLUGIN_STAGE

extends_documentation_fragment:

- vars_plugin_staging

"""

At times a value provided by a vars plugin will contain unsafe values. The utility function

wrap_var

provided by

ansible.utils.unsafe_proxy

should be used to ensure that Ansible handles the variable and value correctly. The use cases for unsafe data is covered in

Unsafe or raw strings

.

from

```
ansible.plugins.vars
import
BaseVarsPlugin
from
ansible.utils.unsafe_proxy
import
wrap_var
class
VarsPlugin
(
BaseVarsPlugin
):
def
get_vars
(
self
,
loader
,
path
,
entities
):
return
dict
(
something_unsafe
=
wrap_var
(
"{{ SOMETHING_UNSAFE }}"
)
)
)
```

For example vars plugins, see the source code for the vars plugins included with Ansible Core

.

See also

[Collection Index](#)

[Browse existing collections, modules, and plugins](#)

[Python API](#)

[Learn about the Python API for task execution](#)

[Developing dynamic inventory](#)

[Learn about how to develop dynamic inventory sources](#)

[Developing modules](#)

[Learn about how to write Ansible modules](#)

[Communication](#)

[Got questions? Need help? Want to share your ideas? Visit the Ansible communication guide](#)

[Adjacent YAML documentation files](#)

[Alternate YAML files as documentation](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_program_flow_modules.html

Ansible module architecture ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Developing
ansible-core
Ansible module architecture
Edit on GitHub
Ansible module architecture
?

If you are working on the
ansible-core

code, writing an Ansible module, or developing an action plugin, you may need to understand how Ansible's program flow executes. If you are just using Ansible Modules in playbooks, you can skip this section.

Types of modules
Action plugins
New-style modules
Python
PowerShell
JSONARGS modules
Non-native want JSON modules
Binary modules
Old-style modules
How modules are executed
Executor/task_executor
The
normal
action plugin
Executor/module_common.py
Assembler frameworks
Module Replacer framework
Ansiballz framework
Passing args
Internal arguments
_ansible_no_log
_ansible_debug
_ansible_diff
_ansible_verbosity
_ansible_selinux_special_fs
_ansible_syslog_facility
_ansible_version
_ansible_module_name

`_ansible_keep_remote_files`

`_ansible_socket`

`_ansible_shell_executable`

`_ansible_tmpdir`

`_ansible_remote_tmp`

Module return values & Unsafe strings

Special considerations

Pipelining

Why pass args over stdin?

AnsibleModule

Argument spec

Dependencies between module options

Declaring check mode support

Adding file options

Types of modules

?

Ansible supports several different types of modules in its code base. Some of these are for backwards compatibility and others are to enable flexibility.

Action plugins

?

Action plugins look like modules to anyone writing a playbook. Usage documentation for most action plugins lives inside a module of the same name. Some action plugins do all the work, with the module providing only documentation. Some action plugins execute modules. The

normal

action plugin executes modules that don't have special action plugins. Action plugins always execute on the control node.

Some action plugins do all their work on the control node. For

example, the

debug

action plugin (which prints text for

the user to see) and the

assert

action plugin (which

tests whether values in a playbook satisfy certain criteria) execute entirely on the control node.

Most action plugins set up some values on the control node, then invoke an

actual module on the managed node that does something with these values. For example, the

template

action plugin takes values from

the user to construct a file in a temporary location on the control node using

variables from the playbook environment. It then transfers the temporary file

to a temporary file on the remote system. After that, it invokes the

copy module

which operates on the remote system to move the file

into its final location, sets file permissions, and so on.

New-style modules

?

All of the modules that ship with Ansible fall into this category. While you can write modules in any language, all official modules (shipped with Ansible) use either Python or PowerShell.

New-style modules have the arguments to the module embedded inside of them in

some manner. Old-style modules must copy a separate file over to the

managed node, which is less efficient as it requires two over-the-wire

connections instead of only one.

Python

?

New-style Python modules use the Ansiballz framework for constructing modules. These modules use imports from `ansible.module_utils` to pull in boilerplate module code, such as argument parsing, formatting of return values as JSON

, and various file operations.

Note

In Ansible, up to version 2.0.x, the official Python modules used the Module Replacer framework. For module authors, Ansiballz framework is largely a superset of Module Replacer framework functionality, so you usually do not need to understand the differences between them.

PowerShell

?

New-style PowerShell modules use the Module Replacer framework for constructing modules. These modules get a library of PowerShell code embedded

in them before being sent to the managed node.

JSONARGS modules

?

These modules are scripts that include the string

```
<<INCLUDE_ANSIBLE_MODULE_JSON_ARGS>>
```

in their body.

This string is replaced with the JSON-formatted argument string. These modules typically set a variable to that value like this:

```
json_arguments
```

```
=
```

```
"""<<INCLUDE_ANSIBLE_MODULE_JSON_ARGS>>"""
```

Which is expanded as:

```
json_arguments
```

```
=
```

```
"""{"param1": "test's quotes", "param2": "  
\"
```

To be or not to be

```
\"
```

```
- Hamlet"}"""
```

Note

Ansible outputs a

JSON

string with bare quotes. Double quotes are

used to quote string values, double quotes inside of string values are

backslash escaped, and single quotes may appear unescaped inside of

a string value. To use JSONARGS, your scripting language must have a way

to handle this type of string. The example uses Python's triple quoted

strings to do this. Other scripting languages may have a similar quote

character that won't be confused by any quotes in the JSON or it may allow you to define your own start-of-quote and end-of-quote characters.

If the language doesn't give you any of these then you'll need to write

a

non-native JSON module

or

Old-style module

instead.

These modules typically parse the contents of

json_arguments

using a JSON

library and then use them as native variables throughout the code.

Non-native want JSON modules

?

If a module has the string

WANT_JSON

in it anywhere, Ansible treats

it as a non-native module that accepts a file name as its only command-line

parameter. The file name is for a temporary file containing a

JSON

string containing the module's parameters. The module needs to open the file,

read and parse the parameters, operate on the data, and print its return data

as a JSON encoded dictionary to stdout before exiting.

These types of modules are self-contained entities. As of Ansible 2.1, Ansible

only modifies them to change a shebang line if present.

See also

Examples of Non-native modules written in ruby are in the

Ansible

for Rubyists

repository.

Binary modules

?

From Ansible 2.2 onwards, modules may also be small binary programs. Ansible

doesn't perform any magic to make these portable to different systems so they

may be specific to the system on which they were compiled or require other

binary runtime dependencies. Despite these drawbacks, you may have

to compile a custom module against a specific binary

library if that's the only way to get access to certain resources.

Binary modules take their arguments and return data to Ansible in the same

way as

want JSON modules

.

See also

One example of a

binary module

written in go.

Old-style modules

?

Old-style modules are similar to

want JSON modules

, except that the file that

they take contains

key=value

pairs for their parameters instead of

JSON

. Ansible decides that a module is old-style when it doesn't have any of the markers that would show that it is one of the other types.

How modules are executed

?

When a user uses

ansible

or

ansible-playbook

, they

specify a task to execute. The task is usually the name of a module along with several parameters to be passed to the module. Ansible takes these values and processes them in various ways before they are finally executed on the remote machine.

Executor/task_executor

?

The TaskExecutor receives the module name and parameters that were parsed from the

playbook

(or from the command-line in the case of

/usr/bin/ansible

). It uses the name to decide whether it is looking

at a module or an

Action Plugin

. If it is

a module, it loads the

Normal Action Plugin

and passes the name, variables, and other information about the task and play to that Action Plugin for further processing.

The

normal

action plugin

?

The

normal

action plugin executes the module on the remote host. It is the primary coordinator of much of the work to actually execute the module on the managed machine.

It loads the appropriate connection plugin for the task, which then transfers or executes as needed to create a connection to that host.

It adds any internal Ansible properties to the module's parameters (for instance, the ones that pass along

no_log

to the module).

It works with other plugins (connection, shell, become, other action plugins) to create any temporary files on the remote machine and cleans up afterwards.

It pushes the module and module parameters to the remote host, although the

module_common

code described in the next section decides which format those will take.

It handles any special cases regarding modules (for example, async

execution, or complications around Windows modules that must have the same names as Python modules, so that

internal calling of modules from other Action Plugins work.)

Much of this functionality comes from the

BaseAction

class,

which lives in

plugins/action/___init___py

. It uses the

Connection

and

Shell

objects to do its work.

Note

When

tasks

are run with the

async:

parameter, Ansible

uses the

async

Action Plugin instead of the

normal

Action Plugin

to invoke it. That program flow is currently not documented. Read the

source for information on how that works.

Executor/module_common.py

?

Code in

executor/module_common.py

assembles the module

to be shipped to the managed node. The module is first read in, then examined

to determine its type:

PowerShell

and

JSON-args modules

are passed through

Module Replacer

.

New-style

Python modules

are assembled by

Ansiballz framework

.

Non-native-want-JSON

,

Binary modules

, and

Old-Style modules

aren't touched by either of these and pass through unchanged.

After the assembling step, one final

modification is made to all modules that have a shebang line. Ansible checks

whether the interpreter in the shebang line has a specific path configured with

an

ansible_\$X_interpreter

inventory variable. If it does, Ansible

substitutes that path for the interpreter path given in the module. After this, Ansible returns the complete module data and the module type to the

Normal Action

which continues execution of the module.

Assembler frameworks

?

Ansible supports two assembler frameworks: Ansiballz and the older Module Replacer.

Module Replacer framework

?

The Module Replacer framework is the original framework implementing new-style modules, and is still used for PowerShell modules. It is essentially a preprocessor (like the C Preprocessor for those familiar with that programming language). It does straight substitutions of specific substring patterns in the module file. There are two types of substitutions:

Replacements that only happen in the module file. These are public replacement strings that modules can utilize to get helpful boilerplate or access to arguments.

from

ansible.module_utils.MOD_LIB_NAME

import

*

is replaced with the

contents of the

ansible/module_utils/MOD_LIB_NAME.py

These should

only be used with

new-style Python modules

.

#<<INCLUDE_ANSIBLE_MODULE_COMMON>>

is equivalent to

from

ansible.module_utils.basic

import

*

and should also only apply to new-style Python modules.

#

POWERSHELL_COMMON

substitutes the contents of

ansible/module_utils/powershell.ps1

. It should only be used with

new-style Powershell modules

.

Replacements that are used by

ansible.module_utils

code. These are internal replacement patterns. They may be used internally, in the above public replacements, but shouldn't be used directly by modules.

"<<ANSIBLE_VERSION>>"

is substituted with the Ansible version. In

new-style Python modules

under the

Ansiballz framework

framework the proper way is to instead instantiate an

AnsibleModule
and then access the version from
:attr:
AnsibleModule.ansible_version
.
"<<INCLUDE_ANSIBLE_MODULE_COMPLEX_ARGS>>"
is substituted with
a string which is the Python
repr
of the
JSON
encoded module
parameters. Using
repr
on the JSON string makes it safe to embed in
a Python file. In new-style Python modules under the Ansiballz framework
this is better accessed by instantiating an
AnsibleModule
and
then using
AnsibleModule.params
.
<<SELINUX_SPECIAL_FILESYSTEMS>>
substitutes a string which is
a comma-separated list of file systems which have a file system dependent
security context in SELinux. In new-style Python modules, if you really
need this you should instantiate an
AnsibleModule
and then use
AnsibleModule._selinux_special_fs
. The variable has also changed
from a comma-separated string of file system names to an actual python
list of file system names.
<<INCLUDE_ANSIBLE_MODULE_JSON_ARGS>>
substitutes the module
parameters as a JSON string. Care must be taken to properly quote the
string as JSON data may contain quotes. This pattern is not substituted
in new-style Python modules as they can get the module parameters another
way.
The string
syslog.LOG_USER
is replaced wherever it occurs with the
syslog_facility
which was named in
ansible.cfg
or any
ansible_syslog_facility
inventory variable that applies to this host. In
new-style Python modules this has changed slightly. If you really need to
access it, you should instantiate an
AnsibleModule
and then use
AnsibleModule._syslog_facility
to access it. It is no longer the

actual syslog facility and is now the name of the syslog facility. See the documentation on internal arguments for details.

Ansiballz framework

?

The Ansiballz framework was adopted in Ansible 2.1 and is used for all new-style Python modules. Unlike the Module Replacer, Ansiballz uses real Python imports of things in

ansible/module_utils

instead of merely preprocessing the module. It

does this by constructing a zipfile ? which includes the module file, files in

ansible/module_utils

that are imported by the module, and some

boilerplate to pass in the module's parameters. The zipfile is then Base64

encoded and wrapped in a small Python script which decodes the Base64 encoding

and places the zipfile into a temp directory on the managed node. It then

extracts just the Ansible module script from the zip file and places that in

the temporary directory as well. Then it sets the PYTHONPATH to find Python

modules inside of the zip file and imports the Ansible module as the special name,

__main__

.

Importing it as

__main__

causes Python to think that it is executing a script rather than simply

importing a module. This lets Ansible run both the wrapper script and the module code in a single copy of Python on the remote machine.

Note

Ansible wraps the zipfile in the Python script for two reasons:

for compatibility with Python 2.6 which has a less

functional version of Python's

-m

command-line switch.

so that pipelining will function properly. Pipelining needs to pipe the

Python module into the Python interpreter on the remote node. Python

understands scripts on stdin but does not understand zip files.

Prior to Ansible 2.7, the module was executed by a second Python interpreter instead of being

executed inside of the same process. This change was made once Python-2.4 support was dropped

to speed up module execution.

In Ansiballz, any imports of Python modules from the

ansible.module_utils

package trigger inclusion of that Python file

into the zipfile. Instances of

#<<INCLUDE_ANSIBLE_MODULE_COMMON>>

in

the module are turned into

from

ansible.module_utils.basic

import

*

and

ansible/module-utils/basic.py

is then included in the zipfile.

Files that are included from

module_utils
are themselves scanned for
imports of other Python modules from
module_utils
to be included in
the zipfile as well.

Passing args

?

Arguments are passed differently by the two frameworks:

In

Module Replacer framework

, module arguments are turned into a JSON-ified string and substituted into the combined module file.

In

Ansiballz framework

, the JSON-ified string is part of the script which wraps the zipfile. Just before the wrapper script imports the Ansible module as

__main__

, it monkey-patches the private,

_ANSIBLE_ARGS

variable in

basic.py

with the variable values. When a

ansible.module_utils.basic.AnsibleModule

is instantiated, it parses this string and places the args into

AnsibleModule.params

where it can be accessed by the module's other code.

Warning

If you are writing modules, remember that the way we pass arguments is an internal implementation detail: it has changed in the past and will change again as soon as changes to the common module_utils code allow Ansible modules to forgo using

ansible.module_utils.basic.AnsibleModule

. Do not rely on the internal global

_ANSIBLE_ARGS

variable.

Very dynamic custom modules which need to parse arguments before they instantiate an

AnsibleModule

may use

_load_params

to retrieve those parameters.

Although

_load_params

may change in breaking ways if necessary to support

changes in the code, it is likely to be more stable than either the way we pass parameters or the internal global variable.

Note

Prior to Ansible 2.7, the Ansible module was invoked in a second Python interpreter and the arguments were then passed to the script over the script's stdin.

Internal arguments

?

Both

Module Replacer framework

and

Ansiballz framework

send additional arguments to

the Ansible module beyond those which the user specified in the playbook. These additional arguments are internal parameters that help implement global Ansible features. Modules often do not need to know about these explicitly because the features are implemented in `ansible.module_utils.basic`.

However, certain features need support from modules and some knowledge of the internal arguments is useful. The internal arguments in this section are global. If you need to add a local internal argument to a custom module, create an action plugin for that specific module. See

`_original_basename`

in the `copy` action plugin for an example.

`_ansible_no_log`

?

Type:

bool

Set to

True

whenever an argument in a task or play specifies

`no_log`

. Any module that calls the `AnsibleModule.log()`

function handles this action automatically. If you have a module that implements its own logging then you need to check the value of

`_ansible_no_log`

. To access

`_ansible_no_log`

in a module, instantiate the

`AnsibleModule`

utility and then check the value of

`AnsibleModule.no_log`

.

Note

`no_log`

specified in a module?

`argument_spec`

is handled by a different mechanism.

`_ansible_debug`

?

Type:

bool

Operates verbose logging and logging of external commands that a module executes. If the module uses the `AnsibleModule.debug()`

function rather than the

`AnsibleModule.log()`

function then the messages are only logged if you set the

`_ansible_debug`

argument to

True

. To access

`_ansible_debug`

in a module, instantiate the

`AnsibleModule`

utility and access

`AnsibleModule._debug`

. For more details, see

`DEFAULT_DEBUG`

.

`_ansible_diff`

?

Type:

bool

With this parameter you can configure your module to show a unified diff of changes that will be applied to the templated files. To access

`_ansible_diff`

in a module, instantiate the

`AnsibleModule`

utility and access

`AnsibleModule._diff`

. You can also access this parameter using the

`diff`

keyword in your playbook, or the relevant environment variable. For more details, see

[Playbook Keywords](#)

and the

`DIFF_ALWAYS`

configuration option.

`_ansible_verbosity`

?

Type:

int

You can use this argument to control the level (0 for none) of verbosity in logging.

`_ansible_selinux_special_fs`

?

Type:

list

Elements:

strings

This argument provides modules with the names of file systems which should have a special SELinux context. They are used by the

`AnsibleModule`

methods which operate on files (changing attributes, moving, and copying).

Most modules can use the built-in

`AnsibleModule`

methods to manipulate files. To access in a module that needs to know about these special context file systems, instantiate

`AnsibleModule`

and examine the list in

`AnsibleModule._selinux_special_fs`

.

This argument replaces

`ansible.module_utils.basic.SELINUX_SPECIAL_FS`

from

Module Replacer framework

. In the module replacer framework the argument was formatted as a comma-separated string of file system names.

Under the Ansiballz framework it is a list. You can access

`_ansible_selinux_special_fs`

using the corresponding environment variable. For more details, see the

DEFAULT_SELINUX_SPECIAL_FS

configuration option.

New in version 2.1.

`_ansible_syslog_facility`

?

This argument controls which syslog facility the module logs to. Most modules should just use the

`AnsibleModule.log()`

function which will then make use of this. If a module has to use this on its own, it should instantiate the

`AnsibleModule`

method and then retrieve the name of the syslog facility from

`AnsibleModule._syslog_facility`

. The Ansiballz code is less elegant than the

Module Replacer framework

code:

Old module_replacer way

import

syslog

syslog

.

openlog

(

NAME

,

0

,

syslog

.

LOG_USER

)

New Ansiballz way

import

syslog

facility_name

=

module

.

`_syslog_facility`

facility

=

getattr

(

syslog

,

facility_name

,

syslog

.

LOG_USER

)

syslog

.

openlog

(

NAME

,
0

,
facility
)

For more details, see the
DEFAULT_SYSLOG_FACILITY
configuration option.

New in version 2.1.

_ansible_version
?

This argument passes the version of Ansible to the module. To access it, a module should instantiate the
AnsibleModule
method and then retrieve the version from
AnsibleModule.ansible_version

. This replaces
ansible.module_utils.basic.ANSIBLE_VERSION
from
Module Replacer framework

.
New in version 2.1.
_ansible_module_name
?

Type:
str
This argument passes the information to modules about their name. For more details see, the configuration option
DEFAULT_MODULE_NAME

.
_ansible_keep_remote_files
?

Type:
bool
This argument provides instructions that modules must be ready if they need to keep the remote files. For more details,
see the
DEFAULT_KEEP_REMOTE_FILES
configuration option.

_ansible_socket
?
This argument provides modules with a socket for persistent connections. The argument is created using the
PERSISTENT_CONTROL_PATH_DIR
configuration option.

_ansible_shell_executable
?
Type:
bool
This argument ensures that modules use the designated shell executable. For more details, see the
ansible_shell_executable
remote host environment parameter.

_ansible_tmpdir
?
Type:
str
This argument provides instructions to modules that all commands must use the designated temporary directory, if
created. The action plugin designs this temporary directory.

Modules can access this parameter by using the public

`tmpdir`

property. The

`tmpdir`

property will create a temporary directory if the action plugin did not set the parameter.

The directory name is generated randomly, and the the root of the directory is determined by one of these:

`DEFAULT_LOCAL_TMP`

`remote_tmp`

`system_tmpdirs`

As a result, using the

`ansible.cfg`

configuration file to activate or customize this setting will not guarantee that you control the full value.

`_ansible_remote_tmp`

?

The module's

`tmpdir`

property creates a randomized directory name in this directory if the action plugin did not set

`_ansible_tmpdir`

. For more details, see the

`remote_tmp`

parameter of the shell plugin.

Module return values & Unsafe strings

?

At the end of a module's execution, it formats the data that it wants to return as a JSON string and prints the string to its stdout. The normal action plugin receives the JSON string, parses it into a Python dictionary, and returns it to the executor.

If Ansible templated every string return value, it would be vulnerable to an attack from users with access to managed nodes. If an unscrupulous user disguised malicious code as Ansible return value strings, and if those strings were then templated on the control node, Ansible could execute arbitrary code. To prevent this scenario, Ansible marks all strings inside returned data as

Unsafe

, emitting any Jinja2 templates in the strings verbatim, not expanded by Jinja2.

Strings returned by invoking a module through

`ActionPlugin._execute_module()`

are automatically marked as

Unsafe

by the normal action plugin. If another action plugin retrieves information from a module through some other means, it must mark its return data as

Unsafe

on its own.

In case a poorly-coded action plugin fails to mark its results as 'Unsafe', Ansible audits the results again when they are returned to the executor,

marking all strings as

Unsafe

. The normal action plugin protects itself and any other code that it calls with the result data as a parameter. The check inside the executor protects the output of all other action plugins, ensuring that subsequent tasks run by Ansible will not template anything from those results either.

Special considerations

?

Pipelining

?

Ansible can transfer a module to a remote machine in one of two ways:

it can write out the module to a temporary file on the remote host and then

use a second connection to the remote host to execute it with the

interpreter that the module needs

or it can use what's known as pipelining to execute the module by piping it into the remote interpreter's stdin.

Pipelining only works with modules written in Python at this time because Ansible only knows that Python supports this mode of operation. Supporting pipelining means that whatever format the module payload takes before being sent over the wire must be executable by Python through stdin.

Why pass args over stdin?

?

Passing arguments through stdin was chosen for the following reasons:

When combined with

`ANSIBLE_PIPELINING`

, this keeps the module's arguments from

temporarily being saved onto disk on the remote machine. This makes it harder (but not impossible) for a malicious user on the remote machine to steal any sensitive information that may be present in the arguments.

Command line arguments would be insecure as most systems allow unprivileged users to read the full commandline of a process.

Environment variables are usually more secure than the commandline but some systems limit the total size of the environment. This could lead to truncation of the parameters if we hit that limit.

`AnsibleModule`

?

Argument spec

?

The

`argument_spec`

provided to

`AnsibleModule`

defines the supported arguments for a module, as well as their type, defaults and more.

Example

`argument_spec`

:

`module`

=

`AnsibleModule`

(

`argument_spec`

=

`dict`

(

`top_level`

=

`dict`

(

`type`

=

`'dict'`

,

`options`

=

`dict`

(

`second_level`

```

=
dict
(
default
=
True
,
type
=
'bool'
,
)
)
)
))

```

This section will discuss the behavioral attributes for arguments:

type
:
type
allows you to define the type of the value accepted for the argument. The default value for **type** is **str**. Possible values are:

```

str
list
dict
bool
int
float
path
raw
jsonarg
json
bytes
bits

```

The **raw** type, performs no type validation or type casting, and maintains the type of the passed value.

elements
:
elements
works in combination with **type** when **type='list'**

.
elements
can then be defined as **elements='int'** or any other type, indicating that each element of the specified list should be of that type.

default
:
The

default

option allows sets a default value for the argument for the scenario when the argument is not provided to the module.

When not specified, the default value is

None

.

fallback

:

fallback

accepts a

tuple

where the first argument is a callable (function) that will be used to perform the lookup, based on the second argument.

The second argument is a list of values to be accepted by the callable.

The most common callable used is

env_fallback

which will allow an argument to optionally use an environment variable when the argument is not supplied.

Example:

username

=

dict

(

fallback

=

(

env_fallback

,

[

'ANSIBLE_NET_USERNAME'

]))

choices

:

choices

accepts a list of choices that the argument will accept. The types of

choices

should match the

type

.

required

:

required

accepts a boolean, either

True

or

False

that indicates that the argument is required. When not specified,

required

defaults to

False

. This should not be used in combination with

default

.

no_log

:

no_log

accepts a boolean, either

True
or
False
, that indicates explicitly whether or not the argument value should be masked in logs and output.

Note
In the absence of
`no_log`
, if the parameter name appears to indicate that the argument value is a password or passphrase (such as `?admin_password?`), a warning will be shown and the value will be masked in logs but
not
output. To disable the warning and masking for parameters that do not contain sensitive information, set
`no_log`
to
False
.

`aliases`
:
`aliases`
accepts a list of alternative argument names for the argument, such as the case where the argument is
name
but the module accepts
`aliases=['pkg']`
to allow
`pkg`
to be interchangeably with
name
.

Use of aliases can make module interfaces confusing, so we recommend adding them only when necessary. If you are updating argument names to fix a typo or improve the interface, consider moving the old names to
`deprecated_aliases`
rather than keeping them around indefinitely.

`options`
:
`options`
implements the ability to create a `sub-argument_spec`, where the sub options of the top level argument are also validated using the attributes discussed in this section. The example at the top of this section demonstrates use of
`options`
.

`type`
or
`elements`
should be
dict
is this case.
`apply_defaults`
:
`apply_defaults`
works alongside
`options`
and allows the
default
of the sub-options to be applied even when the top-level argument is not supplied.

In the example of the
`argument_spec`
at the top of this section, it would allow

module.params['top_level']['second_level']
to be defined, even if the user does not provide

top_level

when calling the module.

removed_in_version

:

removed_in_version

indicates which version of ansible-core or a collection a deprecated argument will be removed in. Mutually exclusive with

removed_at_date

, and must be used with

removed_from_collection

.

Example:

option

=

{

'type'

:

'str'

,

'removed_in_version'

:

'2.0.0'

,

'removed_from_collection'

:

'testns.testcol'

,

},

removed_at_date

:

removed_at_date

indicates that a deprecated argument will be removed in a minor ansible-core release or major collection release after this date. Mutually exclusive with

removed_in_version

, and must be used with

removed_from_collection

.

Example:

option

=

{

'type'

:

'str'

,

'removed_at_date'

:

'2020-12-31'

,

'removed_from_collection'

:

'testns.testcol'

,

```

},
removed_from_collection
:
Specifies which collection (or ansible-core) deprecates this deprecated argument. Specify
ansible.builtin
for ansible-core, or the collection?s name (format
foo.bar
). Must be used with
removed_in_version
or
removed_at_date
.
deprecated_aliases
:
Deprecates aliases of this argument. Must contain a list or tuple of dictionaries having some the following keys:
name
:
The name of the alias to deprecate. (Required.)
version
:
The version of ansible-core or the collection this alias will be removed in. Either
version
or
date
must be specified.
date
:
The a date after which a minor release of ansible-core or a major collection release will no longer contain this alias..
Either
version
or
date
must be specified.
collection_name
:
Specifies which collection (or ansible-core) deprecates this deprecated alias. Specify
ansible.builtin
for ansible-core, or the collection?s name (format
foo.bar
). Must be used with
version
or
date
.
Examples:
option
=
{
'type'
:
'str'
,
'aliases'
:

```

```
[
  'foo'
,
  'bar'
],
'deprecated_aliases'
:
[
  {
    'name'
    :
    'foo'
    ,
    'version'
    :
    '2.0.0'
    ,
    'collection_name'
    :
    'testns.testcol'
    ,
  },
  {
    'name'
    :
    'foo'
    ,
    'date'
    :
    '2020-12-31'
    ,
    'collection_name'
    :
    'testns.testcol'
    ,
  },
],
},
mutually_exclusive
:
If
options
is specified,
mutually_exclusive
refers to the sub-options described in
options
and behaves as in
Dependencies between module options
-
required_together
:
If
options
is specified,
```

`required_together`
refers to the sub-options described in
`options`
and behaves as in
Dependencies between module options

.
`required_one_of`
:
If
`options`
is specified,
`required_one_of`
refers to the sub-options described in
`options`
and behaves as in
Dependencies between module options

.
`required_if`
:
If
`options`
is specified,
`required_if`
refers to the sub-options described in
`options`
and behaves as in
Dependencies between module options

.
`required_by`
:
If
`options`
is specified,
`required_by`
refers to the sub-options described in
`options`
and behaves as in
Dependencies between module options

.
`context`
:
New in version 2.17.
You can set the value of the
`context`
key to a dict of custom content. This allows you to provide additional context in the argument spec. The content provided is not validated or utilized by the core engine.

Example:

```
option
=
{
'type'
:
'str'
,
```



```

'context'
:
{
'disposition'
:
'/properties/apiType'
,
},
'choices'
:
[
'http'
,
'soap'
],
}

```

Dependencies between module options

?

The following are optional arguments for
AnsibleModule()

:

module

=

AnsibleModule

(

argument_spec

,

mutually_exclusive

=

[

(

'path'

,

'content'

),

],

required_one_of

=

[

(

'path'

,

'content'

),

],

)

mutually_exclusive

:

Must be a sequence (list or tuple) of sequences of strings. Every sequence of strings is a list of option names which are mutually exclusive. If more than one options of a list are specified together, Ansible will fail the module with an error.

Example:

mutually_exclusive

=

[

```
(
'path'
,
'content'
),
(
'repository_url'
,
'repository_filename'
),
],
```

In this example, the options

path

and

content

must not be specified at the same time. Also the options

repository_url

and

repository_filename

must not be specified at the same time. But specifying

path

and

repository_url

is accepted.

To ensure that precisely one of two (or more) options is specified, combine

mutually_exclusive

with

required_one_of

```
.
required_together
:
```

Must be a sequence (list or tuple) of sequences of strings. Every sequence of strings is a list of option names which are must be specified together. If at least one of these options are specified, the other ones from the same sequence must all be present.

Example:

```
required_together
```

```
=
[
(
'file_path'
,
'file_hash'
),
],
```

In this example, if one of the options

file_path

or

file_hash

is specified, Ansible will fail the module with an error if the other one is not specified.

```
required_one_of
```

```
:
```

Must be a sequence (list or tuple) of sequences of strings. Every sequence of strings is a list of option names from which at least one must be specified. If none one of these options are specified, Ansible will fail module execution.

Example:

required_one_of

=

[

(

'path'

,

'content'

),

],

In this example, at least one of

path

and

content

must be specified. If none are specified, execution will fail. Specifying both is explicitly allowed; to prevent this, combine

required_one_of

with

mutually_exclusive

.

required_if

:

Must be a sequence of sequences. Every inner sequence describes one conditional dependency. Every sequence must have three or four values. The first two values are the option's name and the option's value which describes the condition. The further elements of the sequence are only needed if the option of that name has precisely this value.

If you want that all options in a list of option names are specified if the condition is met, use one of the following forms:

(

'option_name'

,

option_value

,

(

'option_a'

,

'option_b'

,

...

)),

(

'option_name'

,

option_value

,

(

'option_a'

,

'option_b'

,

...

),

False

),

If you want that at least one option of a list of option names is specified if the condition is met, use the following form:

(

'option_name'

,

```
option_value
```

```
,  
(  
'option_a'  
,  
'option_b'  
,  
...  
),  
True  
),  
Example:  
required_if  
=  
[  
(  
'state'  
,  
'present'  
,  
(  
'path'  
,  
'content'  
)  
),  
True  
)  
(  
'force'  
,  
True  
,  
(  
'force_reason'  
,  
'force_code'  
))  
],
```

In this example, if the user specifies

state=present

, at least one of the options

path

and

content

must be supplied (or both). To make sure that precisely one can be specified, combine

required_if

with

mutually_exclusive

.

On the other hand, if

force

(a boolean parameter) is set to

true

,

yes
and so on, both
force_reason
and
force_code
must be specified.
required_by

:

Must be a dictionary mapping option names to sequences of option names. If the option name in a dictionary key is specified, the option names it maps to must all also be specified. Note that instead of a sequence of option names, you can also specify one single option name.

Example:

```
required_by
=
{
'force'
:
'force_reason'
,
'path'
:
(
'mode'
,
'owner'
,
'group'
),
},
```

In the example, if

force
is specified,
force_reason
must also be specified. Also, if
path
is specified, then three three options
mode

,
owner
and
group
also must be specified.

Declaring check mode support

?

To declare that a module supports check mode, supply
supports_check_mode=True
to the
AnsibleModule()
call:

```
module
=
AnsibleModule
(
argument_spec
```

```
,
supports_check_mode
=
True
)
```

The module can determine whether it is called in check mode by checking the boolean value `module.check_mode`

. If it evaluates to

`True`

, the module must take care not to do any modification.

If

`supports_check_mode=False`

is specified, which is the default value, the module will exit in check mode with

`skipped=True`

and message

`remote`

`module`

(`<insert`

`module`

`name`

`here>`)

does

not

support

check

mode

.

Adding file options

?

To declare that a module should add support for all common file options, supply

`add_file_common_args=True`

to the

`AnsibleModule()`

call:

`module`

=

`AnsibleModule`

(

`argument_spec`

,

`add_file_common_args`

=

`True`

)

You can find

a list of all file options [here](#)

. It is recommended that you make your

`DOCUMENTATION`

extend the doc fragment

`ansible.builtin.files`

(see

[Documentation fragments](#)

) in this case, to make sure that all these fields are correctly documented.

The helper functions

```

module.load_file_common_arguments()
and
module.set_fs_attributes_if_different()
can be used to handle these arguments for you:
argument_spec
=
{
'path'
:
{
'type'
:
'str'
,
'required'
:
True
,
},
}
module
=
AnsibleModule
(
argument_spec
,
add_file_common_args
=
True
)
changed
=
False
# TODO do something with module.params['path'], like update its contents
# Ensure that module.params['path'] satisfies the file options supplied by the user
file_args
=
module
.
load_file_common_arguments
(
module
.
params
)
changed
=
module
.
set_fs_attributes_if_different
(
file_args
,
changed

```

```
)  
module  
.  
exit_json  
(  
  changed  
  =  
  changed  
)
```

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: https://docs.ansible.com/developing_python_3.html

Ansible and Python 3 ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords

Return Values

Ansible Configuration Settings

Controlling how Ansible behaves: precedence rules

YAML Syntax

Python 3 Support

Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Developer Guide
Ansible and Python 3
Edit on GitHub
Ansible and Python 3
?
The
ansible-core
code runs Python 3 (for specific versions check
Control Node Requirements
Contributors to
ansible-core
and to Ansible Collections should be aware of the tips in this document so that they can write code
that will run on the same versions of Python as the rest of Ansible.
Minimum version of Python 3.x and Python 2.x
Developing Ansible code that supports Python 2 and Python 3
Understanding strings in Python 2 and Python 3
Control node string strategy: the Unicode Sandwich
Unicode Sandwich common borders: places to convert bytes to text in control node code
Reading and writing to files
Filesystem interaction
Interacting with other programs
Module string strategy: Native String
Module_utils string strategy: hybrid
Tips, tricks, and idioms for Python 2/Python 3 compatibility
Use forward-compatibility boilerplate
Prefix byte strings with
b_
Import Ansible's bundled Python
six
library
Handle exceptions with
as
Update octal numbers
String formatting for control node code
Use
str.format()
for Python 2.6 compatibility
Use percent format with byte strings
We do have some considerations depending on the types of Ansible code:

code on the control node - code that runs on the machine where you invoke

`/usr/bin/ansible`

, only needs to support the control node's Python versions.

modules - the code which Ansible transmits to and invokes on the managed machine. Modules need to support the managed node's Python versions, with some exceptions.

shared

`module_utils`

code - the common code that is used by modules to perform tasks and sometimes used by code on the control node.

Shared

`module_utils`

code needs to support the same range of Python as the modules.

However, the three types of code do not use the same string strategy. If you're developing a module or some

`module_utils`

code, be sure to read the section on string strategy carefully.

Minimum version of Python 3.x and Python 2.x

?

See

Control Node Requirements

and

Managed Node Requirements

for the

specific versions supported.

Your custom modules can support any version of Python (or other languages) you want, but the above are the requirements for the code contributed to the Ansible project.

Developing Ansible code that supports Python 2 and Python 3

?

The best place to start learning about writing code that supports both Python 2 and Python 3

is

Lennart Regebro's book: *Porting to Python 3*

.

The book describes several strategies for porting to Python 3. The one we're

using is

to support Python 2 and Python 3 from a single code base

Understanding strings in Python 2 and Python 3

?

Python 2 and Python 3 handle strings differently, so when you write code that supports Python 3

you must decide what string model to use. Strings can be an array of bytes (like in C) or

they can be an array of text. Text is what we think of as letters, digits,

numbers, other printable symbols, and a small number of unprintable symbols?

(control codes).

In Python 2, the two types for these (

`str`

for bytes and

`unicode`

for text) are often used interchangeably. When dealing only

with ASCII characters, the strings can be combined, compared, and converted

from one type to another automatically. When non-ASCII characters are

introduced, Python 2 starts throwing exceptions due to not knowing what encoding

the non-ASCII characters should be in.

Python 3 changes this behavior by making the separation between bytes (

`bytes`

)

and text (

`str`

) more strict. Python 3 will throw an exception when trying to combine and compare the two types. The programmer has to explicitly convert from one type to the other to mix values from each.

In Python 3 it is immediately apparent to the programmer when code is mixing the byte and text types inappropriately, whereas in Python 2, code that mixes those types may work until a user causes an exception by entering non-ASCII input.

Python 3 forces programmers to proactively define a strategy for working with strings in their program so that they don't mix text and byte strings unintentionally.

Ansible uses different strategies for working with strings in the code on the control node, in

:ref:

modules <module_string_strategy>

, and in

module_utils

code.

Control node string strategy: the Unicode Sandwich

?

Until recently

ansible-core

supported Python 2.x and followed this strategy, known as the Unicode Sandwich (named after Python 2's

unicode

text type). For Unicode Sandwich we know that

at the border of our code and the outside world (for example, file and network IO, environment variables, and some library calls) we are going to receive bytes.

We need to transform these bytes into text and use that throughout the internal portions of our code. When we have to send those strings back out to the outside world we first convert the text back into bytes.

To visualize this, imagine a ?sandwich? consisting of a top and bottom layer of bytes, a layer of conversion between, and all text type in the center.

For compatibility reasons you will see a bunch of custom functions we developed (

to_text

/

to_bytes

/

to_native

)

and while Python 2 is not a concern anymore we will continue to use them as they apply for other cases that make dealing with unicode problematic.

While we will not be using it most of it anymore, the documentation below is still useful for those developing modules that still need to support both Python 2 and 3 simultaneously.

Unicode Sandwich common borders: places to convert bytes to text in control node code

?

This is a partial list of places where we have to convert to and from bytes when using the Unicode Sandwich string strategy. It is not exhaustive but it gives you an idea of where to watch for problems.

Reading and writing to files

?

In Python 2, reading from files yields bytes. In Python 3, it can yield text.

To make code that's portable to both we don't make use of Python 3's ability to yield text but instead do the conversion explicitly ourselves. For example:

from

ansible.module_utils.common.text.converters

import

to_text

```

with
open
(
'filename-with-utf8-data.txt'
,
'rb'
)
as
my_file
:
b_data
=
my_file
.
read
()
try
:
data
=
to_text
(
b_data
,
errors
=
'surrogate_or_strict'
)
except
UnicodeError
:
# Handle the exception gracefully -- usually by displaying a good
# user-centric error message that can be traced back to this piece
# of code.
pass

```

Note

Much of Ansible assumes that all encoded text is UTF-8. At some point, if there is demand for other encodings we may change that, but for now it is safe to assume that bytes are UTF-8.

Writing to files is the opposite process:

```

from
ansible.module_utils.common.text.converters
import
to_bytes
with
open
(
'filename.txt'
,
'wb'
)
as
my_file
:

```

```
my_file
```

```
.
```

```
write
```

```
(
```

```
to_bytes
```

```
(
```

```
some_text_string
```

```
))
```

Note that we don't have to catch

UnicodeError

here because we're

transforming to UTF-8 and all text strings in Python can be transformed back to UTF-8.

Filesystem interaction

?

Dealing with file names often involves dropping back to bytes because on UNIX-like systems file names are bytes. On Python 2, if we pass a text string to these functions, the text string will be converted to a byte string inside of the function and a traceback will occur if non-ASCII characters are present. In Python 3, a traceback will only occur if the text string can't be decoded in the current locale, but it is still good to be explicit and have code which works on both versions:

```
import
```

```
os.path
```

```
from
```

```
ansible.module_utils.common.text.converters
```

```
import
```

```
to_bytes
```

```
filename
```

```
=
```

```
u
```

```
'/var/tmp/?????.txt'
```

```
f
```

```
=
```

```
open
```

```
(
```

```
to_bytes
```

```
(
```

```
filename
```

```
),
```

```
'wb'
```

```
)
```

```
mtime
```

```
=
```

```
os
```

```
.
```

```
path
```

```
.
```

```
getmtime
```

```
(
```

```
to_bytes
```

```
(
```

```
filename
```

```
))
```

```
b_filename
```

```
=
```

```
os
```

```
.
```

```
path
```

```
.
```

```
expandvars
```

```
(
```

```
to_bytes
```

```
(
```

```
filename
```

```
))
```

```
if
```

```
os
```

```
.
```

```
path
```

```
.
```

```
exists
```

```
(
```

```
to_bytes
```

```
(
```

```
filename
```

```
)):
```

```
pass
```

When you are only manipulating a filename as a string without talking to the filesystem (or a C library which talks to the filesystem) you can often get away without converting to bytes:

```
import
```

```
os.path
```

```
os
```

```
.
```

```
path
```

```
.
```

```
join
```

```
(
```

```
u
```

```
'/var/tmp/café'
```

```
,
```

```
u
```

```
'????'
```

```
)
```

```
os
```

```
.
```

```
path
```

```
.
```

```
split
```

```
(
```

```
u
```

```
'/var/tmp/café/????'
```

```
)
```

On the other hand, if the code needs to manipulate the file name and also talk to the filesystem, it can be more convenient to transform to bytes right away and manipulate in bytes.

Warning

Make sure all variables passed to a function are the same type.

If you're working with something like

```
os.path.join()
```

which takes

multiple strings and uses them in combination, you need to make sure that

all the types are the same (either all bytes or all text). Mixing

bytes and text will cause tracebacks.

Interacting with other programs

?

Interacting with other programs goes through the operating system and

C libraries and operates on things that the UNIX kernel defines. These

interfaces are all byte-oriented so the Python interface is byte oriented as

well. On both Python 2 and Python 3, byte strings should be given to Python's

subprocess library and byte strings should be expected back from it.

One of the main places in Ansible's control node code that we interact with

other programs is the connection plugins?

```
exec_command
```

methods. These

methods transform any text strings they receive in the command (and arguments

to the command) to execute into bytes and return stdout and stderr as byte strings

Higher level functions (like action plugins?

```
_low_level_execute_command
```

```
)
```

transform the output into text strings.

Module string strategy: Native String

?

In modules we use a strategy known as Native Strings. This makes things

easier on the community members who maintain so many of Ansible's

modules, by not breaking backwards compatibility by

mandating that all strings inside of modules are text and converting between

text and bytes at the borders.

Native strings refer to the type that Python uses when you specify a bare

string literal:

```
"This is a native string"
```

In Python 2, these are byte strings. In Python 3 these are text strings. Modules should be

coded to expect bytes on Python 2 and text on Python 3.

Module_utils string strategy: hybrid

?

In

```
module_utils
```

code we use a hybrid string strategy. Although Ansible's

```
module_utils
```

code is largely like module code, some pieces of it are

used by the control node as well. So it needs to be compatible with modules

and with the control node's assumptions, particularly the string strategy.

The module_utils code attempts to accept native strings as input

to its functions and emit native strings as their output.

In

```
module_utils
```

```
code:
```

Functions

must

accept string parameters as either text strings or byte strings.

Functions may return either the same type of string as they were given or the native string type for the Python version

they are run on.

Functions that return strings

must

document whether they return strings of the same type as they were given or native strings.

Module-utils functions are therefore often very defensive in nature.

They convert their string parameters into text (using

`ansible.module_utils.common.text.converters.to_text`

)

at the beginning of the function, do their work, and then convert

the return values into the native string type (using

`ansible.module_utils.common.text.converters.to_native`

)

or back to the string type that their parameters received.

Tips, tricks, and idioms for Python 2/Python 3 compatibility

?

Use forward-compatibility boilerplate

?

Use the following boilerplate code at the top of all python files

to make certain constructs act the same way on Python 2 and Python 3:

Make coding more python3-ish

from

`__future__`

import

(

`absolute_import`

,

`division`

,

`print_function`

)

`__metaclass__`

=

`type`

`__metaclass__`

=

`type`

makes all classes defined in the file into new-style

classes without explicitly inheriting from

`object`

.

The

`__future__`

imports do the following:

`absolute_import`

:

Makes imports look in

`sys.path`

for the modules being

imported, skipping the directory in which the module doing the importing

lives. If the code wants to use the directory in which the module doing

the importing, there's a new dot notation to do so.

`division`

:

Makes division of integers always return a float. If you need to

find the quotient use

x

//

y

instead of

x

/

y

.

print_function

:

Changes

print

from a keyword into a function.

See also

PEP 0328: Absolute Imports

PEP 0238: Division

PEP 3105: Print function

Prefix byte strings with

b_

?

Since mixing text and bytes types leads to tracebacks we want to be clear about what variables hold text and what variables hold bytes. We do this by prefixing any variable holding bytes with

b_

. For example:

filename

=

u

'/var/tmp/café.txt'

b_filename

=

to_bytes

(

filename

)

with

open

(

b_filename

)

as

f

:

data

=

f

.

read

()

We do not prefix the text strings instead because we only operate on byte strings at the borders, so there are fewer variables that need bytes than text.

Import Ansible's bundled Python

six
library
?
The third-party Python

six
library exists
to help projects create code that runs on both Python 2 and Python 3. Ansible
includes a version of the library in module_utils so that other modules can use it
without requiring that it is installed on the remote system. To make use of
it, import it like this:

```
from  
ansible.module_utils  
import
```

six
Note

Ansible can also use a system copy of six
Ansible will use a system copy of six if the system copy is a later
version than the one Ansible bundles.

Handle exceptions with

```
as  
?
```

In order for code to function on Python 2.6+ and Python 3, use the
new exception-catching syntax which uses the

```
as  
keyword:
```

```
try  
:  
a  
=  
2  
/  
0  
except  
ValueError
```

```
as  
e  
:  
module  
.  
fail_json
```

```
(  
msg  
=  
"Tried to divide by zero:
```

```
%s  
"
```

```
%  
e  
)
```

```
Do  
not
```

use the following syntax as it will fail on every version of Python 3:

```
try:  
    a = 2/0
```

except ValueError, e:

```
    module.fail_json(msg="Tried to divide by zero: %s" % e)
```

Update octal numbers

?

In Python 2.x, octal literals could be specified as

0755

. In Python 3,

octals must be specified as

0o755

.

String formatting for control node code

?

Use

```
str.format()
```

for Python 2.6 compatibility

?

Starting in Python 2.6, strings gained a method called

```
format()
```

to put

strings together. However, one commonly used feature of

```
format()
```

wasn't

added until Python 2.7, so you need to remember not to use it in Ansible code:

Does not work in Python 2.6!

```
new_string
```

```
=
```

```
"Dear
```

```
{}
```

```
, Welcome to
```

```
{}
```

```
"
```

.

```
format
```

```
(
```

```
    username
```

```
,
```

```
    location
```

```
)
```

Use this instead

```
new_string
```

```
=
```

```
"Dear
```

```
{0}
```

```
, Welcome to
```

```
{1}
```

```
"
```

.

```
format
```

```
(
```

```
    username
```

```
,
```

```
    location
```

```
)
```

Both of the format strings above map positional arguments of the

`format()`

method into the string. However, the first version doesn't work in

Python 2.6. Always remember to put numbers into the placeholders so the code is compatible with Python 2.6.

See also

Python documentation on format strings:

[format strings in 2.6](#)

[format strings in 3.x](#)

[Use percent format with byte strings](#)

?

In Python 3.5 and later, byte strings do not have a

`format()`

method. However, it

does have support for the older, percent-formatting.

`b_command_line`

=

`b`

`'ansible-playbook --become-user`

`%s`

`-K`

`%s`

`,`

`%`

`(`

`user`

`,`

`playbook_file`

`)`

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: <https://docs.ansible.com/ecosystem.html>

Ansible Ecosystem | Ansible Documentation

Ansible documentation

Toggle navigation

Products

Blog

Community

Webinars and training

Try it now

Toggle navigation

Join the community

Users

Developers

Maintainers

Ansible core

Ansible ecosystem

Red Hat Ansible Automation Platform

Ansible Ecosystem

The Ansible ecosystem contains a wide range of open source projects managed by Red Hat and a vibrant community of contributors.

Got thoughts or feedback on this site? We want to hear from you!

Join us in the

Ansible Forum

or open a

GitHub issue

in the docsite repository.

Awesome Ansible

A collaborative curated list of awesome Ansible resources, tools, roles, tutorials and other related content.

Awesome Ansible list

Ansible collections

Ansible collections offer distributions of playbooks, roles, modules, and plugins.

Collection index

Find out how to use collections

Learn how to contribute to collections

Index of all modules and plugins

Ansible AWX

AWX provides a web-based user interface, REST API, and task engine built on top of Ansible.

Ansible AWX documentation

AWX Operator

Ansible AWX Operator offers built-in intelligence and operational best practices for deploying on Kubernetes environments.

AWX Operator documentation

Ansible Builder

Ansible Builder lets you create Execution Environments, which are container images that act as Ansible control nodes.

Ansible Builder documentation

Ansible Compat

Compat is a Python package that assists with compatibility between different Ansible releases, starting at version 2.9.

Ansible Compat documentation

Ansible Core

Ansible Core is the language and runtime that powers automation. It also provides command-line tools such as Ansible Test.

Ansible Core documentation

Ansible Test documentation

Ansible Creator

Ansible Creator is a Command-Line Interface (CLI) tool designed for effortlessly scaffolding all your Ansible content.

[Ansible Creator documentation](#)

Ansible Development Environment

A pip-like install for Ansible collections.

[Ansible Development Environment documentation](#)

Ansible Development Tools

Ansible Development Tools (ADT) streamlines the setup and usage of several tools for creating Ansible content.

[Ansible Development Tools documentation](#)

Event-Driven Ansible Server

Event-Driven Ansible Server offers scalable and flexible automation that can subscribe to a wide variety of event sources.

[Event-Driven Ansible Server documentation](#)

Edge Automation

Edge provides tooling and collections to run automation jobs on device endpoints at the very edge of your infrastructure.

Osbuild Composer Collection

Common Industrial Protocol (CIP) Collection

FDO Collection

MicroShift Collection

Galaxy NG

Galaxy NG jumpstarts automation projects with Ansible community content.

[Galaxy NG documentation](#)

Ansible Lint

Lint improves code quality through proven best practices, patterns, and behaviors so that your Ansible content results in reliable and consistent automation.

[Ansible Lint documentation](#)

Molecule

Molecule helps you develop and test Ansible roles.

[Molecule documentation](#)

Ansible Navigator

Ansible Navigator is a command-line tool for creating, reviewing, and troubleshooting Ansible content.

[Ansible Navigator documentation](#)

ansible-pylibssh

ansible-pylibssh provides Python bindings for Ansible with the libssh project.

[ansible-pylibssh documentation](#)

Ansible Pytest

Enables the use of Ansible in tests as well as the use of pytest as a collection unit test runner, and exposes molecule scenarios using a pytest fixture.

[Ansible Pytest documentation](#)

Ansible Rulebook

Ansible Rulebook is a command-line tool that listens to events so your automation can react when software or system states change.

[Ansible Rulebook documentation](#)

Ansible Runner

Ansible Runner provides a stable and consistent interface abstraction to Ansible.

[Ansible Runner documentation](#)

Ansible SDK

Ansible SDK is a toolkit that lets you harness the power and simplicity of Ansible automation directly from your applications.

[Ansible SDK documentation](#)

Ansible Sign

Ansible Sign is a utility for signing and verifying Ansible content.

[Ansible Sign documentation](#)

Tox Ansible

Tox Ansible is a utility designed to simplify the testing of Ansible content collections.

[Tox Ansible documentation](#)

[Ansible VS Code Extension](#)

The VS Code extension adds Ansible language support to Visual Studio Code and OpenVSX compatible editors.

[Ansible VS Code Extension documentation](#)

[Ansible community package](#)

The Ansible community package consists of ansible-core and a set of Ansible collections published as the Python ``ansible`` package, in tradition of the Ansible 2.9 and earlier "batteries included" releases.

[PyPI page for Ansible community package](#)

[Documentation of package build process](#)

[Source code of package build](#)

[Source code of the Antsibull build tool](#)

[Antsibull Changelog](#)

A changelog generator used by ansible-core and Ansible collections.

[Documentation](#)

[Source code](#)

[Antsibull Docs](#)

Tooling for building documentation for Ansible collections, ansible-core, and the Ansible community package.

[Documentation](#)

[Source code](#)

[Python library for parsing Ansible markup](#)

[TypeScript library for parsing Ansible markup](#)

[CC BY-SA 4.0 |](#)

[Privacy policy |](#)

[Sponsored by](#)

Content from: <https://docs.ansible.com/index.html>

[Ansible Documentation](#)

[Ansible documentation](#)

[Toggle navigation](#)

[Products](#)

[Blog](#)

[Community](#)

[Webinars and training](#)

[Try it now](#)

[Toggle navigation](#)

[Join the community](#)

[Users](#)

[Developers](#)

[Maintainers](#)

[Ansible core](#)

[Ansible ecosystem](#)

[Red Hat Ansible Automation Platform](#)

[Ansible community documentation](#)

Ansible offers open-source automation that is simple, flexible, and powerful.

Got thoughts or feedback on this site? We want to hear from you!

Join us in the

[Ansible Forum](#)

or open a

[GitHub issue](#)

in the docsite repository.

[Quicklinks](#)

[Ansible package docs home](#)

[Collection index](#)

[Modules and plugins index](#)

[Get started with Ansible](#)

[Understand the fundamentals of Ansible automation](#)

[Install the Ansible package](#)

[Run your first ad hoc command in a few easy steps](#)

[Users](#)

[Start writing Ansible playbooks](#)

[Learn about Ansible modules](#)

[Build inventory files to manage multiple hosts](#)

[Continue the Ansible user journey](#)

[Developers](#)

[Set up your development environment](#)

[Learn how Ansible works](#)

[Write custom modules or plugins](#)

[Continue the Ansible developer journey](#)

[Maintainers](#)

[Review community maintainer responsibilities](#)

[Understand Ansible contributor paths](#)

[Explore ways to grow community](#)

[Continue the Ansible community maintainer journey](#)

[CC BY-SA 4.0](#) |

[Privacy policy](#) |

[Sponsored by](#)

Content from: <https://docs.ansible.com/maintainers.html>

Maintainers | Ansible Documentation

Ansible documentation

Toggle navigation

Products

Blog

Community

Webinars and training

Try it now

Toggle navigation

Join the community

Users

Developers

Maintainers

Ansible core

Ansible ecosystem

Red Hat Ansible Automation Platform

Maintainers

Ansible community maintainers are trusted contributors who oversee project lifecycle and overall health.

Got thoughts or feedback on this site? We want to hear from you!

Join us in the

Ansible Forum

or open a

GitHub issue

in the docsite repository.

Quicklinks

News for maintainers

Community topics

Package inclusion requests

Learn about community maintainer responsibilities

Review community maintainer responsibilities

Backport merged pull requests to stable branches

Regularly release stable versions

Look after collection documentation

Expand community around a collection

Explore ways to grow community

Maintain good contributor documentation

Understand Ansible contributor paths

Get collections included in the Ansible package

Understand the inclusion process

Review other inclusion requests

Submit your collection for inclusion

Participate in cross-project governance

Discuss and vote on community topics

Join the Ansible community steering committee

CC BY-SA 4.0 |

Privacy policy |

Sponsored by

Content from: <https://docs.ansible.com/modules.html>

Using Ansible modules and plugins ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Using Ansible modules and plugins

Introduction to modules

Boolean variables

Module maintenance and support

Rejecting modules

Working with plugins

Modules and plugins index

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins

Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Using Ansible modules and plugins
Edit on GitHub
Using Ansible modules and plugins
?

Note

Making Open Source More Inclusive

Red Hat is committed to replacing problematic language in our code, documentation, and web properties. We are beginning with these four terms: master, slave, blacklist, and whitelist. We ask that you open an issue or pull request if you come upon a term that we have missed. For more details, see our CTO Chris Wright's message

.

Welcome to the Ansible guide for working with modules, plugins, and collections.

Ansible modules are units of code that can control system resources or execute system commands.

Ansible provides a module library that you can execute directly on remote hosts or through playbooks.

You can also write custom modules.

Similar to modules are plugins, which are pieces of code that extend core Ansible functionality.

Ansible uses a plugin architecture to enable a rich, flexible, and expandable feature set.

Ansible ships with several plugins and lets you easily use your own plugins.

Introduction to modules

Boolean variables

Module maintenance and support

Maintenance

Issue Reporting

Support

Rejecting modules

Working with plugins

Action plugins

Become plugins

Cache plugins

Callback plugins

Cliconf plugins

Connection plugins

Docs fragments

[Filter plugins](#)

[Httpapi plugins](#)

[Inventory plugins](#)

[Lookup plugins](#)

[Modules](#)

[Module utilities](#)

[Netconf plugins](#)

[Shell plugins](#)

[Strategy plugins](#)

[Terminal plugins](#)

[Test plugins](#)

[Vars plugins](#)

[Modules and plugins index](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: <https://docs.ansible.com/platform.html>

Red Hat Ansible Automation Platform | Ansible Documentation

Ansible documentation

Toggle navigation

Products

Blog

Community

Webinars and training

Try it now

Toggle navigation

Join the community

Users

Developers

Maintainers

Ansible core

Ansible ecosystem

Red Hat Ansible Automation Platform

Red Hat Ansible Automation Platform

Ansible Automation Platform centralizes automation infrastructure, adds enterprise-grade capabilities, and provides supported collections.

Got thoughts or feedback on this site? We want to hear from you!

Join us in the

Ansible Forum

or open a

GitHub issue

in the docsite repository.

Red Hat Ansible Automation Platform

Red Hat Ansible Automation Platform provides everything needed to create, execute, and manage automation in a single subscription. From execution environments to certified collections to automation analytics, discover the features and benefits of Ansible Automation Platform. This now includes Ansible controller 4.5 documentation, except the CLI and Release notes listed below.

Visit the Red Hat Ansible Automation Platform documentation

Release notes

View the changes made since the last release of Ansible Automation Controller.

Visit the release notes

Automation Controller CLI

Learn how to use the command-line client for AWX and Ansible Automation Controller.

Visit the CLI guide

Automation Controller documentation archive

Find documentation for older versions of Automation Controller.

Visit the Automation Controller documentation archive

CC BY-SA 4.0 |

Privacy policy |

Sponsored by

Content from: https://docs.ansible.com/playbooks_vault.html

Protecting sensitive data with Ansible vault ? Ansible Community Documentation

Blog

Ansible community forum

Documentation

Ansible Community Documentation

Ansible

Select version:

latest

11

devel

Search docs:

Ansible getting started

Getting started with Ansible

Getting started with Execution Environments

Installation, Upgrade & Configuration

Installation Guide

Ansible Porting Guides

Using Ansible

Building Ansible inventories

Using Ansible command line tools

Using Ansible playbooks

Protecting sensitive data with Ansible vault

Ansible Vault

Managing vault passwords

Encrypting content with Ansible Vault

Using encrypted variables and files

Configuring defaults for using encrypted content

When are encrypted files made visible?

Format of files encrypted with Ansible Vault

Using Ansible modules and plugins

Using Ansible collections

Using Ansible on Windows, BSD, and z/OS UNIX

Ansible tips and tricks

Contributing to Ansible

Ansible Community Guide

Ansible Collections Contributor Guide

ansible-core Contributors Guide

Advanced Contributor Guide

Ansible documentation style guide

Extending Ansible

Developer Guide

Common Ansible Scenarios

Legacy Public Cloud Guides

Network Automation

Network Getting Started

Network Advanced Topics

Network Developer Guide

Ansible Galaxy

Galaxy User Guide

Galaxy Developer Guide

Reference & Appendices

Collection Index

Indexes of all modules and plugins
Playbook Keywords
Return Values
Ansible Configuration Settings
Controlling how Ansible behaves: precedence rules
YAML Syntax
Python 3 Support
Interpreter Discovery
Releases and maintenance
Testing Strategies
Sanity Tests
Frequently Asked Questions
Glossary
Ansible Reference: Module Utilities
Special Variables
Red Hat Ansible Automation Platform
Ansible Automation Hub
Logging Ansible output
Roadmaps
Ansible Roadmap
ansible-core Roadmaps
Ansible
Protecting sensitive data with Ansible vault
Edit on GitHub
Protecting sensitive data with Ansible vault
?

Note

Making Open Source More Inclusive

Red Hat is committed to replacing problematic language in our code, documentation, and web properties. We are beginning with these four terms: master, slave, blacklist, and whitelist. We ask that you open an issue or pull request if you come upon a term that we have missed. For more details, see our CTO Chris Wright's message

.

Welcome to the Ansible vault documentation.

Ansible vault provides a way to encrypt and manage sensitive data such as passwords.

This guide introduces you to Ansible vault and covers the following topics:

Managing vault passwords.

Encrypting content and files with Ansible vault.

Using encrypted variables and files.

Ansible Vault

Managing vault passwords

Choosing between a single password and multiple passwords

Managing multiple passwords with vault IDs

Storing and accessing vault passwords

Encrypting content with Ansible Vault

Encrypting individual variables with Ansible Vault

Encrypting files with Ansible Vault

Using encrypted variables and files

Passing a single password

Passing vault IDs

Passing multiple vault passwords

Using

--vault-id

without a vault ID

[Configuring defaults for using encrypted content](#)

[Setting a default vault ID](#)

[Setting a default password source](#)

[When are encrypted files made visible?](#)

[Format of files encrypted with Ansible Vault](#)

[Ansible Vault payload format 1.1 - 1.2](#)

[Previous](#)

[Next](#)

© Copyright Ansible project contributors.

Last updated on Oct 08, 2025.

Content from: <https://docs.ansible.com/users.html>

[Users | Ansible Documentation](#)

[Ansible documentation](#)

[Toggle navigation](#)

[Products](#)

[Blog](#)

[Community](#)

[Webinars and training](#)

[Try it now](#)

[Toggle navigation](#)

[Join the community](#)

[Users](#)

[Developers](#)

[Maintainers](#)

[Ansible core](#)

[Ansible ecosystem](#)

[Red Hat Ansible Automation Platform](#)

[Users](#)

Automate the management of remote systems and control their desired state.

Got thoughts or feedback on this site? We want to hear from you!

Join us in the

[Ansible Forum](#)

or open a

[GitHub issue](#)

in the docsite repository.

[Quicklinks](#)

[YAML syntax](#)

[Playbook variables](#)

[Playbook conditionals](#)

[Create automation](#)

[Start writing Ansible playbooks](#)

[Learn about Ansible modules](#)

[Build inventories](#)

[Build inventory files to manage multiple hosts](#)

[Use dynamic inventories](#)

[Organize automation projects](#)

[Use roles to structure the automation project](#)

[Use Ansible execution environments](#)

[Get started with execution environments](#)

[Build execution environments](#)

[Use Ansible tooling](#)

[Create and test playbooks with Ansible Navigator](#)

[Use Ansible Lint to validate playbooks](#)

[Install Molecule to develop and test Ansible roles](#)

[Use Ansible with Visual Studio Code and OpenVSX compatible editors](#)

[Find automation content](#)

[Start exploring Ansible Galaxy](#)

[Install and use roles](#)

[Install and use collections](#)

[Share automation content](#)

[Submit roles to an existing collection](#)

[Create a new collection](#)

[Upload a collection to Ansible Galaxy](#)

Schedule and run automation jobs

Execute automation jobs on demand

Schedule automation jobs

Use execution environments with AWX jobs

[CC BY-SA 4.0](#) |

[Privacy policy](#) |

Sponsored by

